

Generative adversarial networks

Christopher Beckham

Montréal Institute for Learning Algorithms

github: christopher-beckham

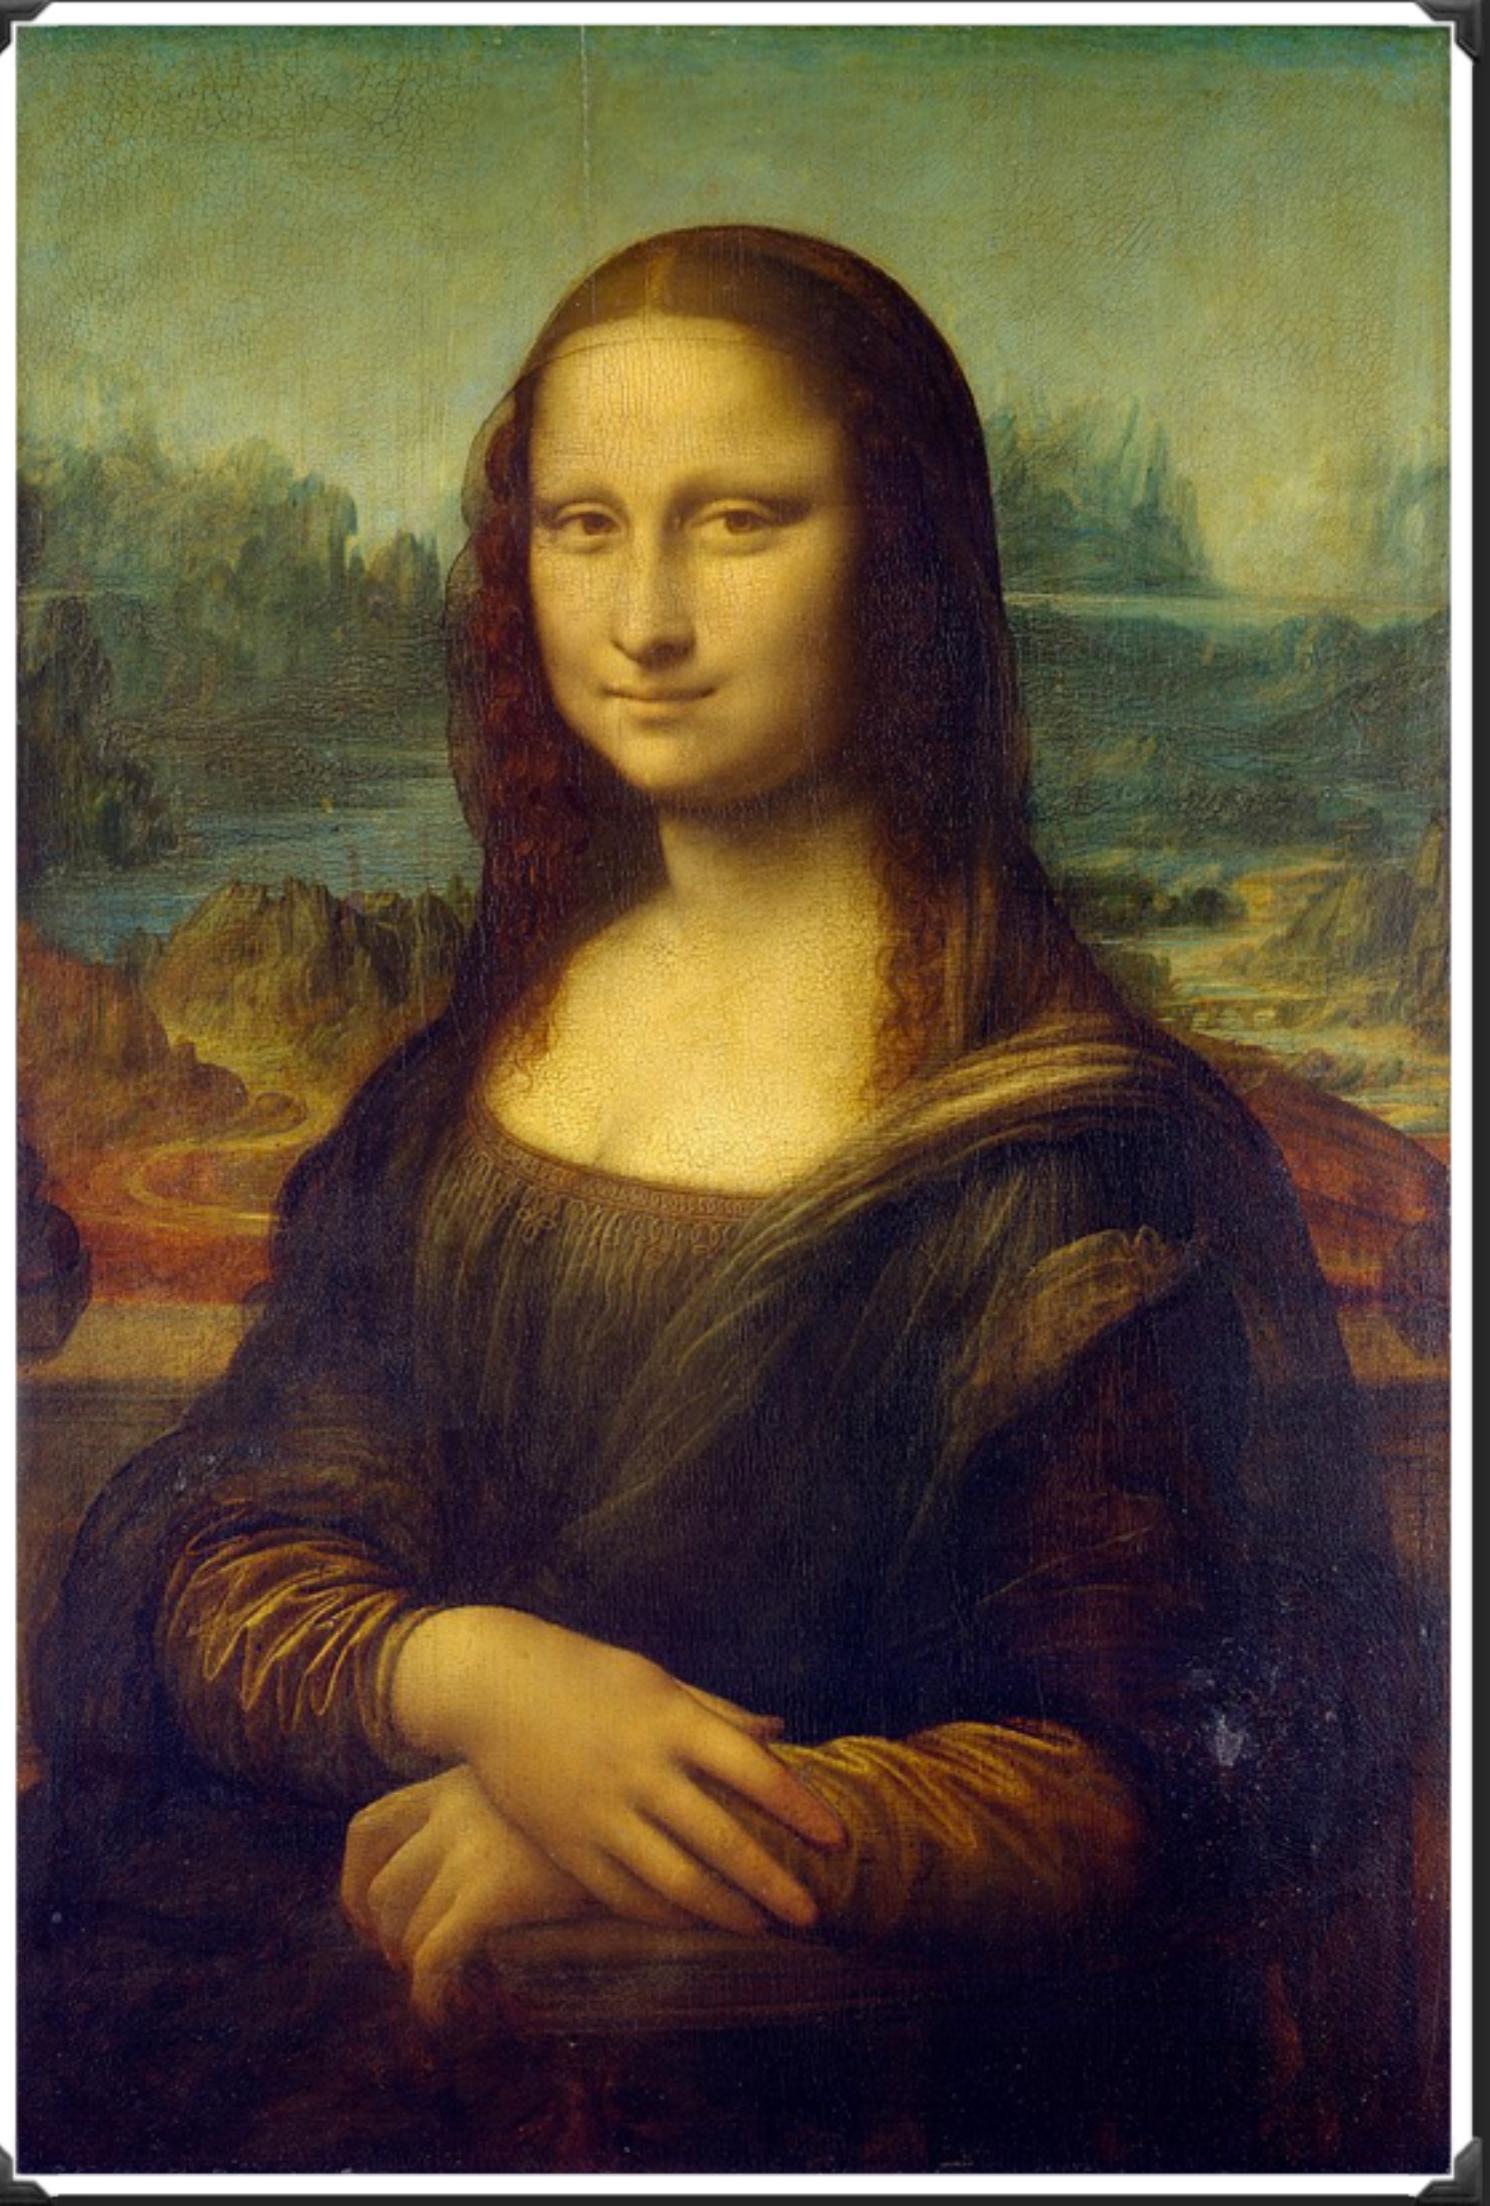
twitter: chris_j_beckham



"What I cannot create, I do
not understand"

- Richard Feynman

https://en.wikipedia.org/wiki/File:Mona_Lisa,_by_Leonardo_da_Vinci,_from_C2RMF_retoucheda.jpg
https://static.giantbomb.com/uploads/scale_medium/0/4604/967633-739px_taz.png
<http://poussiecat.eklablog.com/great-mouse-detective-c26082364>



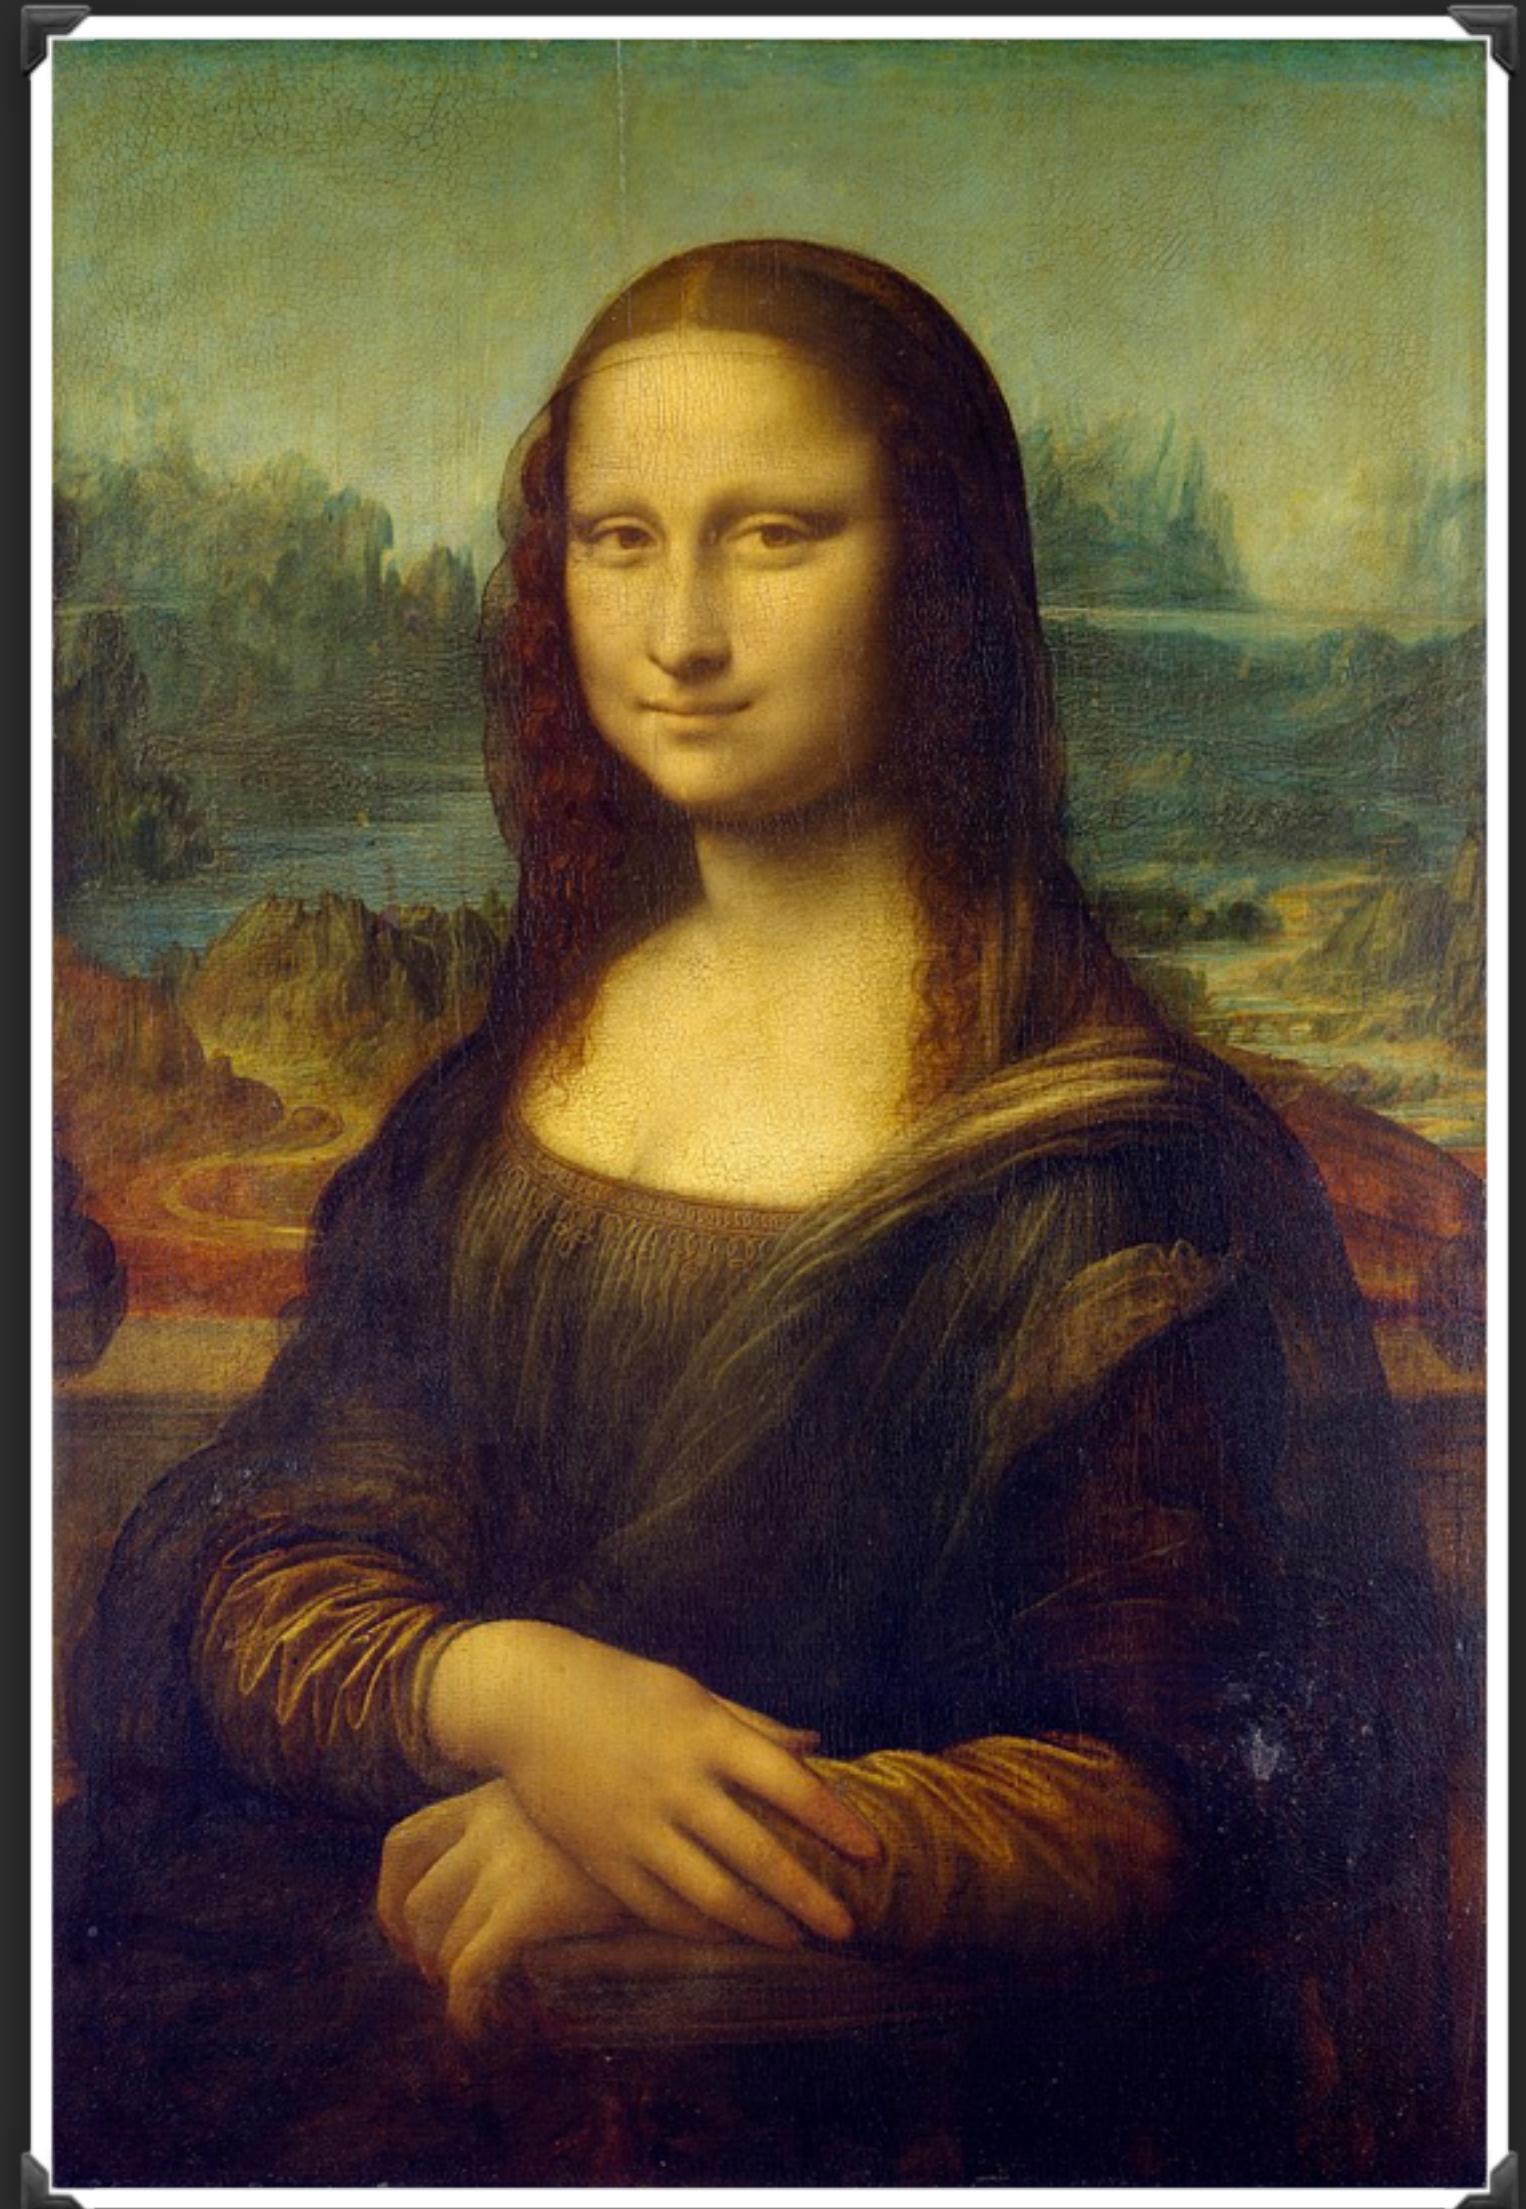
https://en.wikipedia.org/wiki/File:Mona_Lisa,_by_Leonardo_da_Vinci,_from_C2RMF_retouch.jpg

https://static.giantbomb.com/uploads/scale_medium/0/4604/967633-739px_taz.png

<http://poussiecat.eklablog.com/great-mouse-detective-c26082364>



discriminator



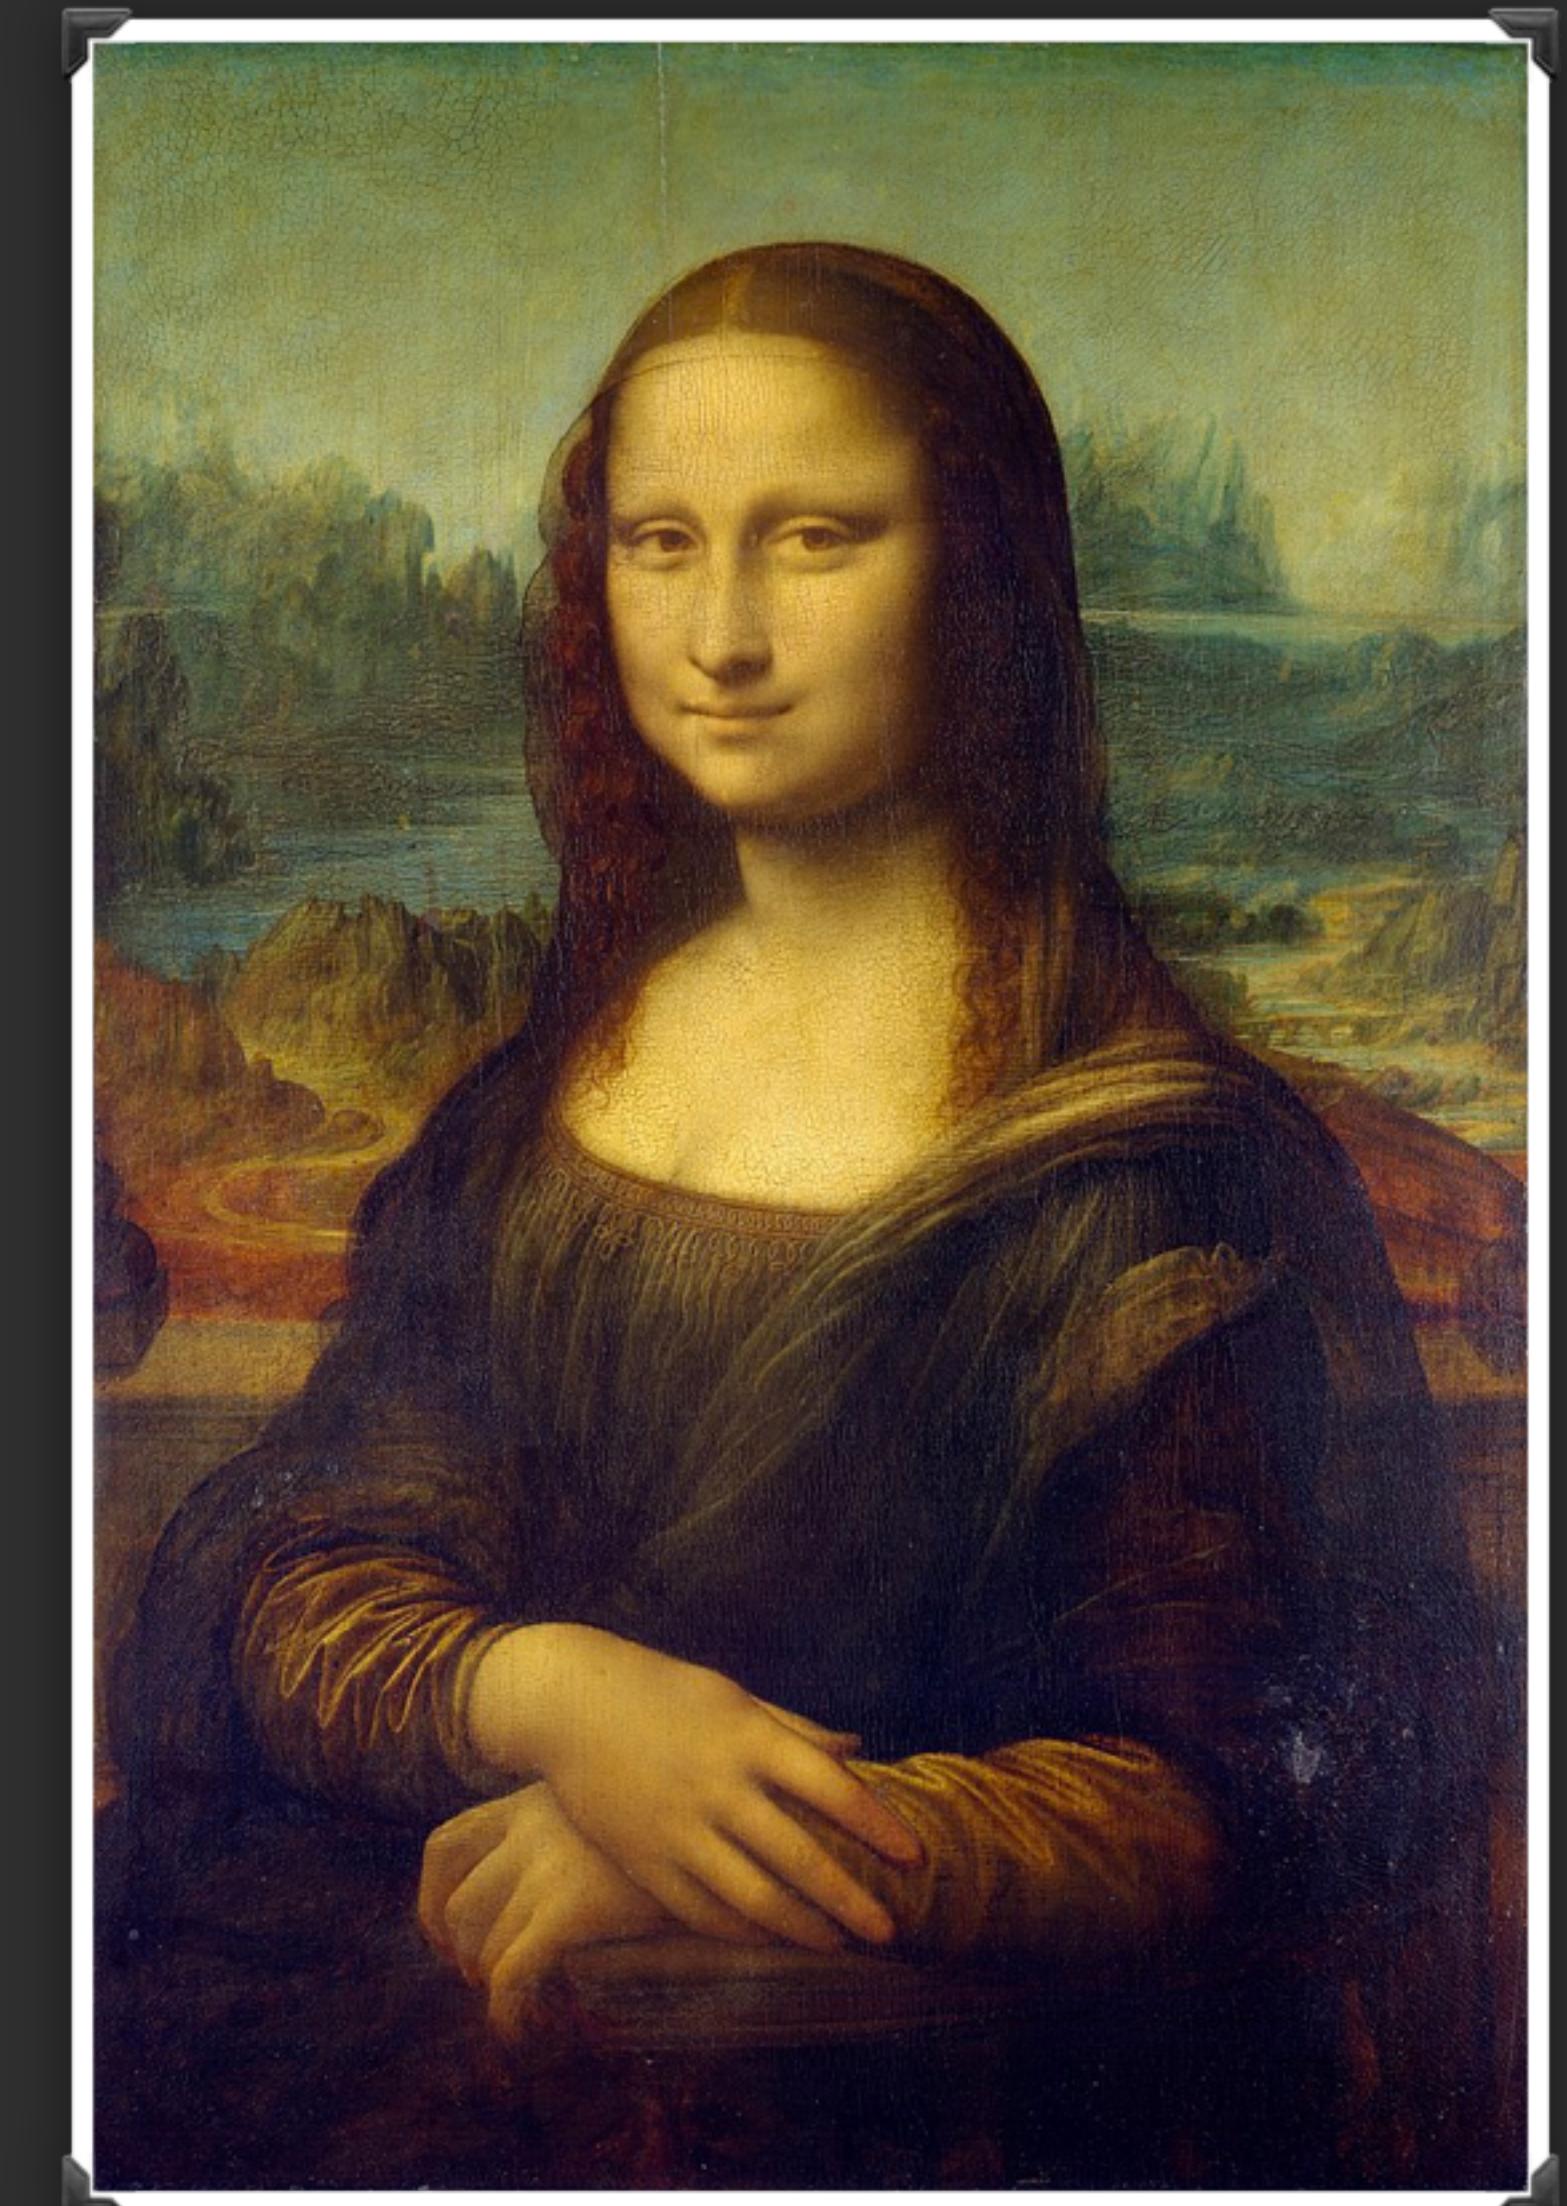
https://en.wikipedia.org/wiki/File:Mona_Lisa,_by_Leonardo_da_Vinci,_from_C2RMF_retoucheda.jpg

https://static.giantbomb.com/uploads/scale_medium/0/4604/967633-739px_taz.png

<http://poussiecat.eklablog.com/great-mouse-detective-c26082364>



discriminator



https://en.wikipedia.org/wiki/File:Mona_Lisa,_by_Leonardo_da_Vinci,_from_C2RMF_retoucheda.jpg

https://static.giantbomb.com/uploads/scale_medium/0/4604/967633-739px_taz.png

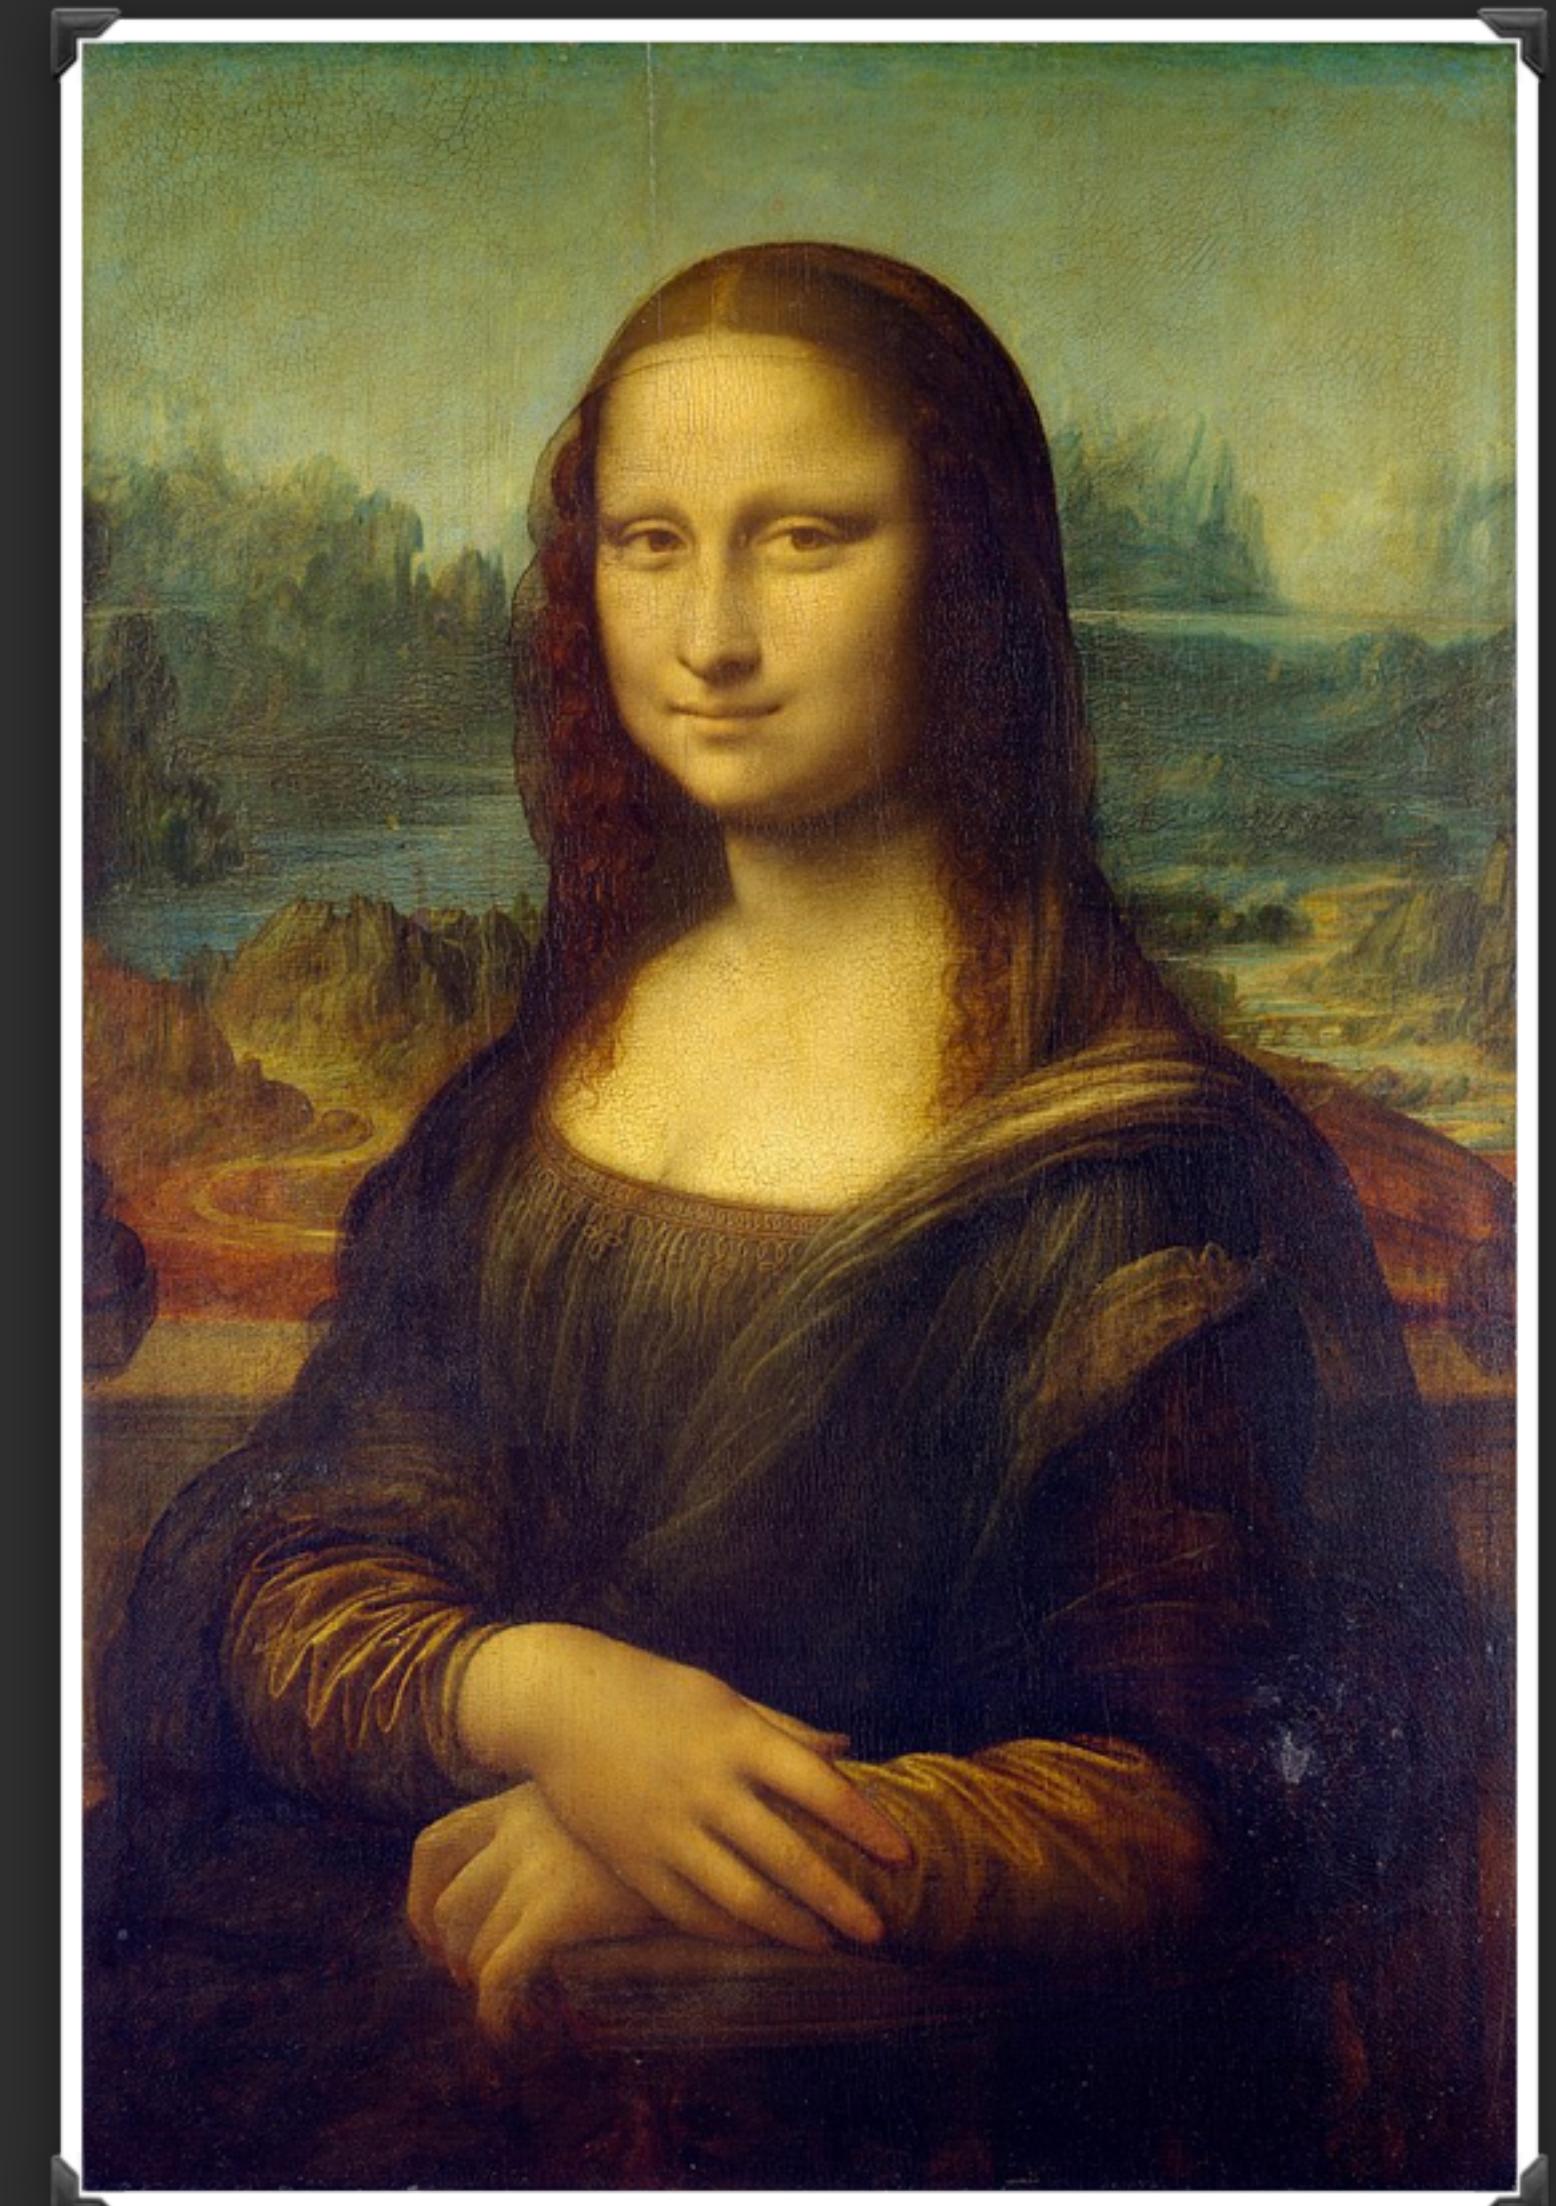
<http://poussiecat.eklablog.com/great-mouse-detective-c26082364>



discriminator



generator



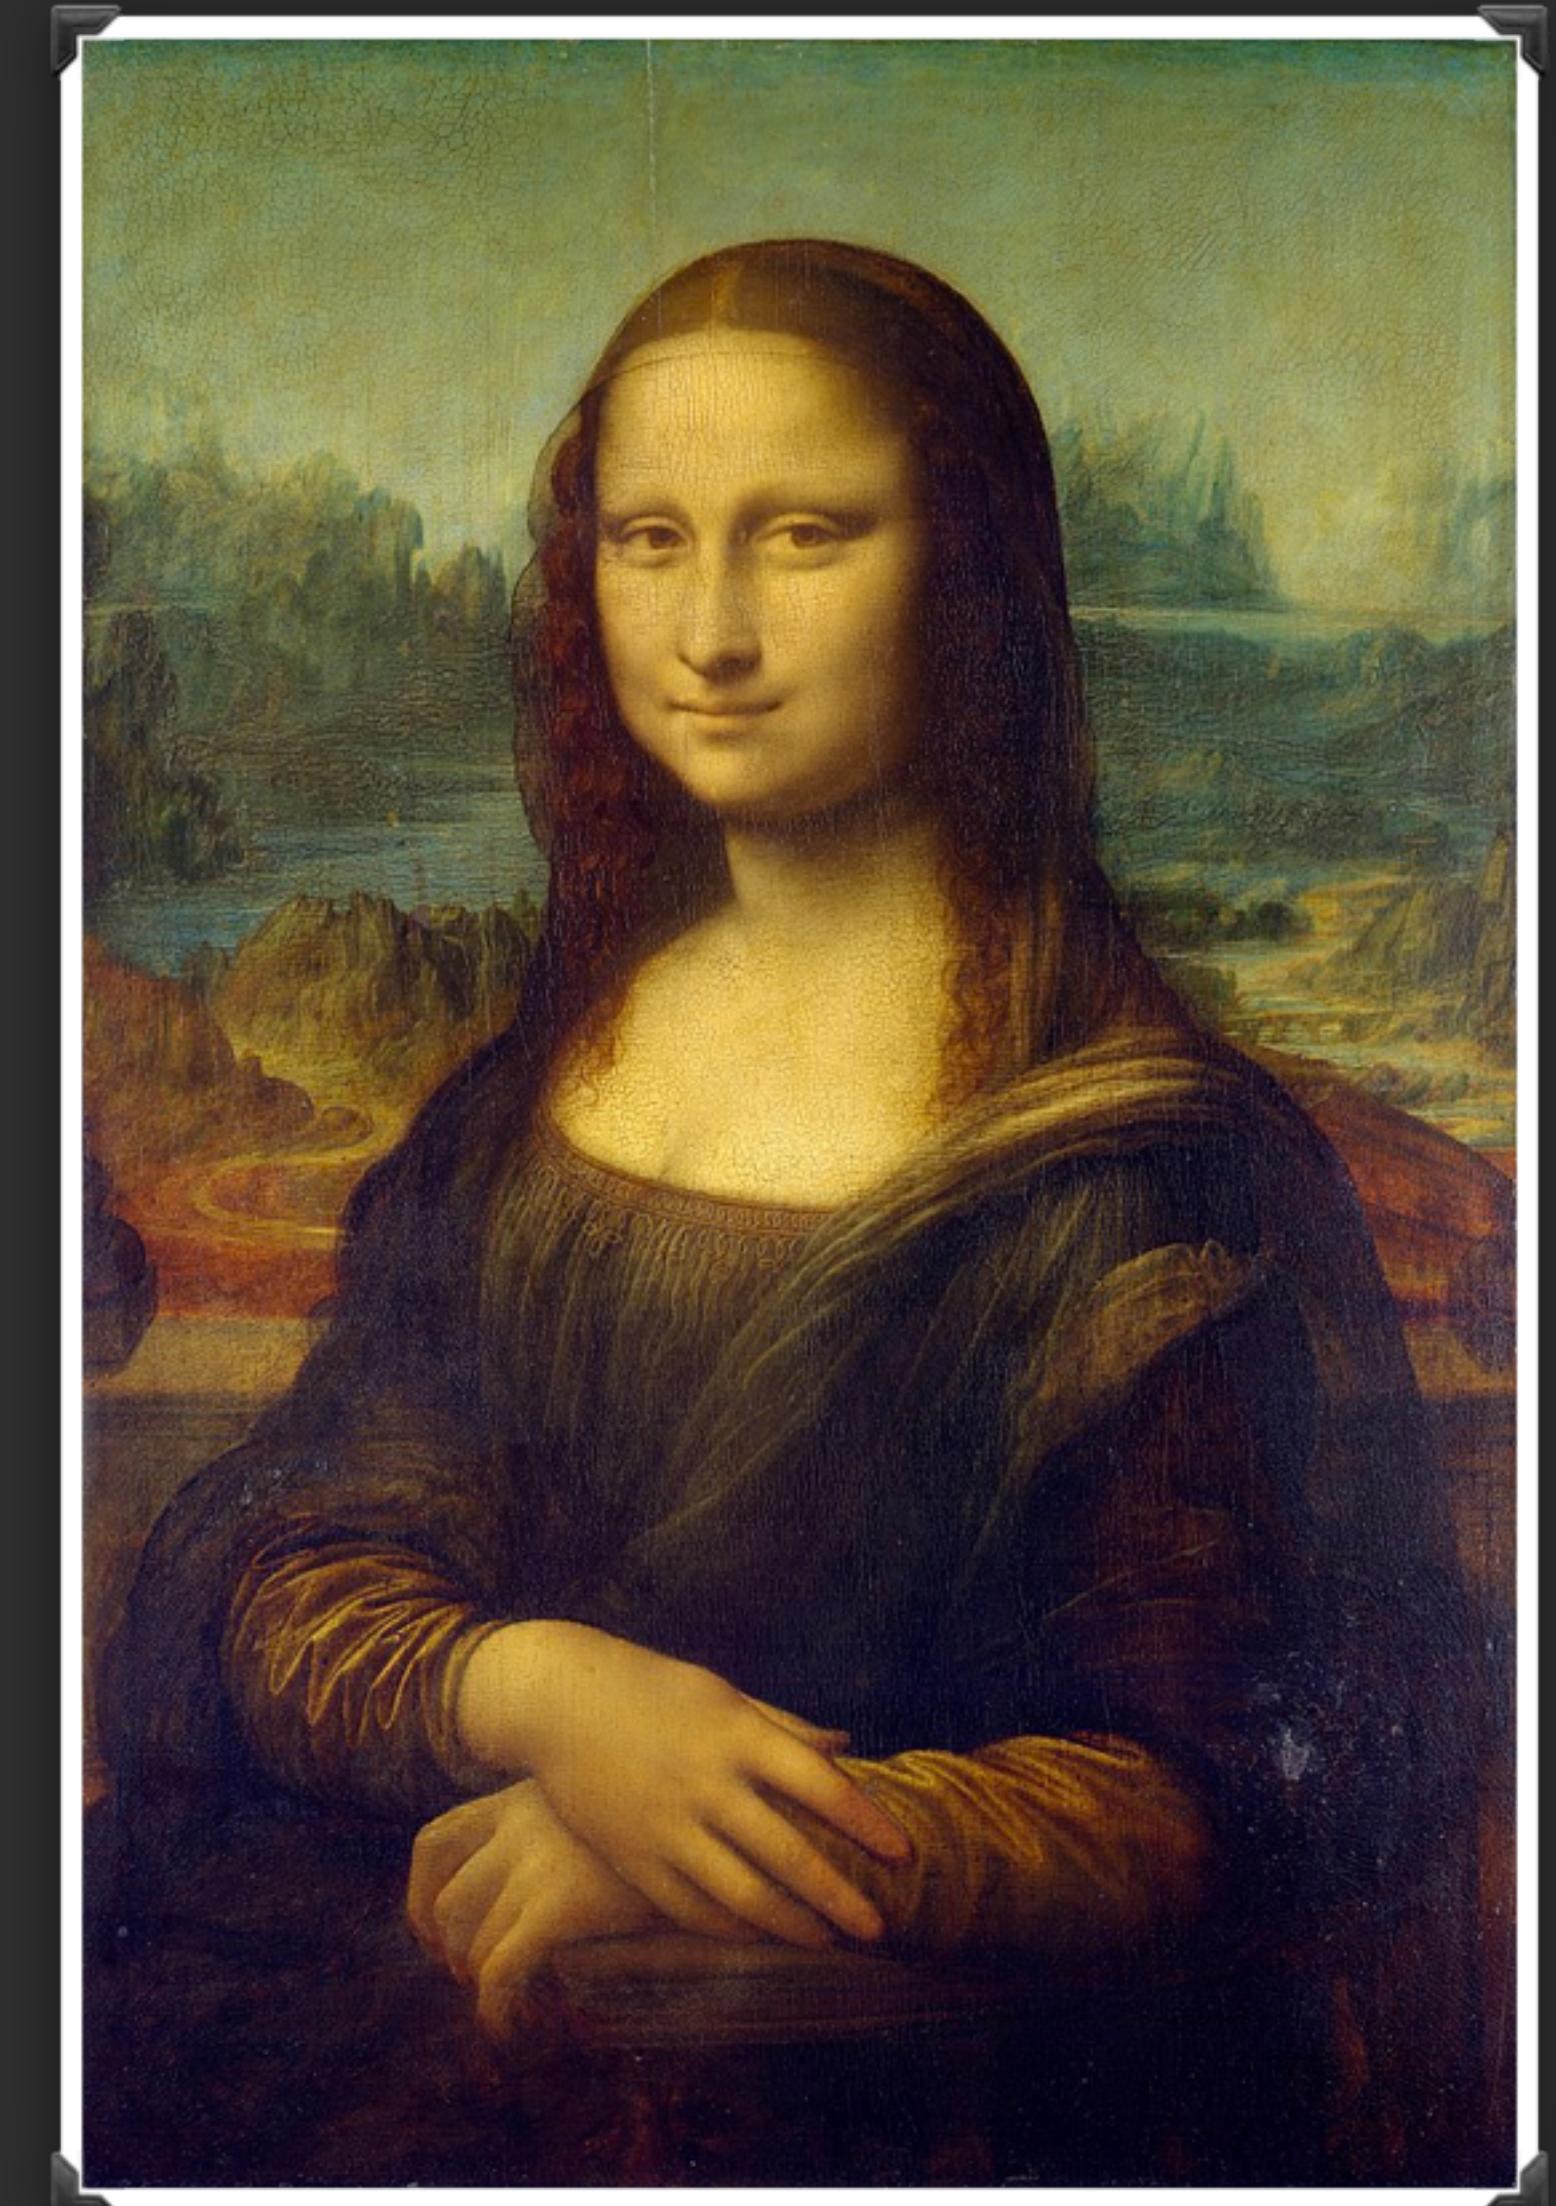
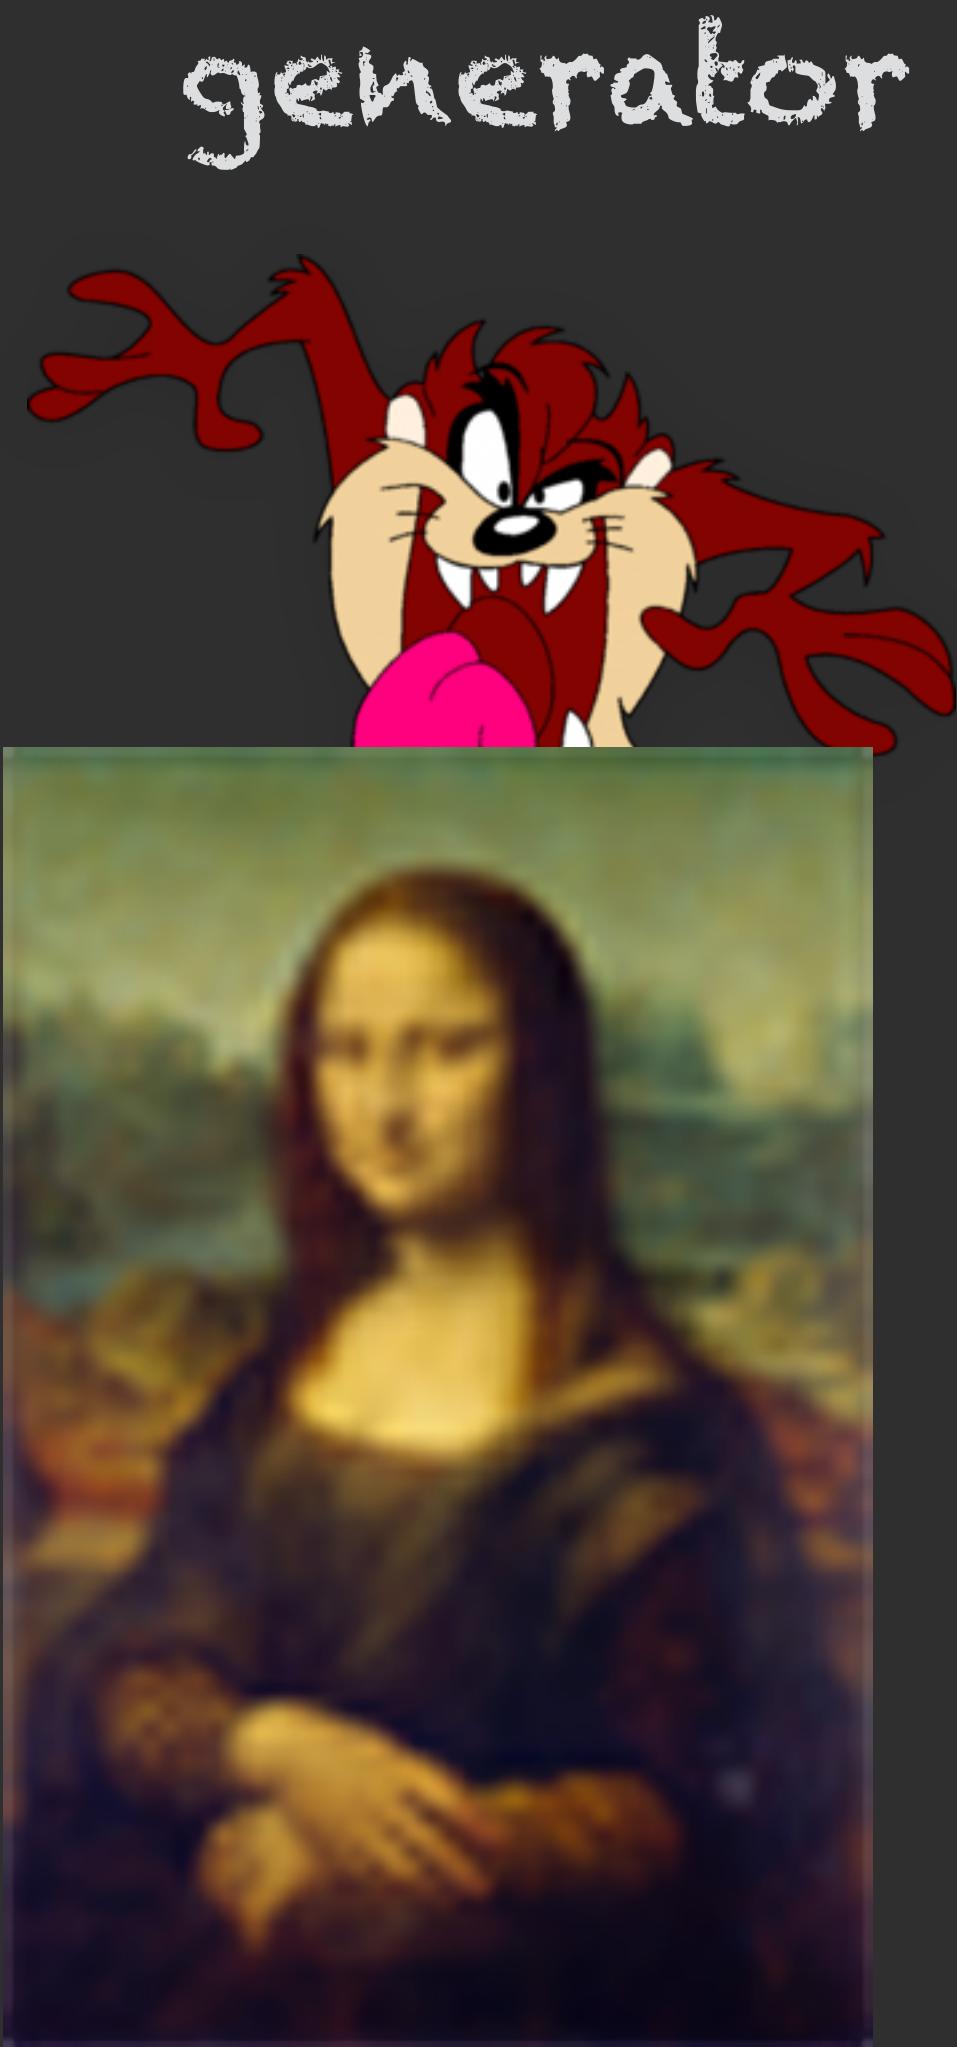
https://en.wikipedia.org/wiki/File:Mona_Lisa,_by_Leonardo_da_Vinci,_from_C2RMF_retoucheda.jpg

https://static.giantbomb.com/uploads/scale_medium/0/4604/967633-739px_taz.png

<http://poussiecat.eklablog.com/great-mouse-detective-c26082364>

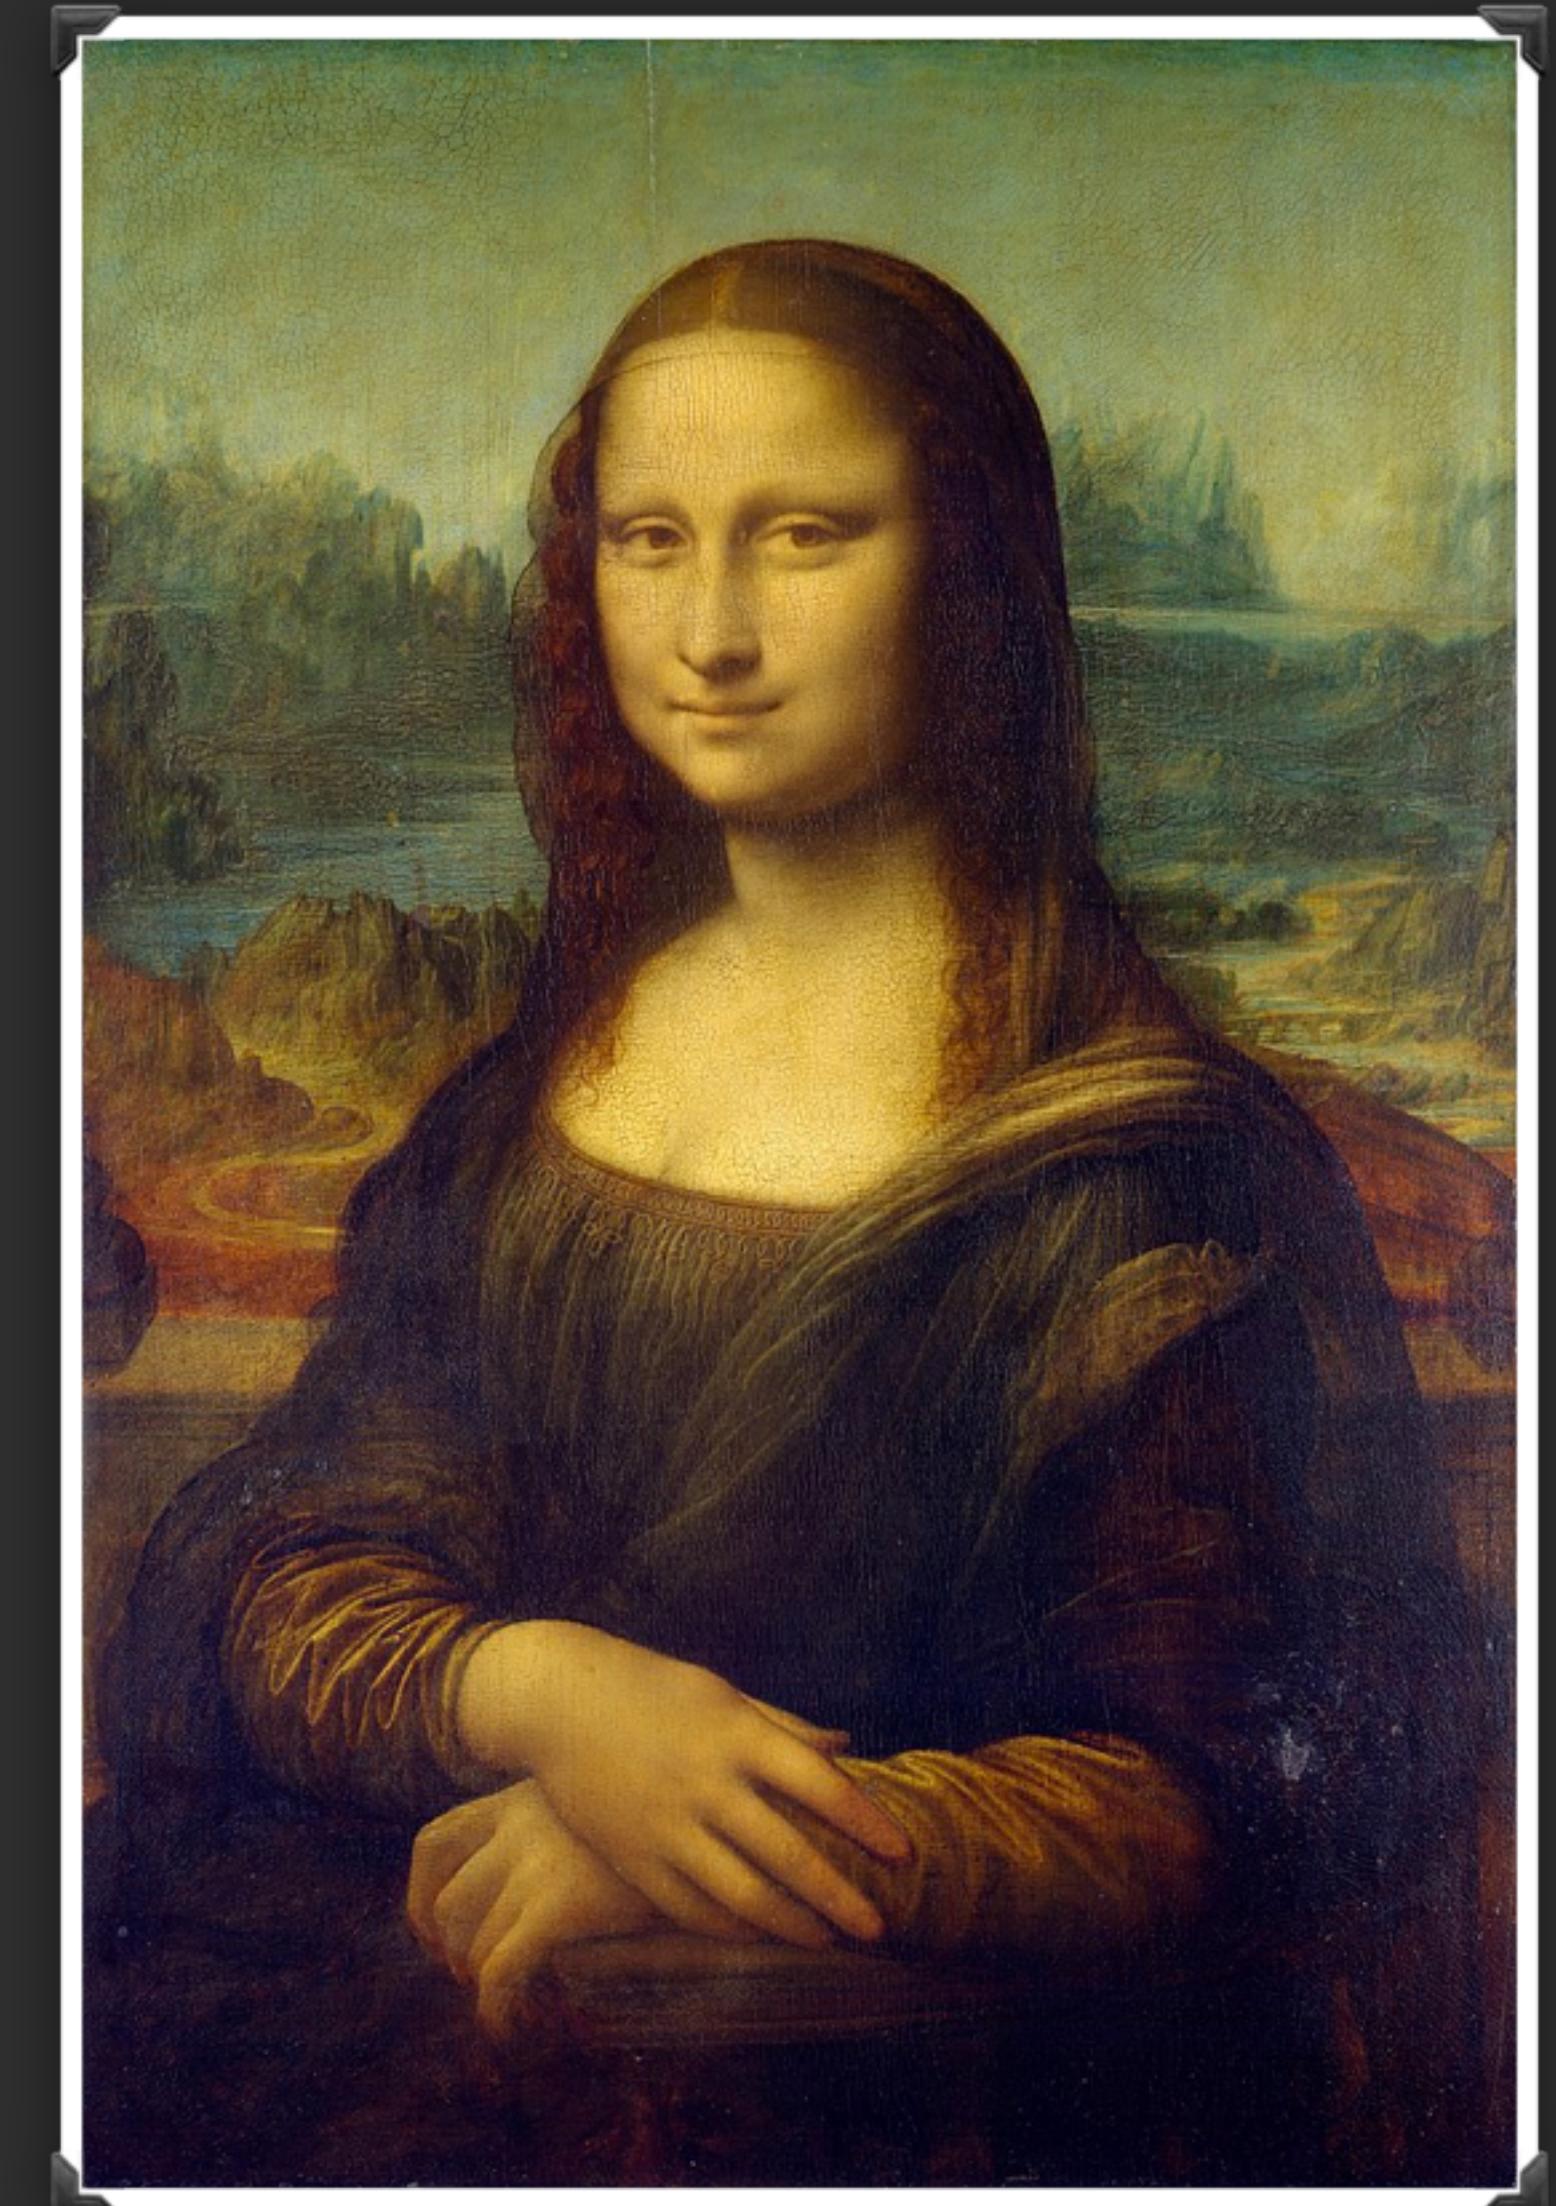
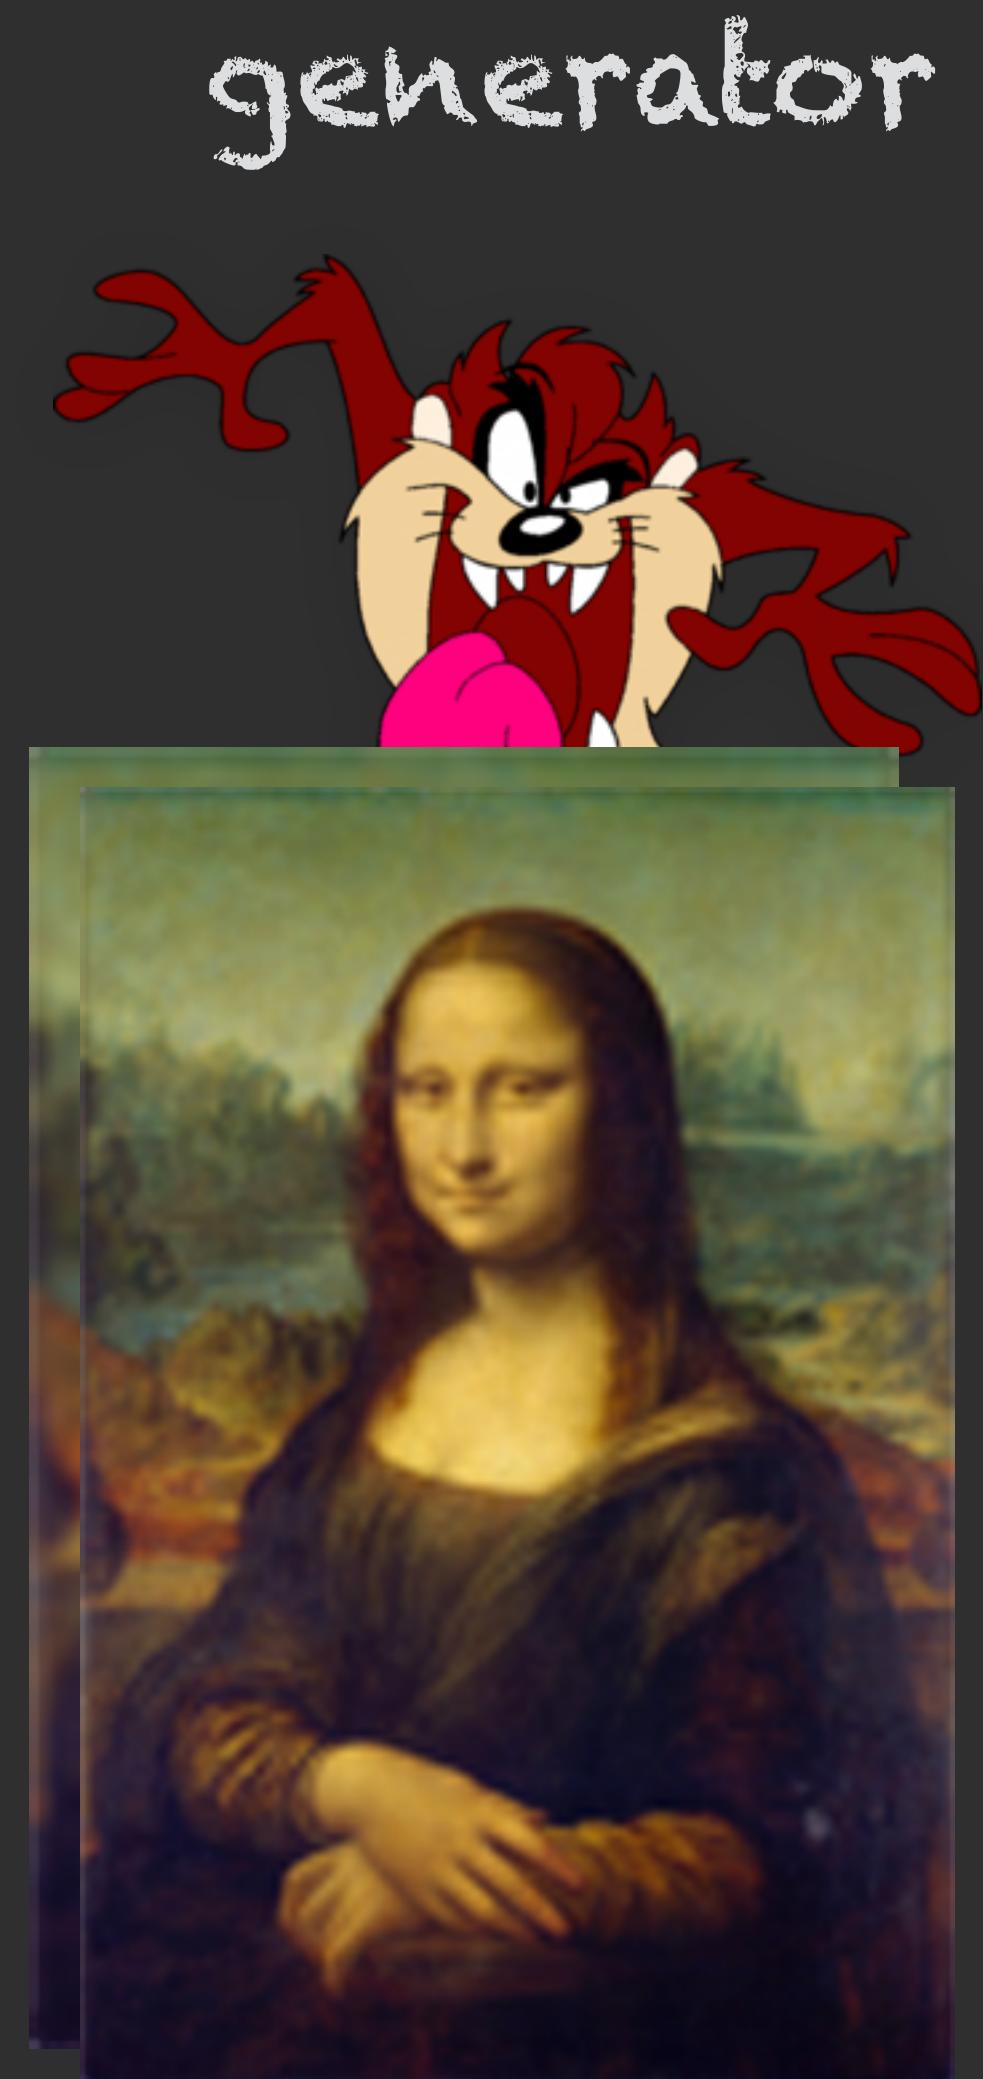


discriminator





discriminator



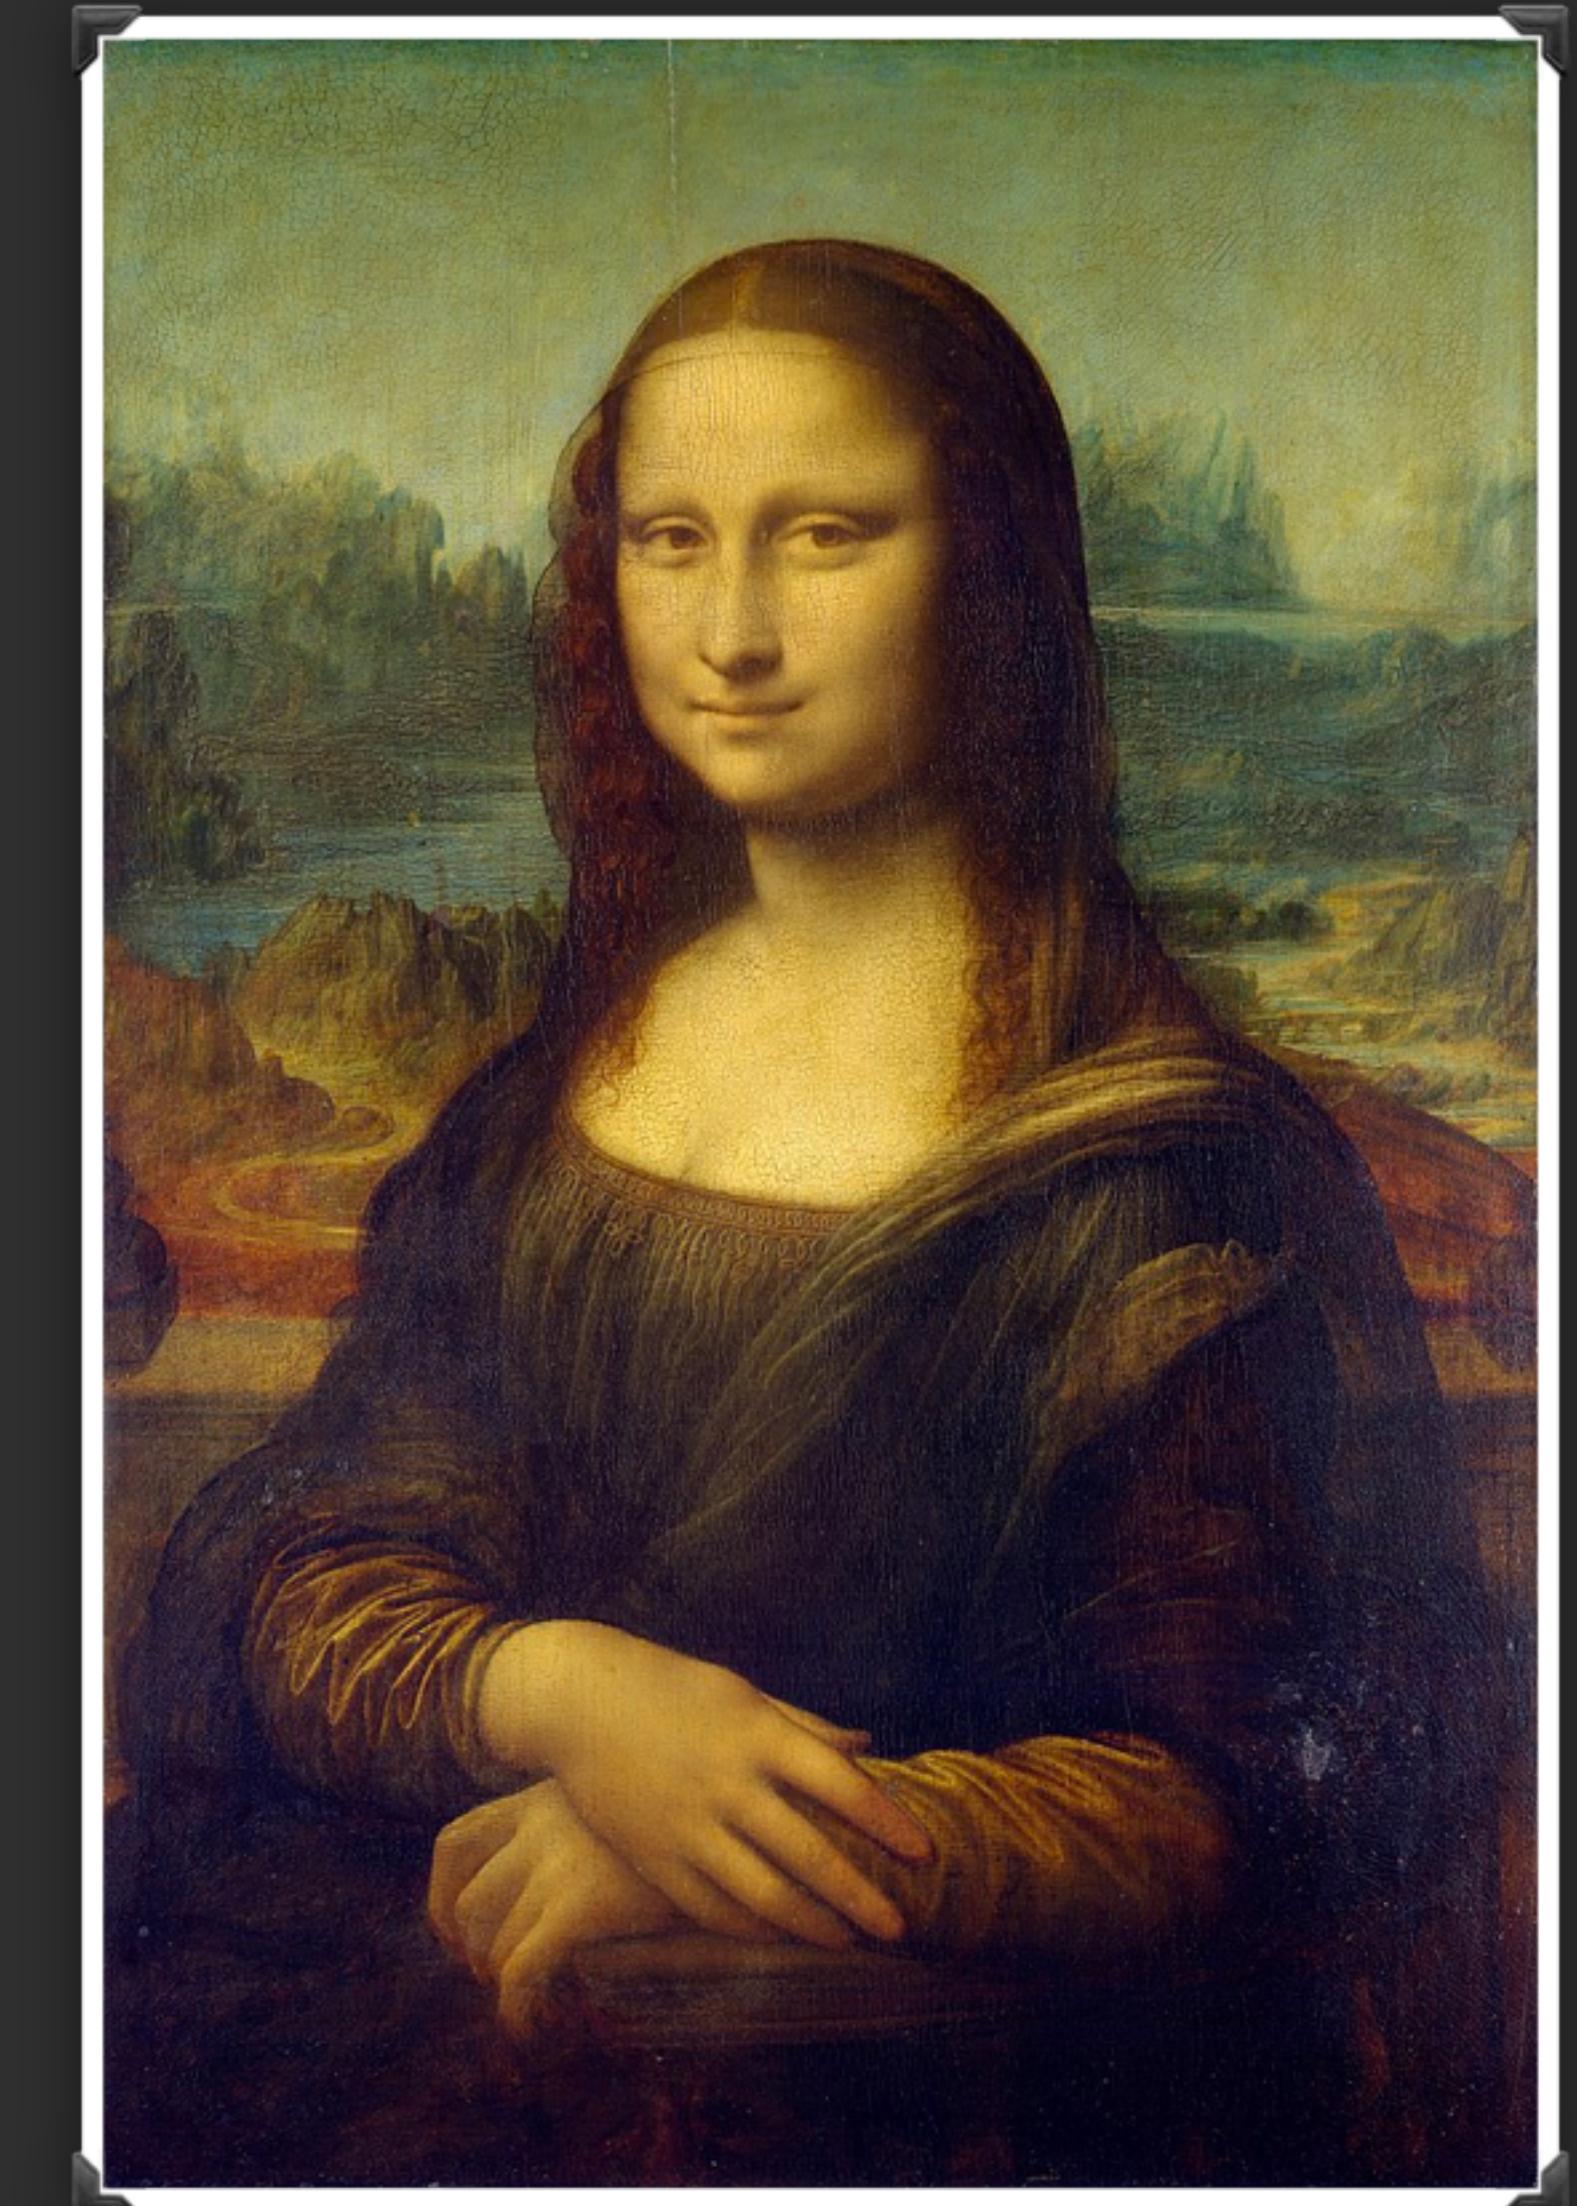
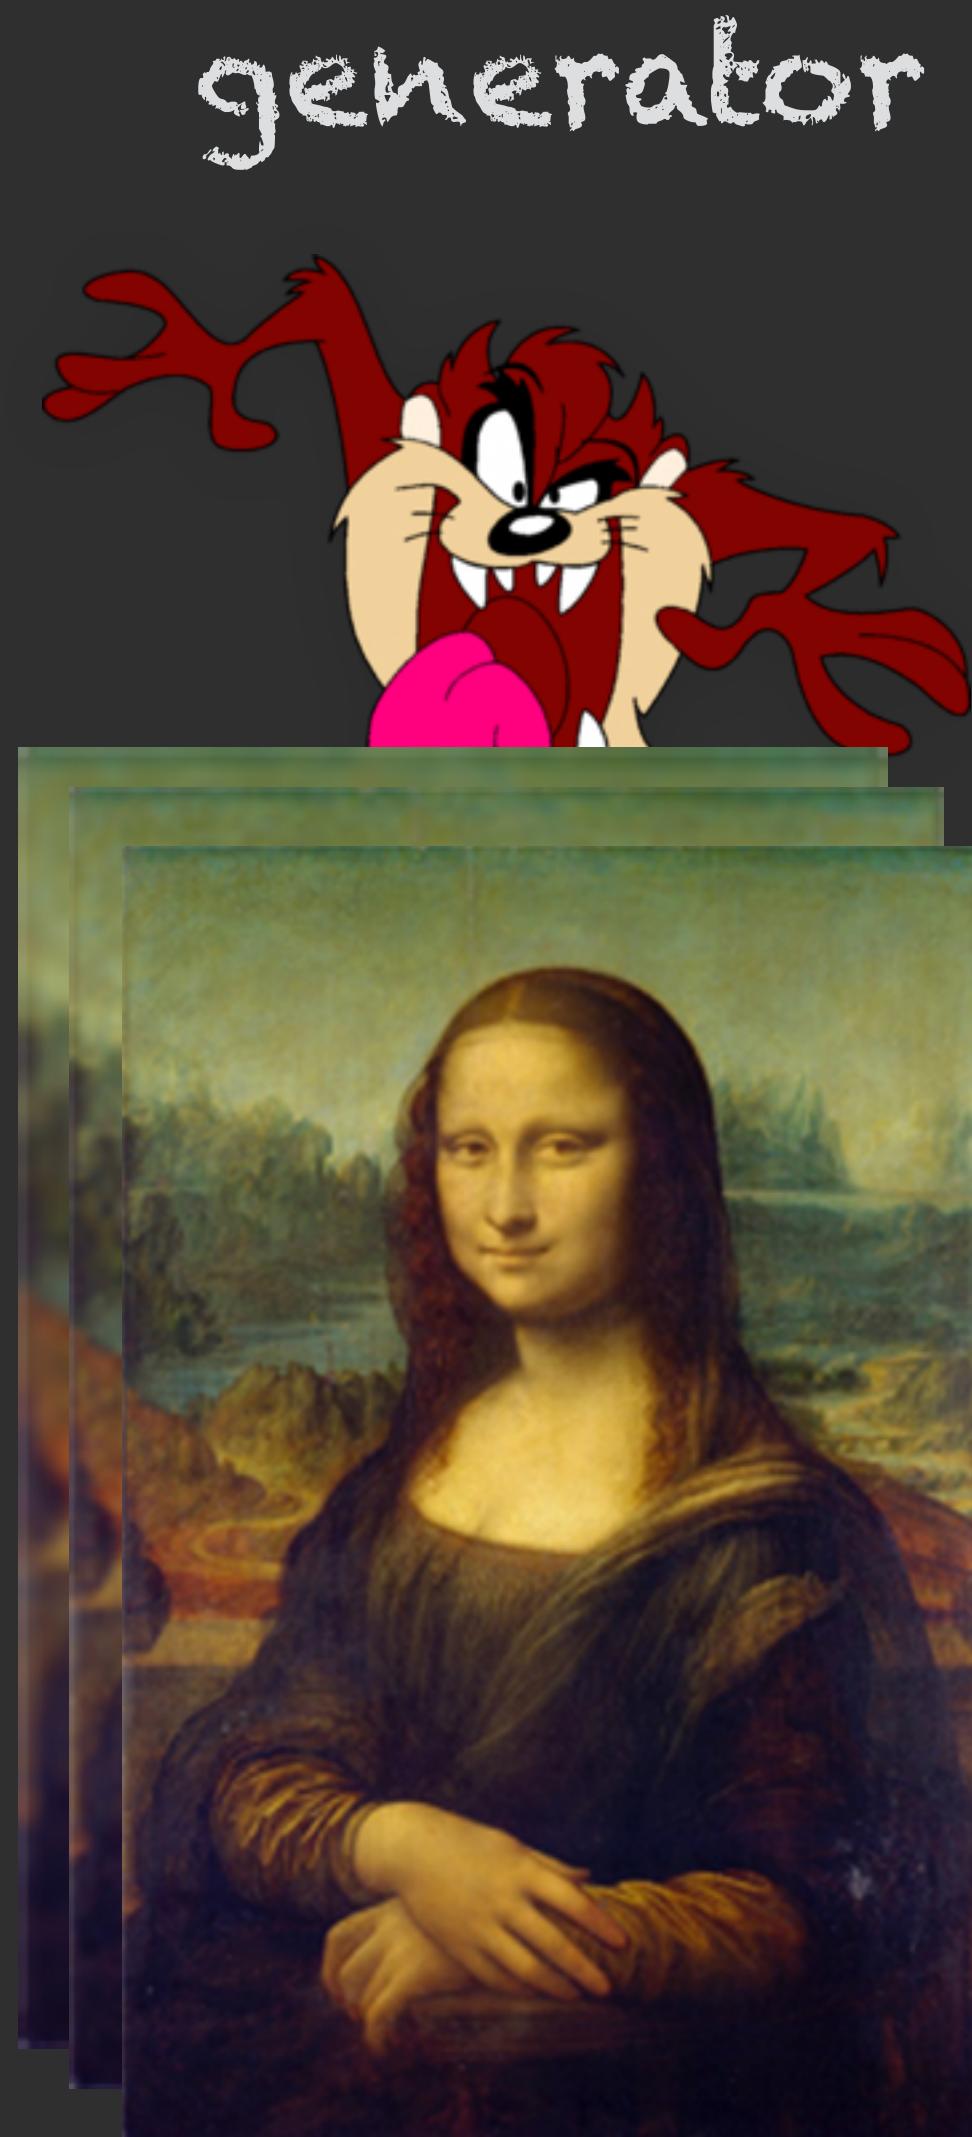
https://en.wikipedia.org/wiki/File:Mona_Lisa,_by_Leonardo_da_Vinci,_from_C2RMF_retouch.jpg

https://static.giantbomb.com/uploads/scale_medium/0/4604/967633-739px_taz.png

<http://poussiecat.eklablog.com/great-mouse-detective-c26082364>



discriminator



https://en.wikipedia.org/wiki/File:Mona_Lisa,_by_Leonardo_da_Vinci,_from_C2RMF_retouchd.jpg

https://static.giantbomb.com/uploads/scale_medium/0/4604/967633-739px_taz.png

<http://poussiecat.eklablog.com/great-mouse-detective-c26082364>

What are generative adversarial networks?

What are generative adversarial networks?

- GANs are a class of algorithms used in unsupervised learning in which the goal is to learn (implicitly) the data distribution

What are generative adversarial networks?

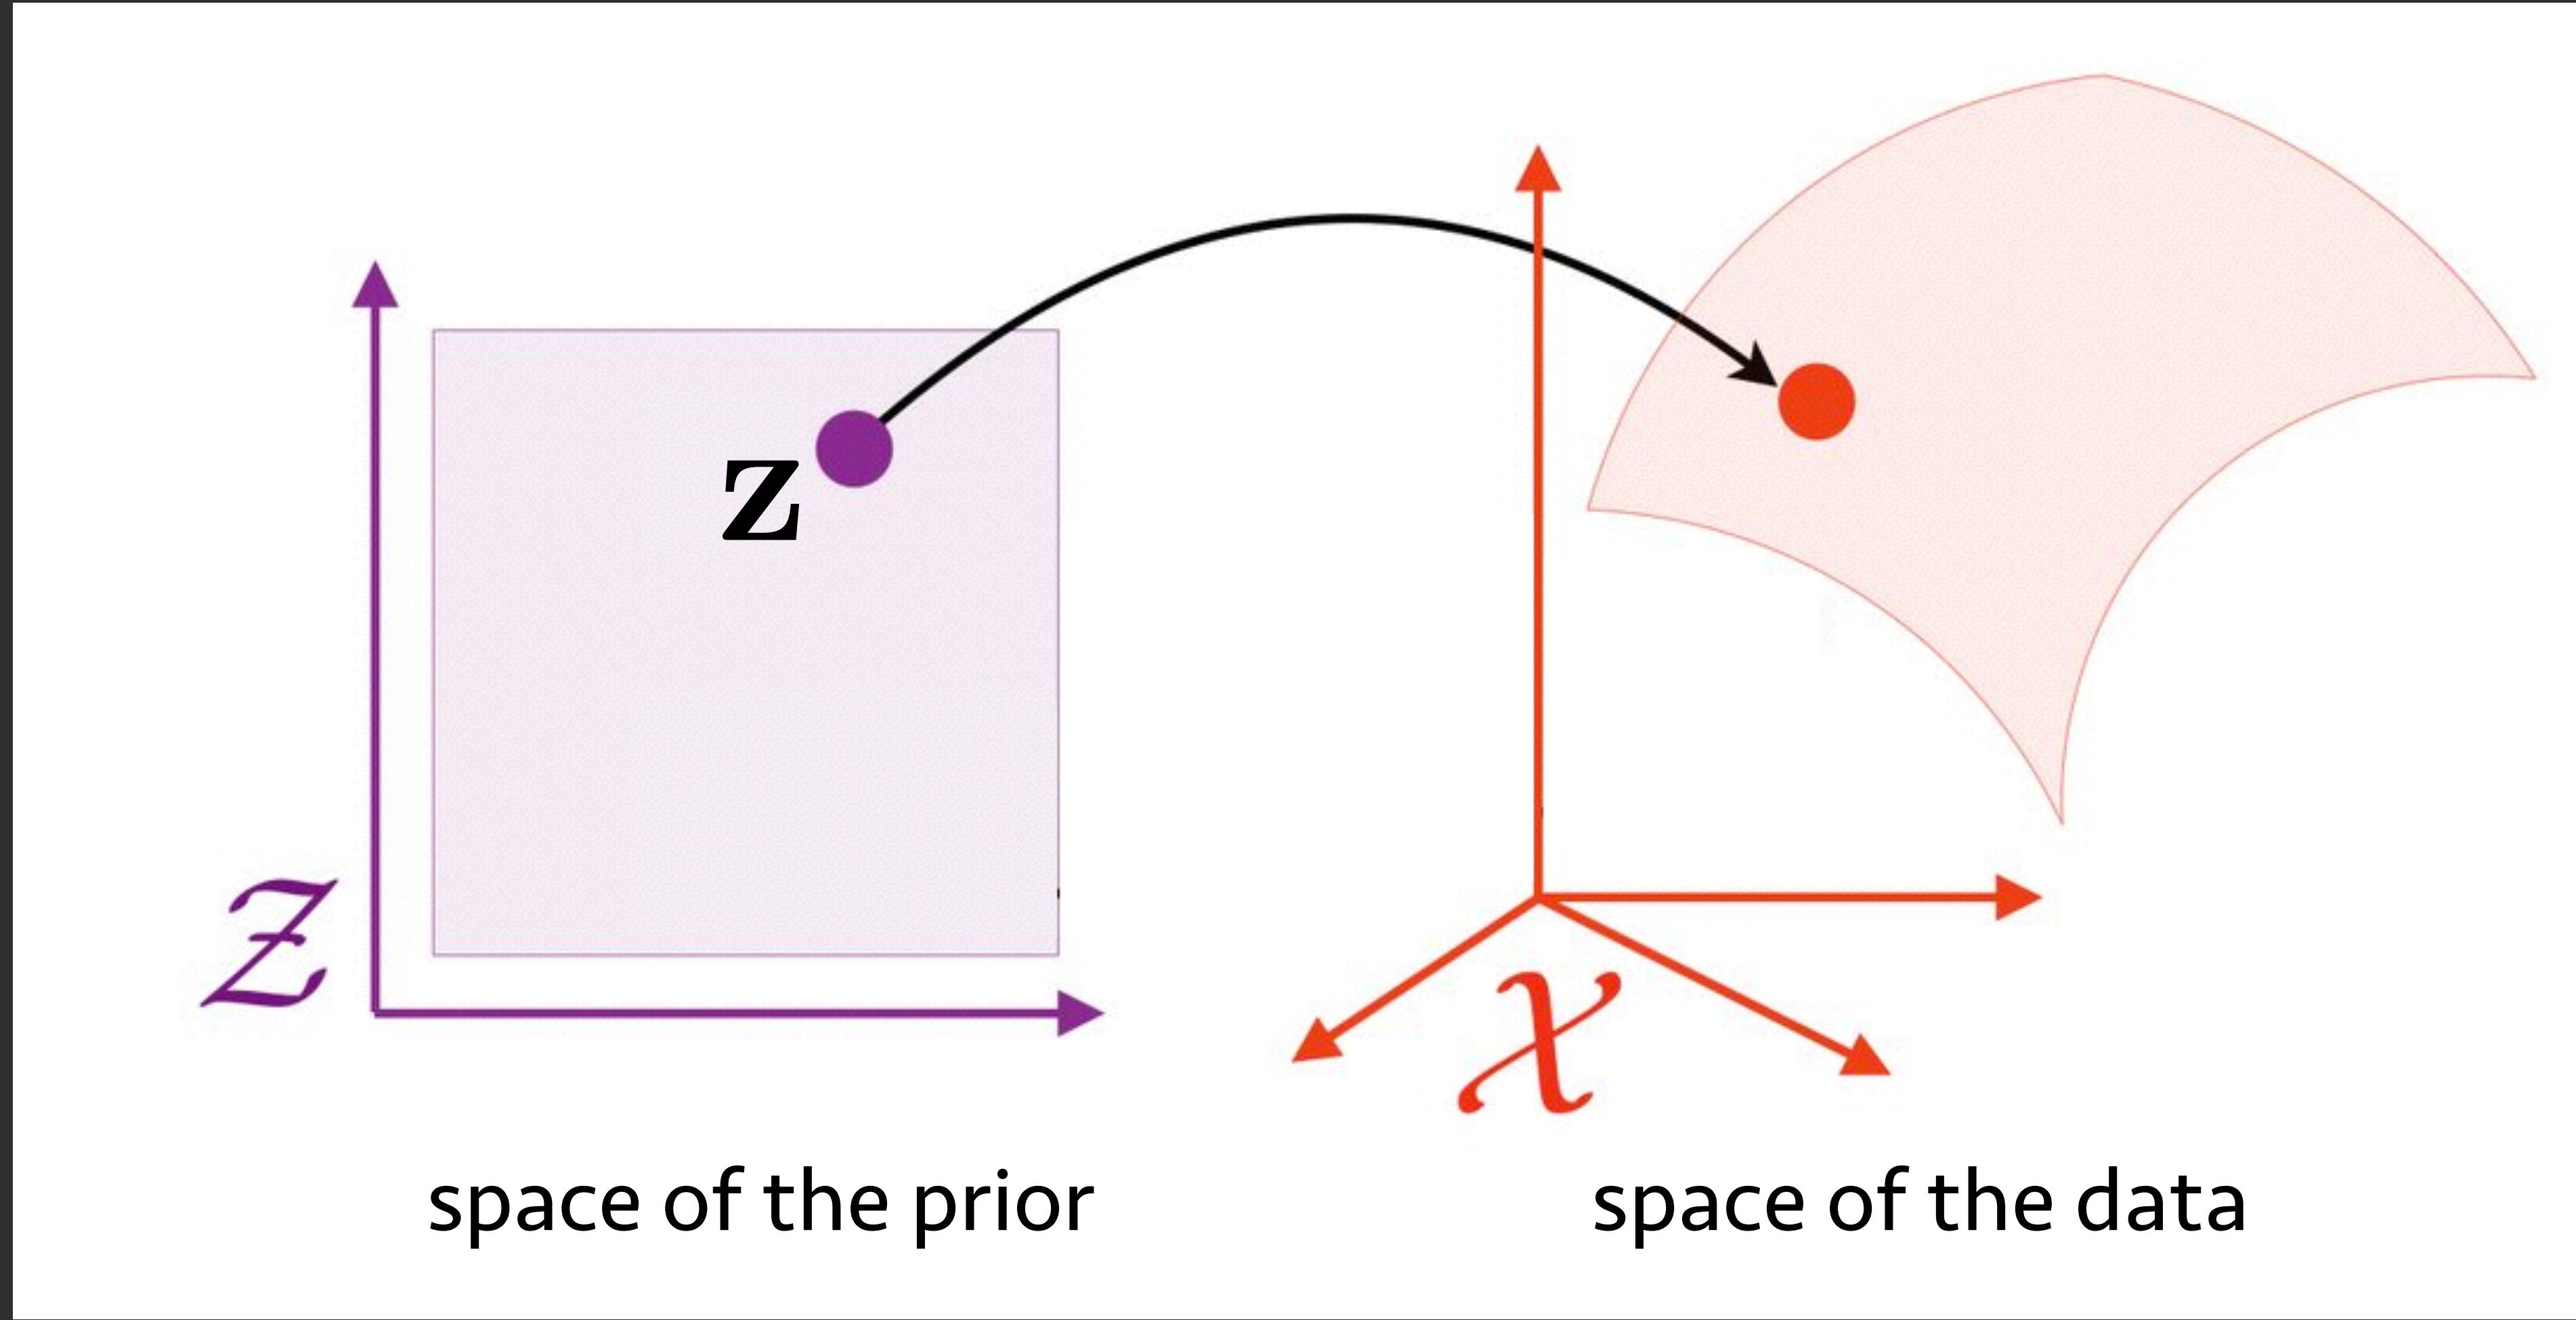
- GANs are a class of algorithms used in unsupervised learning in which the goal is to learn (implicitly) the data distribution
- This is done through two adversaries (generator and discriminator) which compute against each other in a zero-sum game

What are generative adversarial networks?

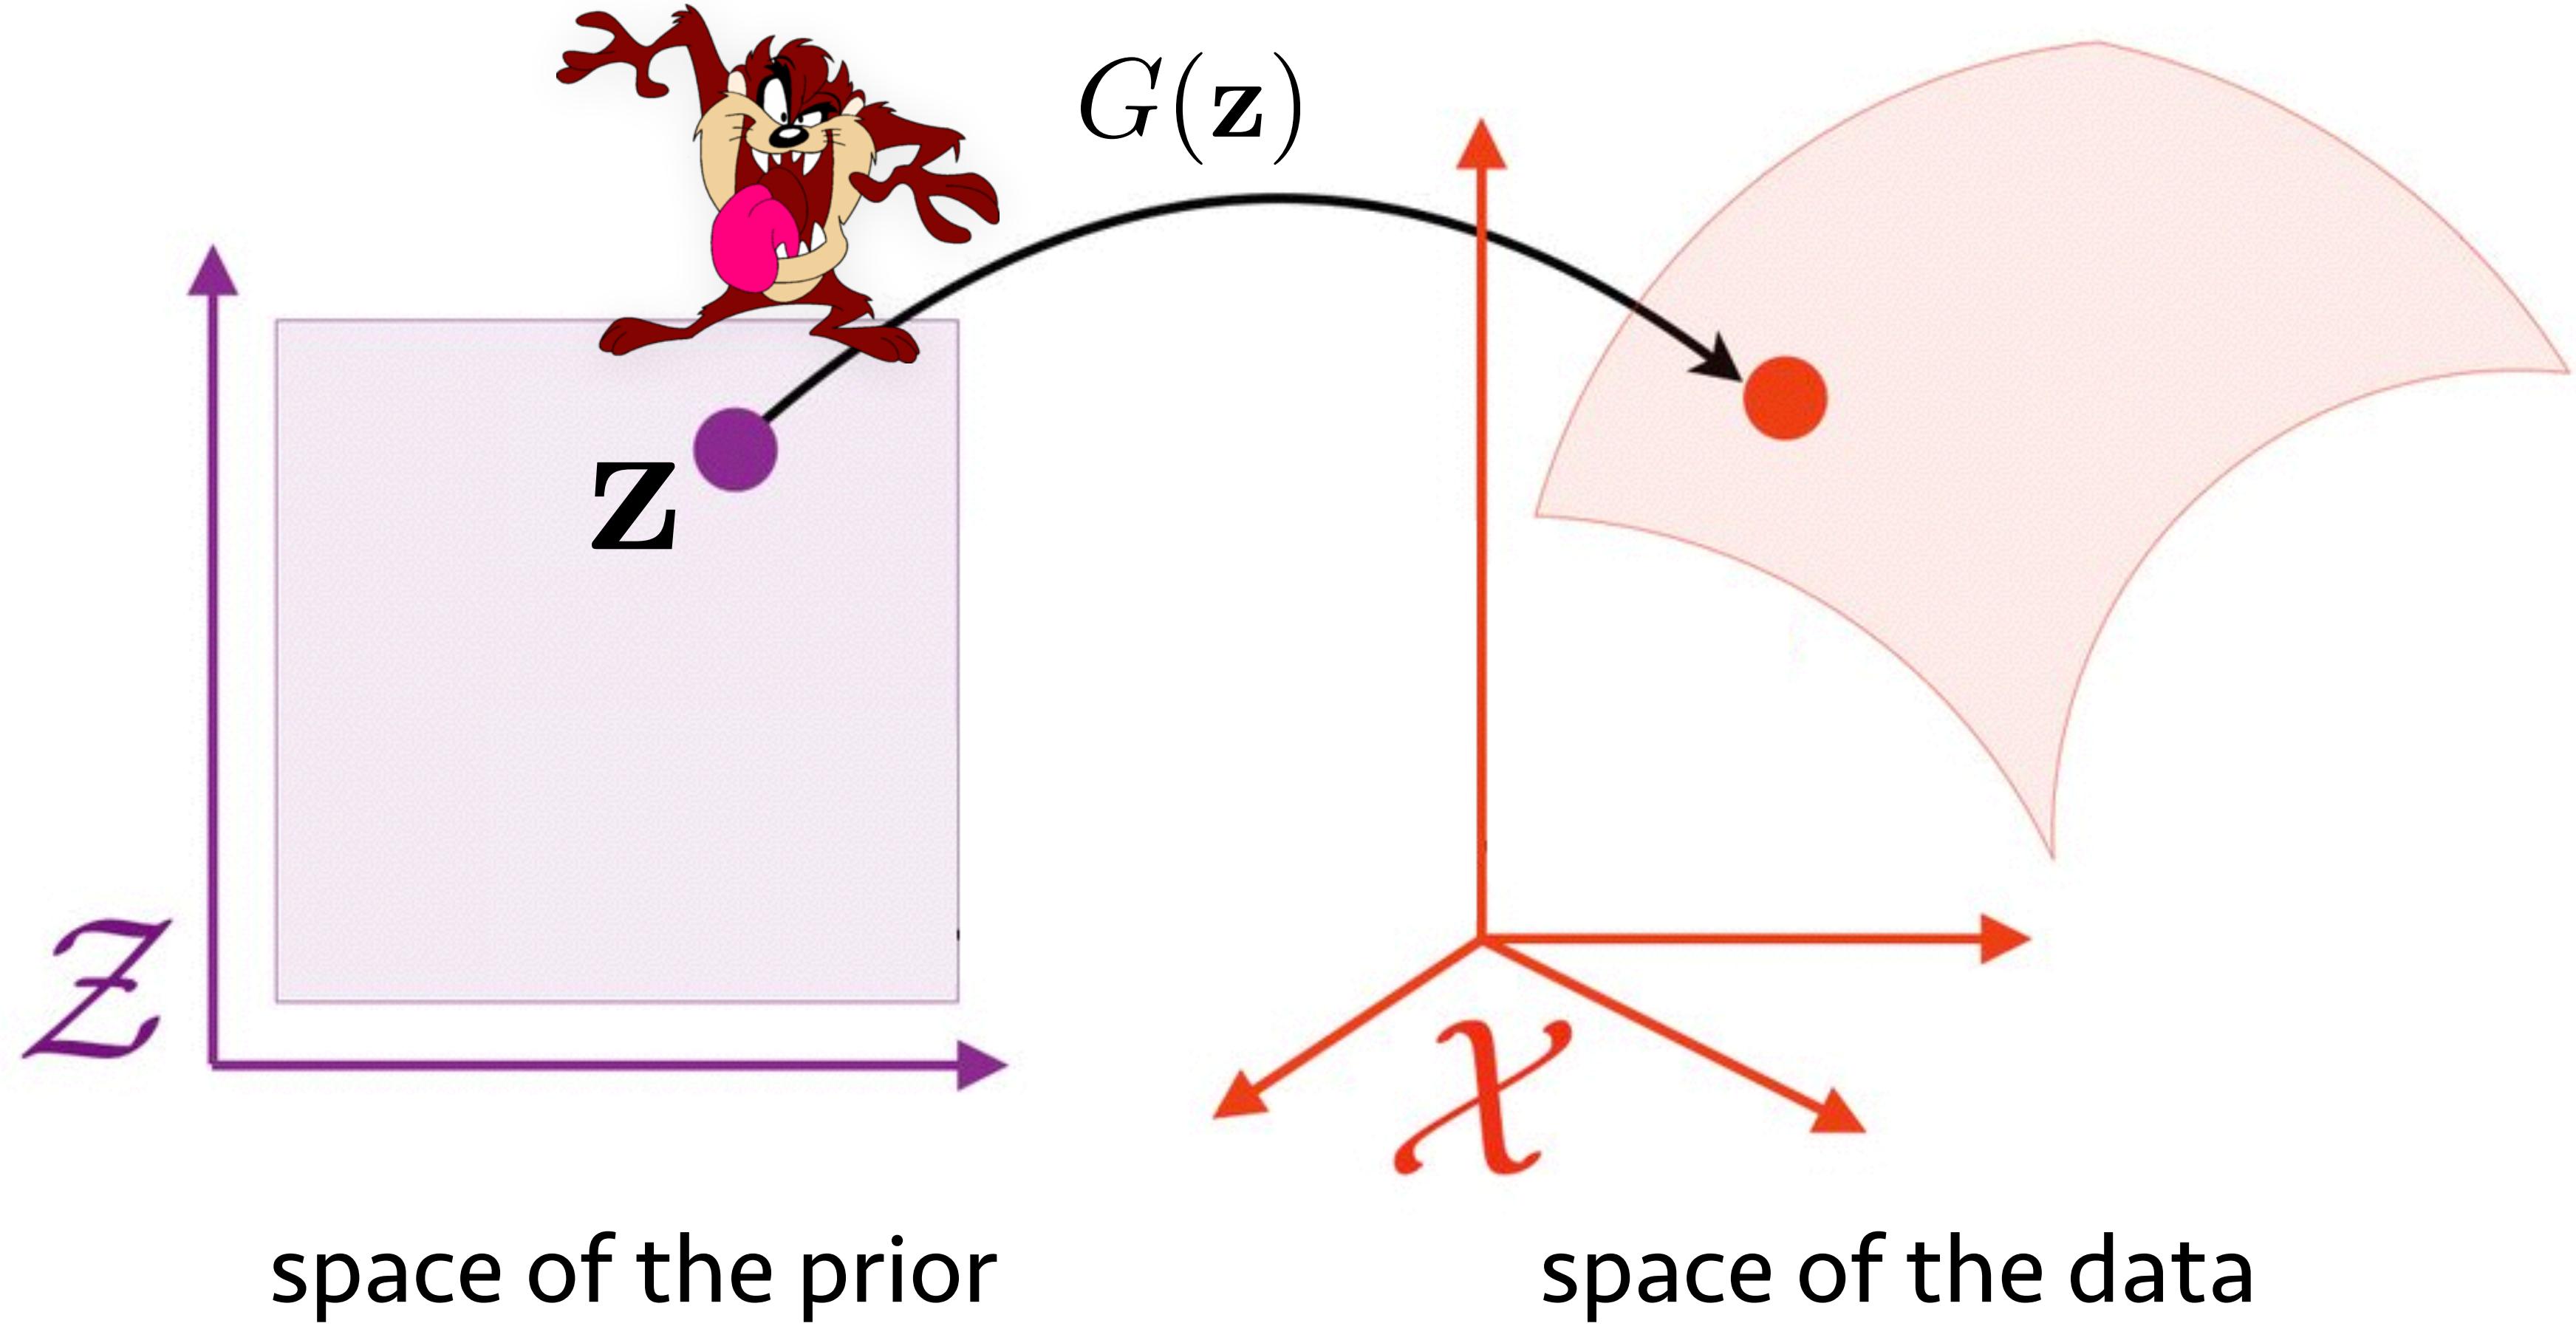
- GANs are a class of algorithms used in unsupervised learning in which the goal is to learn (implicitly) the data distribution
- This is done through two adversaries (generator and discriminator) which compute against each other in a zero-sum game
- Both networks compute against each other until Nash equilibrium

What are generative adversarial networks?

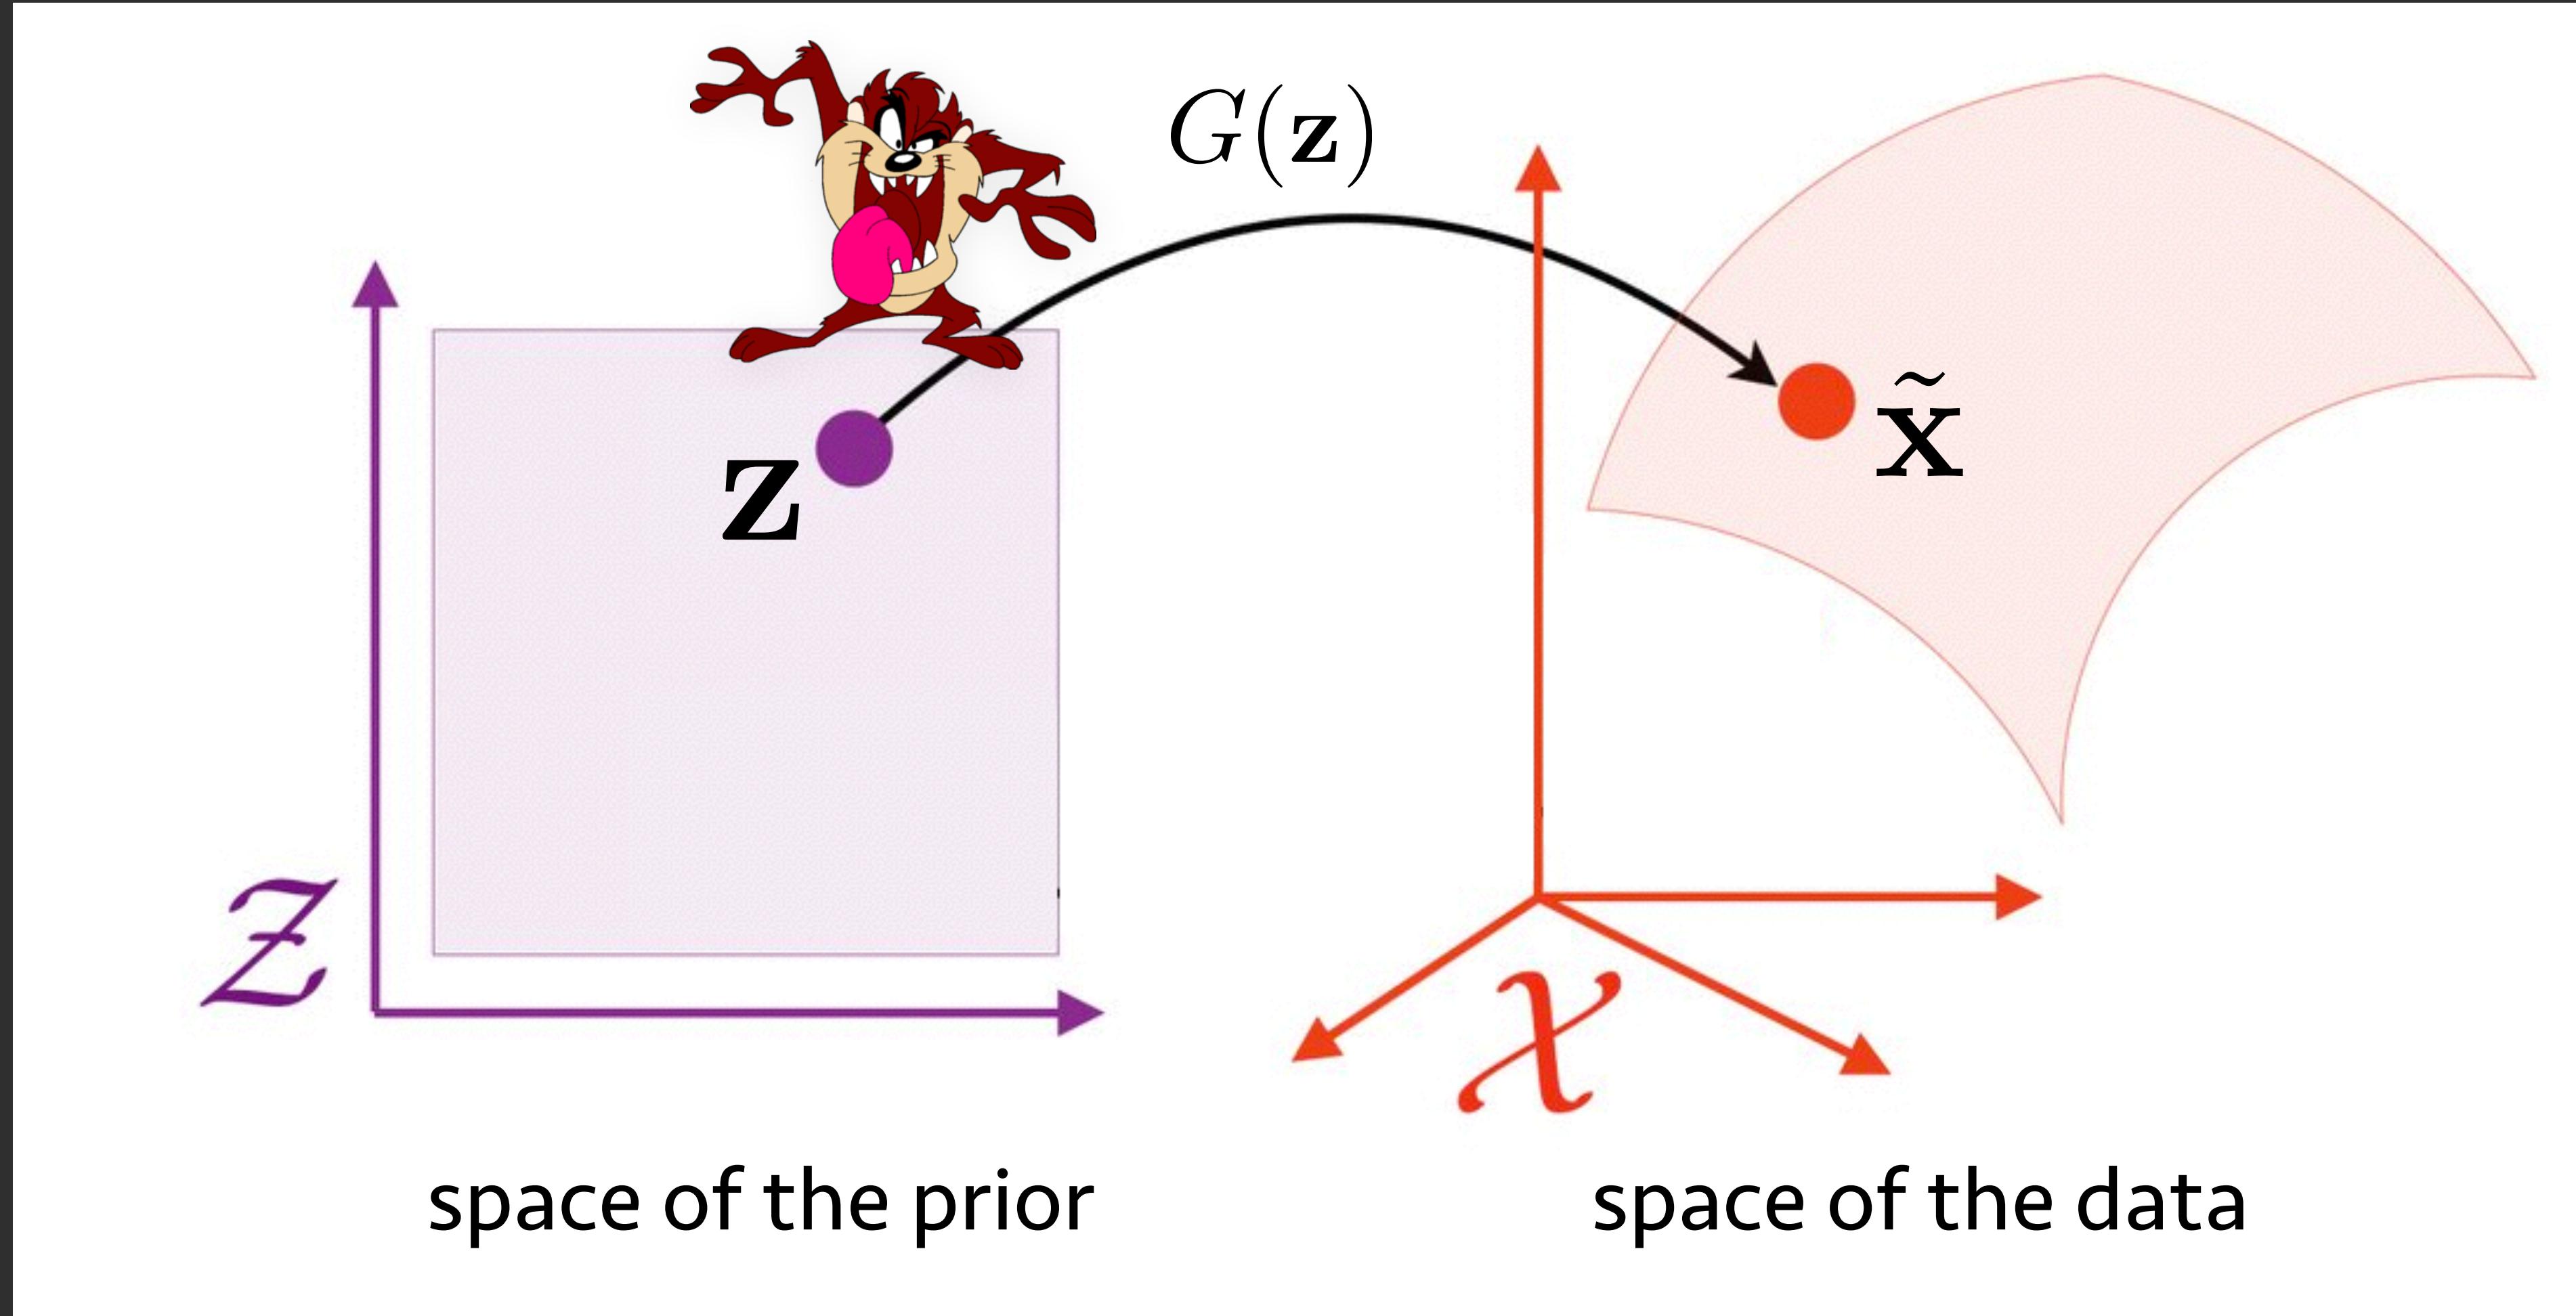
- GANs are a class of algorithms used in unsupervised learning in which the goal is to learn (implicitly) the data distribution
- This is done through two adversaries (generator and discriminator) which compute against each other in a zero-sum game
- Both networks compute against each other until Nash equilibrium
- Theoretically, the game that GANs play is equivalent to minimising some measure of divergence between two distributions (real and fake)



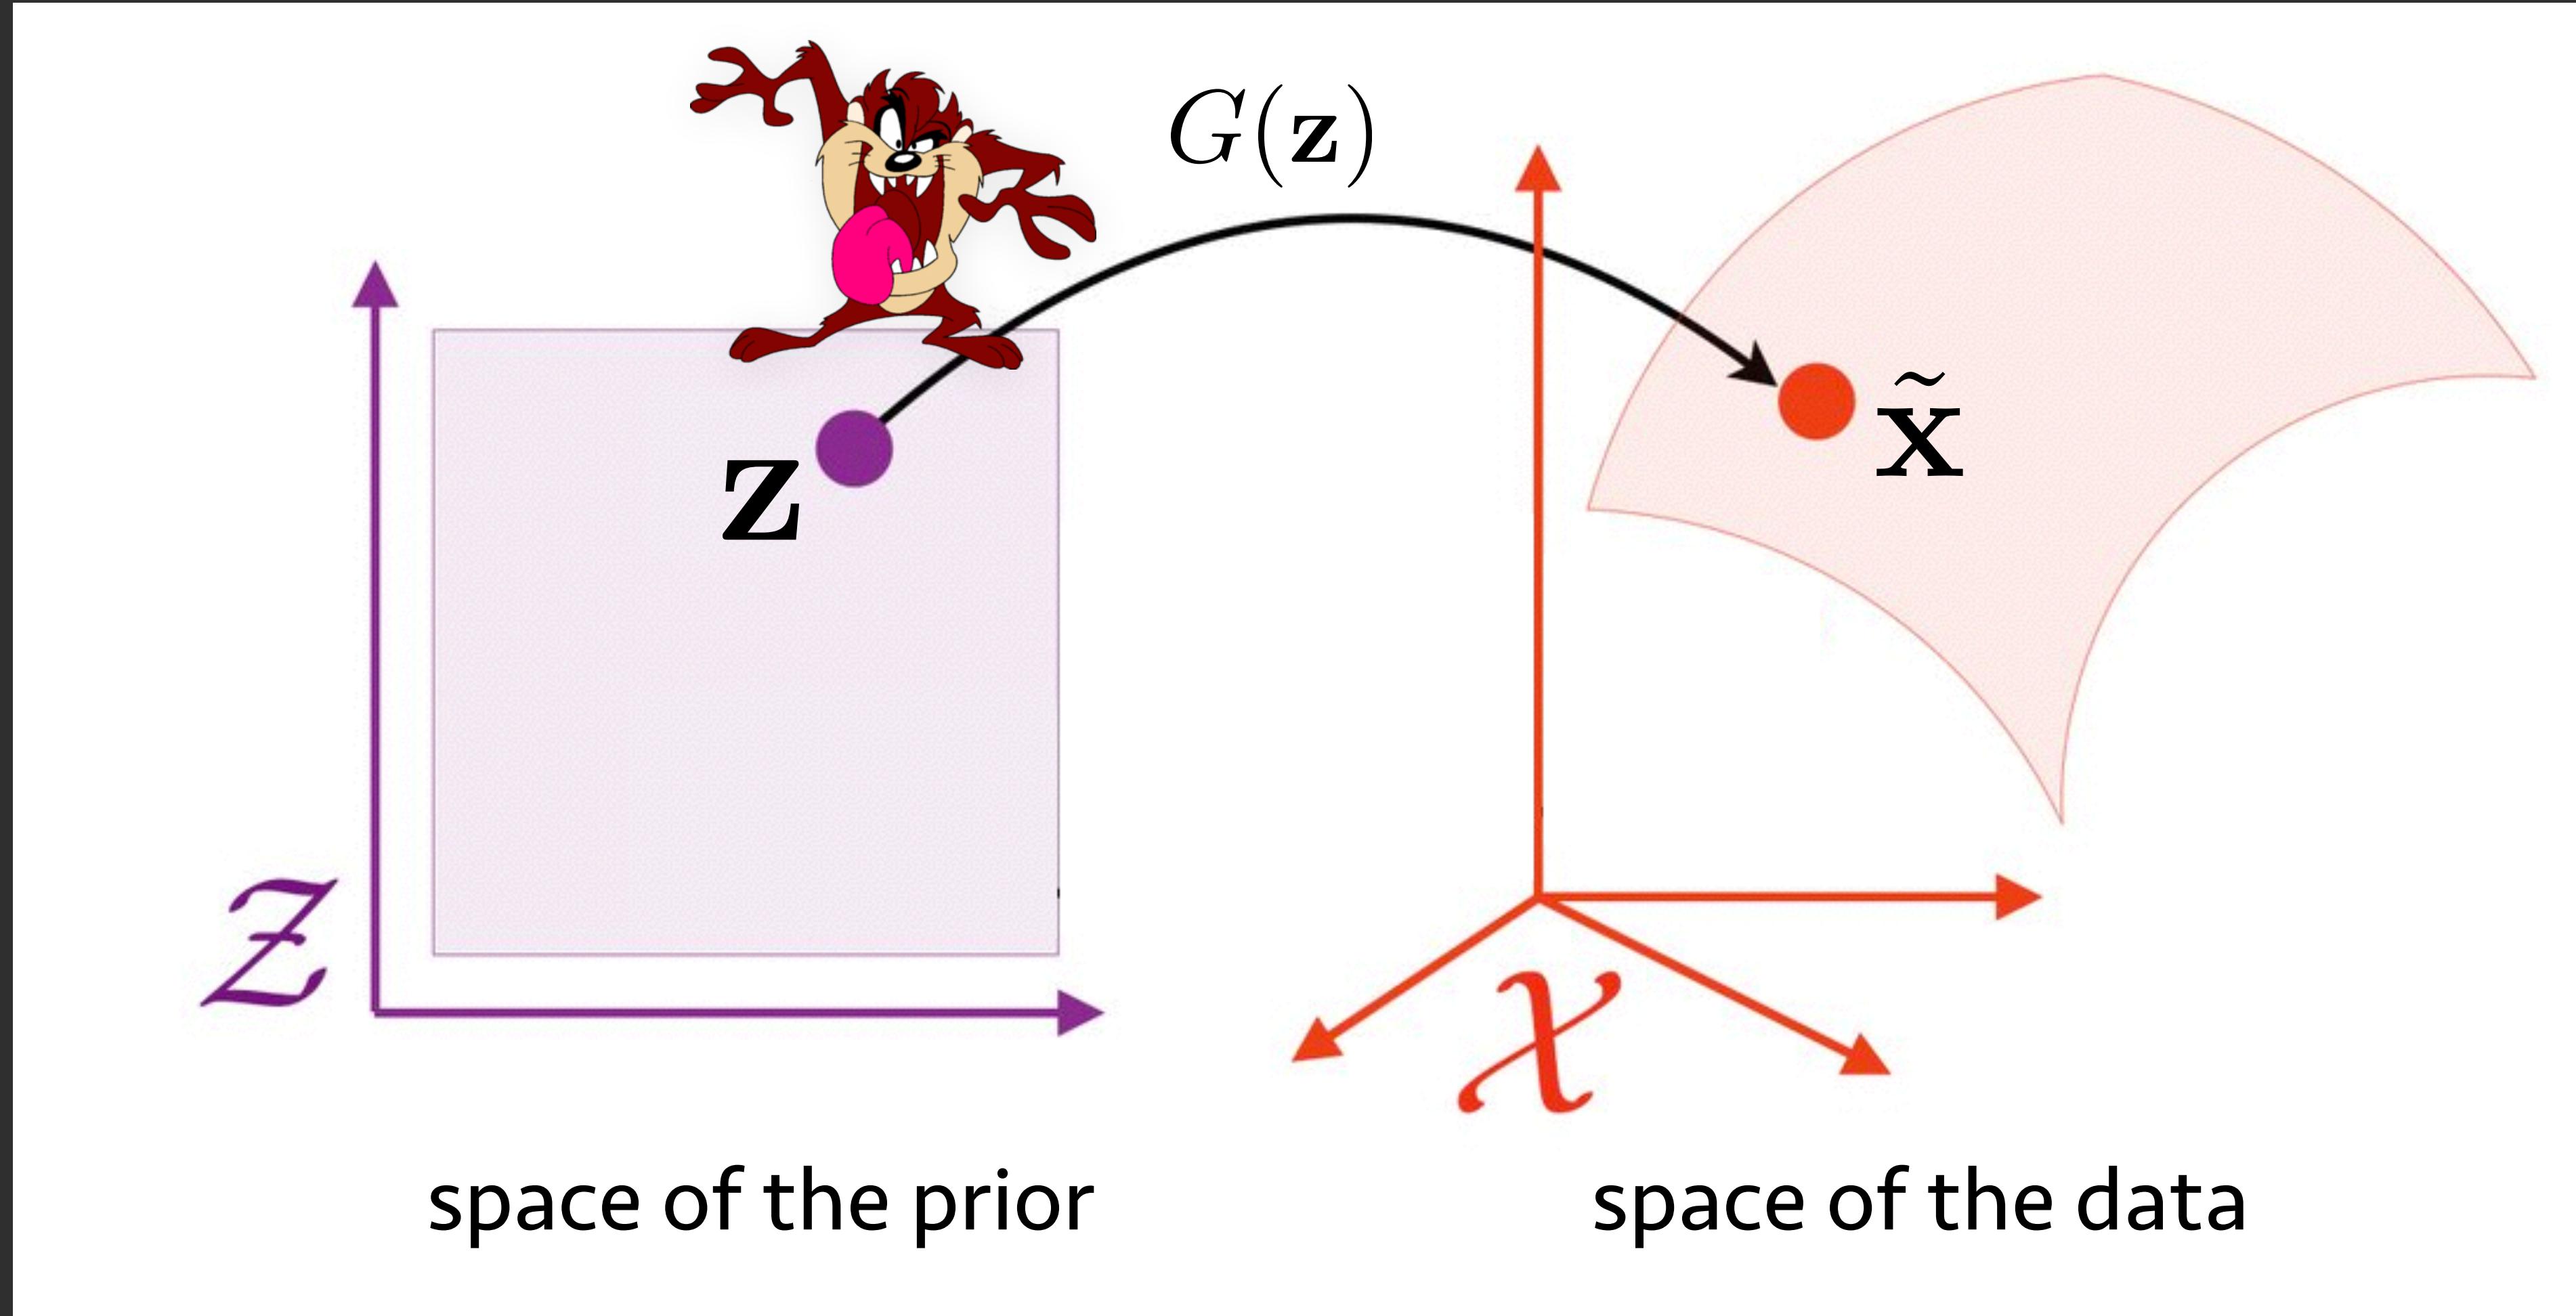
<https://twitter.com/gabrielpeyre/status/1000264262118445057>



<https://twitter.com/gabrielpeyre/status/1000264262118445057>

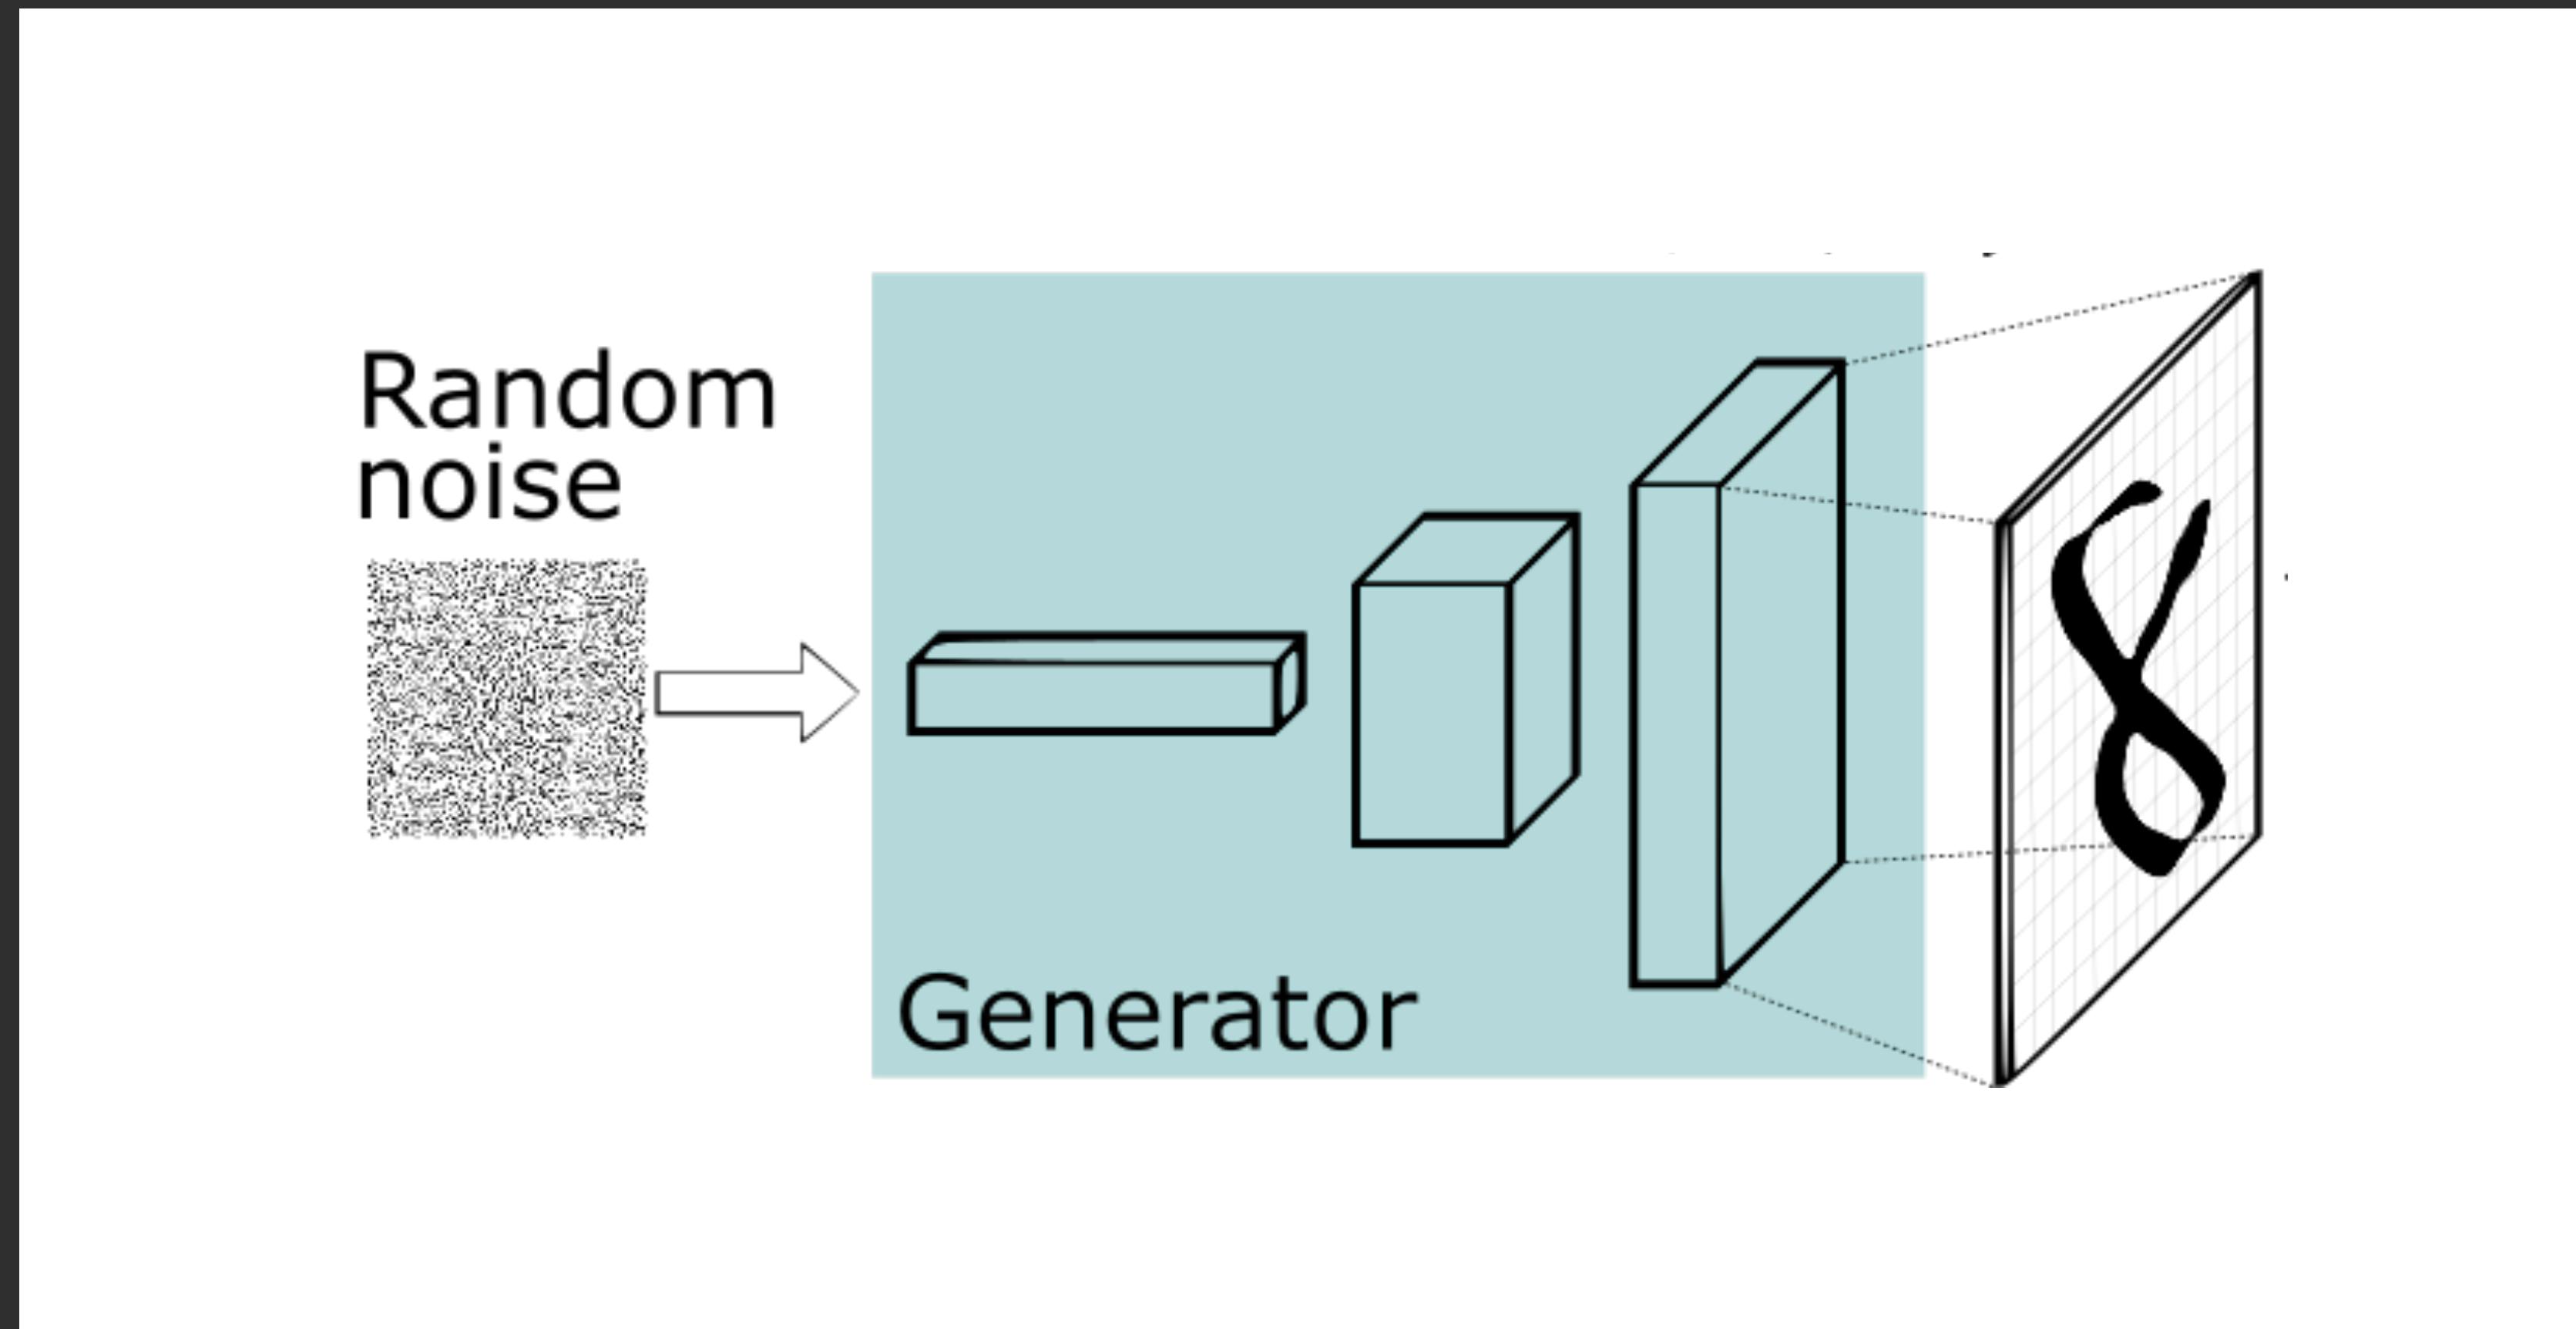


<https://twitter.com/gabrielpeyre/status/1000264262118445057>



<https://twitter.com/gabrielpeyre/status/1000264262118445057>

We can model arbitrary distributions by sampling from a simple distribution (prior) and mapping the sample through a sufficiently powerful function (generator)

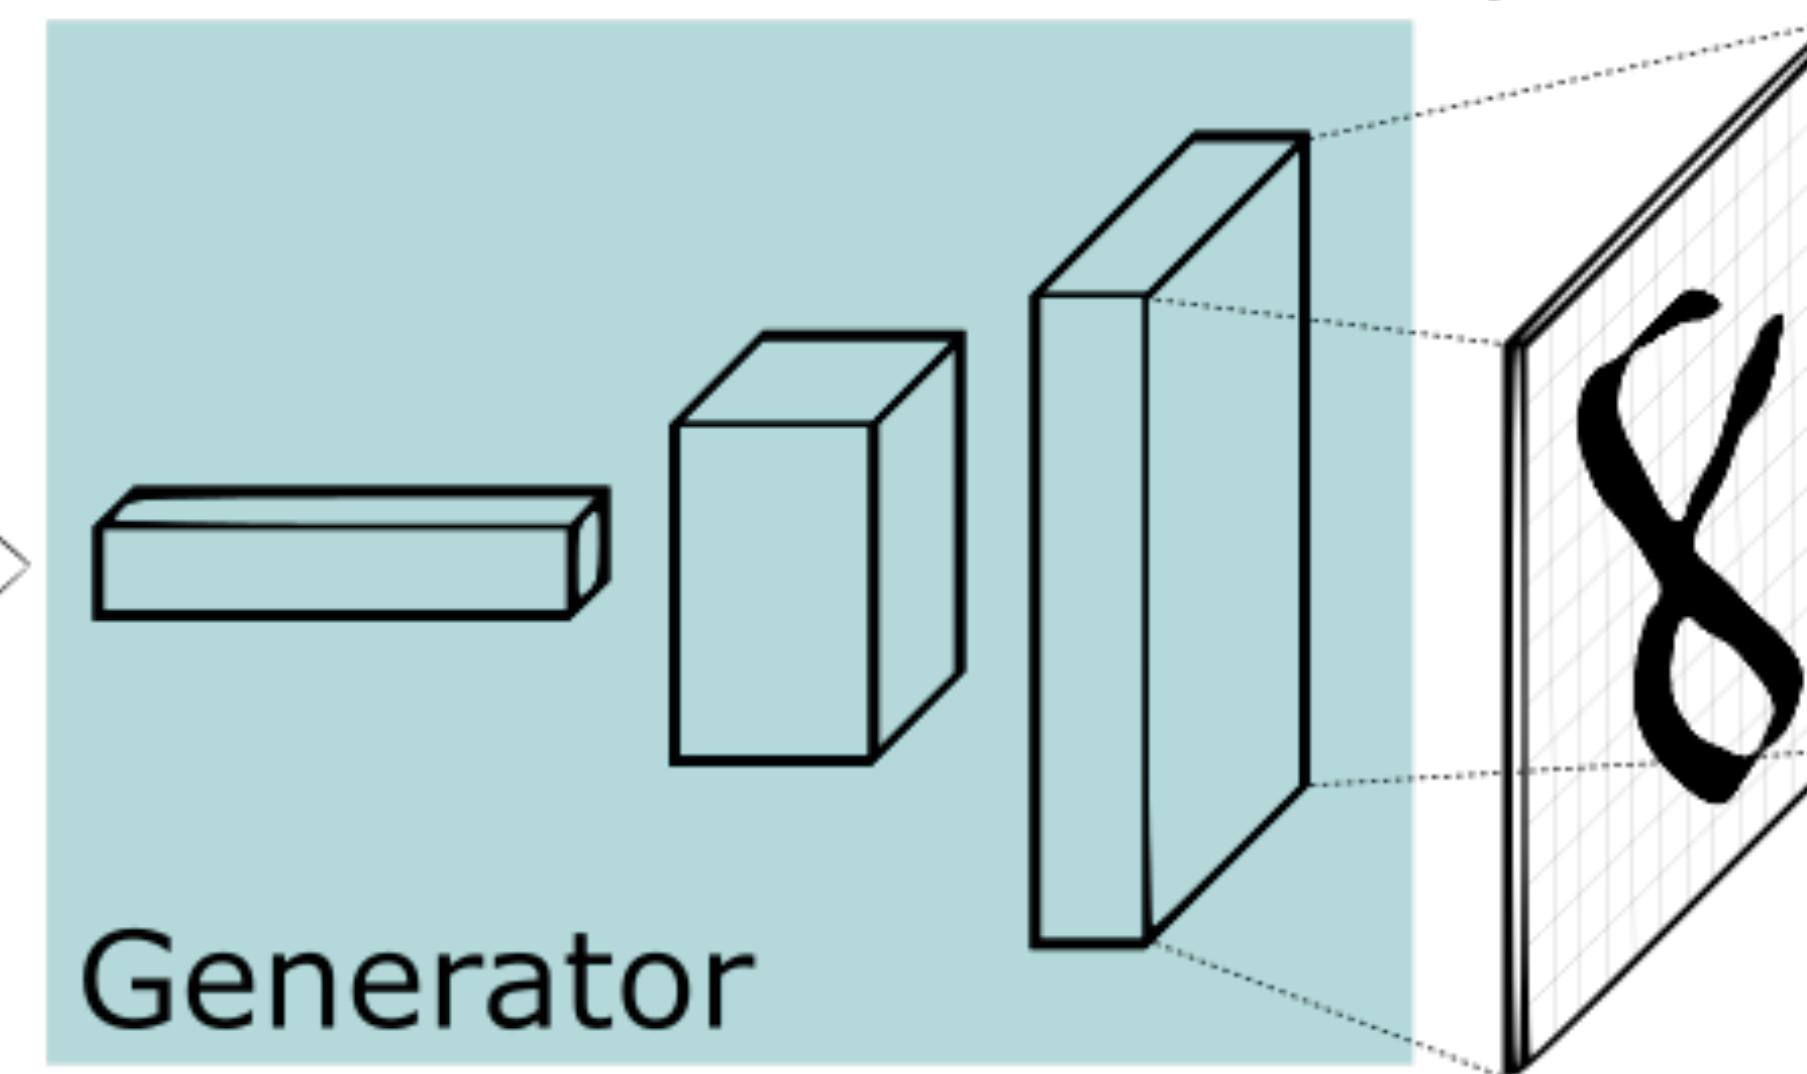


<https://deeplearning4j.org/generative-adversarial-network>

Random
noise



$\mathbf{z} \sim p(\mathbf{z})$

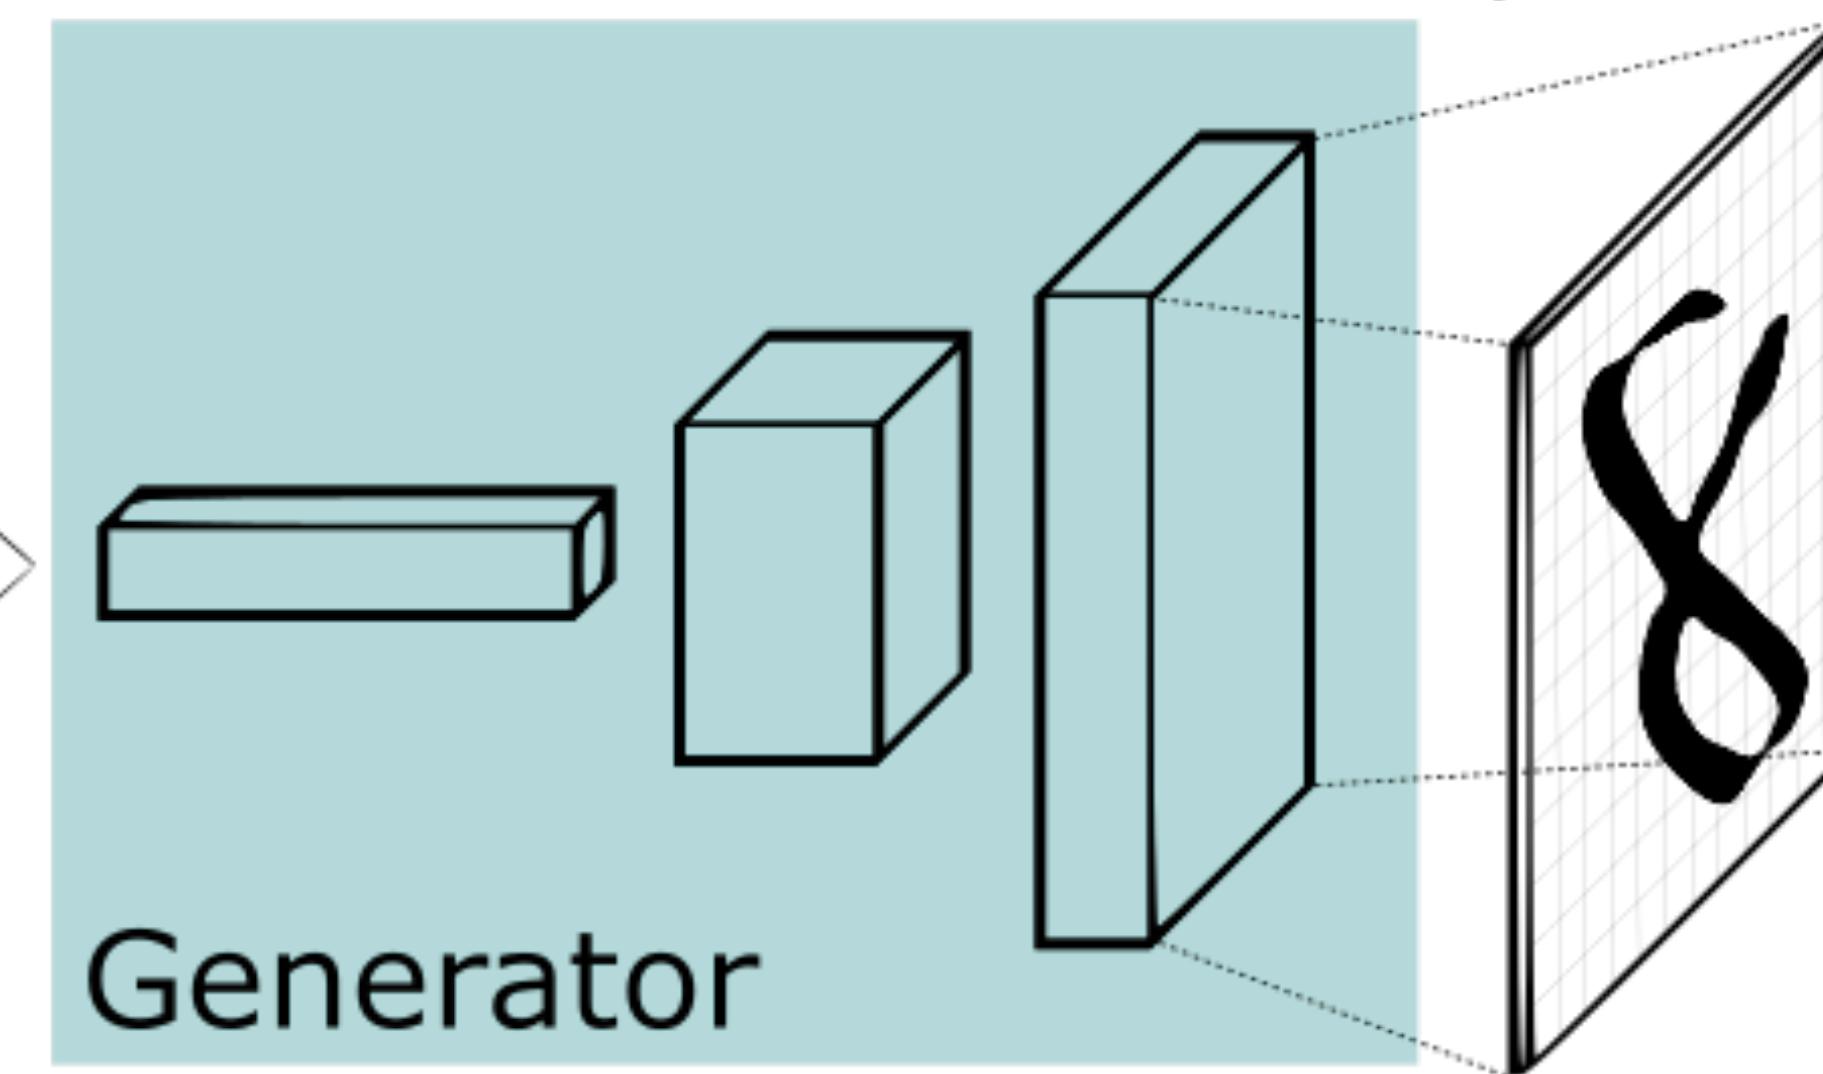


<https://deeplearning4j.org/generative-adversarial-network>

Random
noise

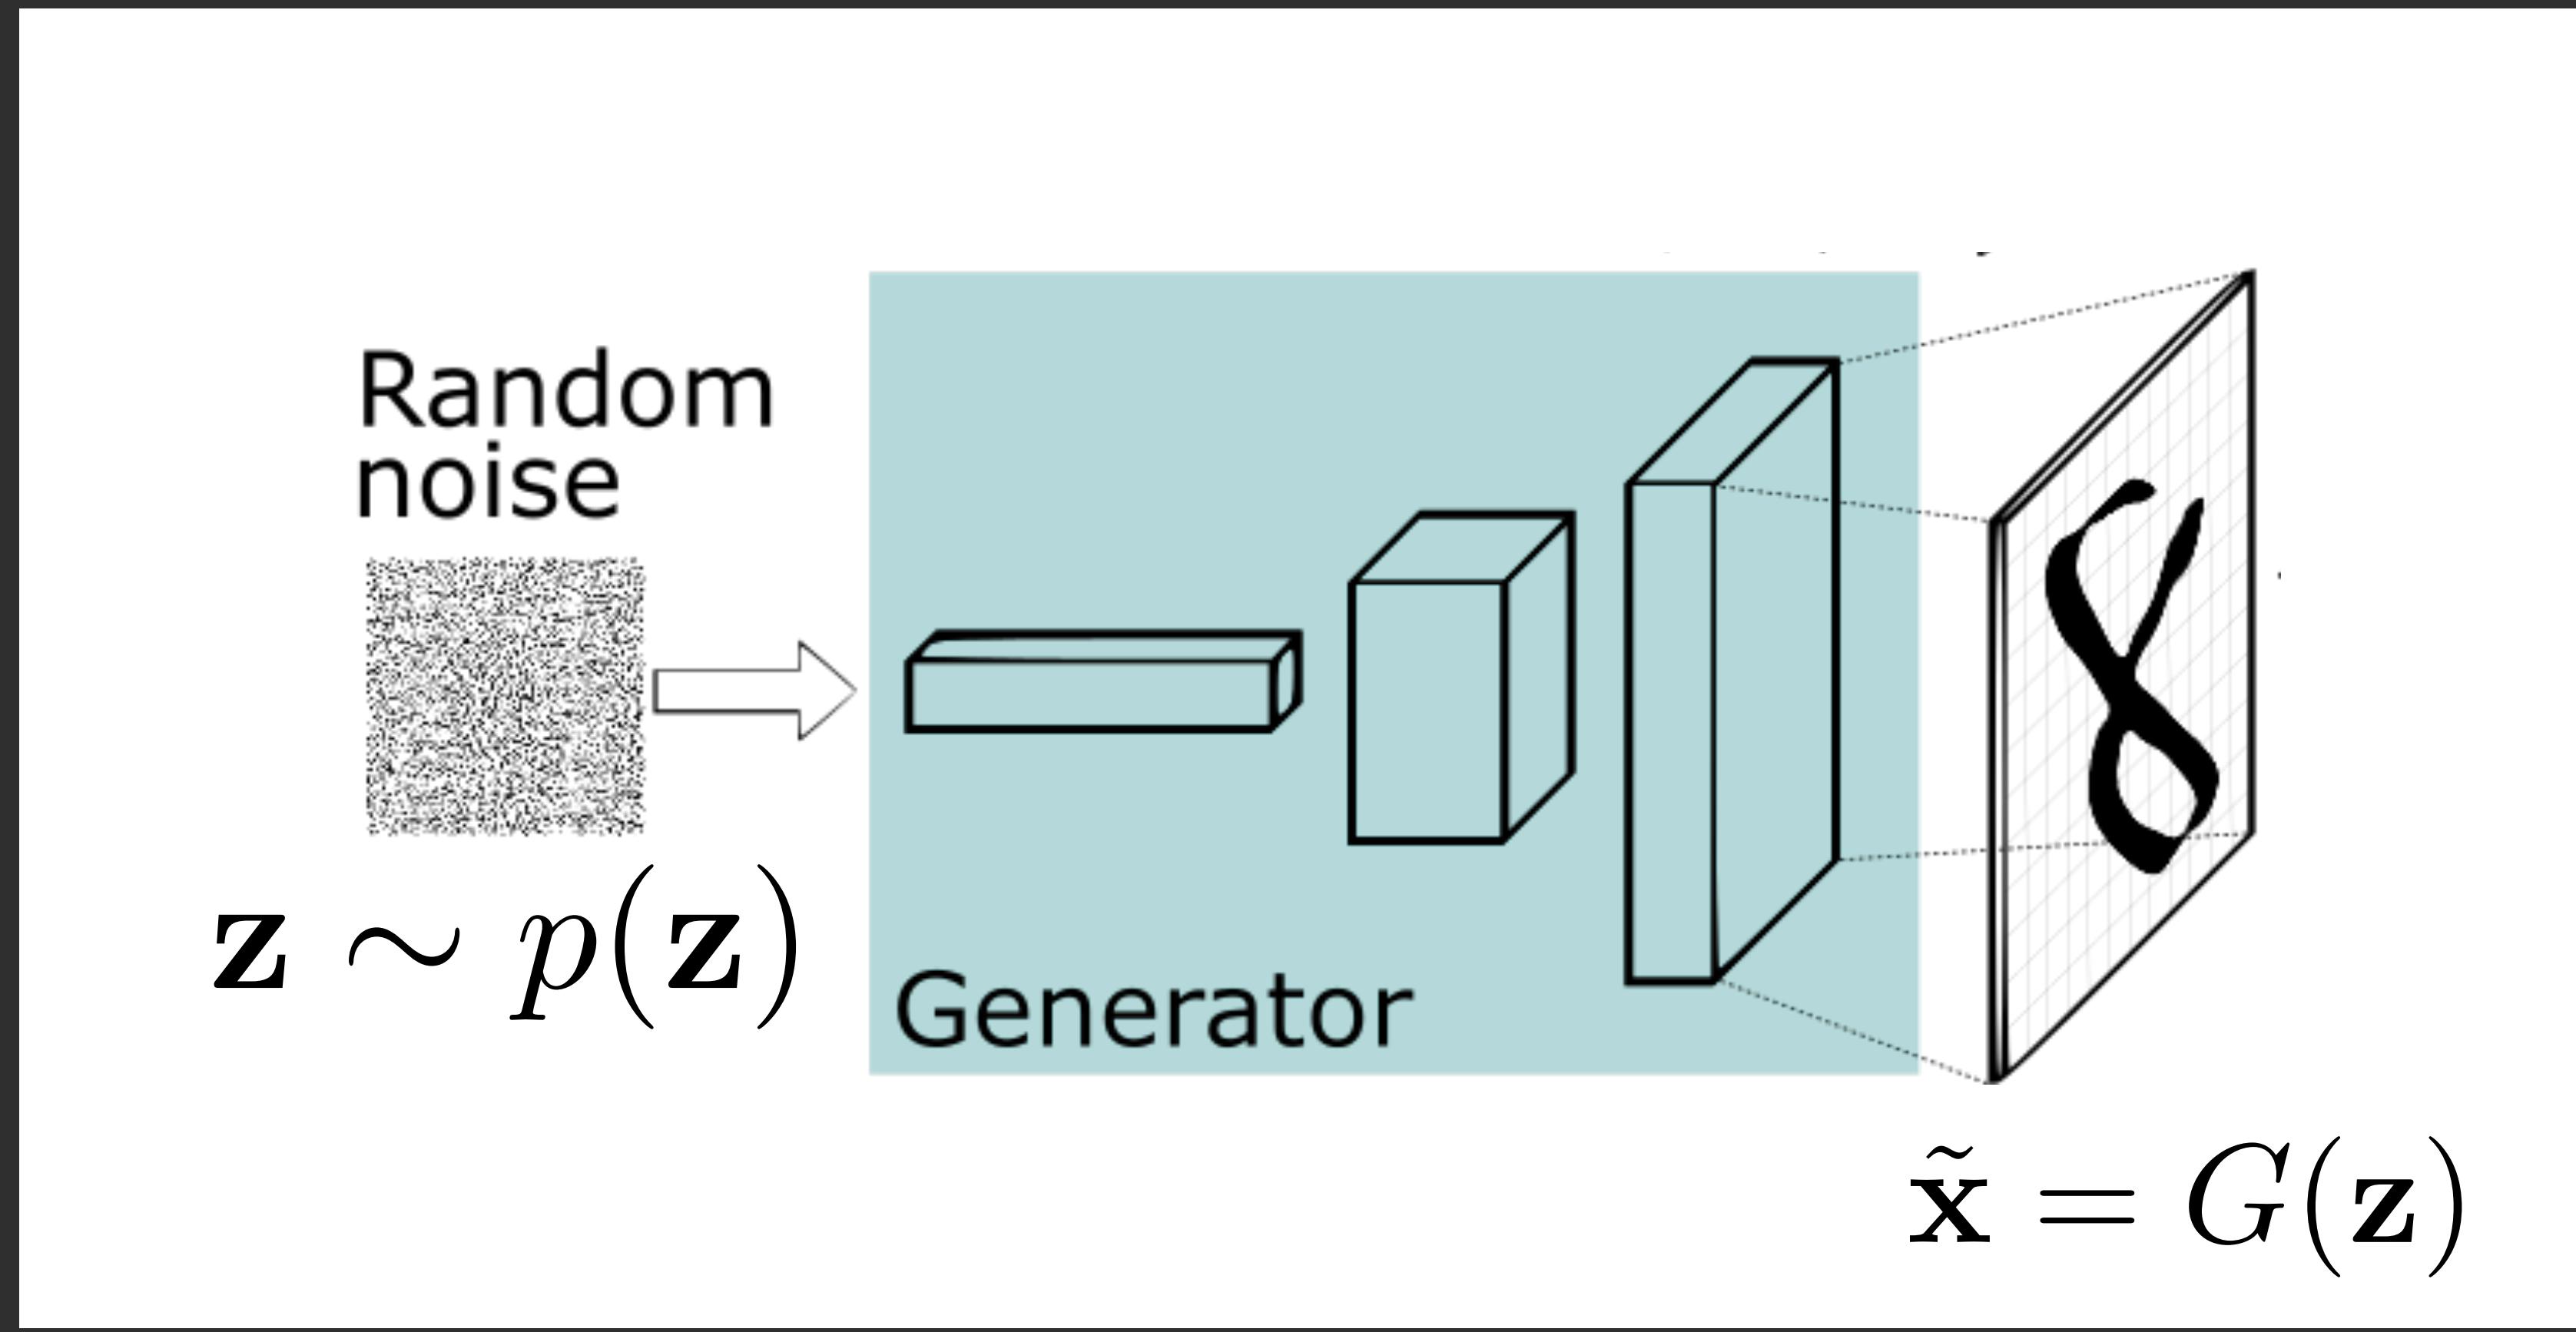


$\mathbf{z} \sim p(\mathbf{z})$



$$\tilde{\mathbf{x}} = G(\mathbf{z})$$

<https://deeplearning4j.org/generative-adversarial-network>

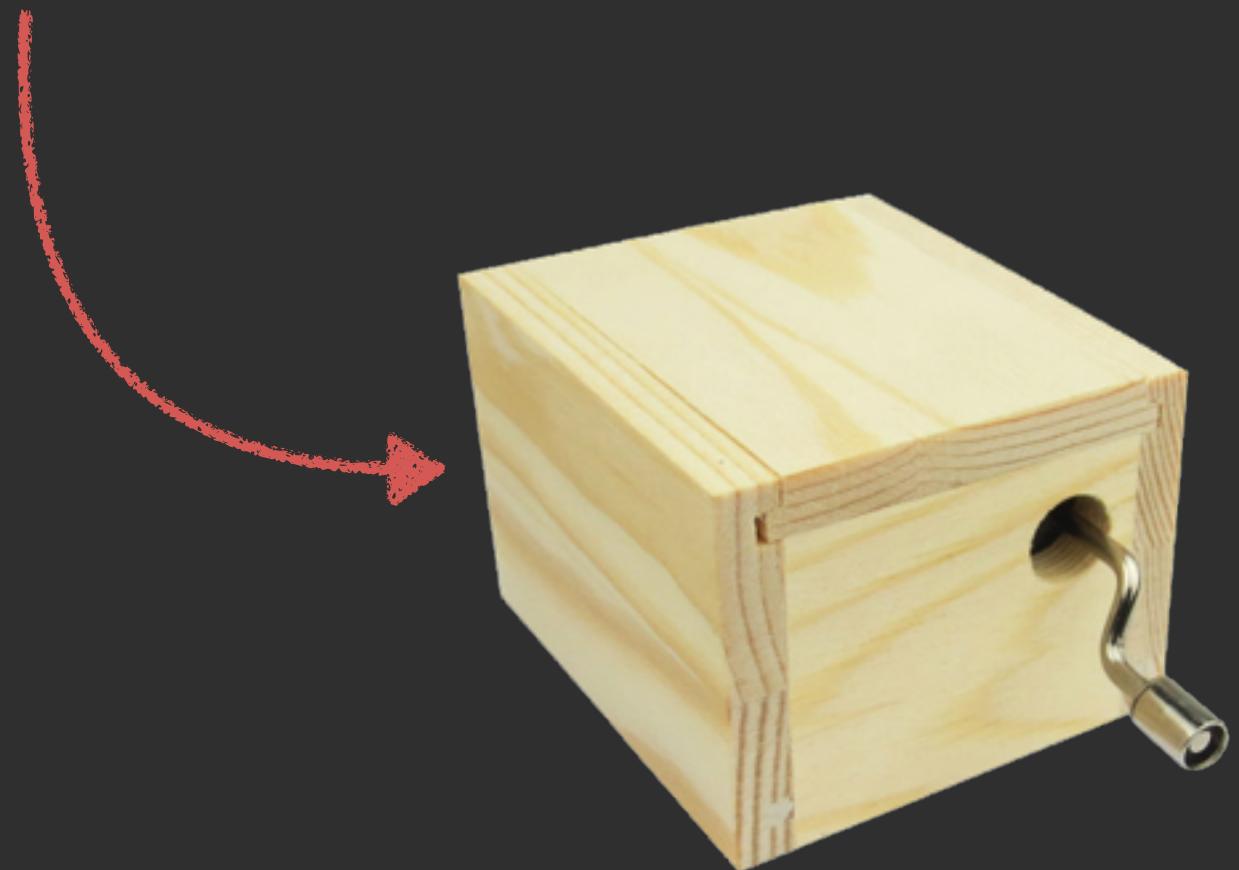


<https://deeplearning4j.org/generative-adversarial-network>

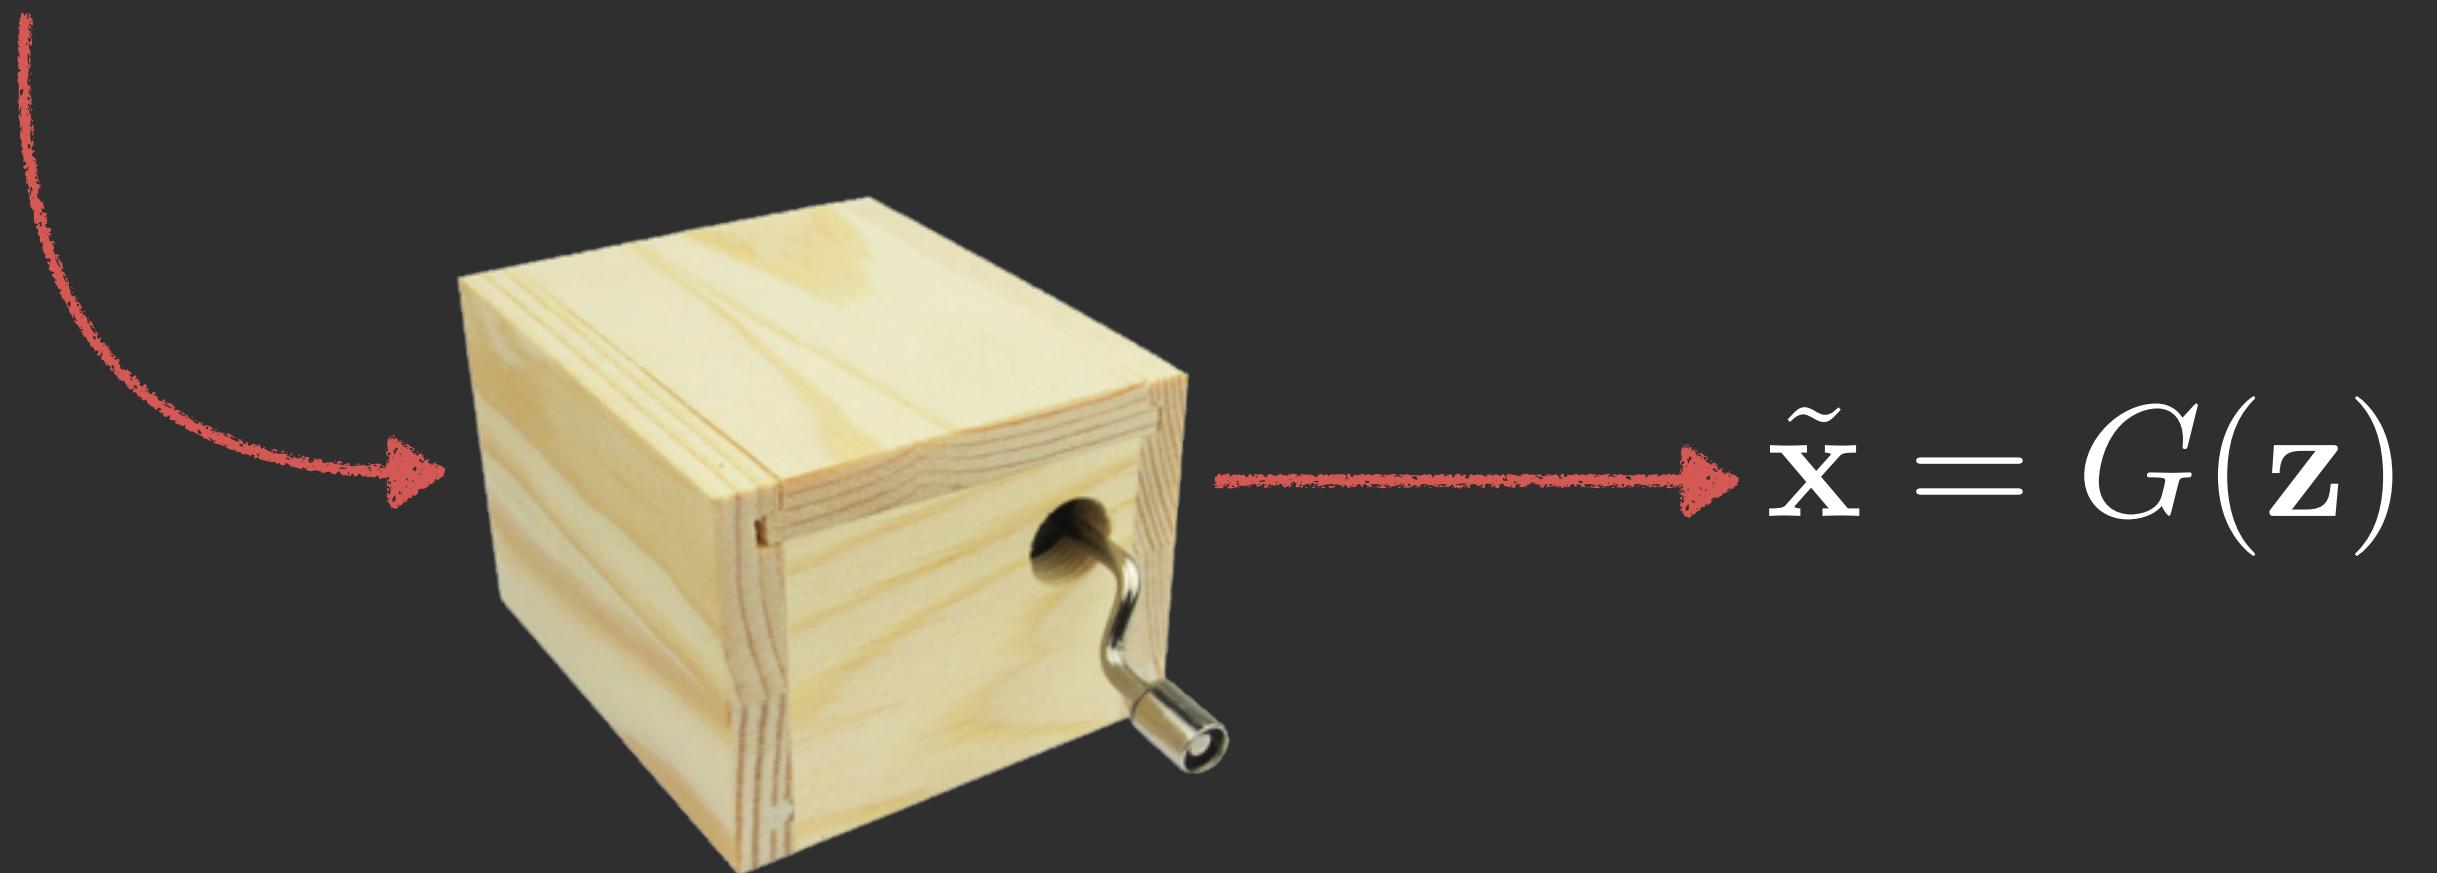
The generator takes noise (a source of stochasticity) and transforms it into something that could plausibly come from the real data

$$\mathbf{z} \sim p(\mathbf{z})$$

$$\mathbf{z} \sim p(\mathbf{z})$$



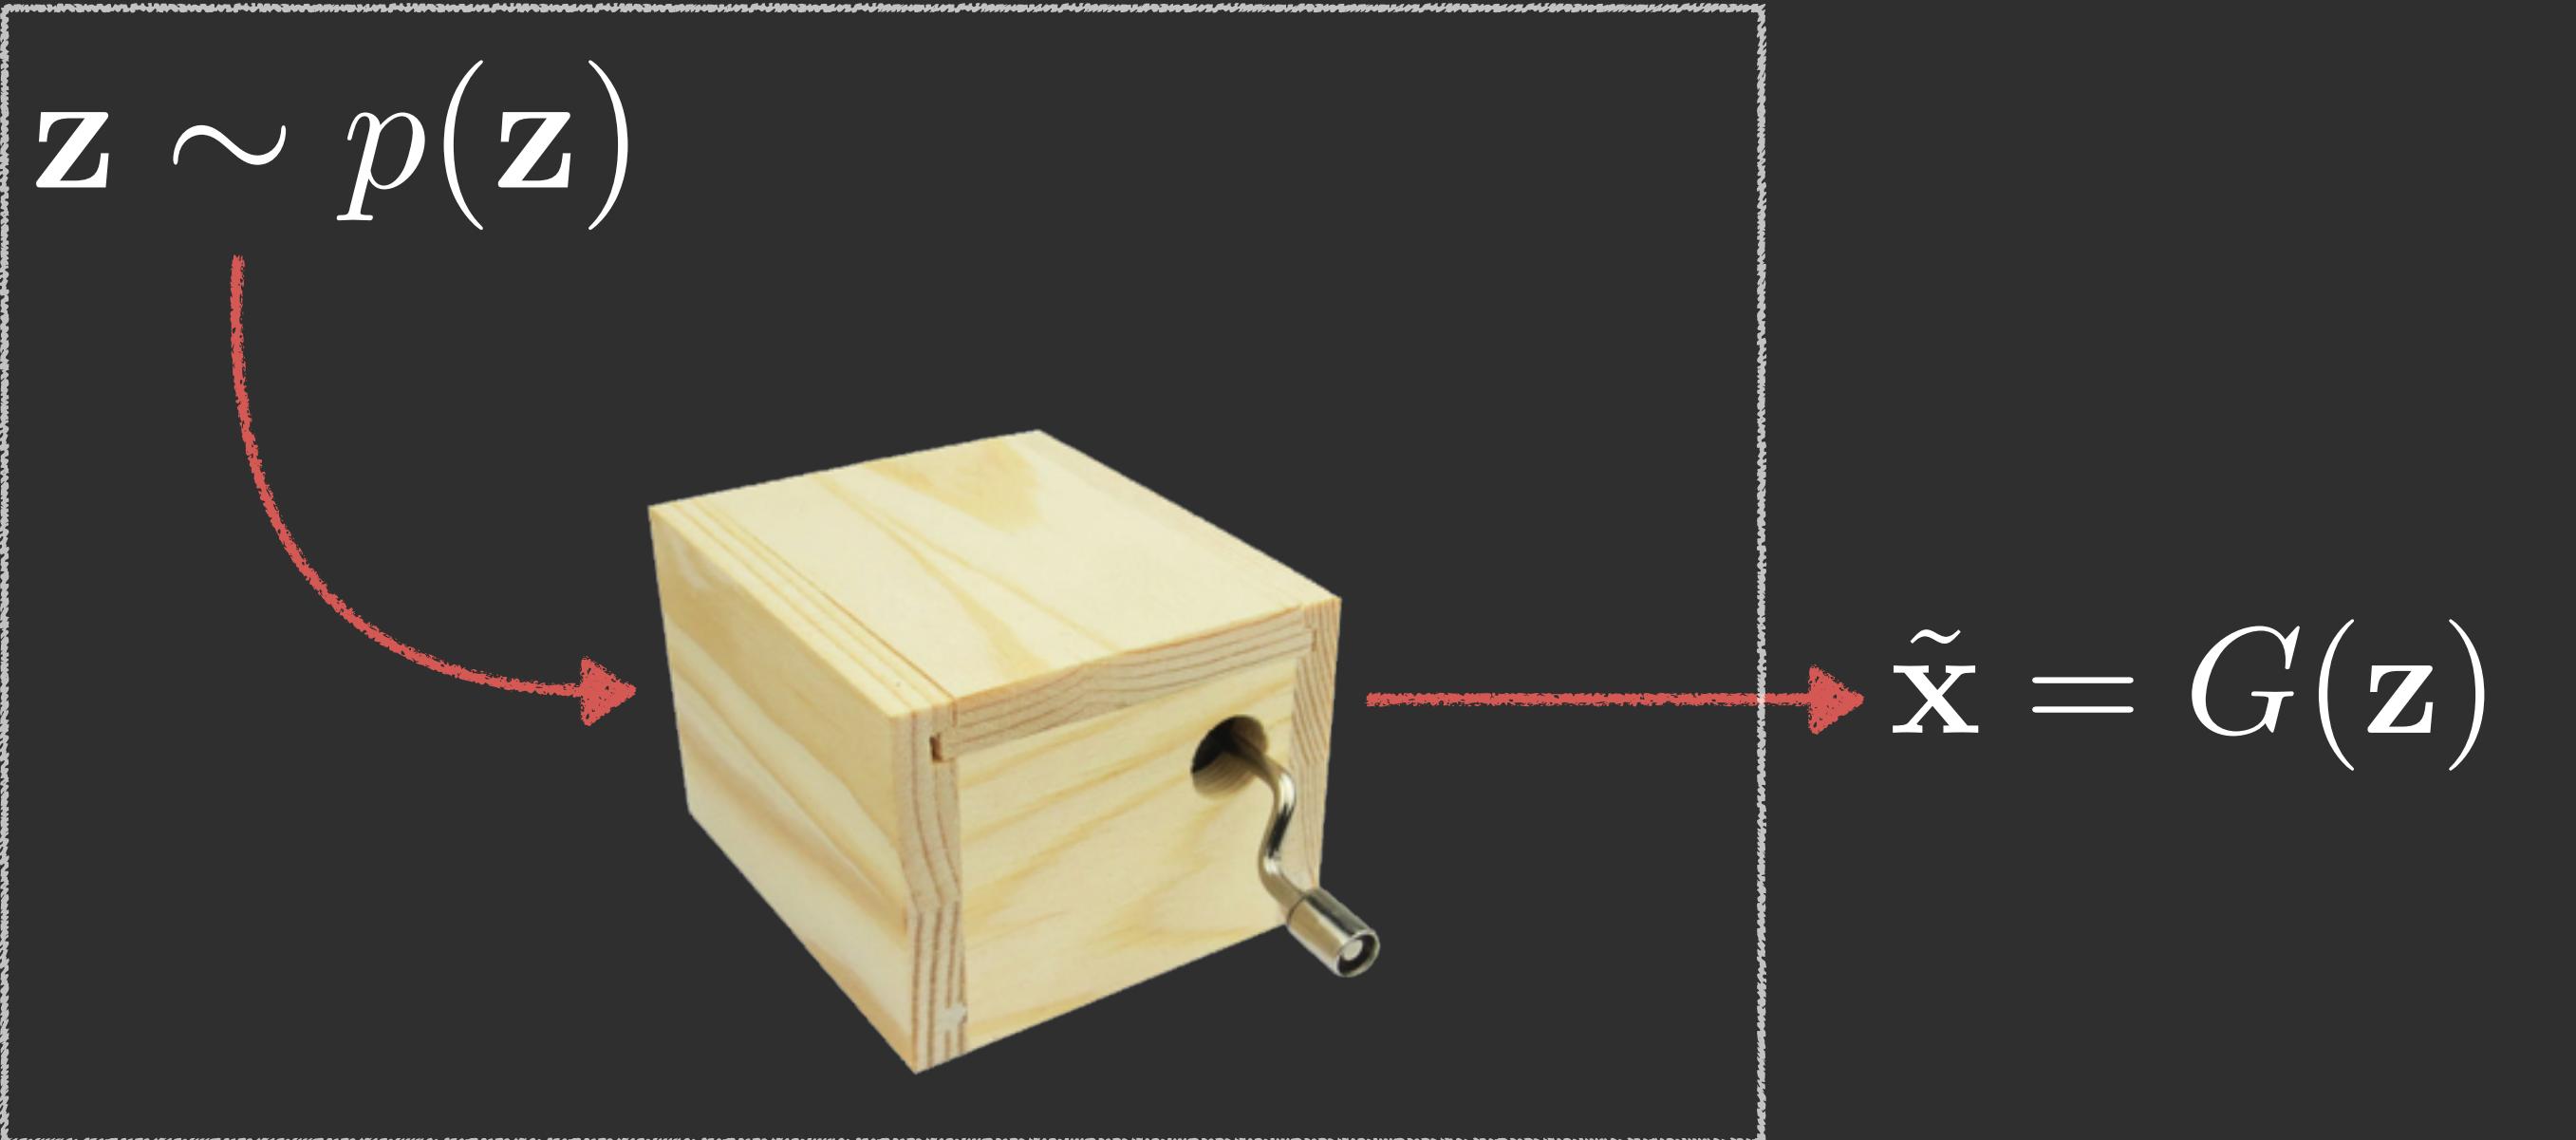
$$\mathbf{z} \sim p(\mathbf{z})$$



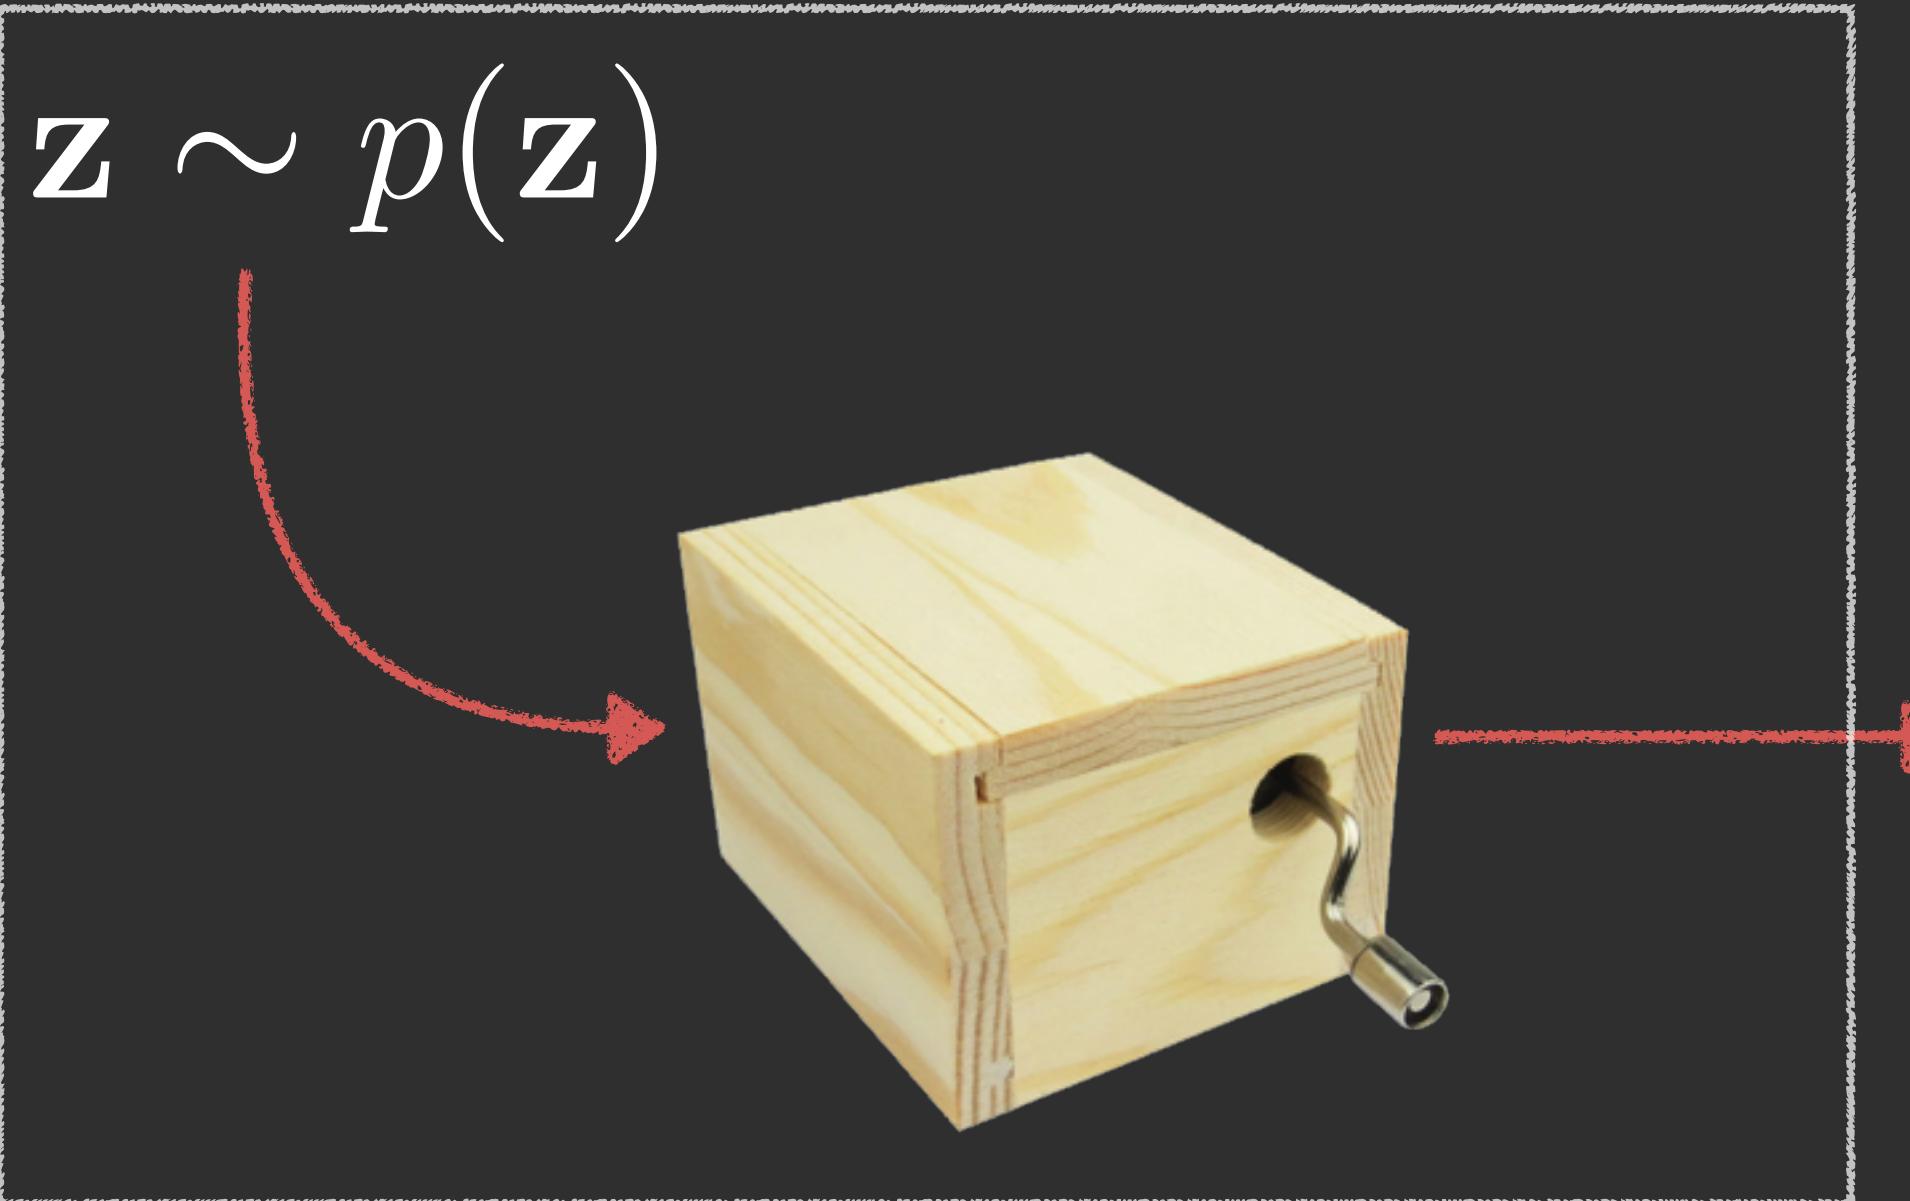
<https://www.etsy.com/listing/216241958/musical-wooden-box-hand-crank-music-box>

<https://www.mathworks.com/matlabcentral/answers/180778-plotting-a-3d-gaussian-function-using-surf>

$$\tilde{\mathbf{x}} \sim p_g(\mathbf{x})$$



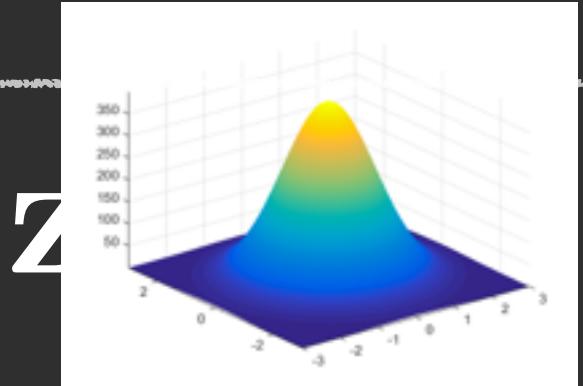
$$\tilde{\mathbf{x}} \sim p_g(\mathbf{x})$$



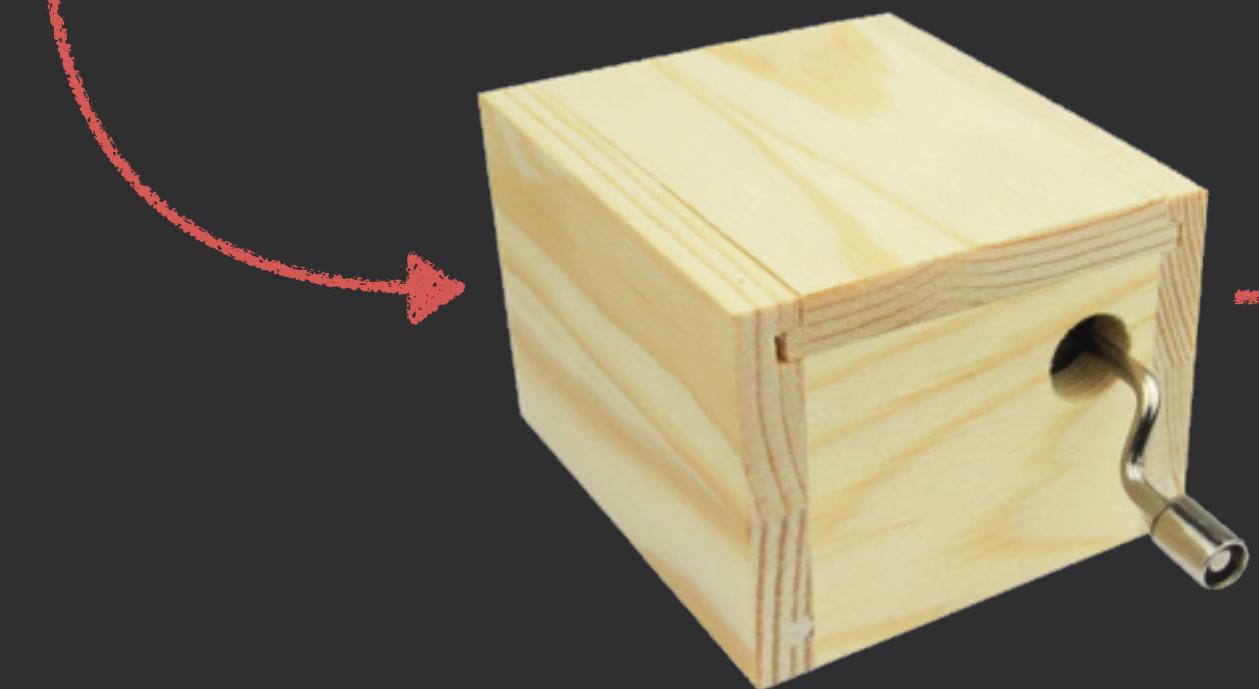
$$\mathbf{x} \sim p_d(\mathbf{x})$$

$$\tilde{\mathbf{x}} = G(\mathbf{z})$$

$$\tilde{\mathbf{x}} \sim p_g(\mathbf{x})$$



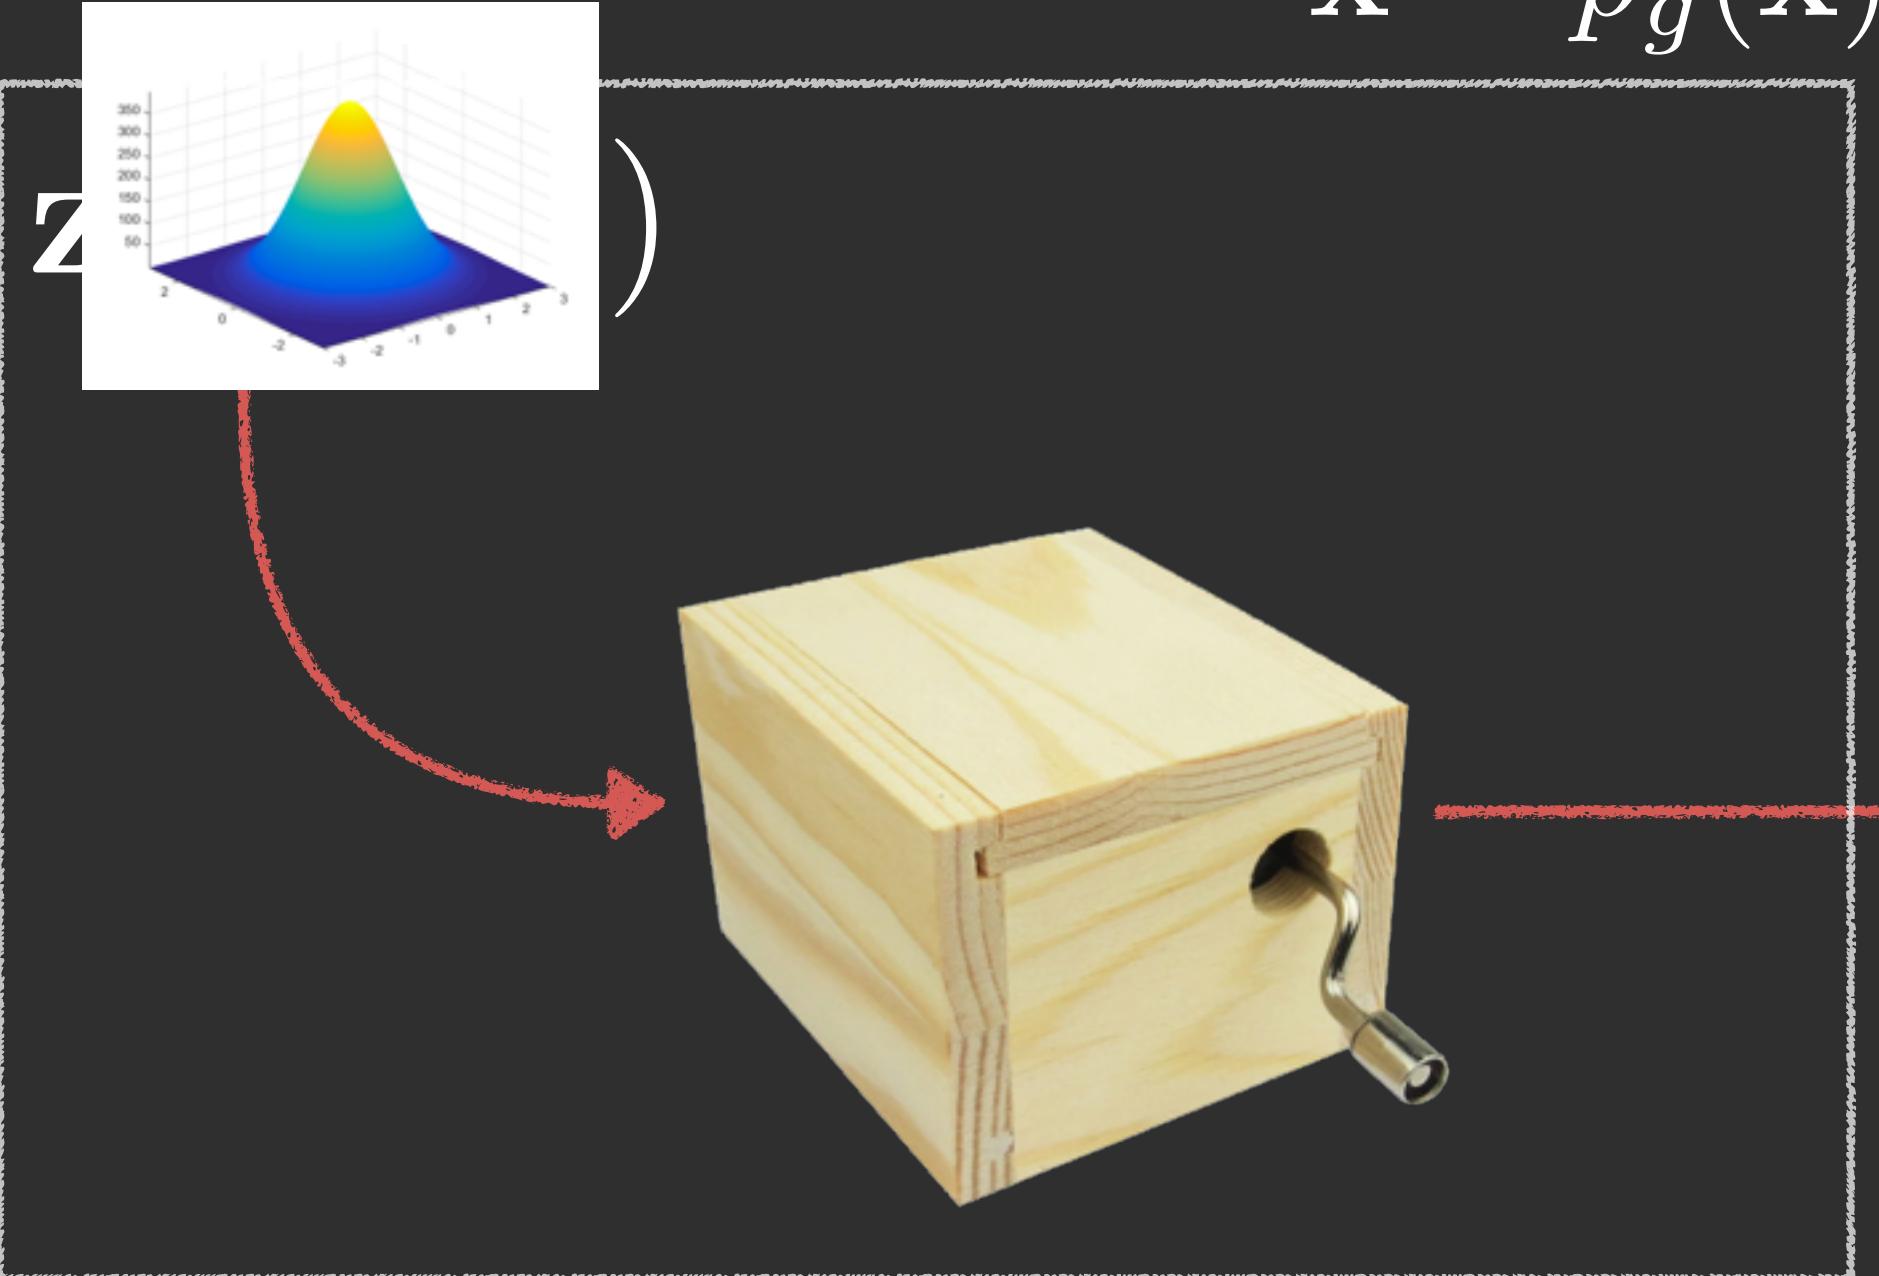
)



3 3 3 3 1 8 6 2
0 9 7 6 4 1 7 9
t 1 8 8 8 8 3 1

1 8 3 1 5 5 8 1
6 1 1 8 7 6 0 7
4 8 6 4 2 3 0 3

$$\tilde{\mathbf{x}} \sim p_g(\mathbf{x})$$

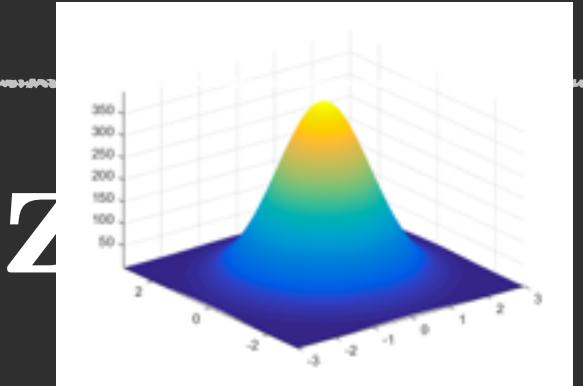


1 8 3 1 5 5 8 1
6 1 1 8 7 6 0 7
4 8 6 4 2 0 0 3

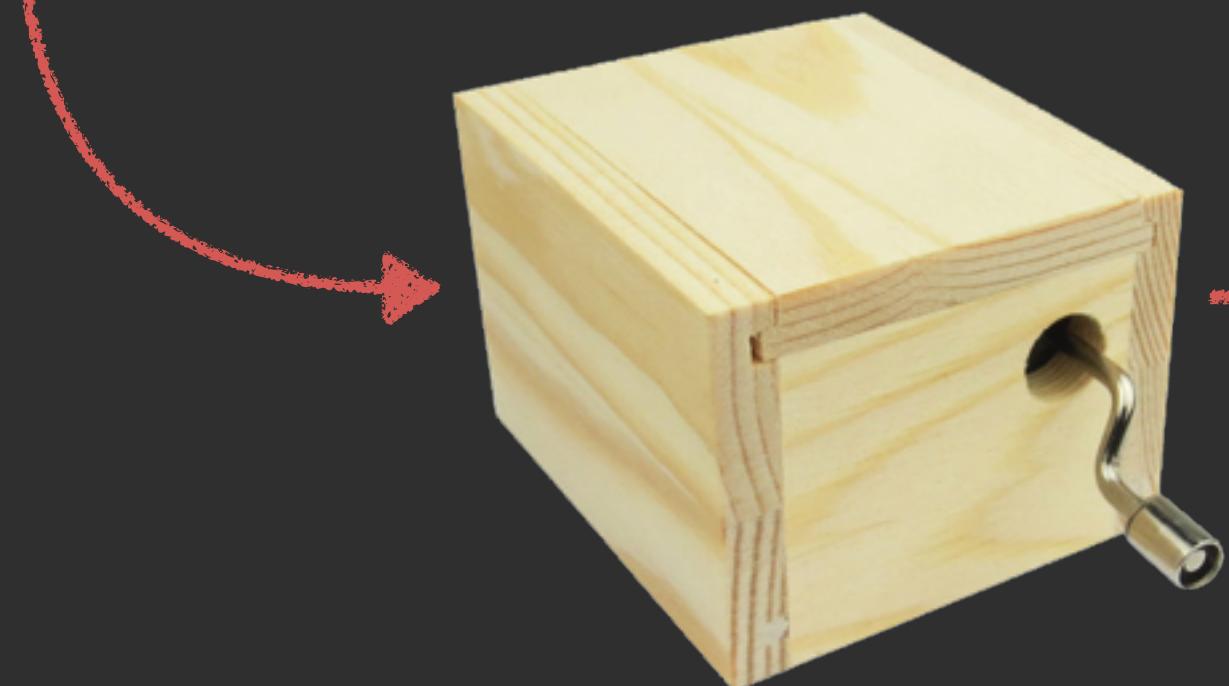
3 3 3 3 1 8 6 0
0 3 7 6 4 1 7 9
t 1 8 8 8 8 3 1

(+)

$$\tilde{\mathbf{x}} \sim p_g(\mathbf{x})$$



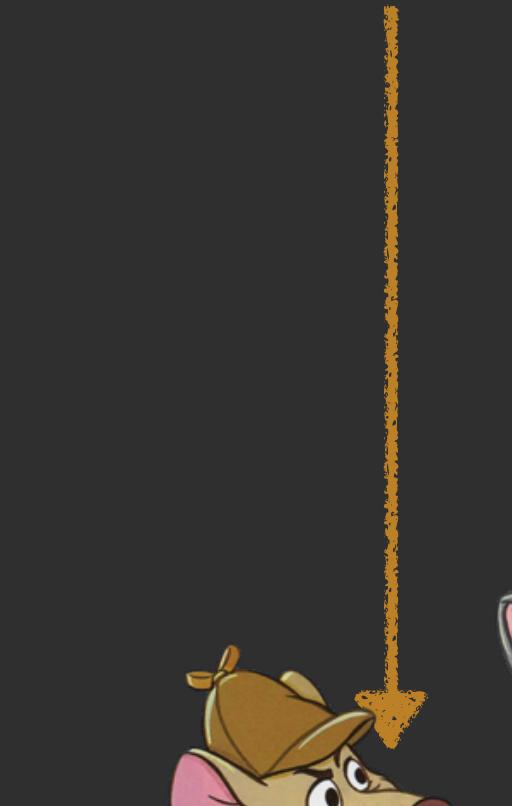
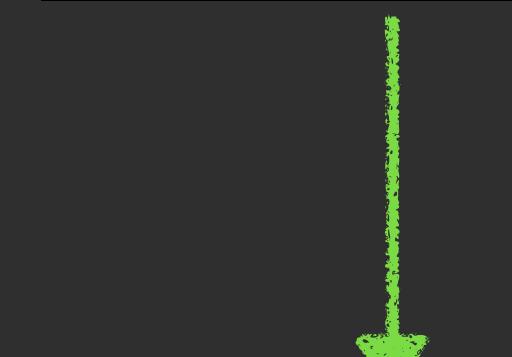
)



3 3 3 3 1 8 6 2
0 3 7 6 4 1 7 3
1 1 8 8 8 2 3 1

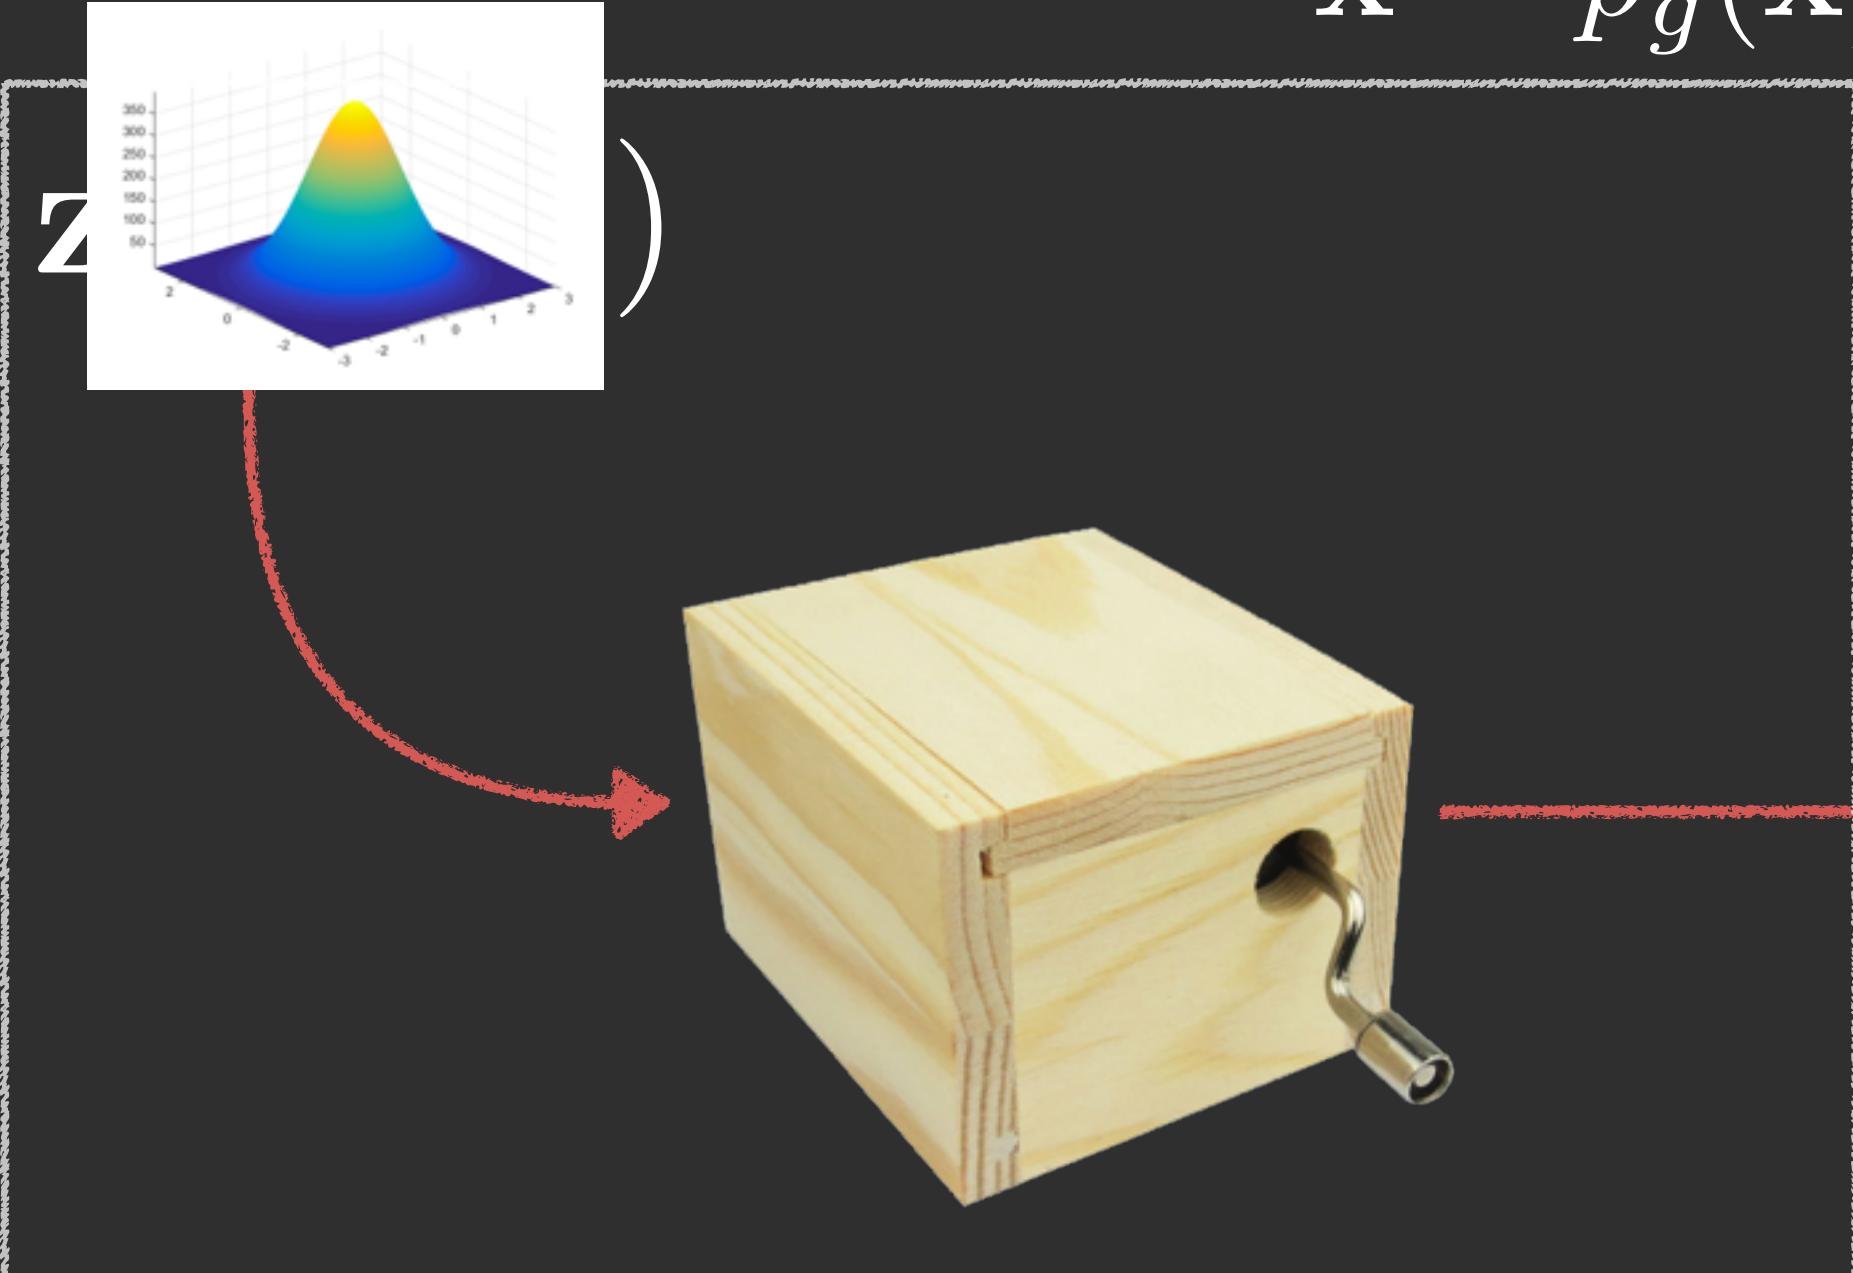
(+)

1 8 3 1 5 5 8 1
6 1 1 8 7 6 0 7
4 8 6 4 2 3 0 3



10% real
50% real
4% real
88% real
....

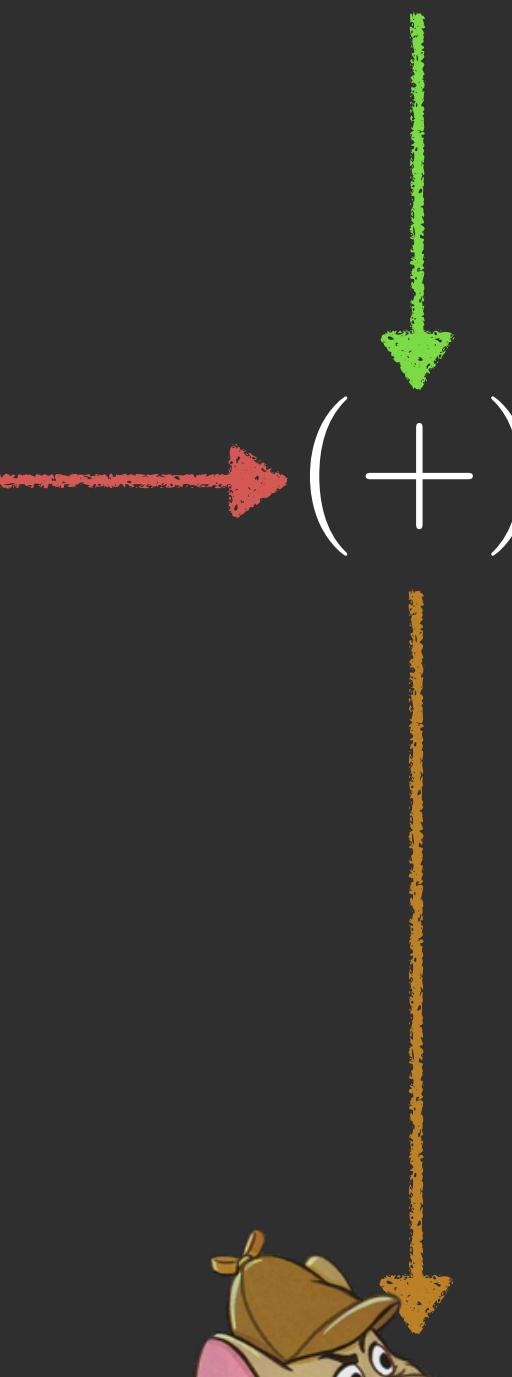
$$\tilde{\mathbf{x}} \sim p_g(\mathbf{x})$$



1 3 3 3 1 8 6 2
0 3 7 6 4 1 7 3
1 1 8 8 8 2 3 1

(+)

1 8 3 1 5 5 8 1
6 1 1 8 7 6 0 7
4 8 6 4 2 3 0 3



How does the generator try to fool the discriminator?

GAN algorithm

GAN algorithm

- Let D denote the discriminator, with $D(\mathbf{x})$ outputting a probability

GAN algorithm

- Let D denote the discriminator, with $D(\mathbf{x})$ outputting a probability
- D wants to do the following:

GAN algorithm

- Let D denote the discriminator, with $D(\mathbf{x})$ outputting a probability
- D wants to do the following:
 - For real \mathbf{x} , update its params so as to minimise $L(D(\mathbf{x}), 1)$

GAN algorithm

- Let D denote the discriminator, with $D(\mathbf{x})$ outputting a probability
- D wants to do the following:
 - For real \mathbf{x} , update its params so as to minimise $L(D(\mathbf{x}), 1)$
 - For $G(z)$, update its params so as to minimise $L(D(G(z)), 0)$

GAN algorithm

- Let D denote the discriminator, with $D(\mathbf{x})$ outputting a probability
- D wants to do the following:
 - For real \mathbf{x} , update its params so as to minimise $L(D(\mathbf{x}), 1)$
 - For $G(z)$, update its params so as to minimise $L(D(G(z)), 0)$
- G wants to do the following:

GAN algorithm

- Let D denote the discriminator, with $D(\mathbf{x})$ outputting a probability
- D wants to do the following:
 - For real \mathbf{x} , update its params so as to minimise $L(D(\mathbf{x}), 1)$
 - For $G(z)$, update its params so as to minimise $L(D(G(z)), 0)$
- G wants to do the following:
 - Update its params so as to minimise $L(D(G(z)), 1)$

GAN algorithm

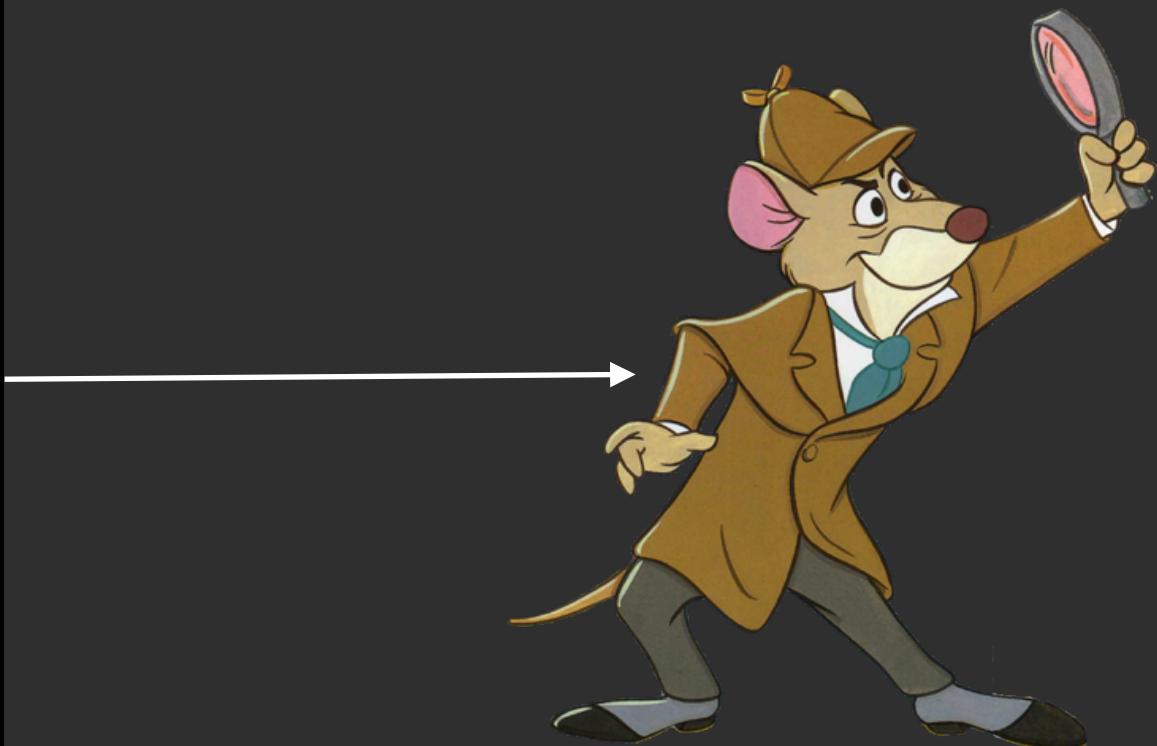
- Let D denote the discriminator, with $D(\mathbf{x})$ outputting a probability
- D wants to do the following:
 - For real \mathbf{x} , update its params so as to minimise $L(D(\mathbf{x}), 1)$
 - For $G(z)$, update its params so as to minimise $L(D(G(z)), 0)$
- G wants to do the following:
 - Update its params so as to minimise $L(D(G(z)), 1)$
 - Denote $L(y_1, y_2)$ to be squared error: $(y_1 - y_2)^2$

GAN algorithm

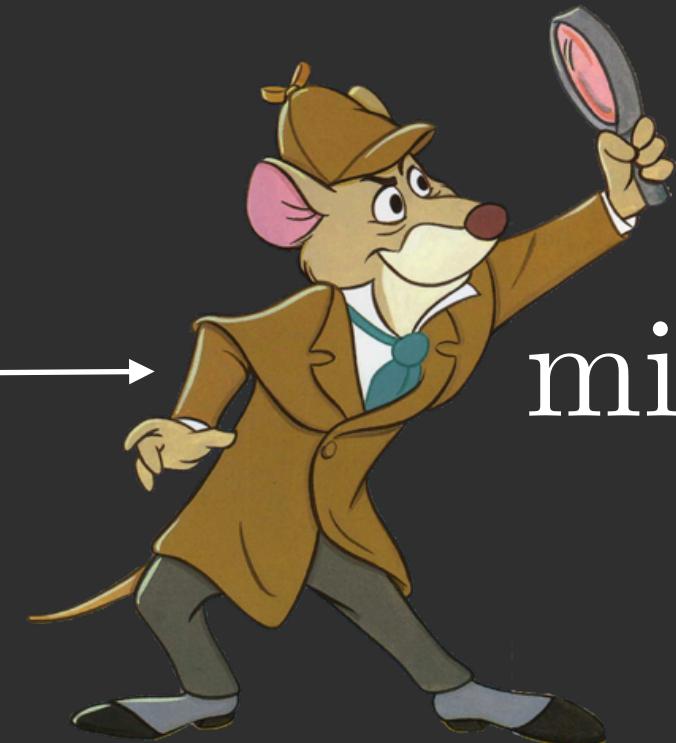
- Let D denote the discriminator, with $D(\mathbf{x})$ outputting a probability
- D wants to do the following:
 - For real \mathbf{x} , update its params so as to minimise $L(D(\mathbf{x}), 1)$
 - For $G(z)$, update its params so as to minimise $L(D(G(z)), 0)$
- G wants to do the following:
 - Update its params so as to minimise $L(D(G(z)), 1)$
- Denote $L(y_1, y_2)$ to be squared error: $(y_1 - y_2)^2$
- LSGAN (Mao et al)

五
〇
七
)

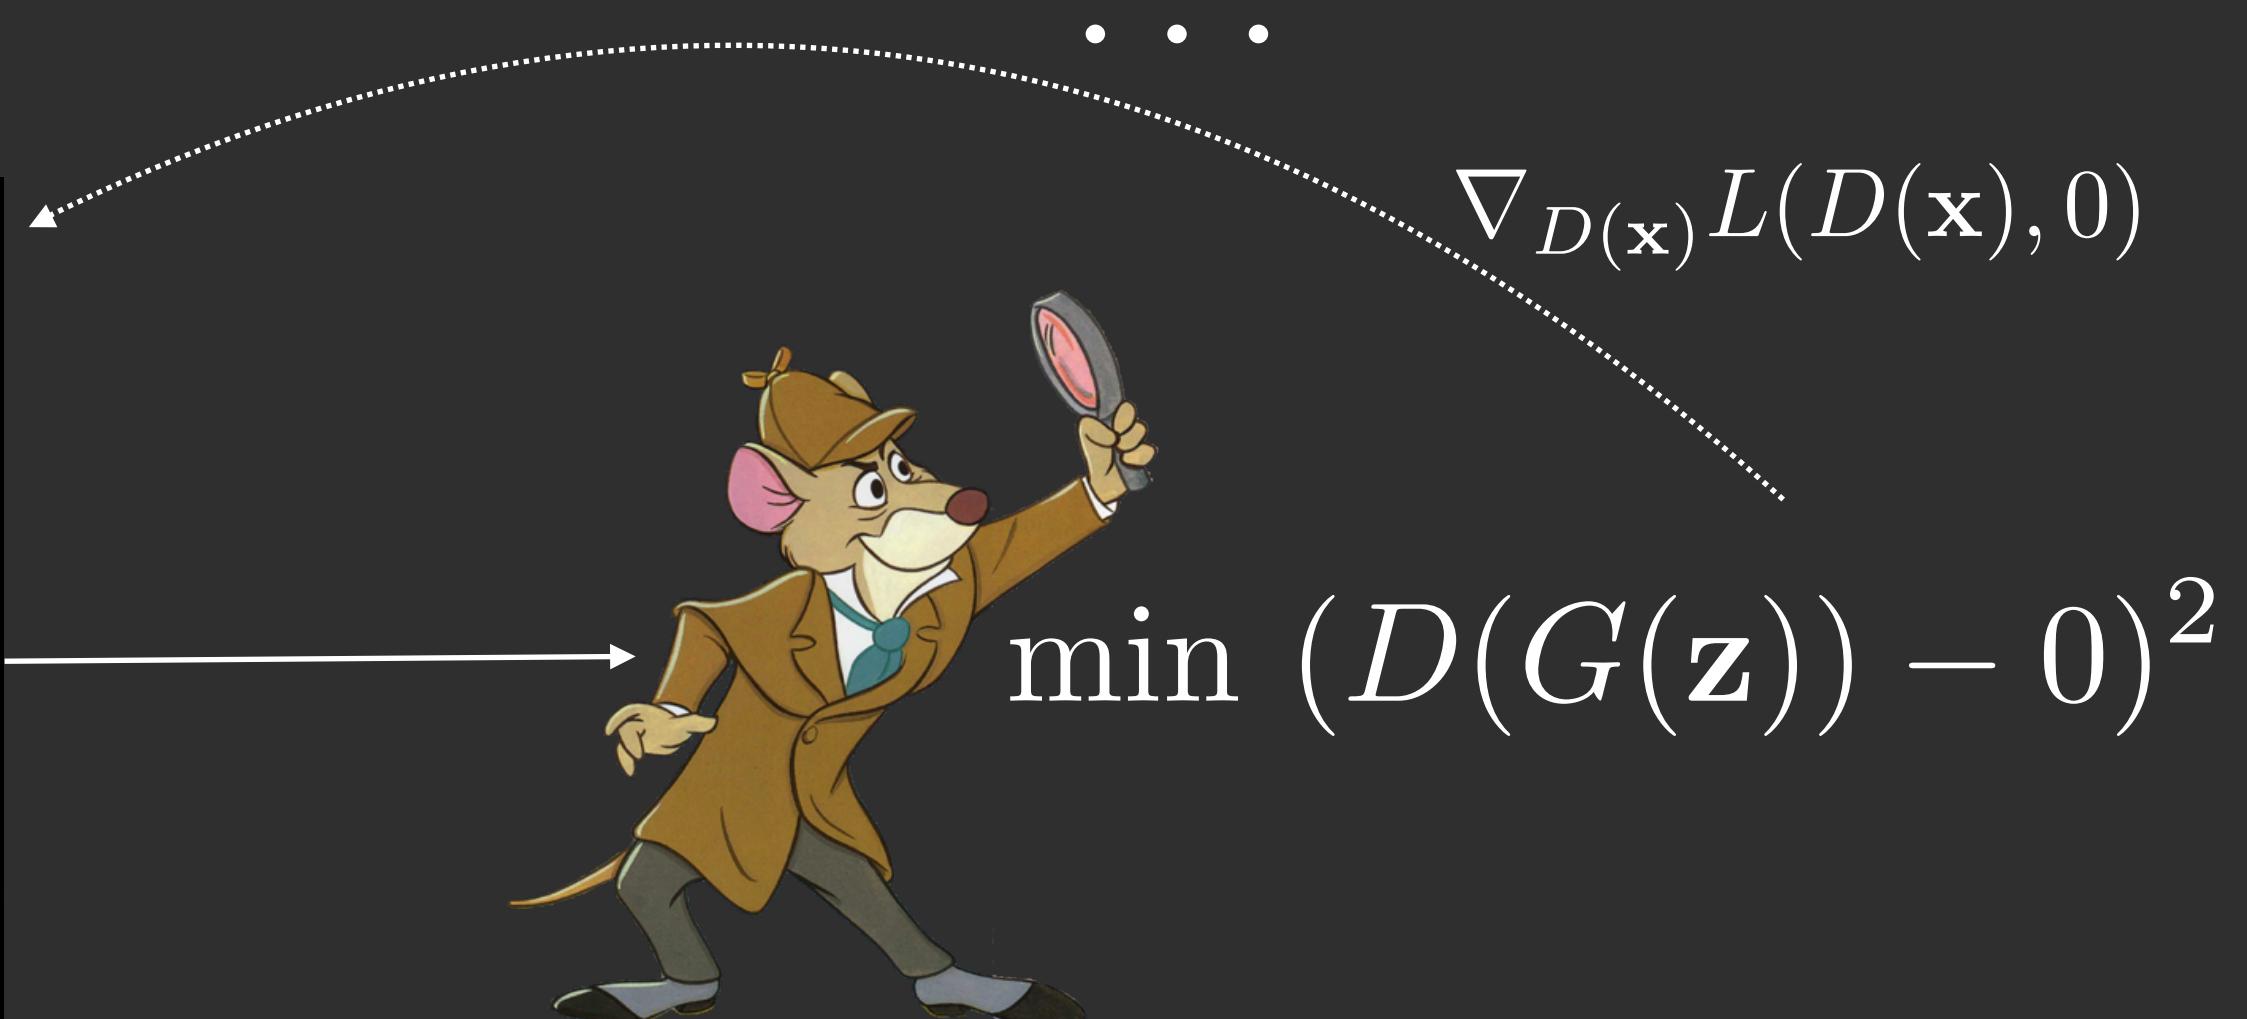
3
0
7
)



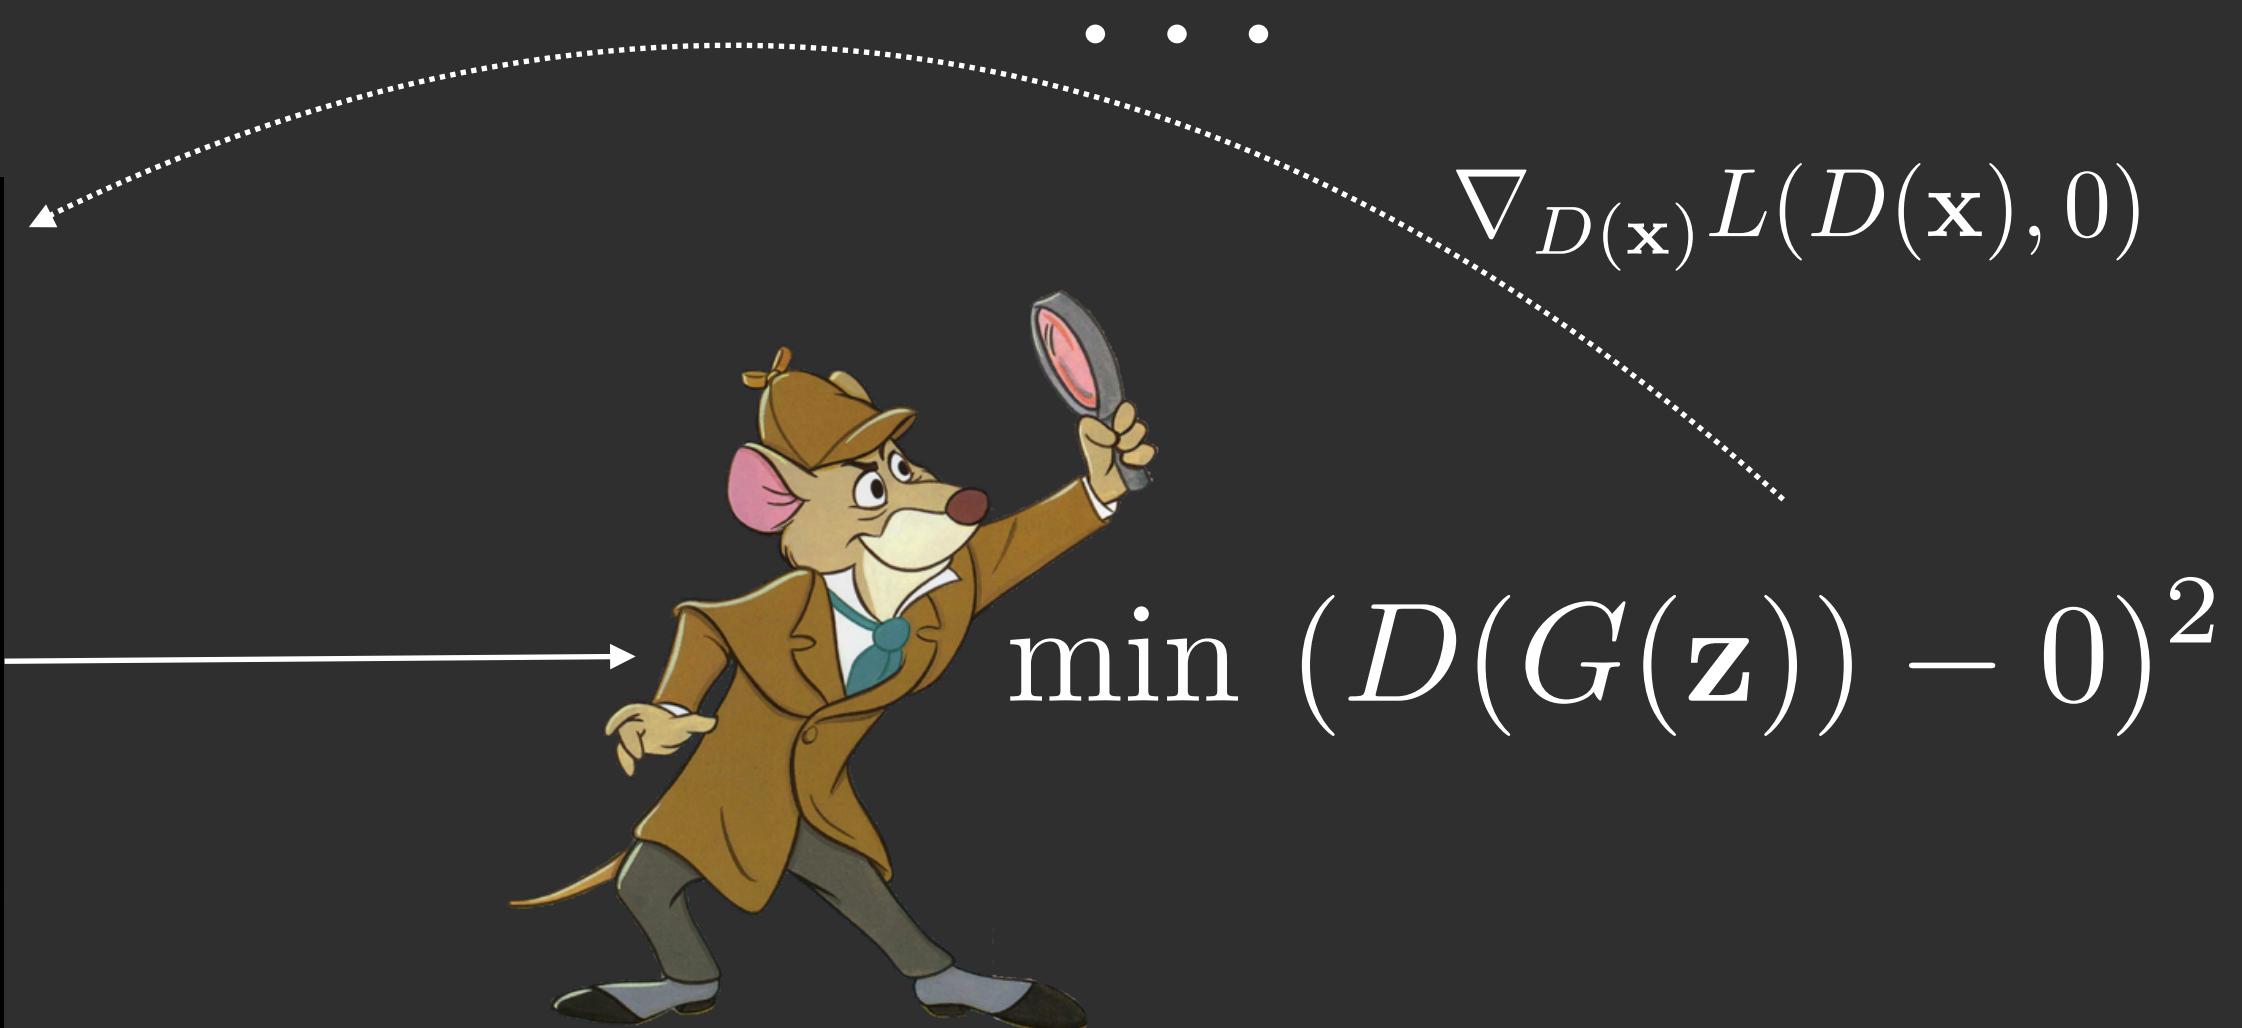
3
0
7
}



$$\min (D(G(\mathbf{z})) - 0)^2$$

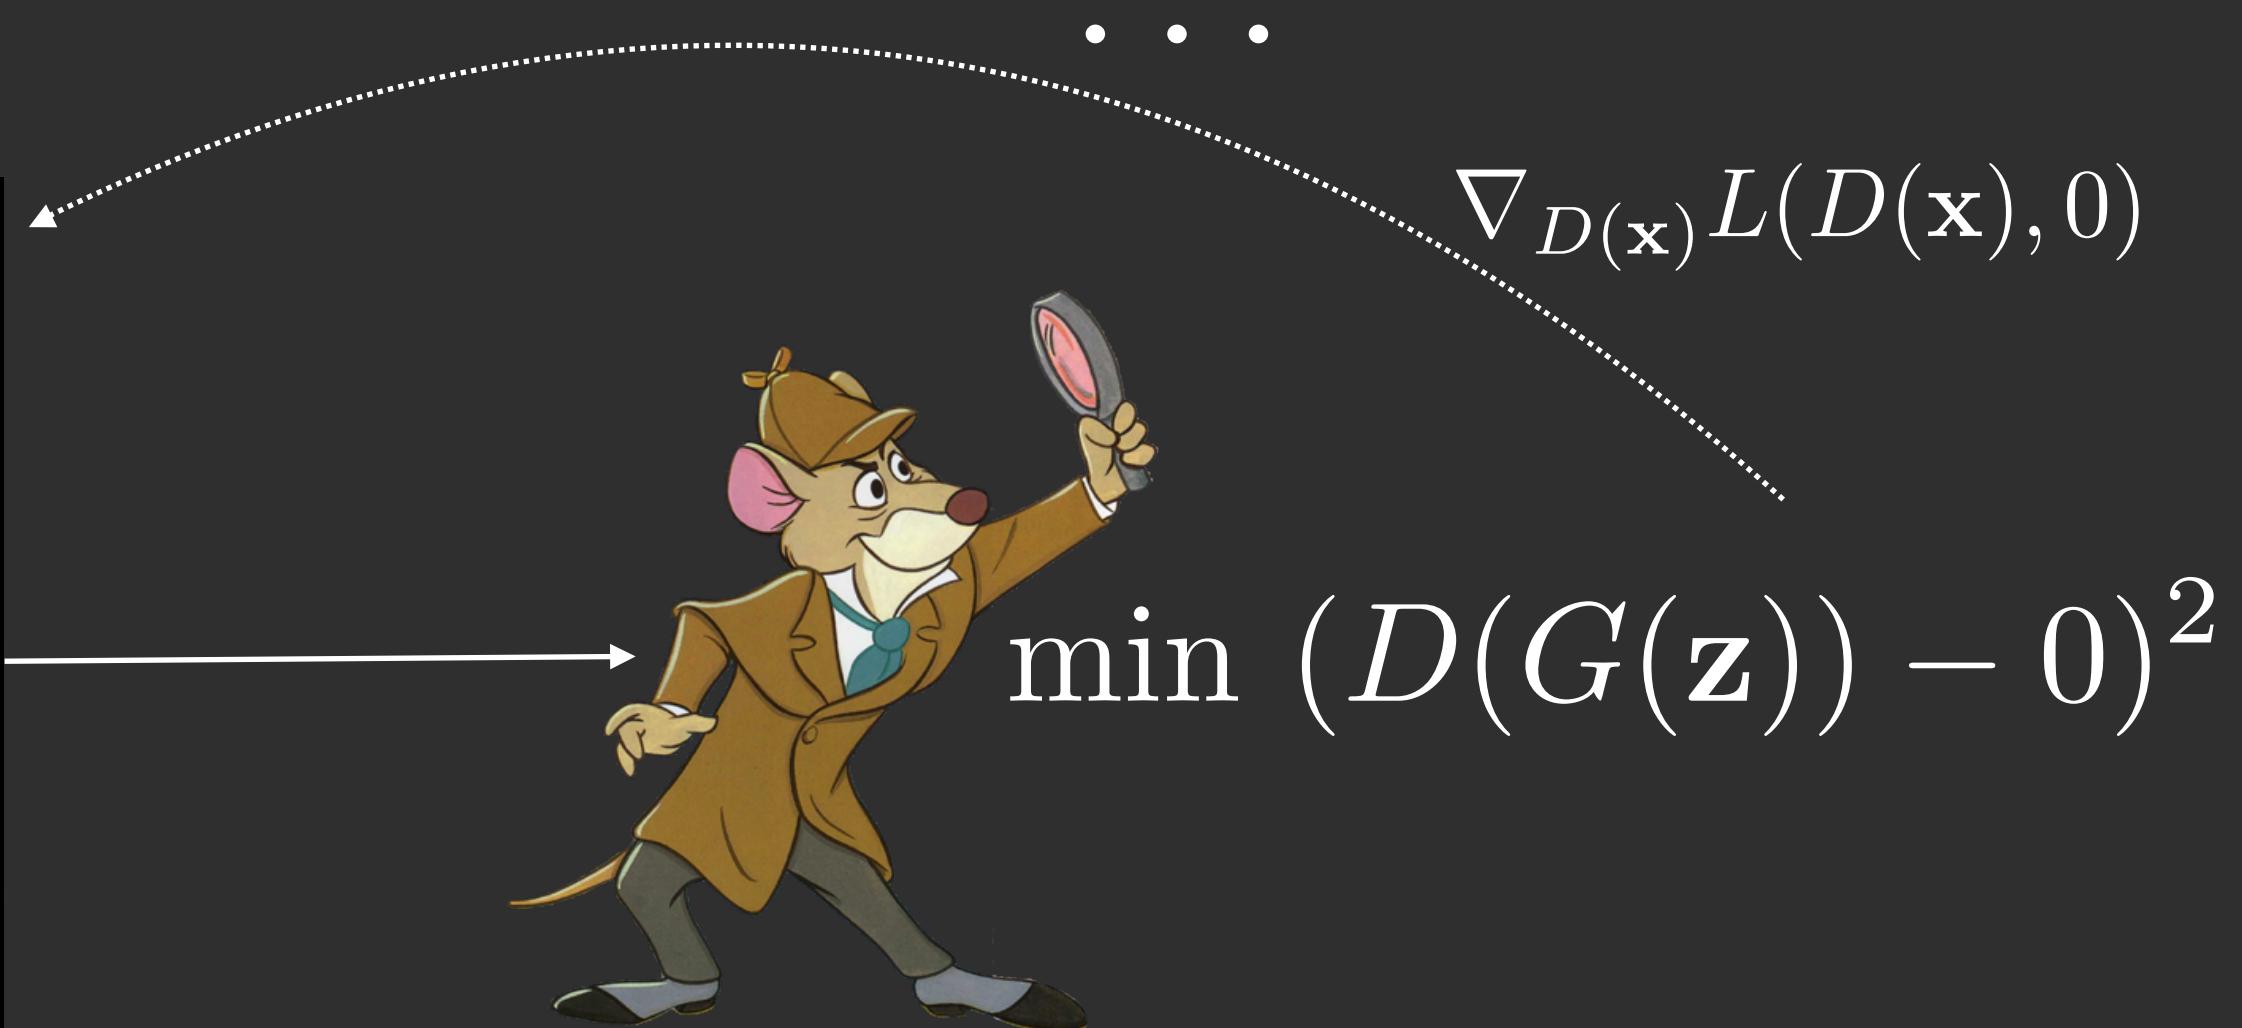


3
0
7
)



6
4
3
B

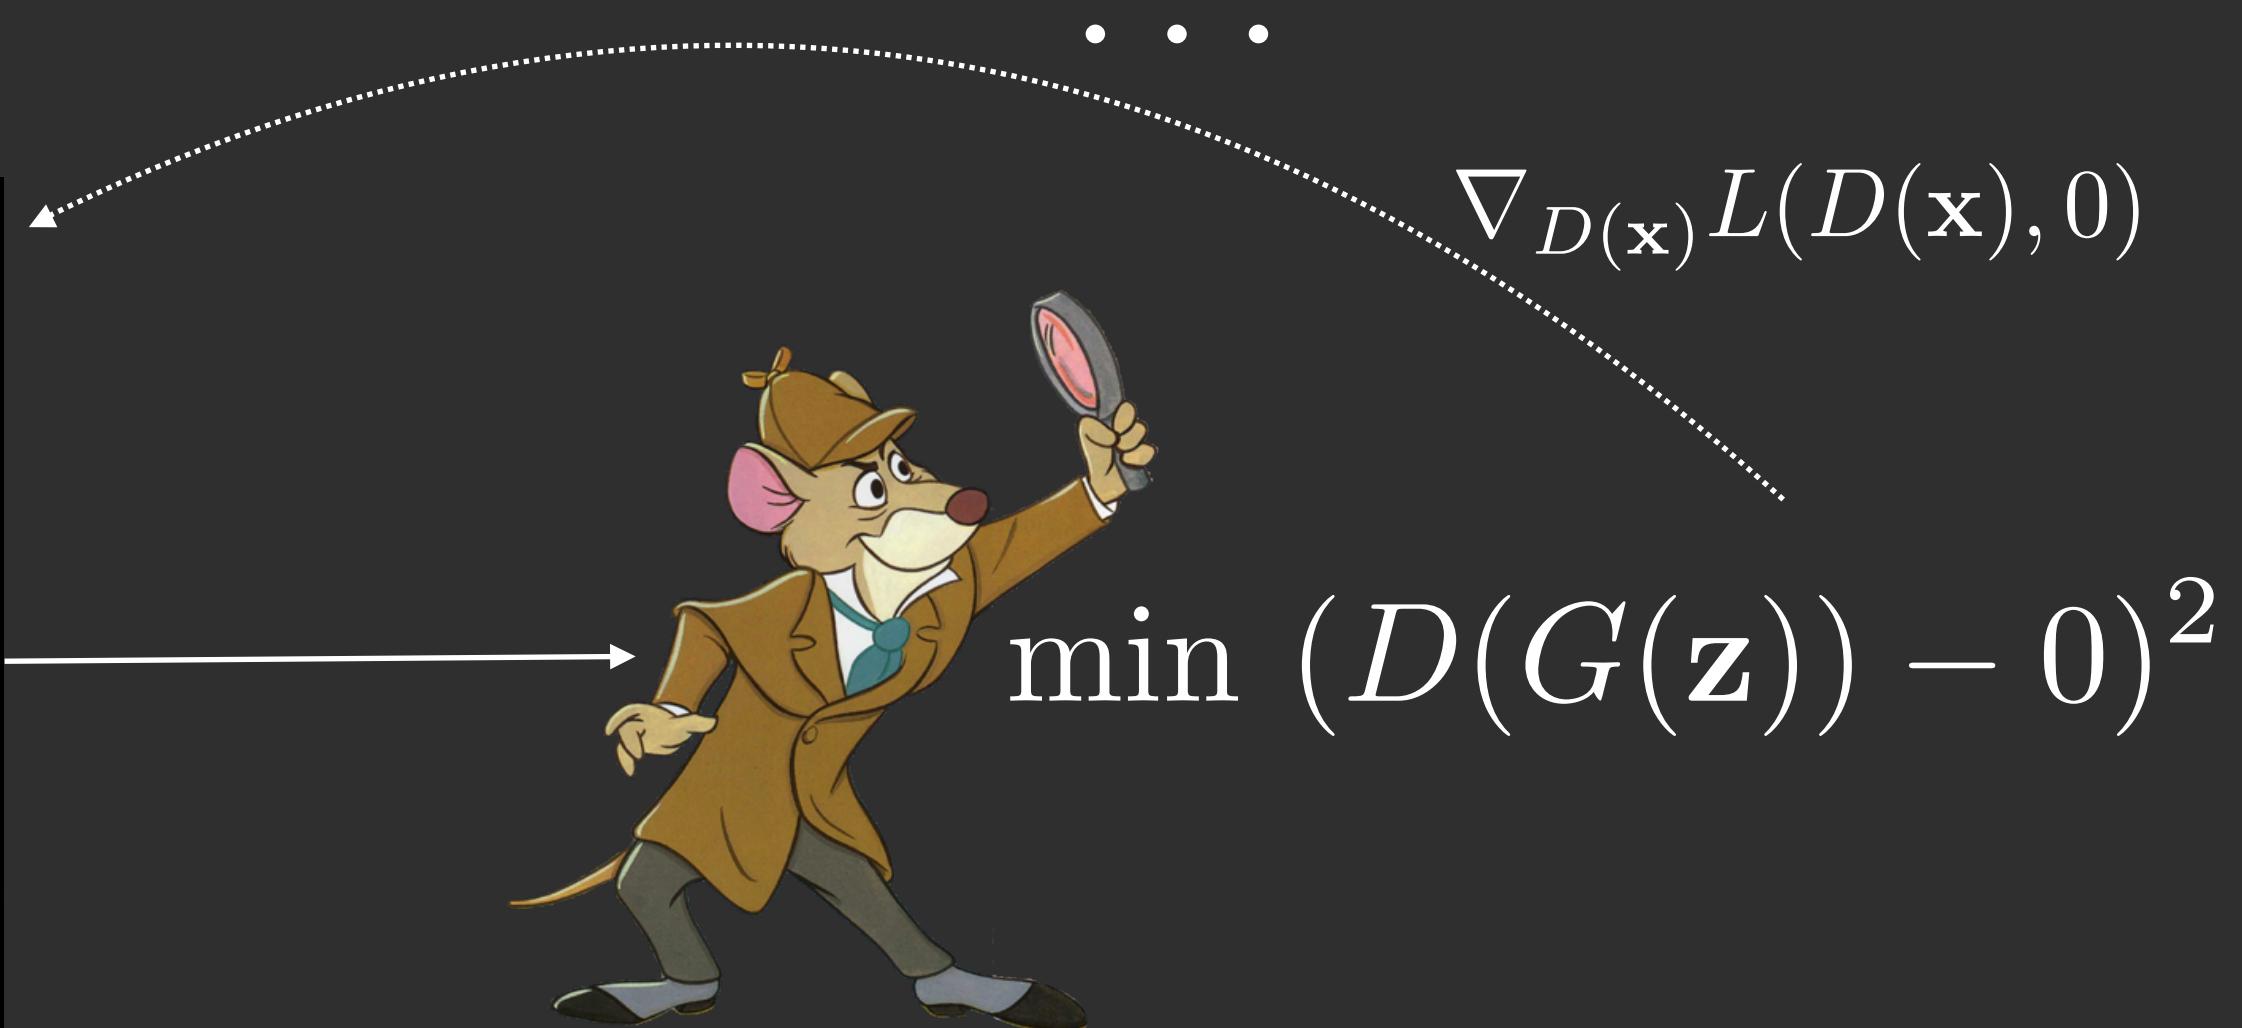
3
0
7
)



6
4
3
B



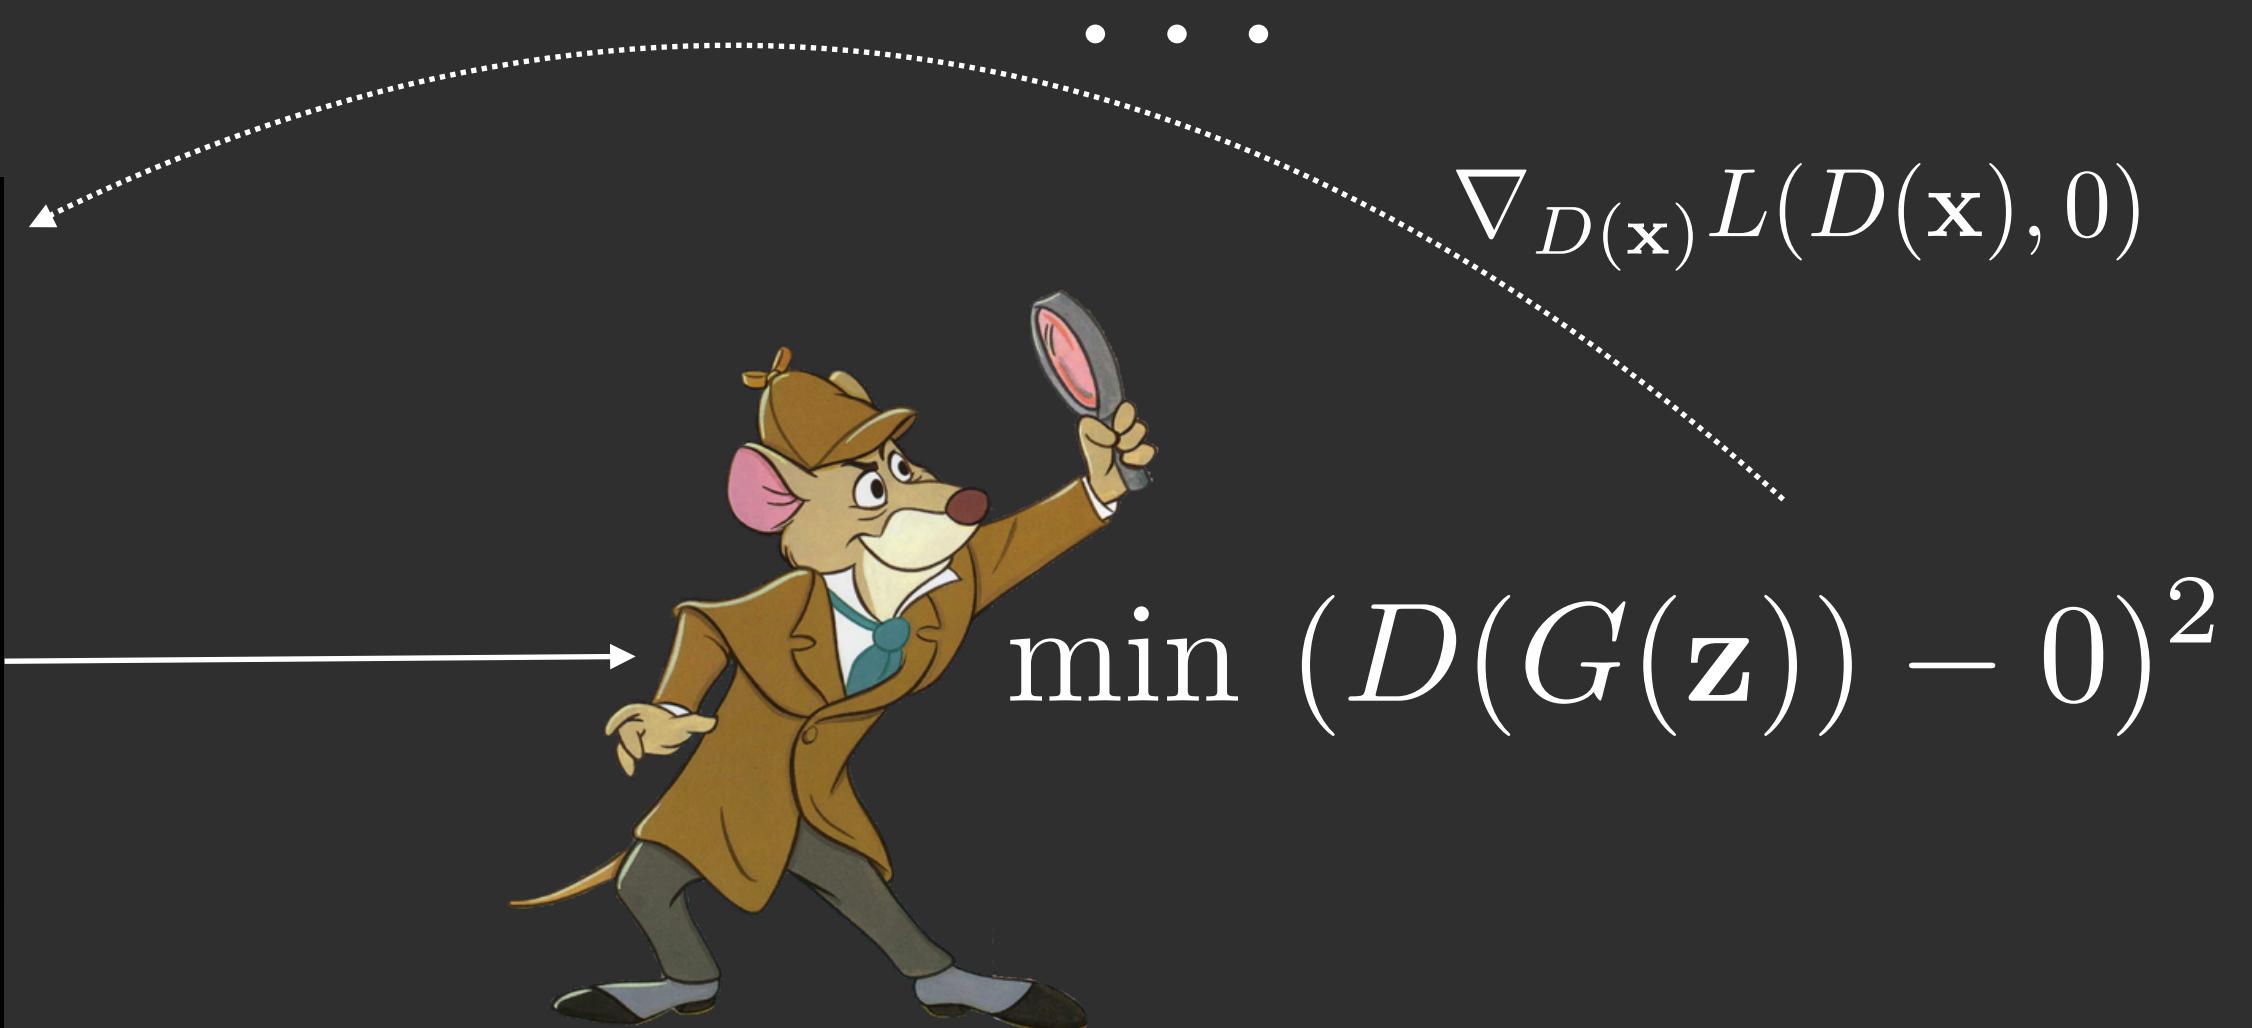
3
0
7
)



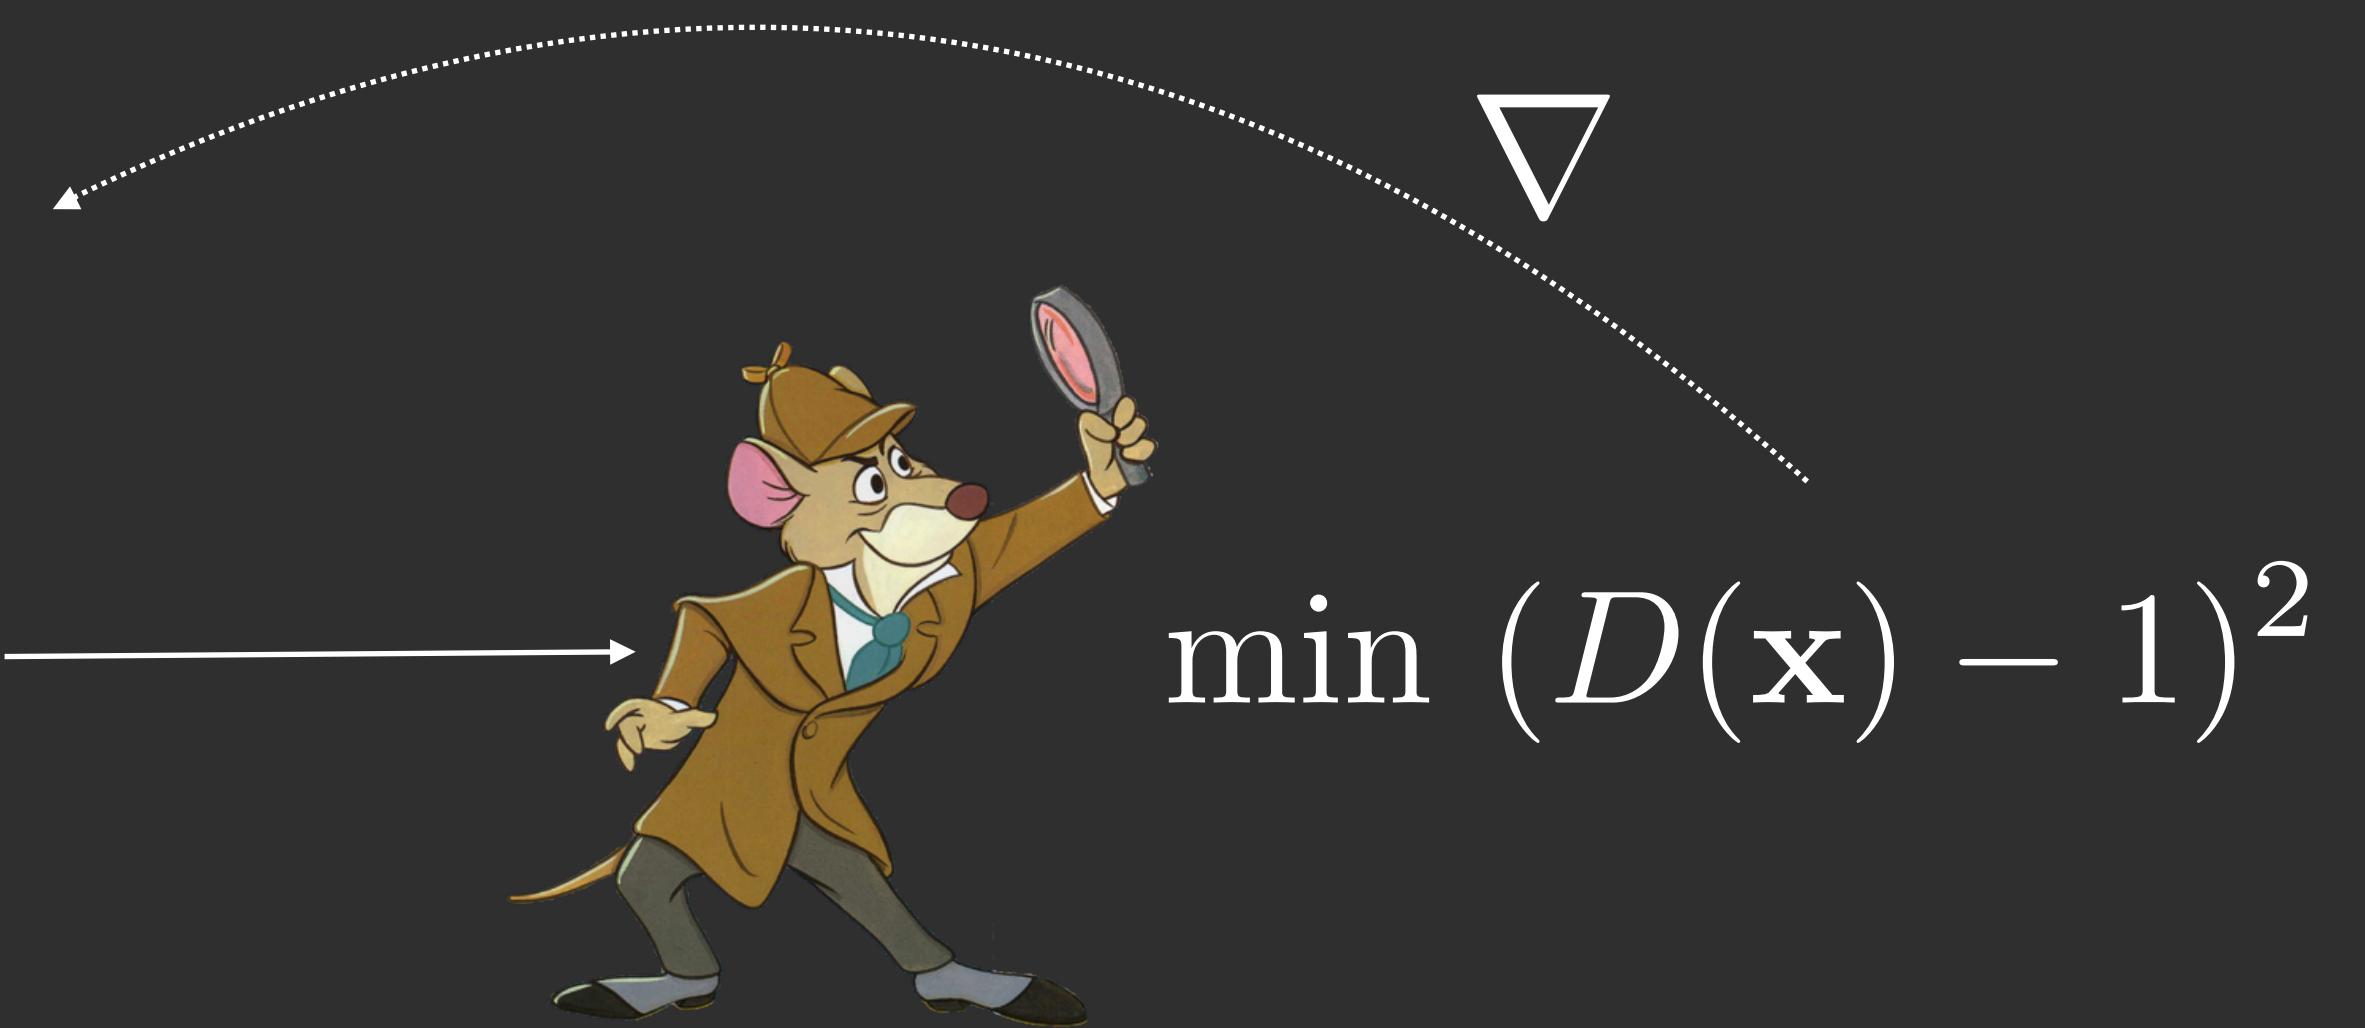
6
4
3
B

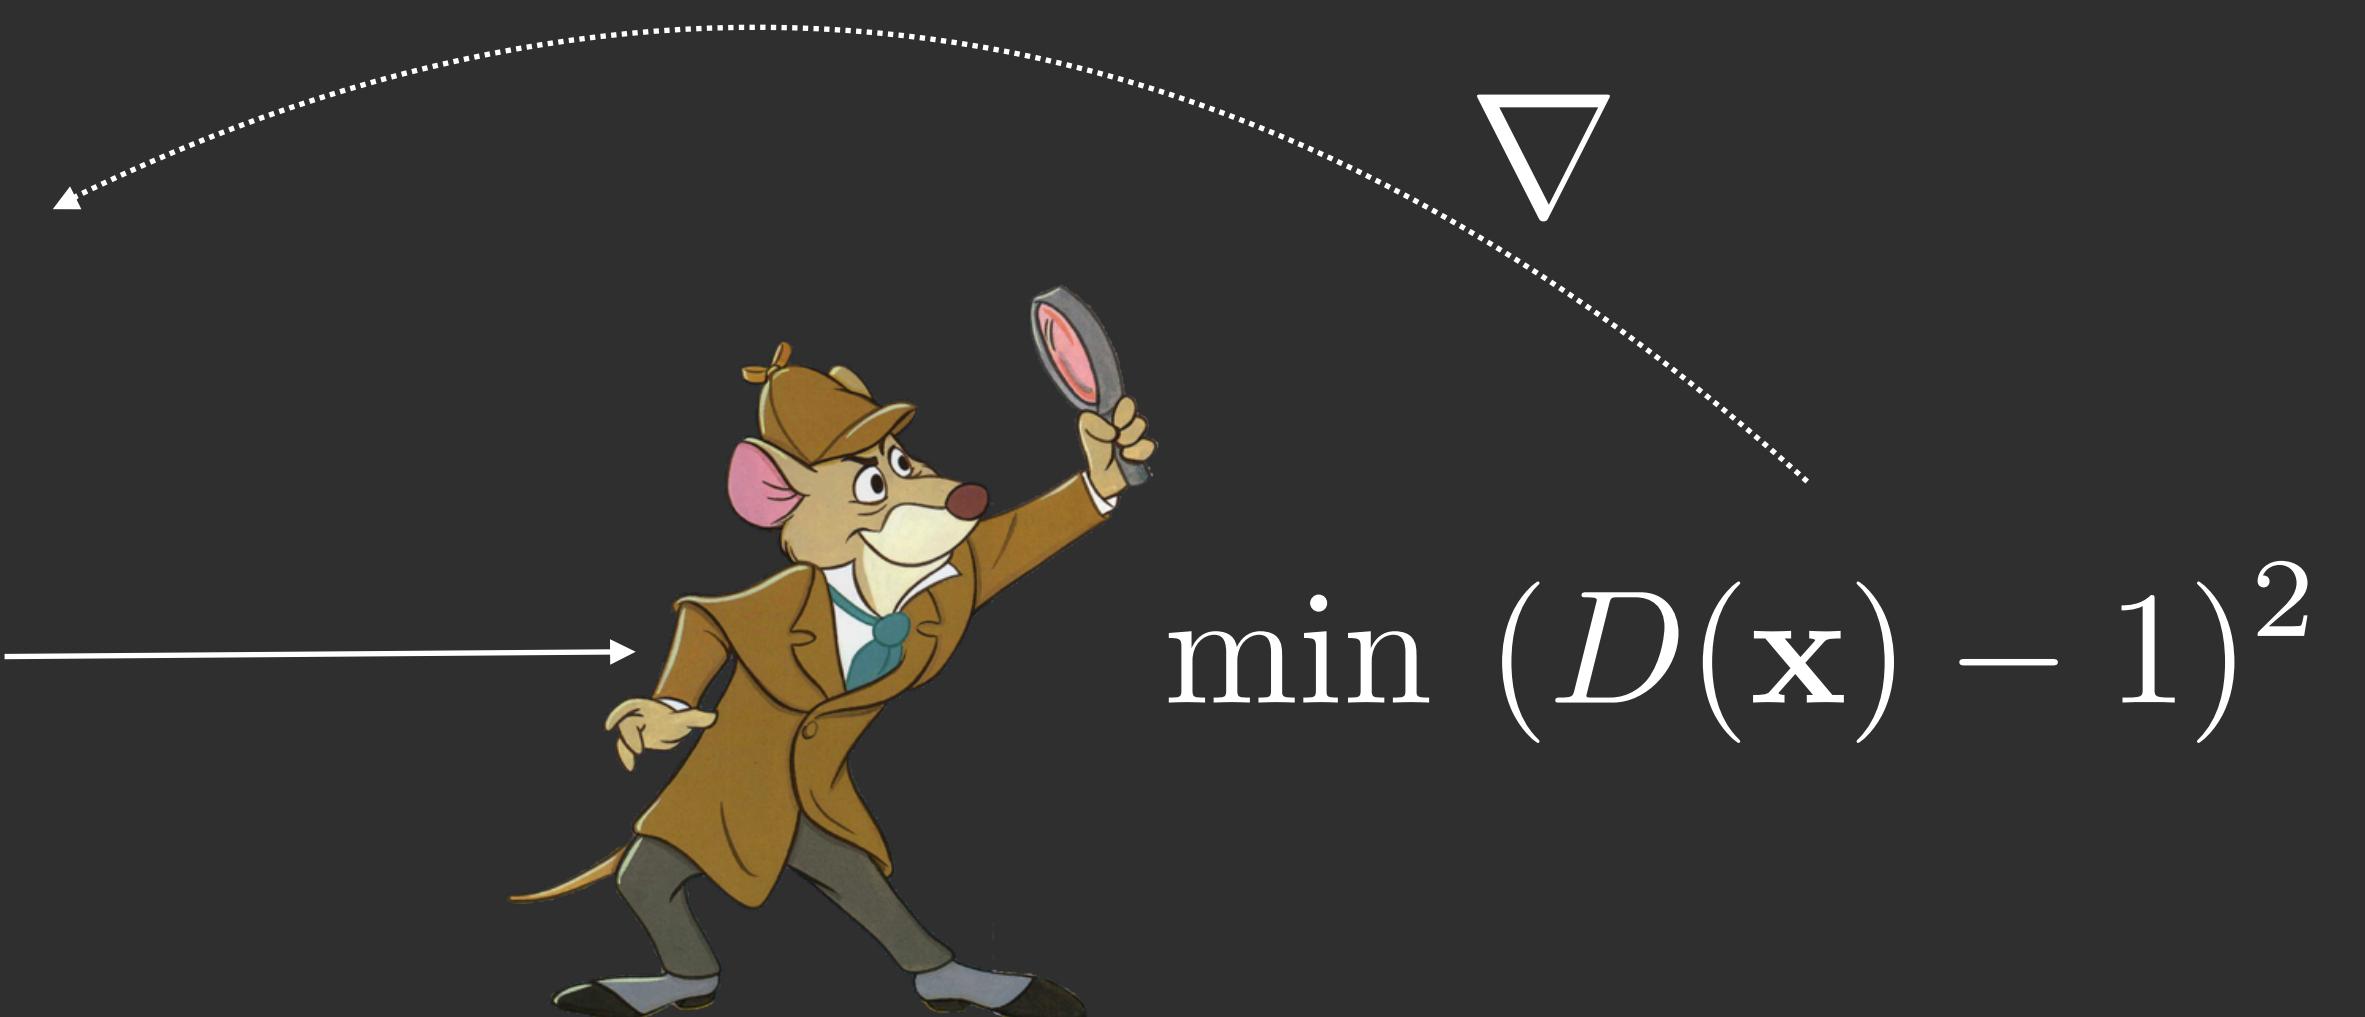
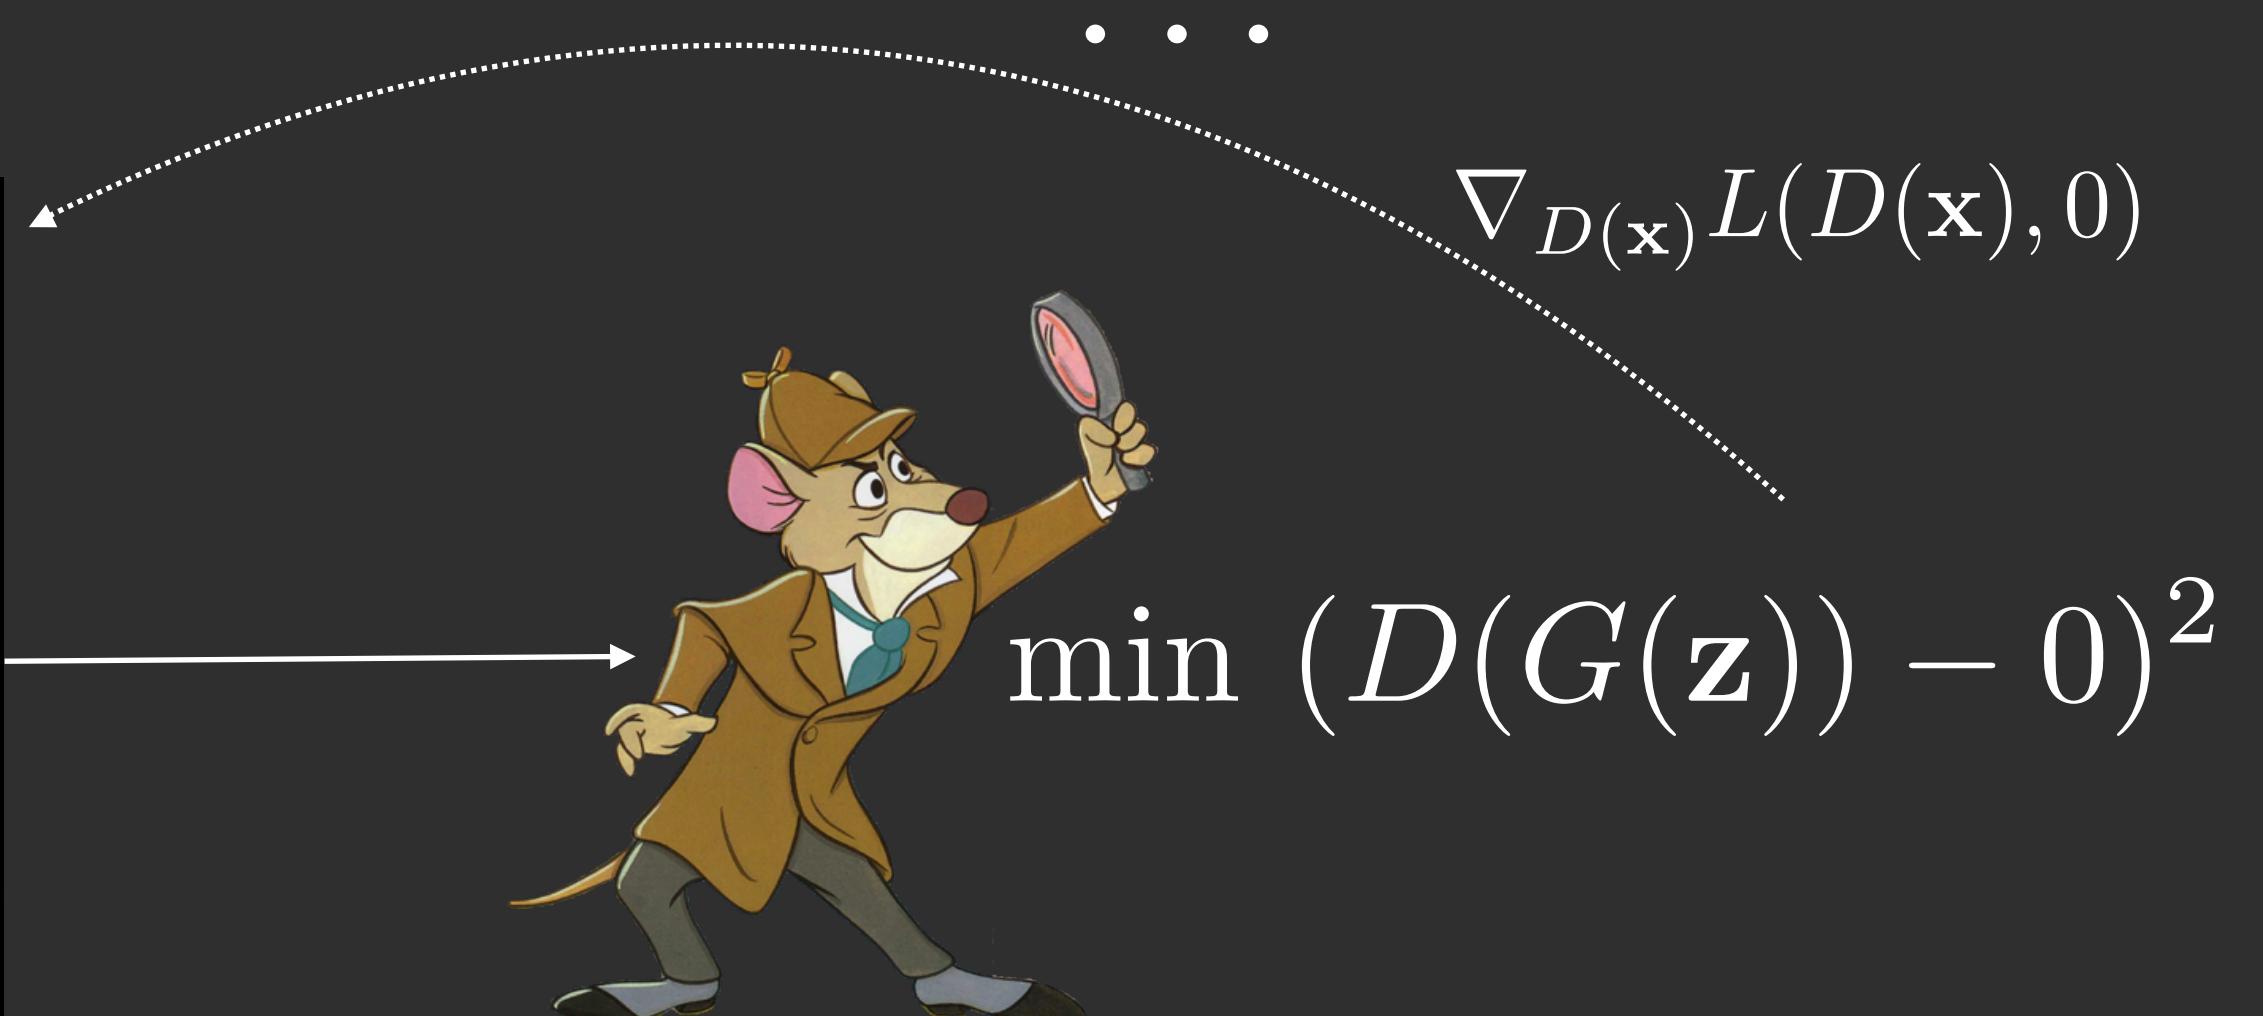


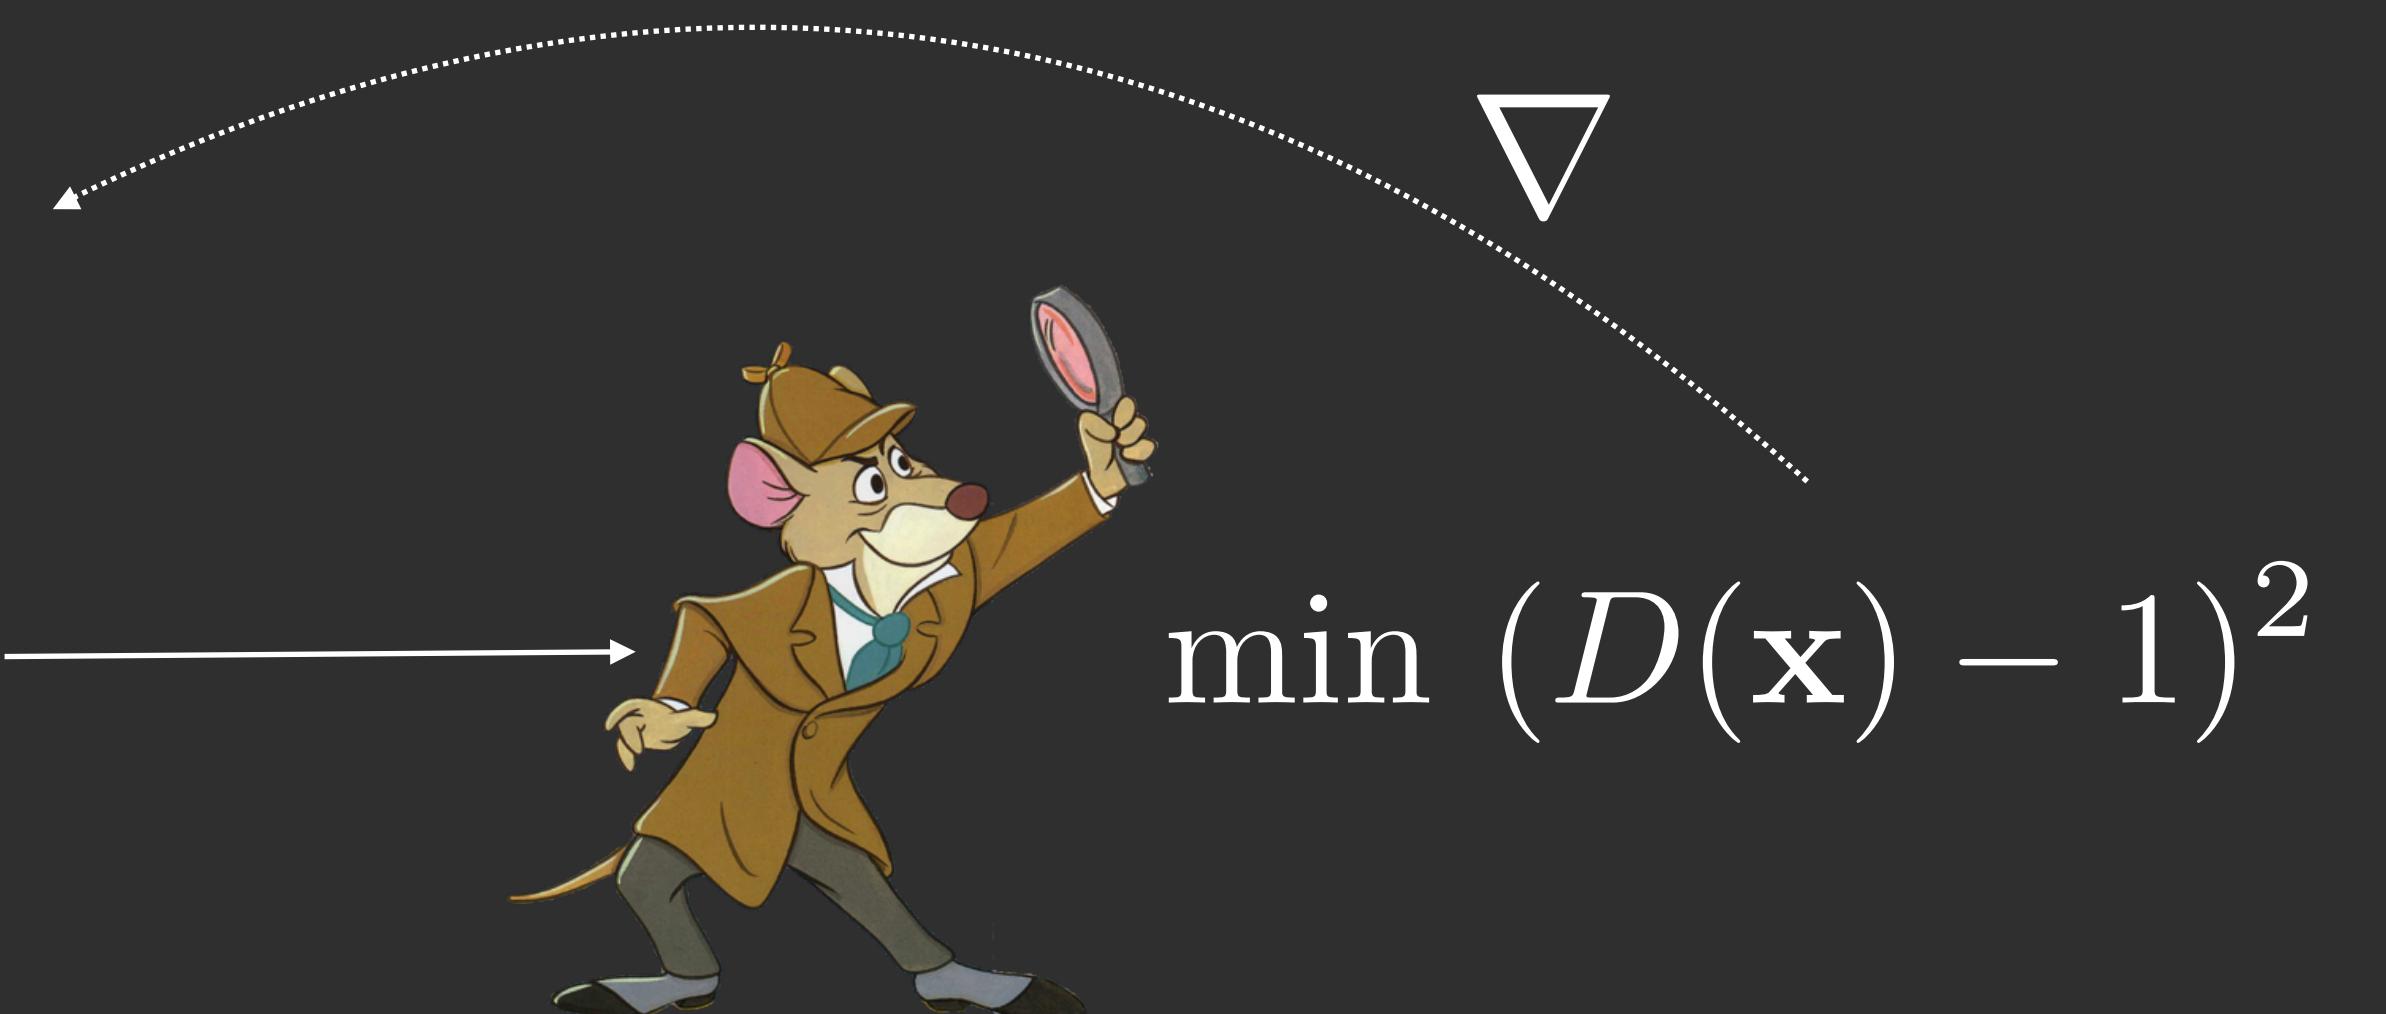
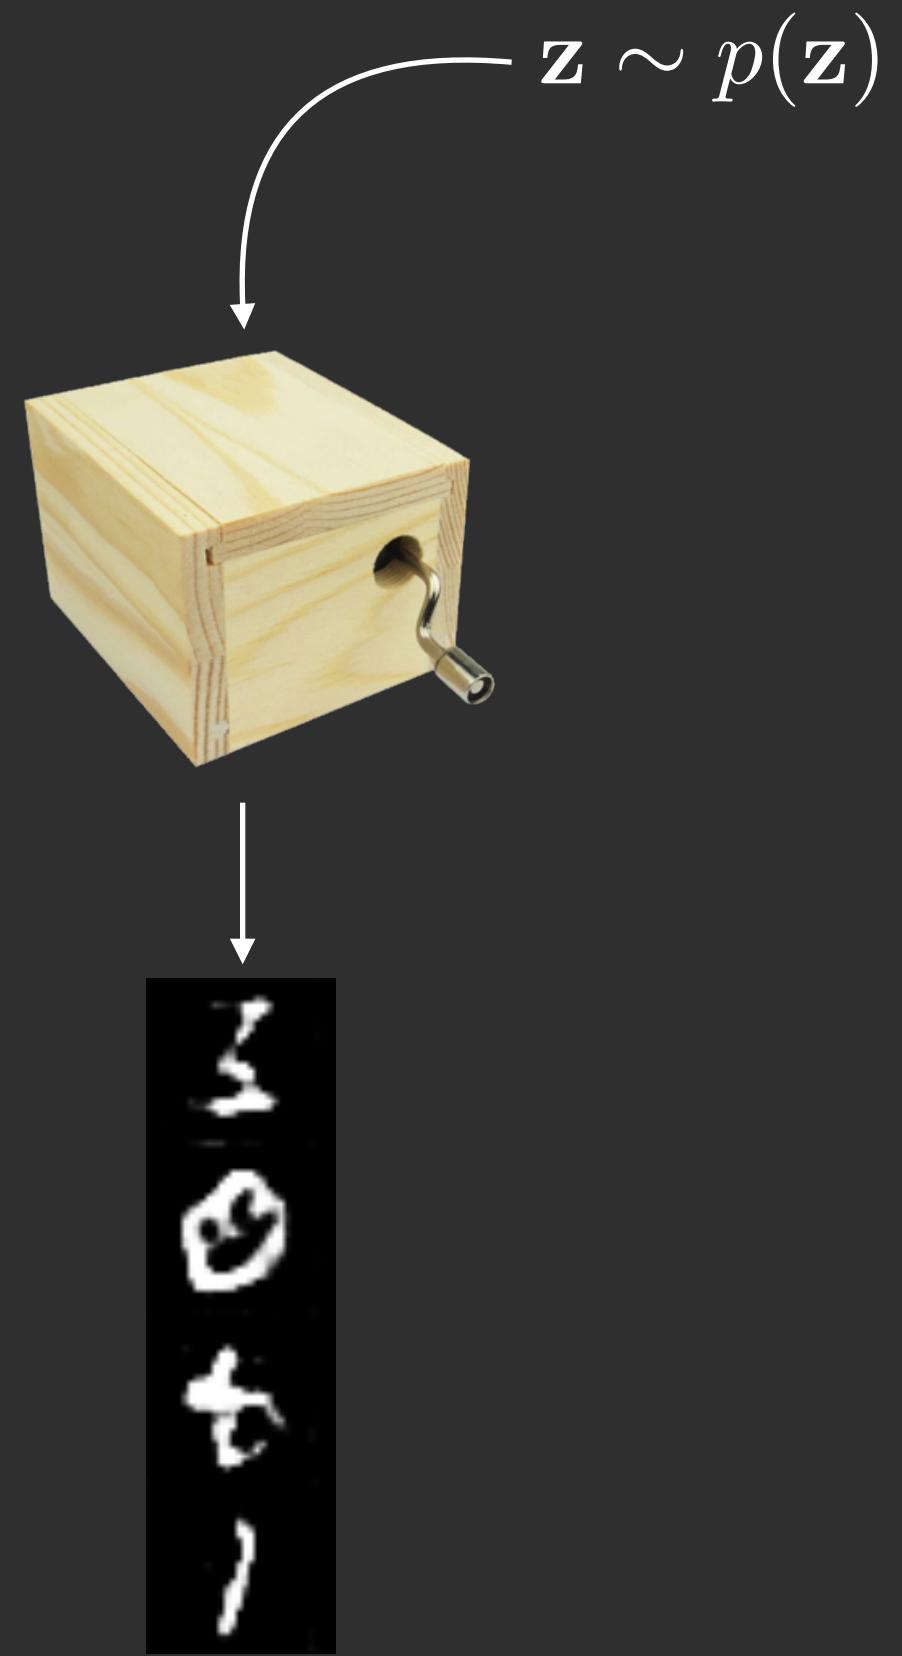
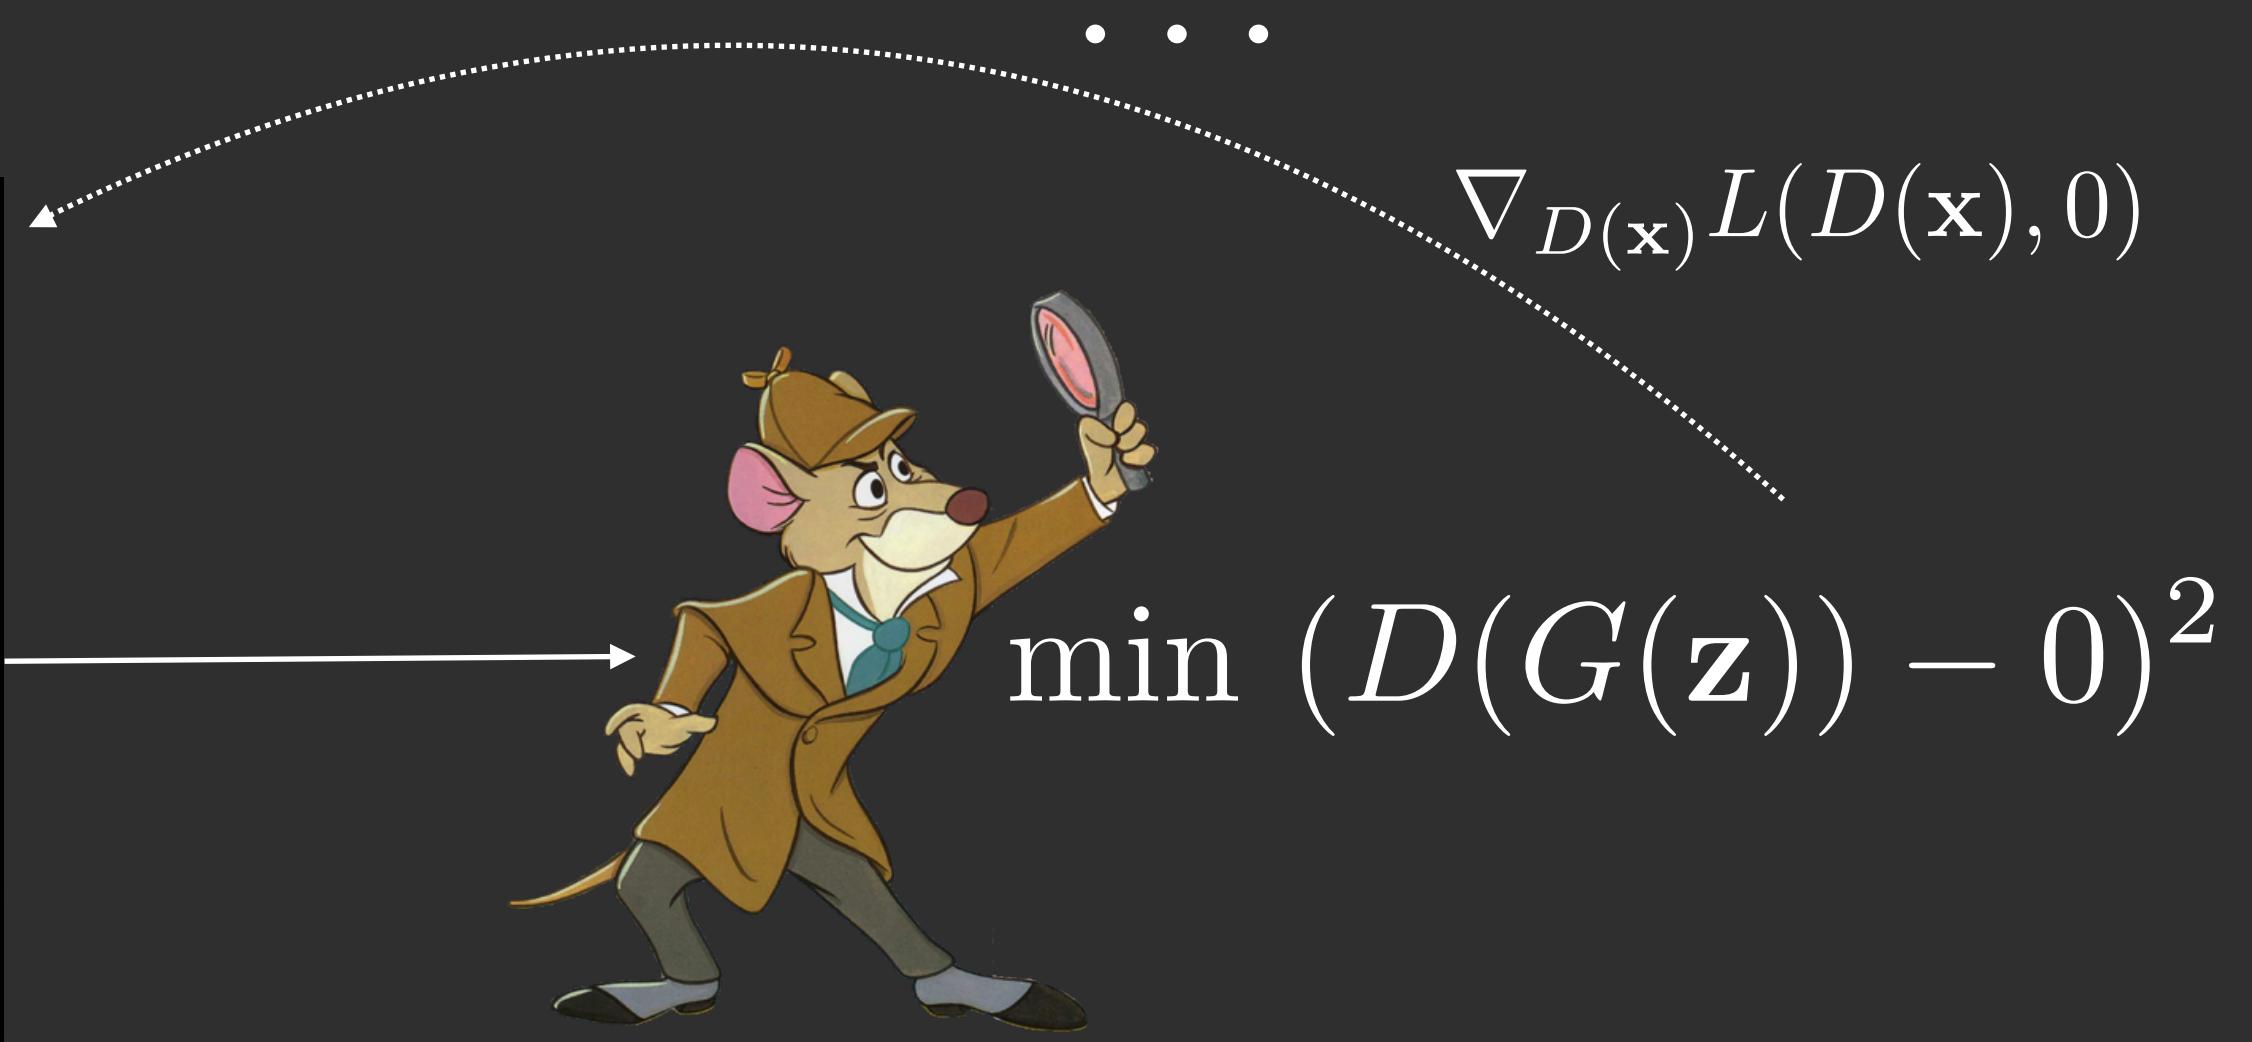
3
0
7
)

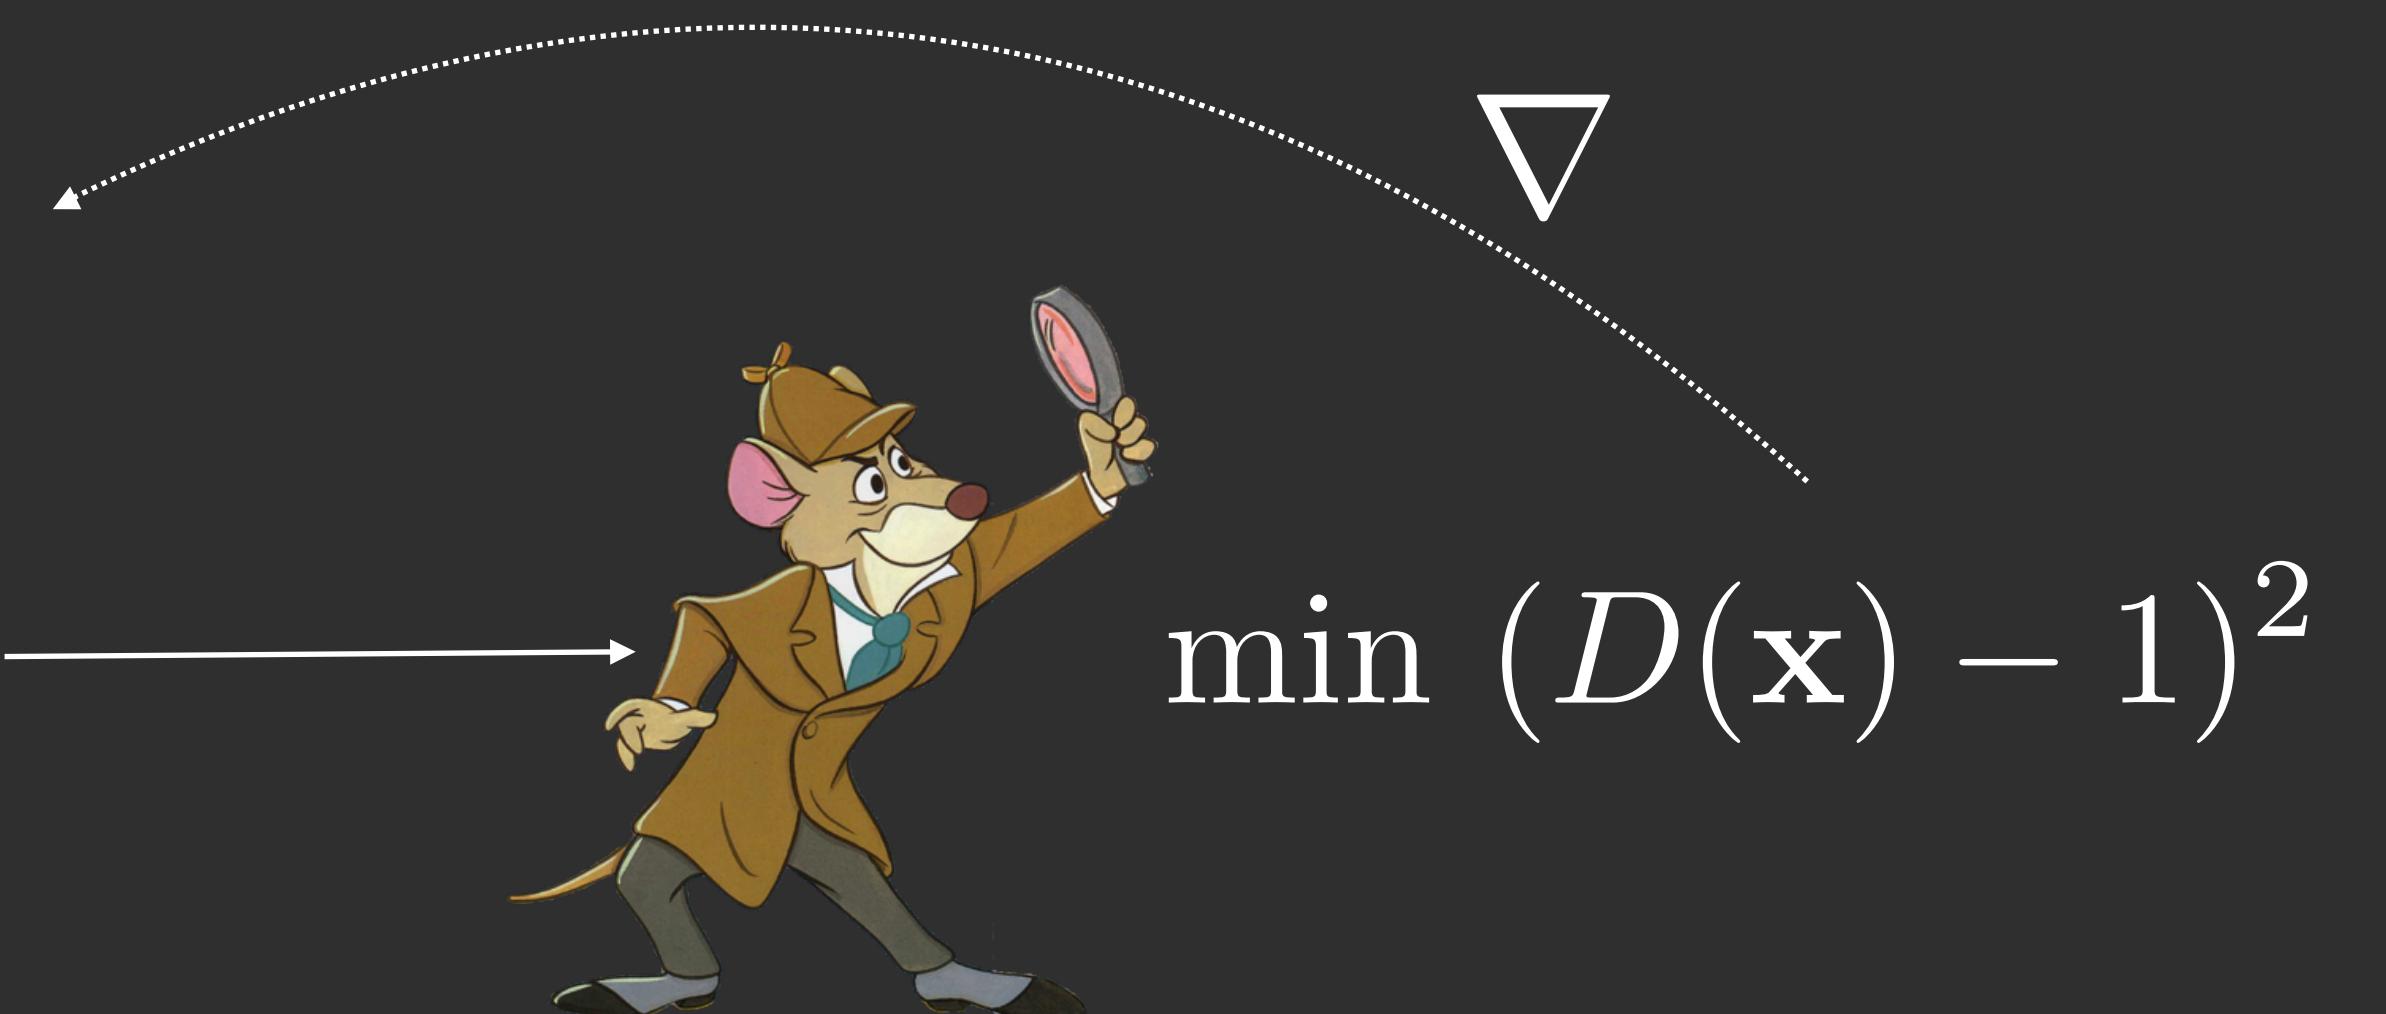
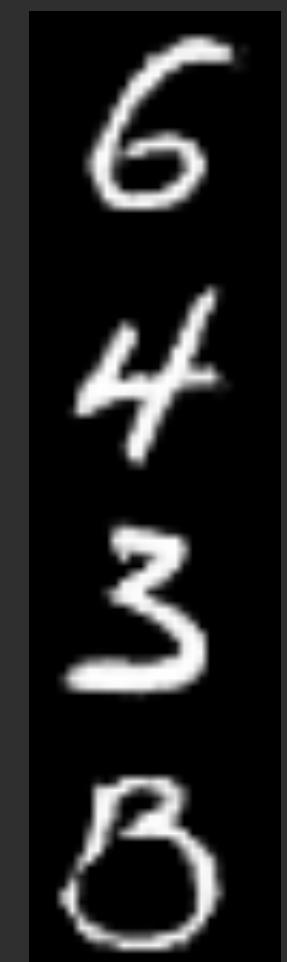
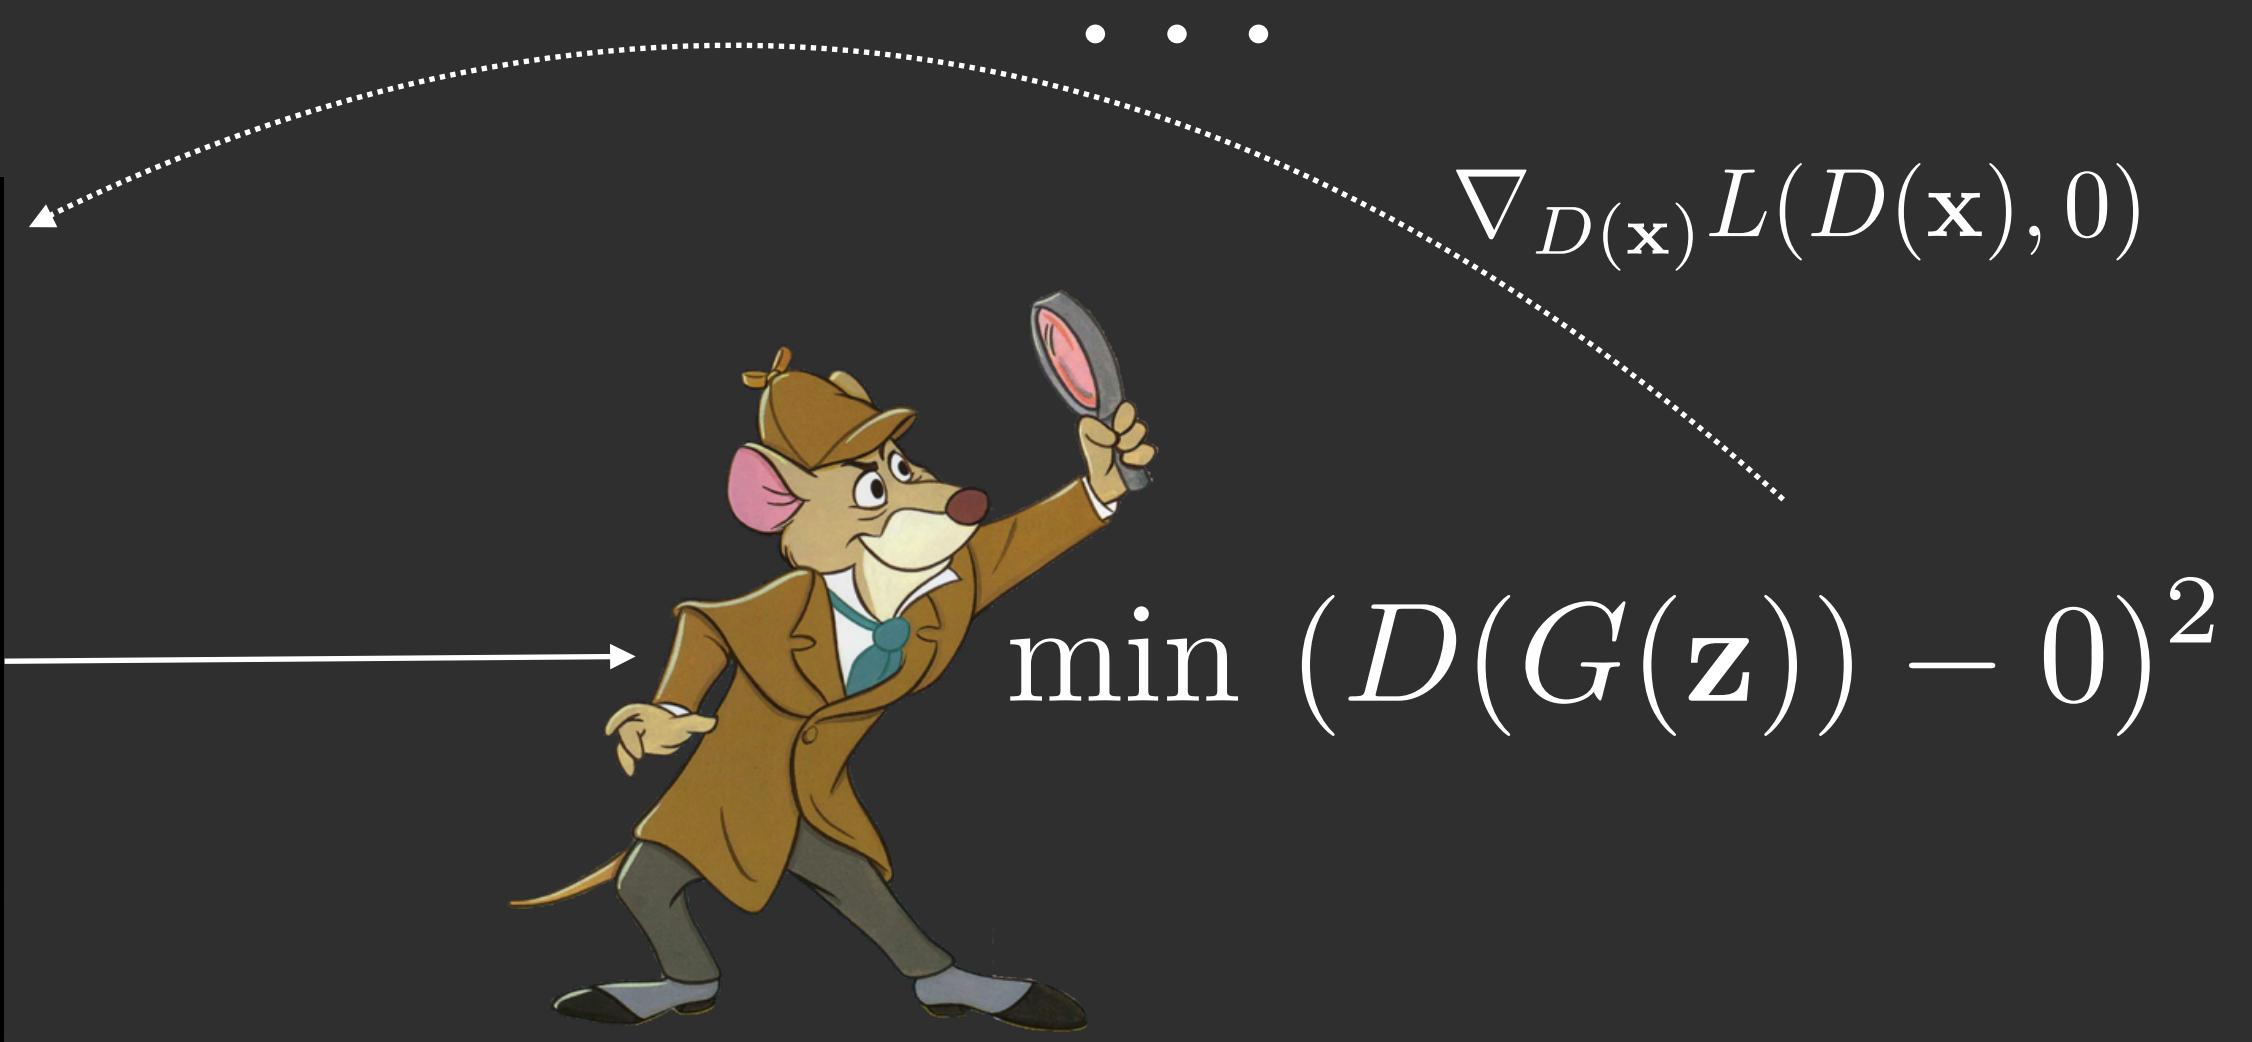


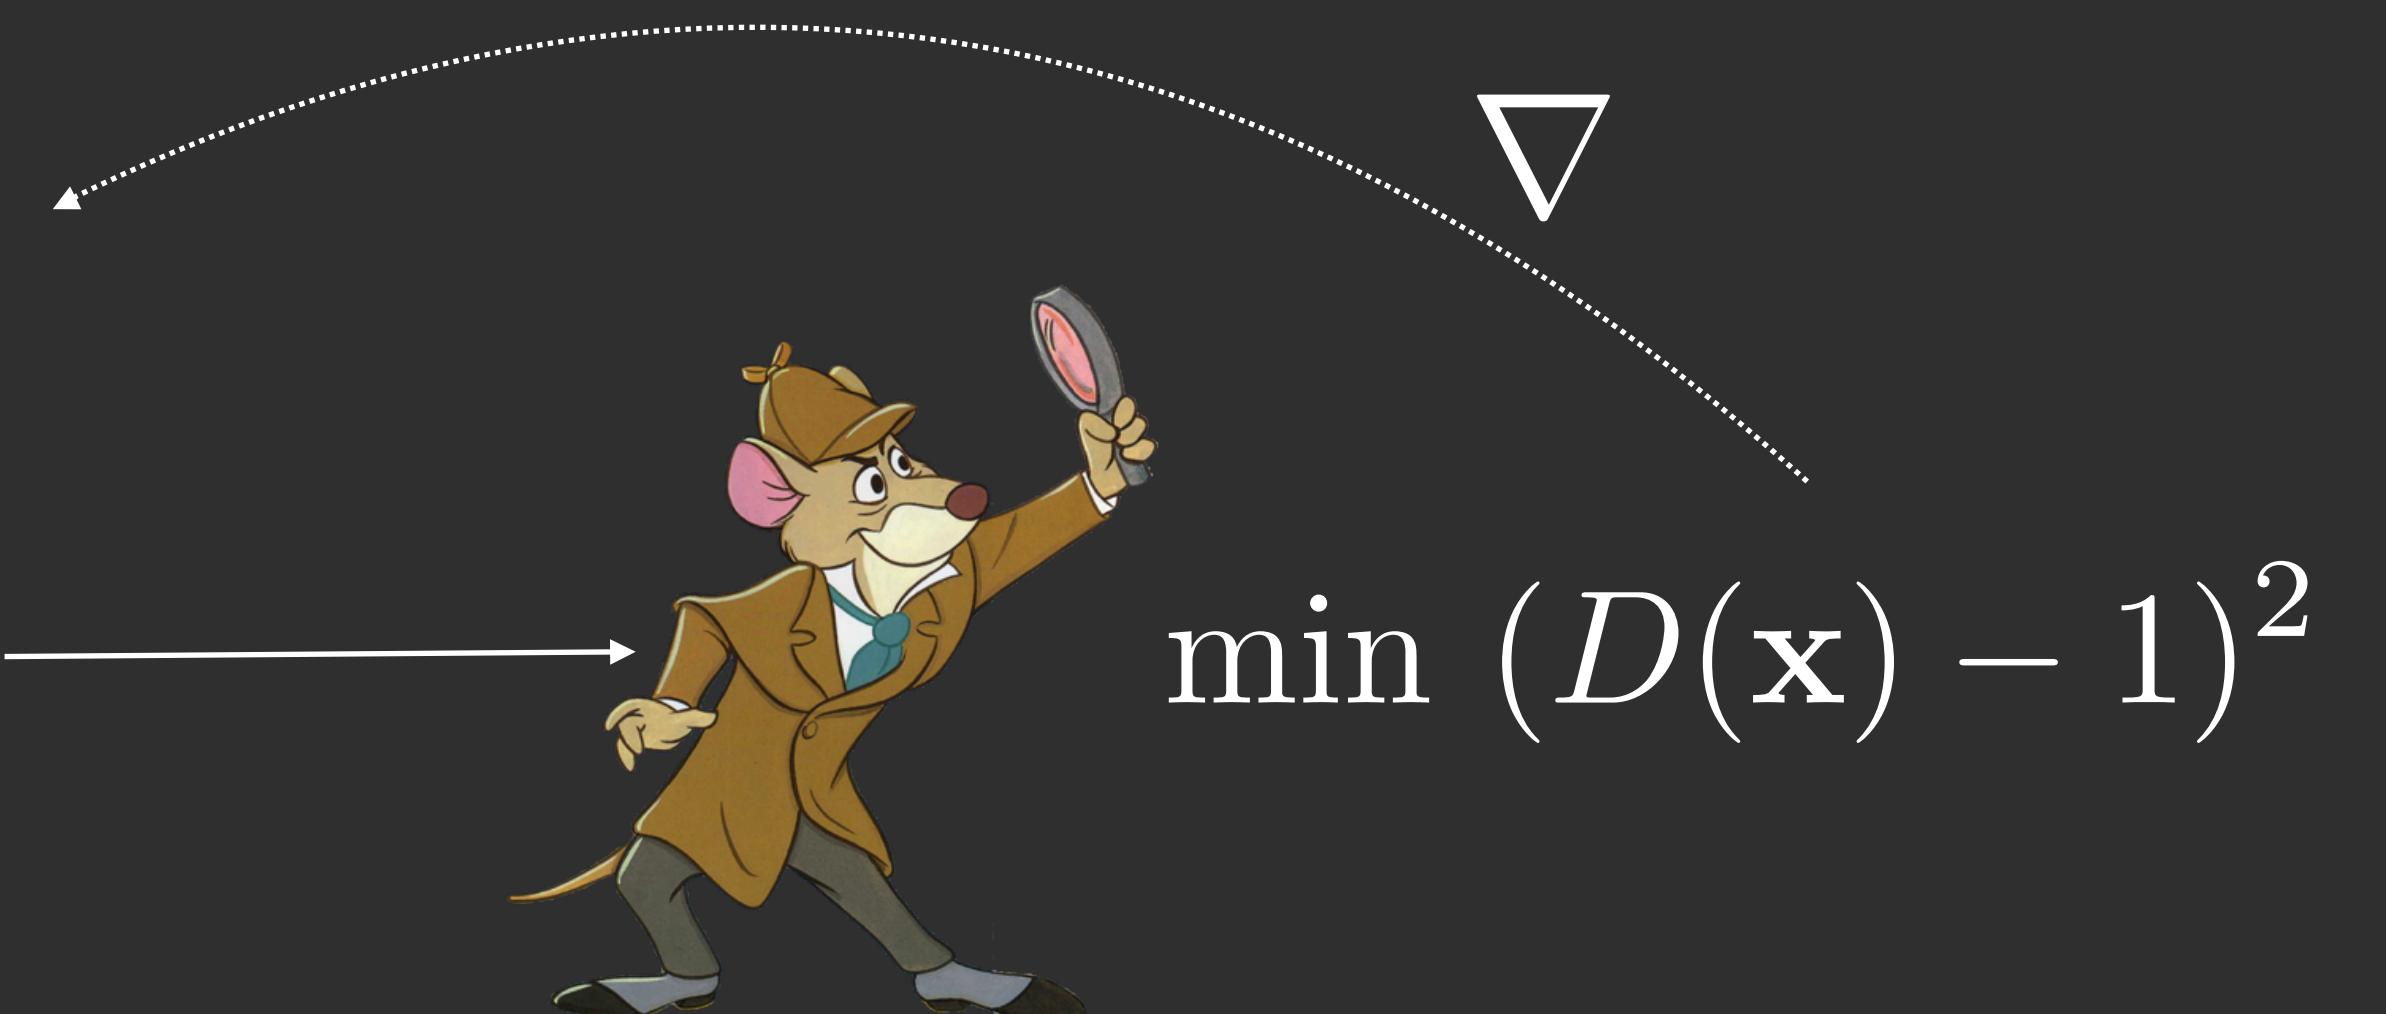
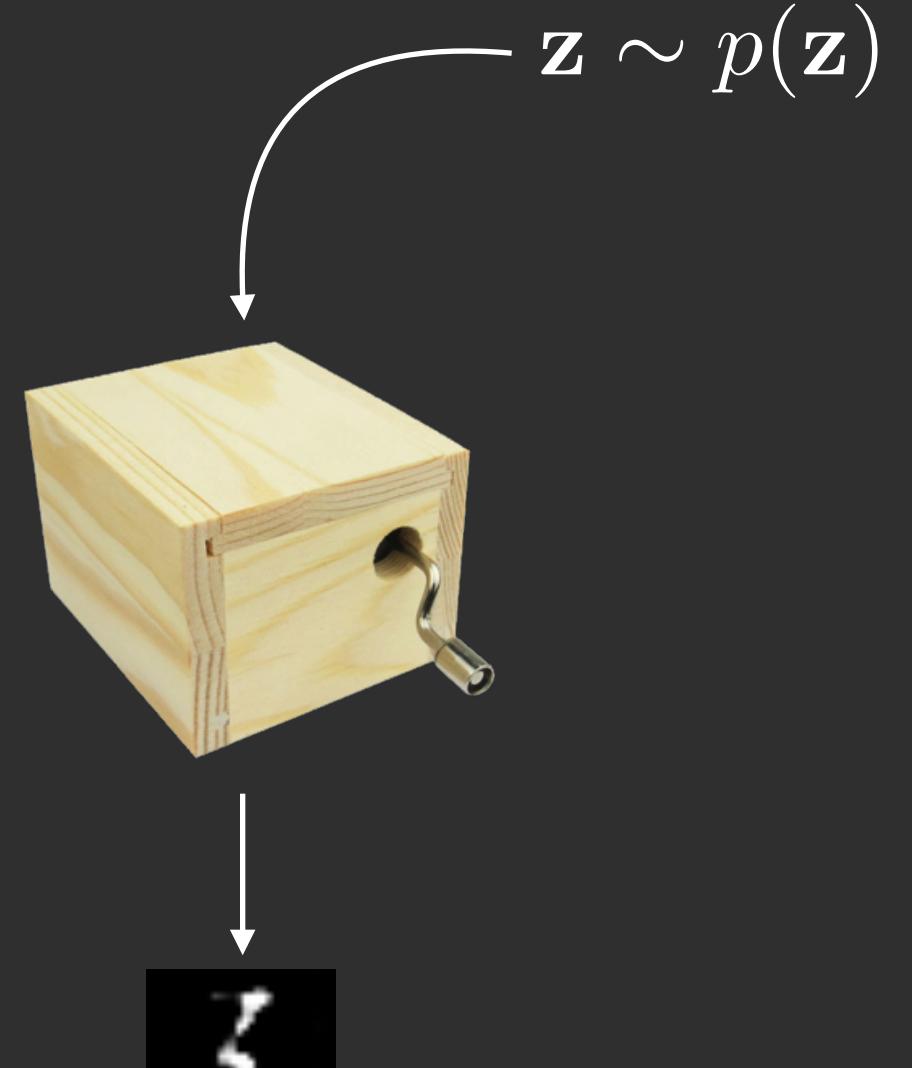
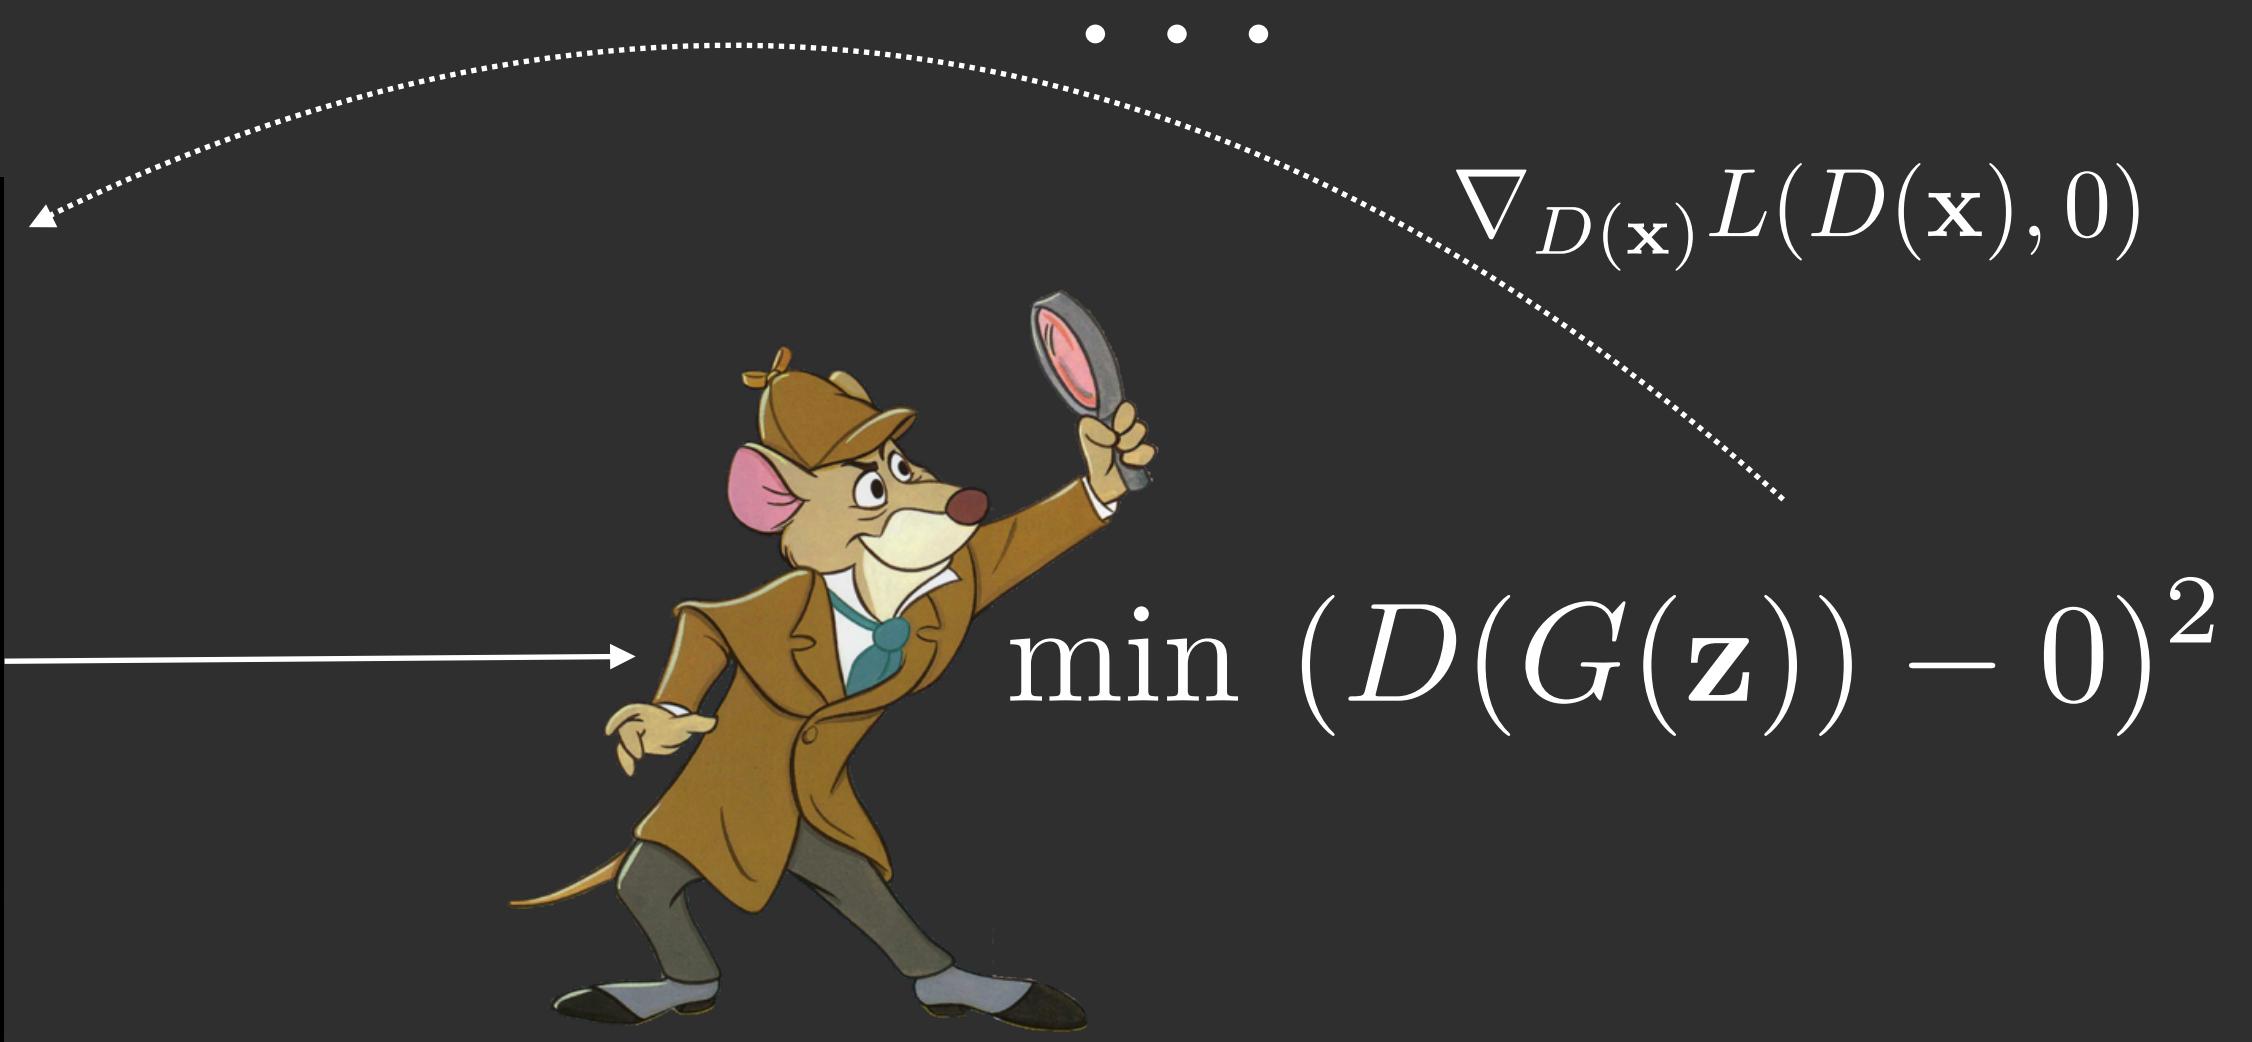
6
4
3
B

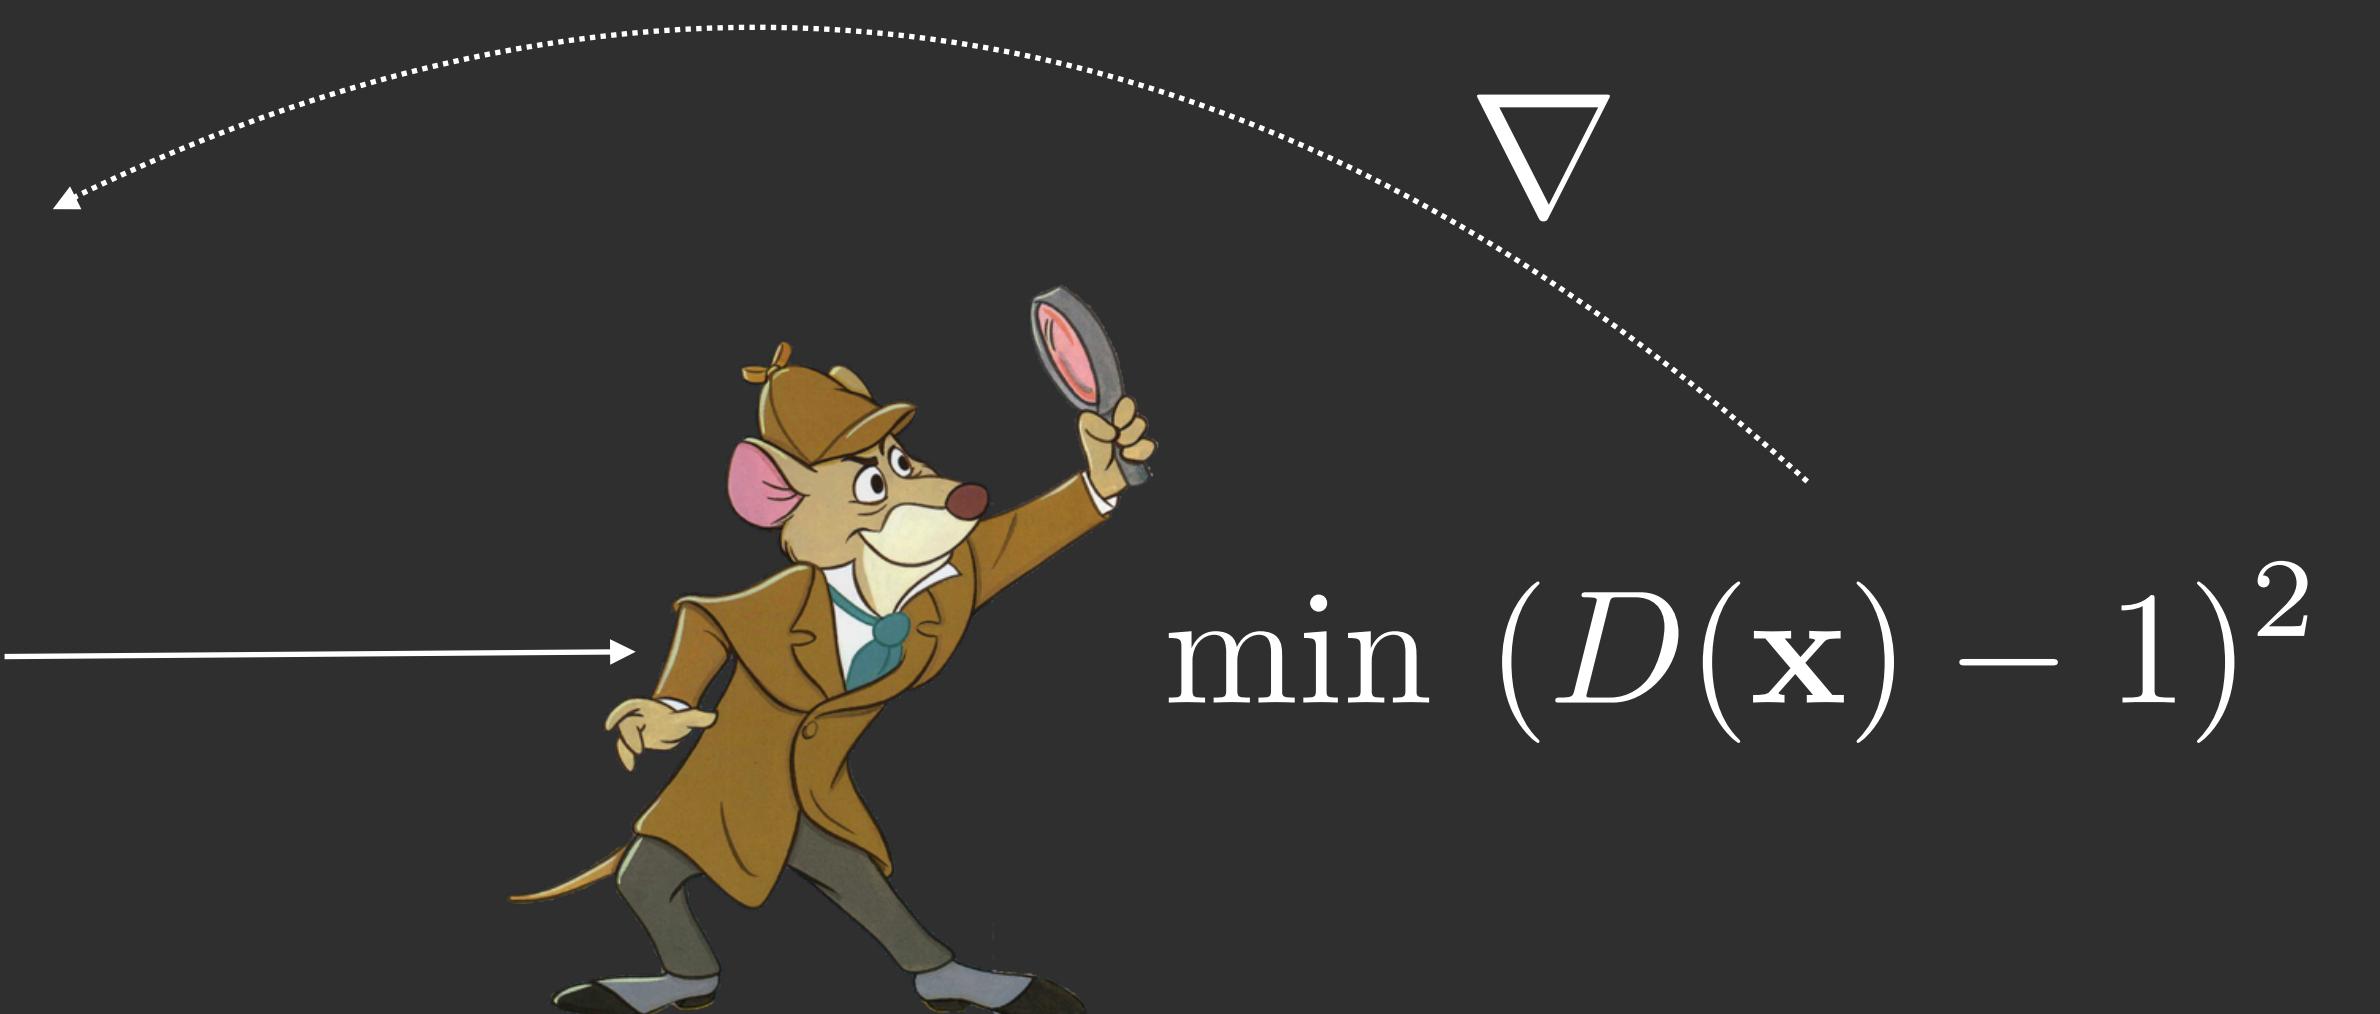
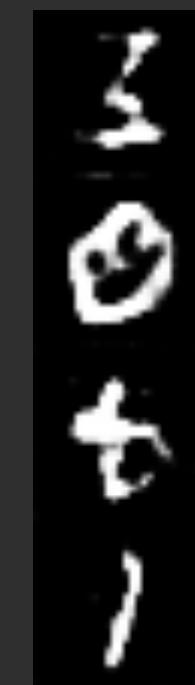
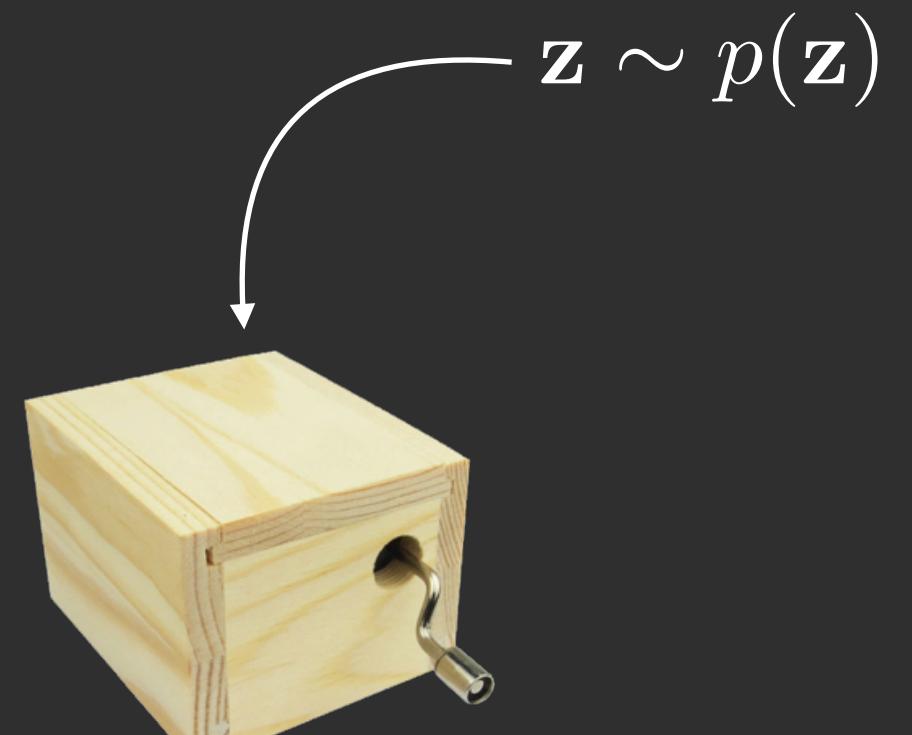
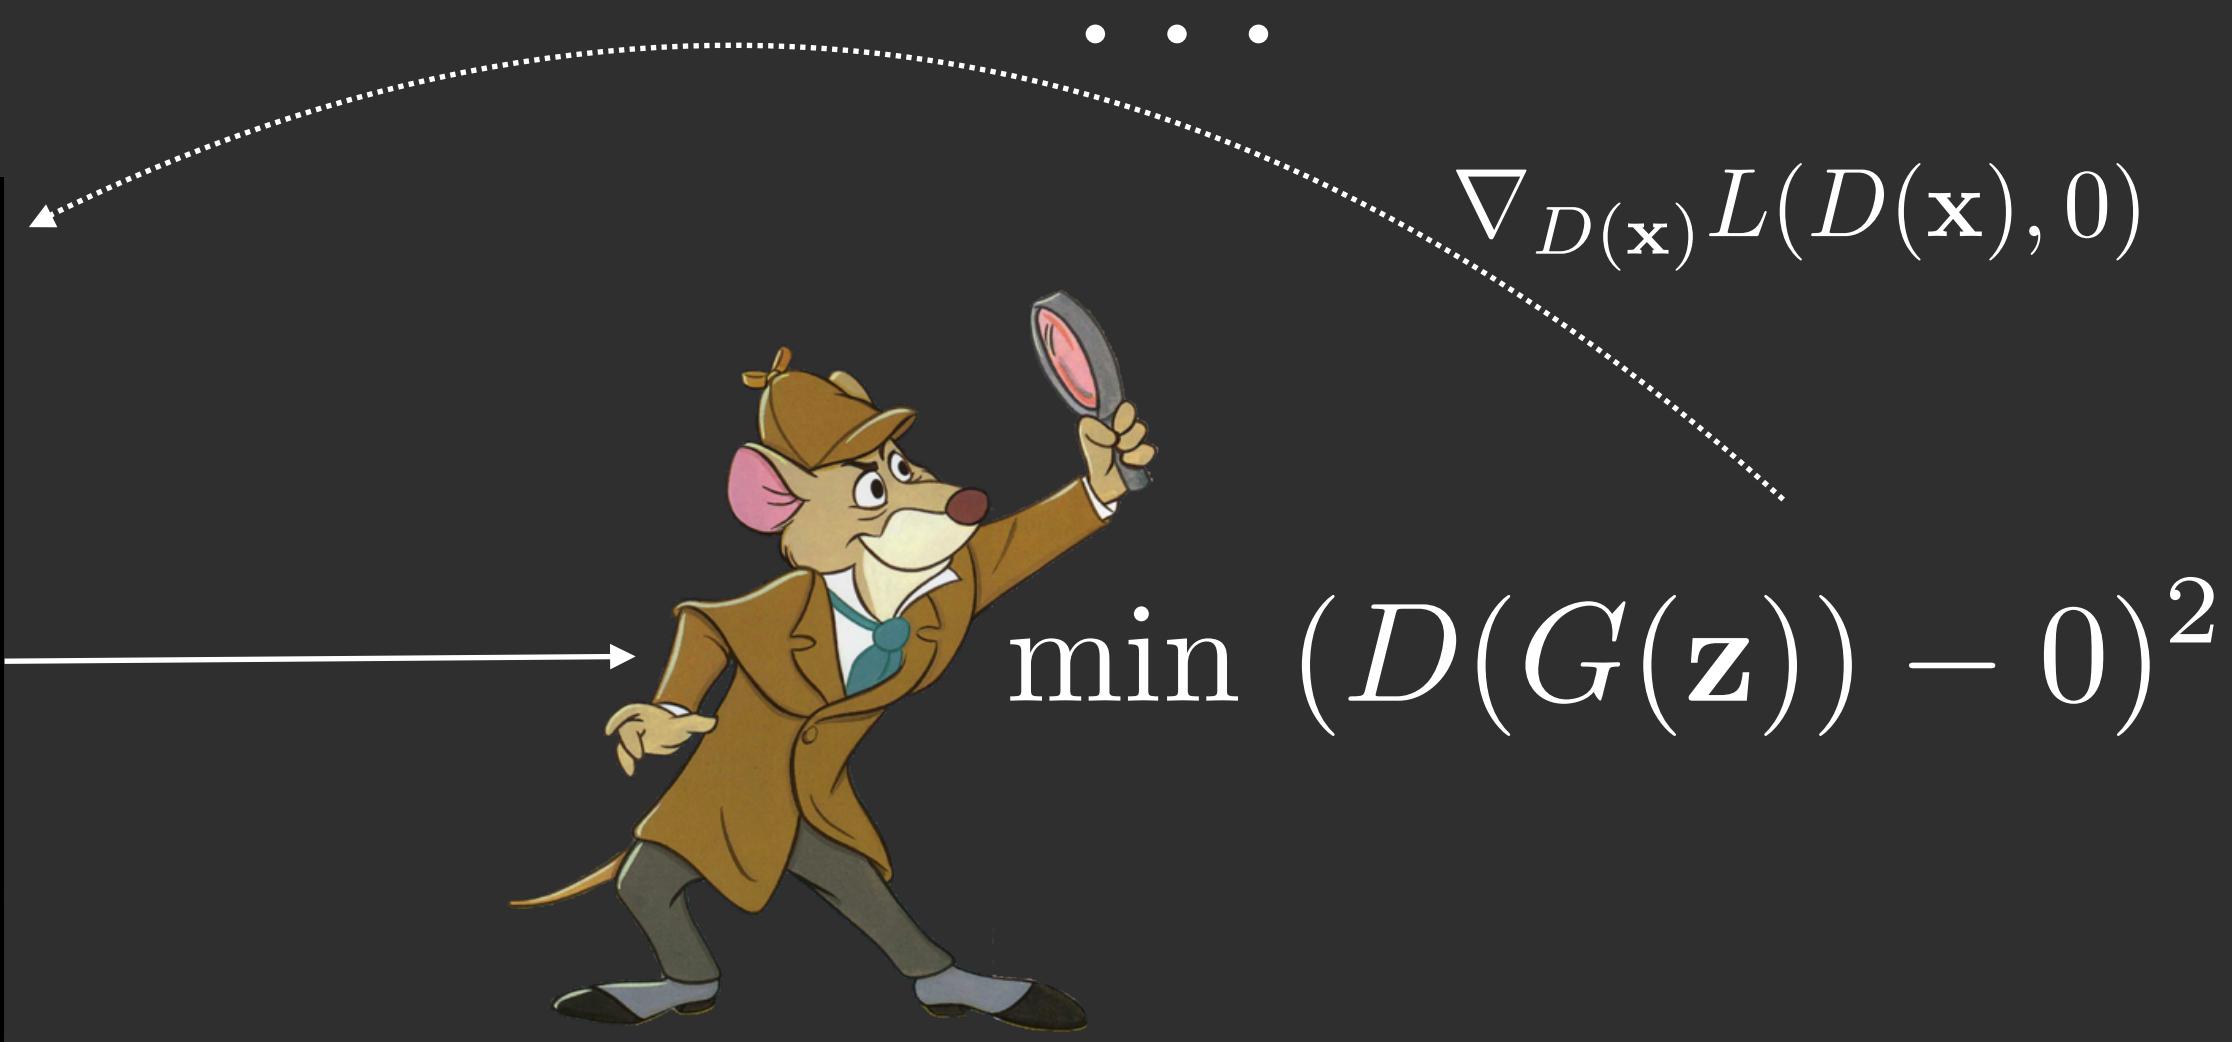


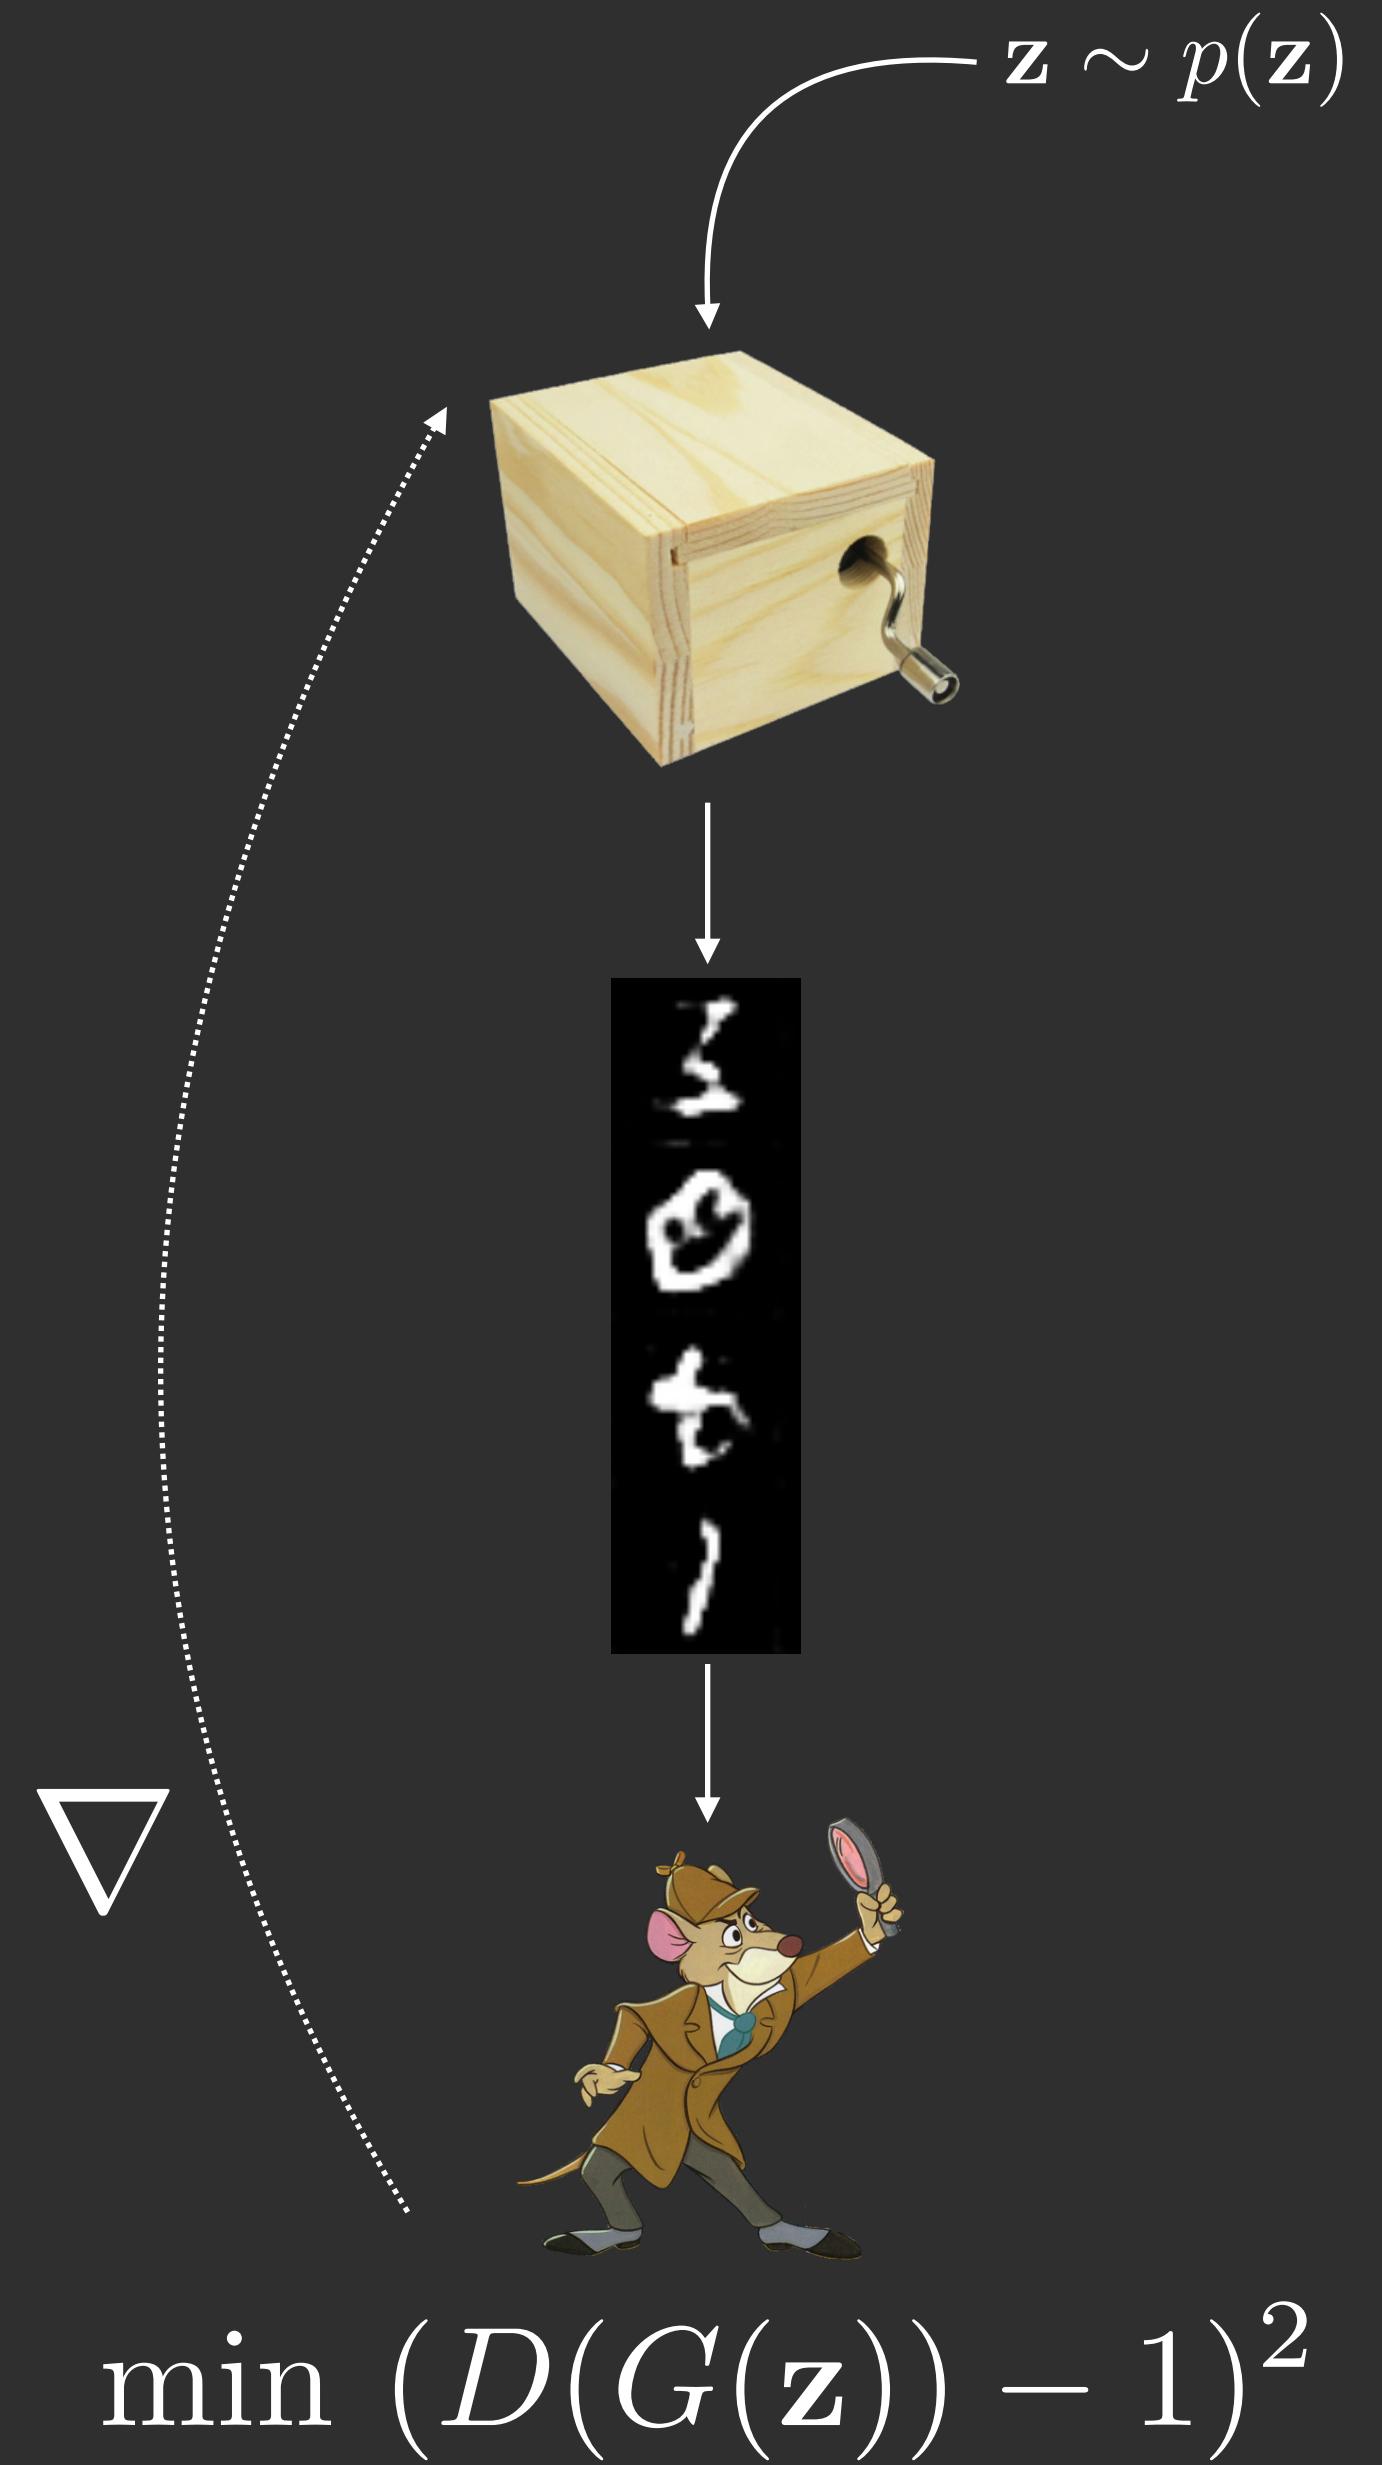
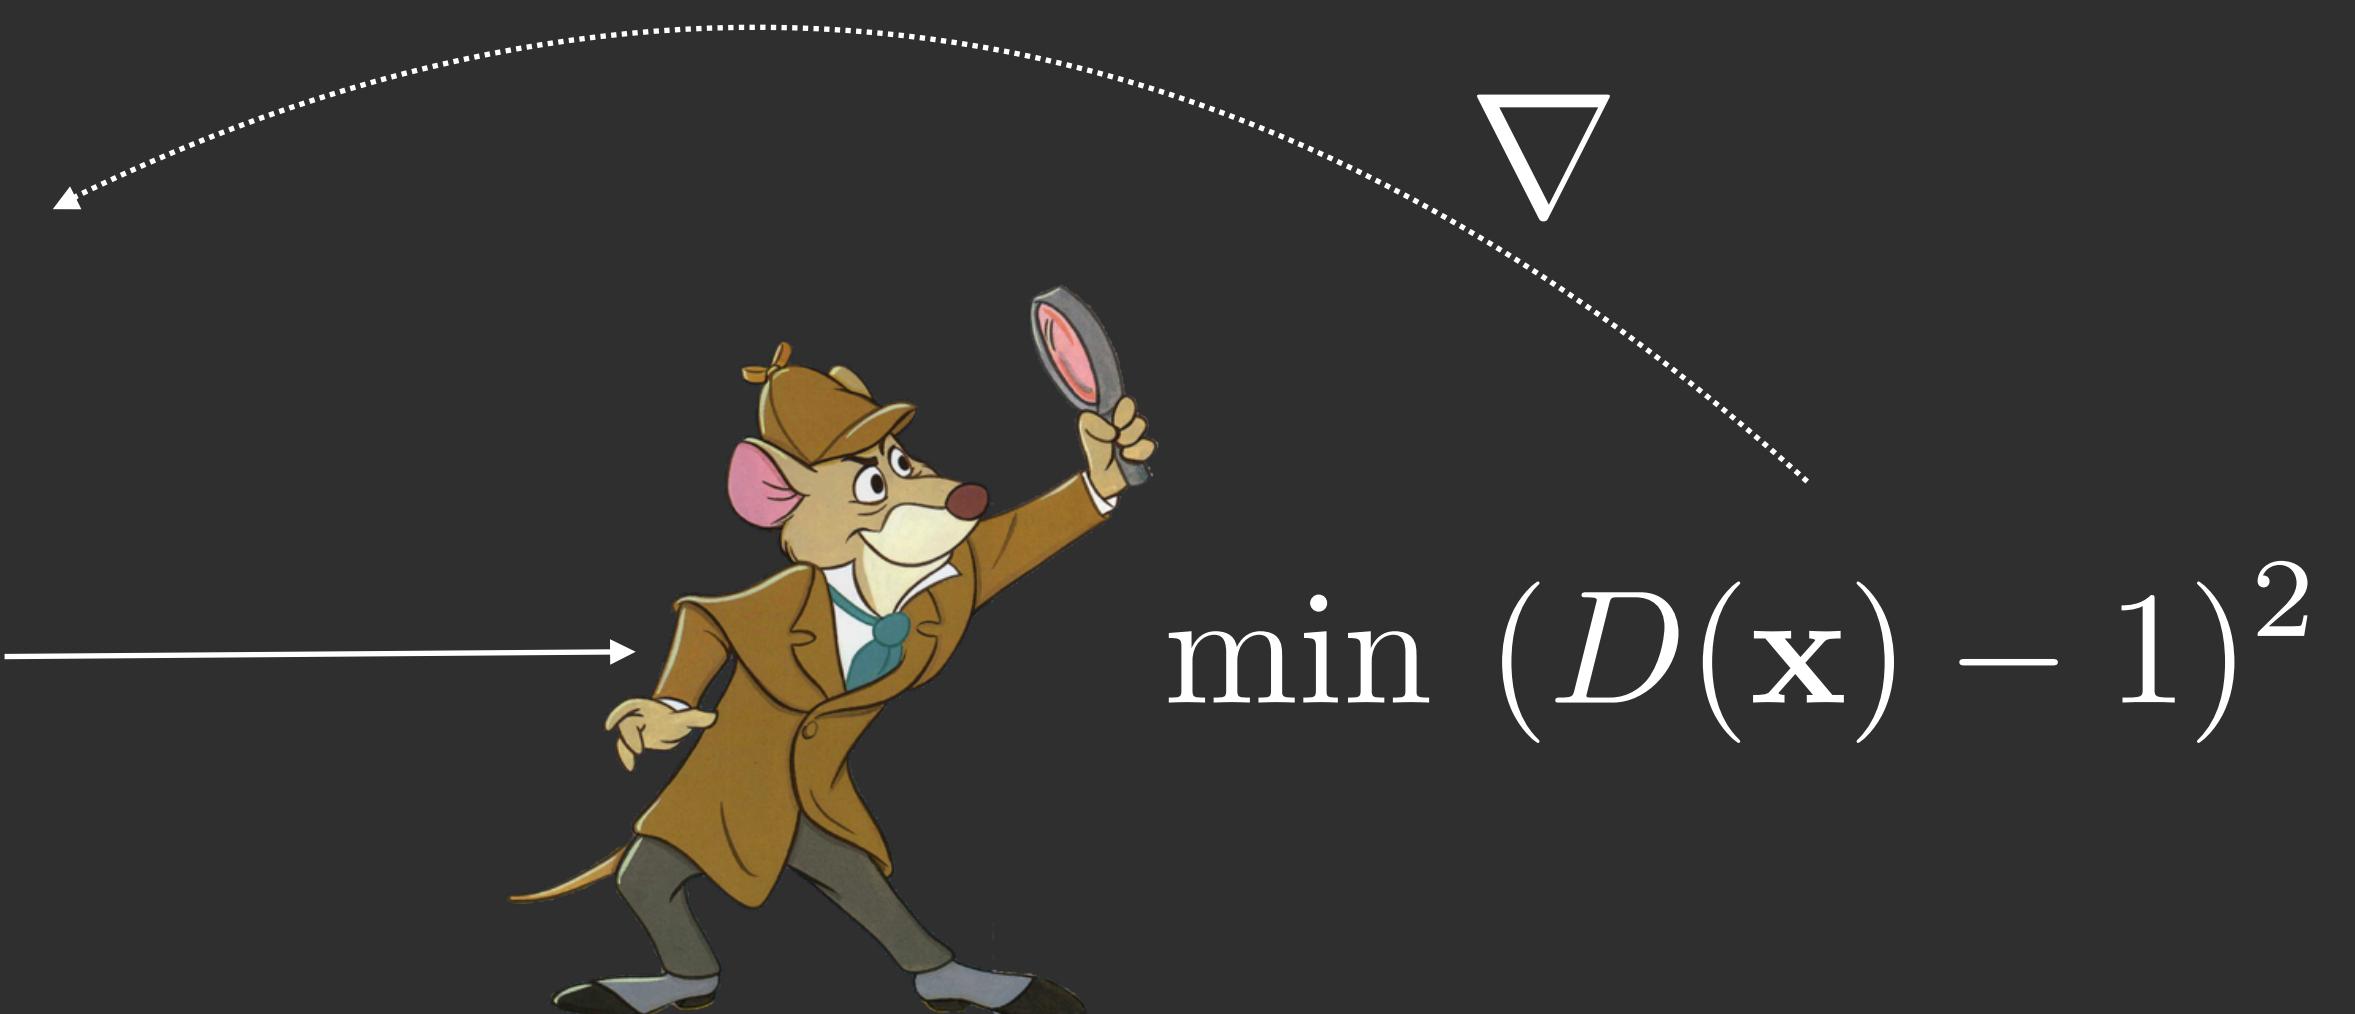
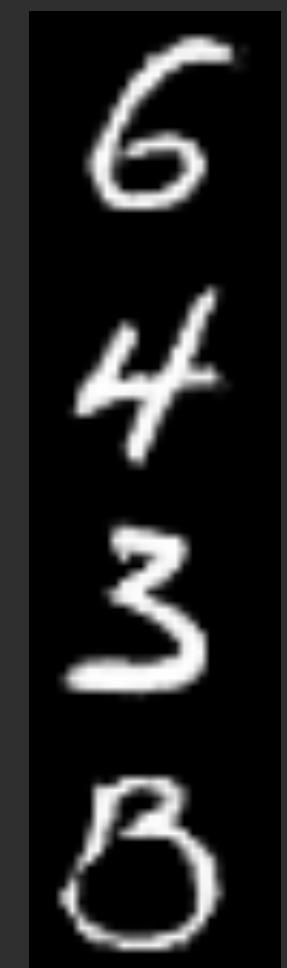
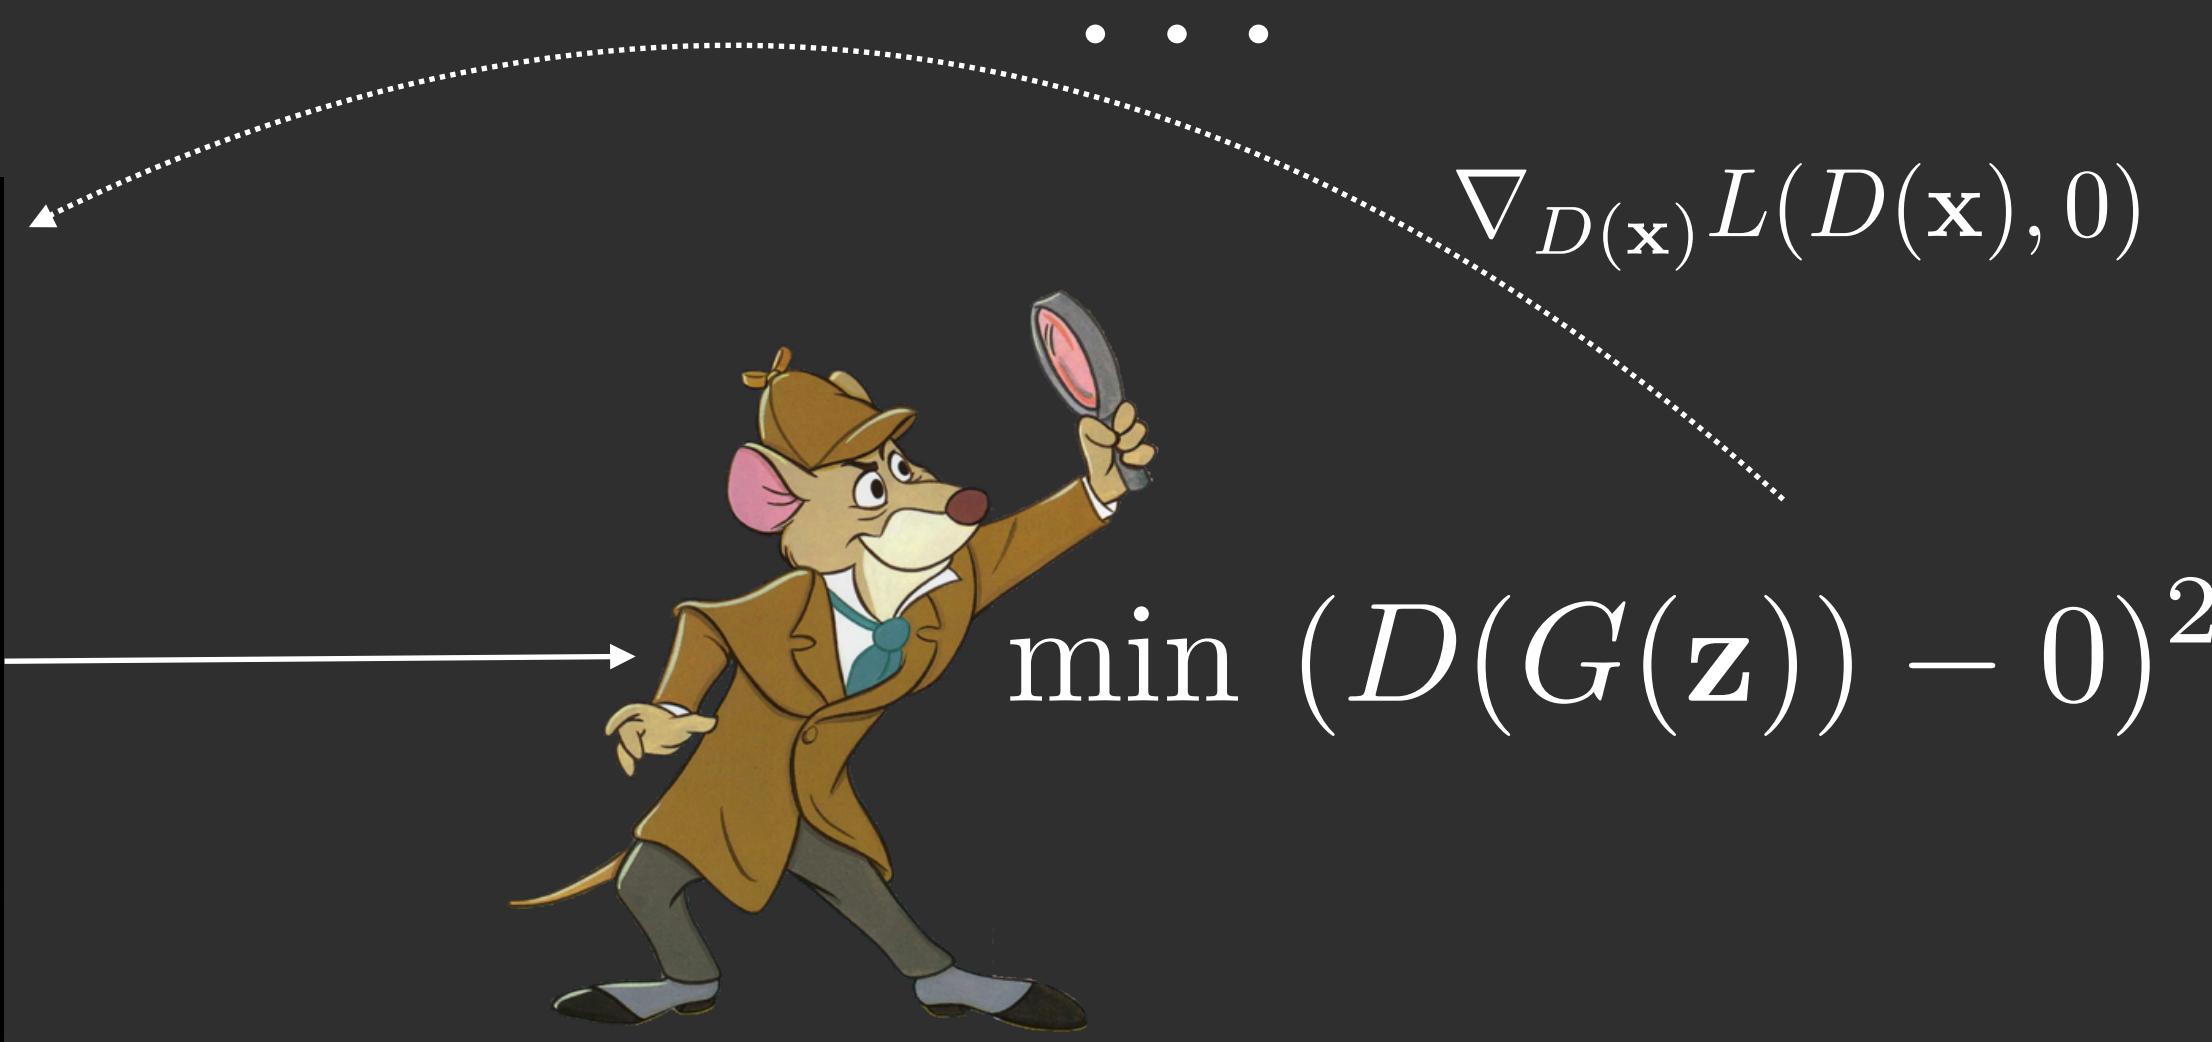









$$\min (D(G(\mathbf{z})) - 1)^2$$



What's the use of learning to generate data?

What's the use of learning to generate data?

You may want to actually generate / augment data!



<https://towardsdatascience.com/artistic-style-transfer-b7566a216431>

GANs to generate / augment data

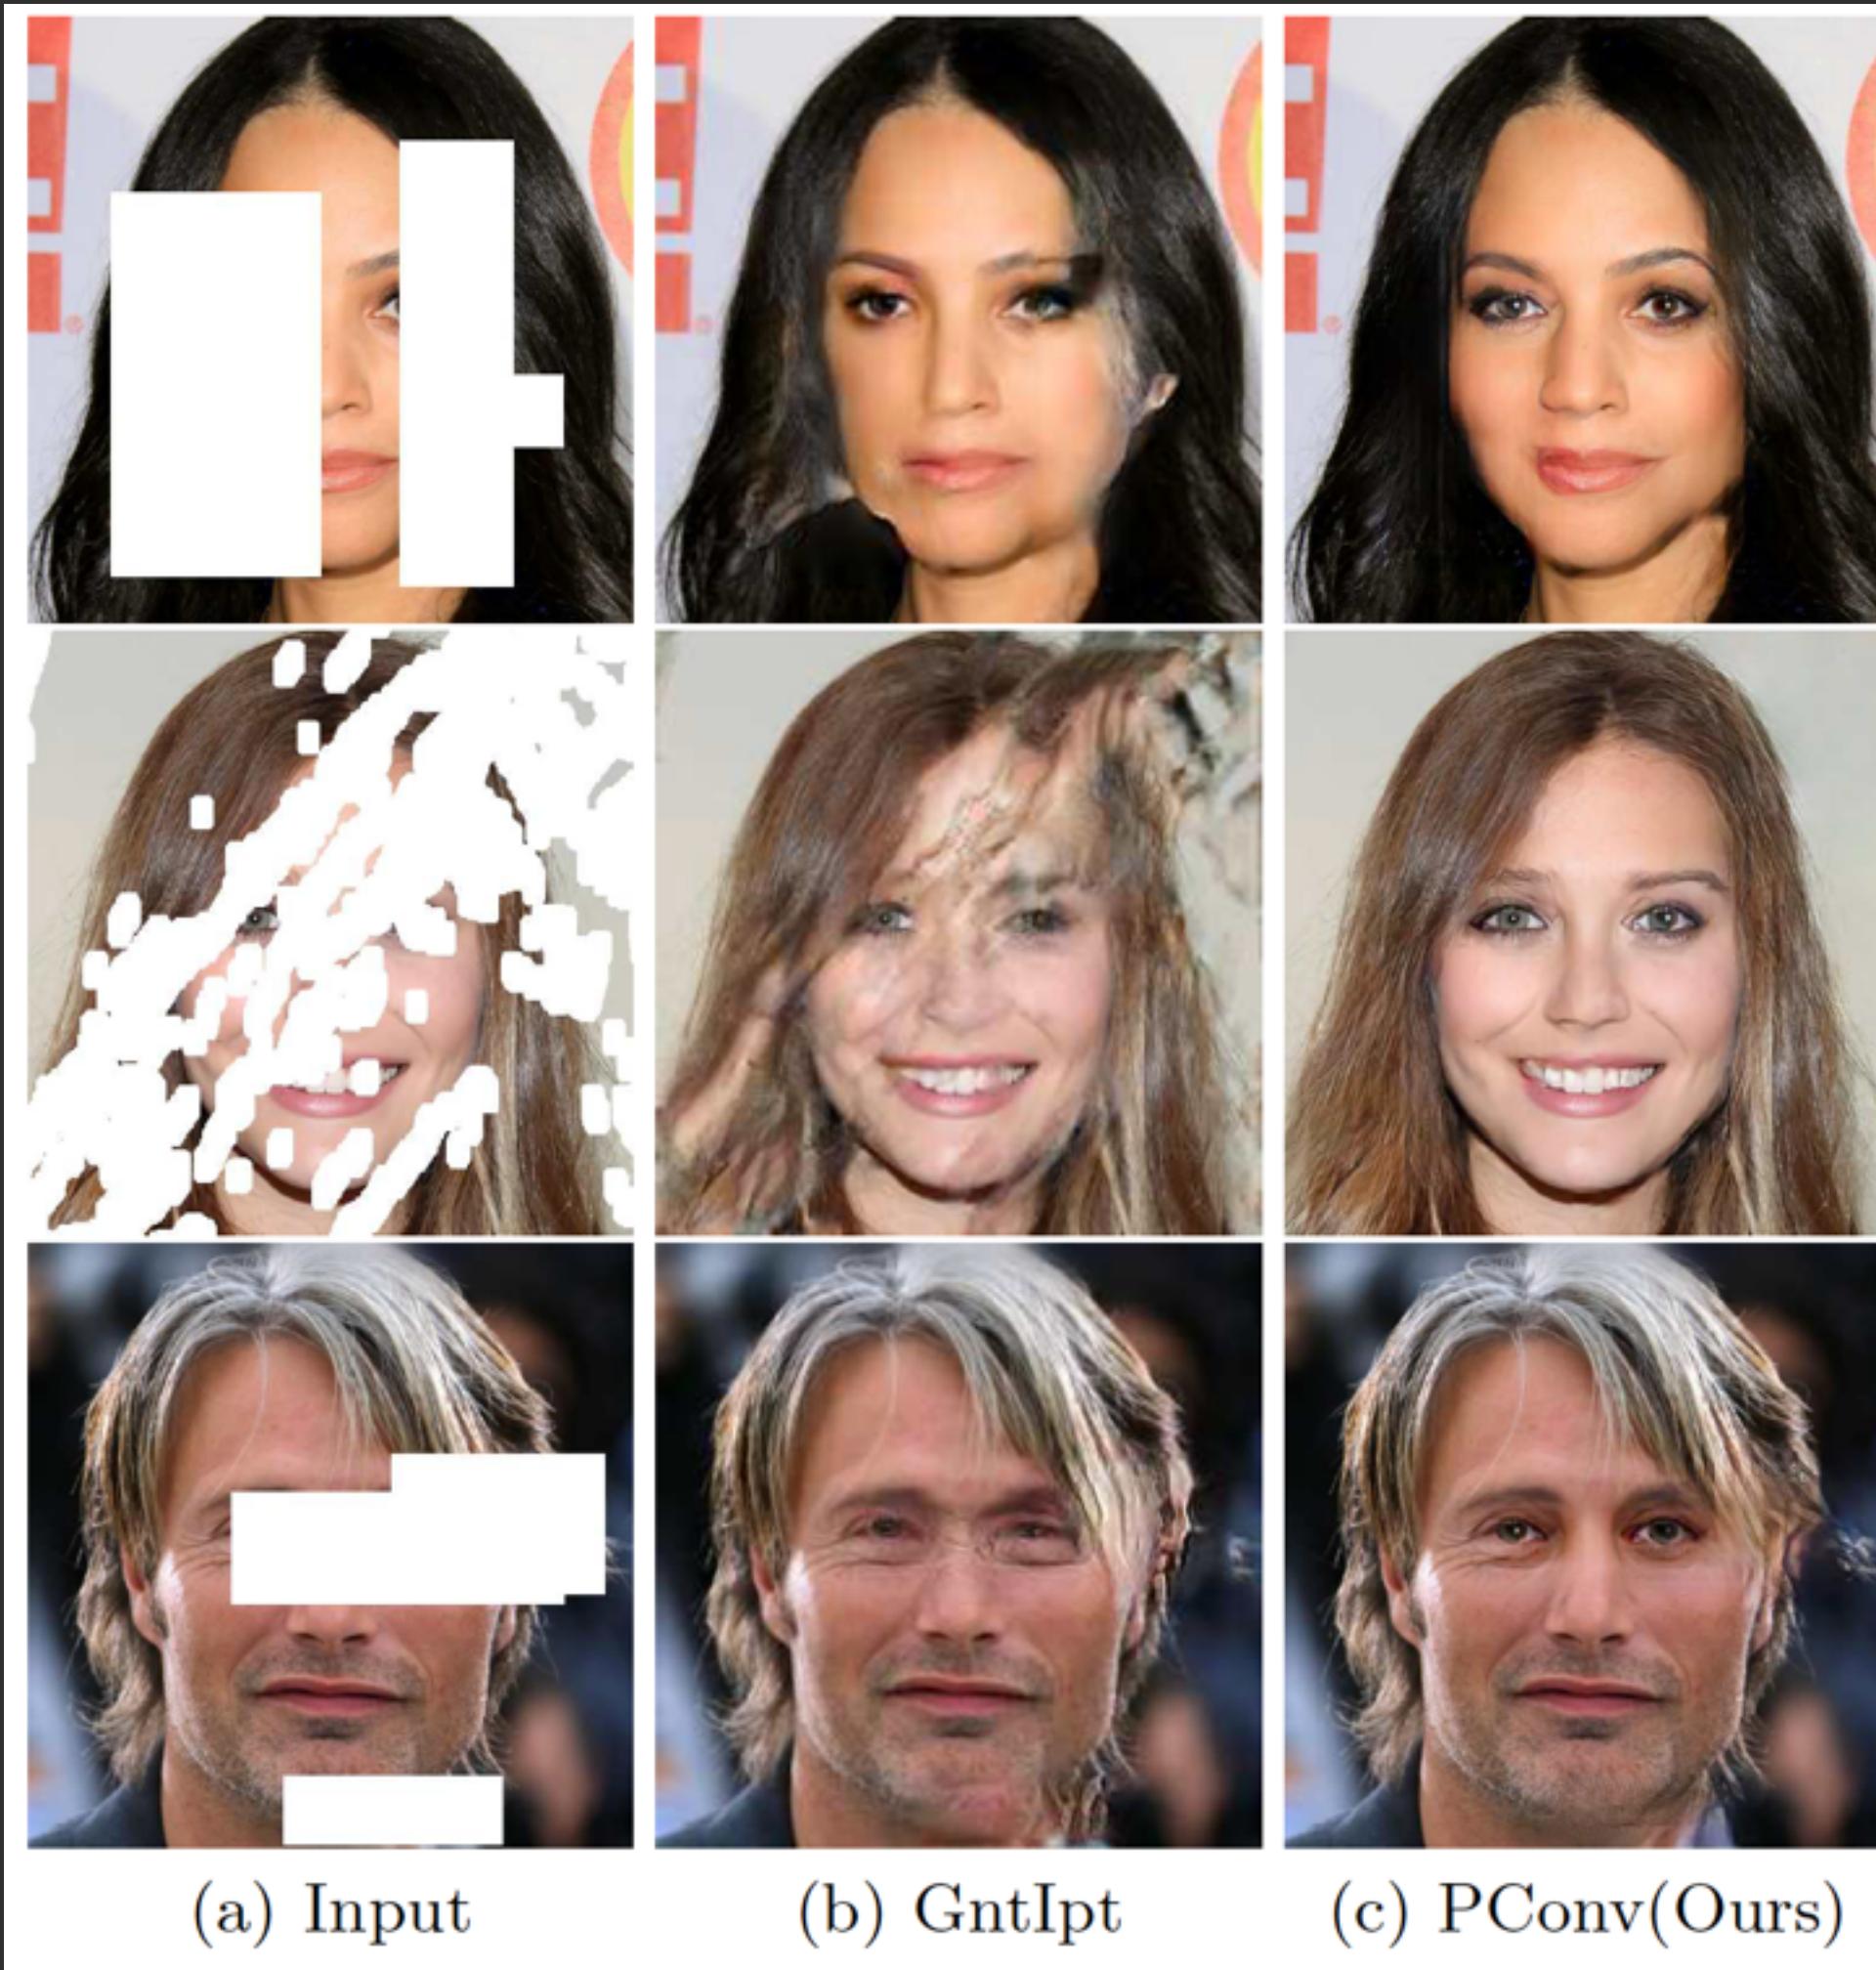
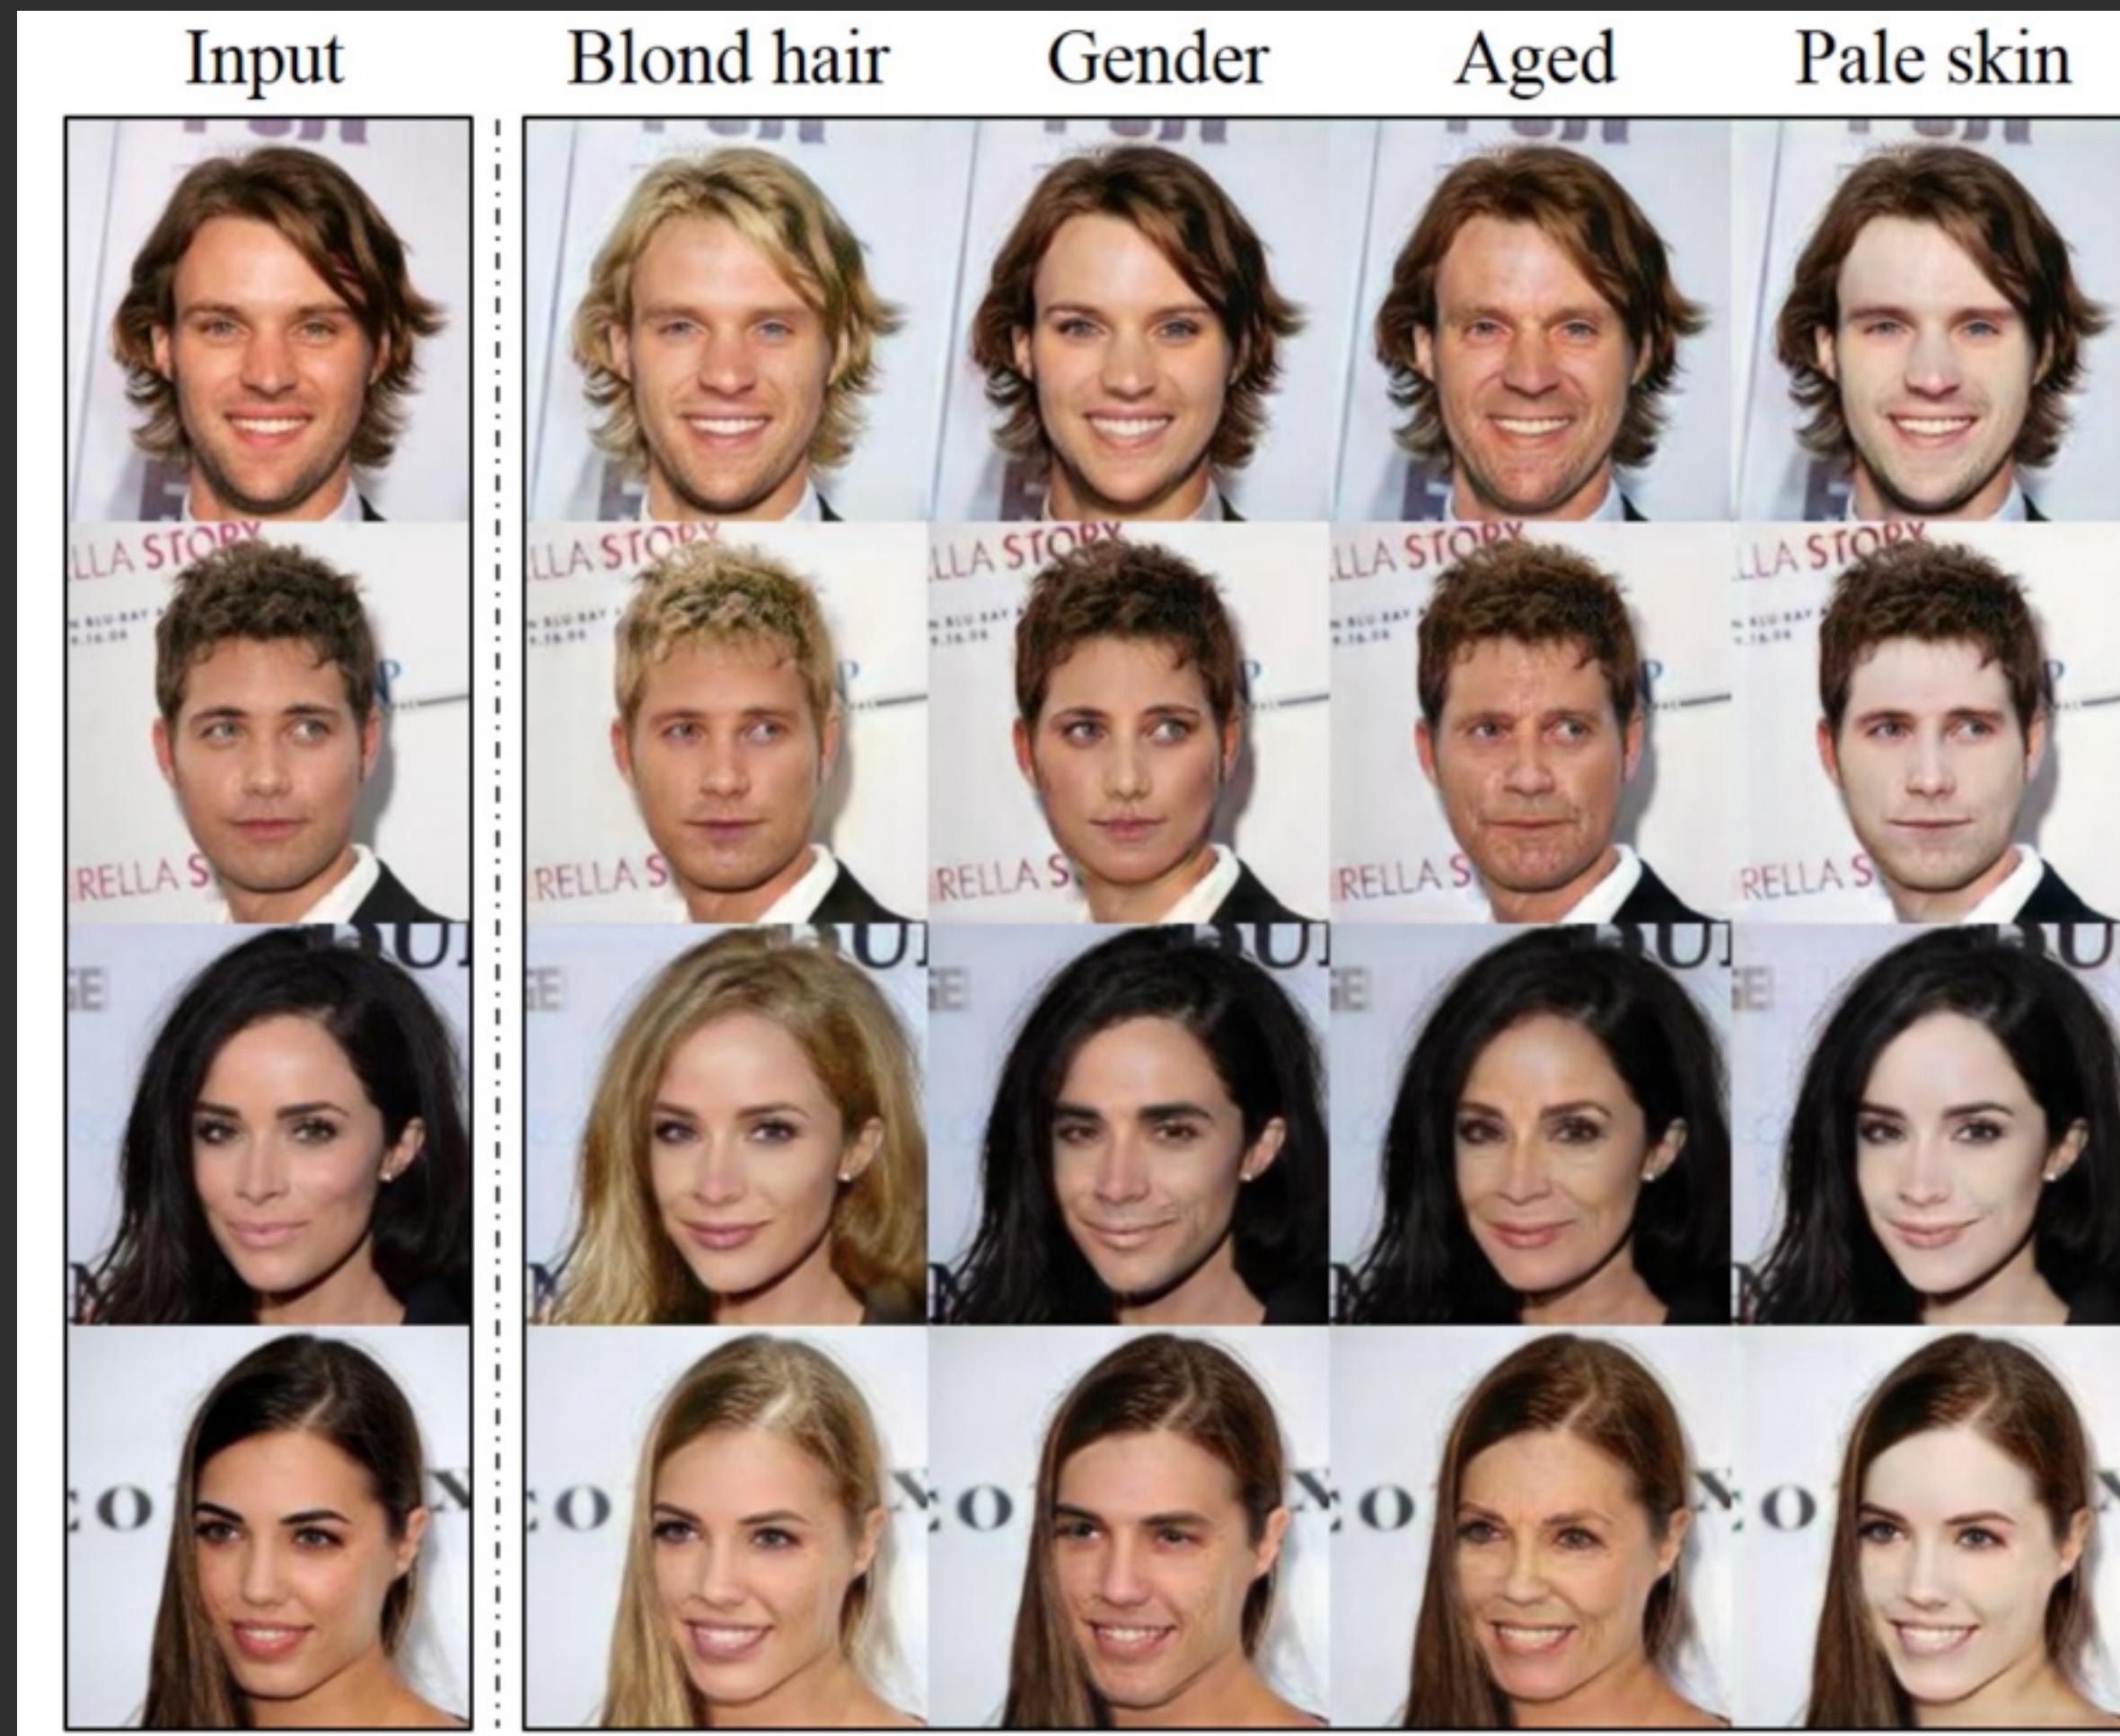


Image Inpainting for Irregular Holes Using Partial Convolutions (Liu et al, 2018)

GANs to generate / augment data



StarGAN: unified generative adversarial networks for multi-domain image-to-image translation
(Choi et al, 2017)

What's the use of learning to generate data?

Learning to generate data may help downstream tasks



What I cannot create,
I do not understand

GANs to aid downstream tasks

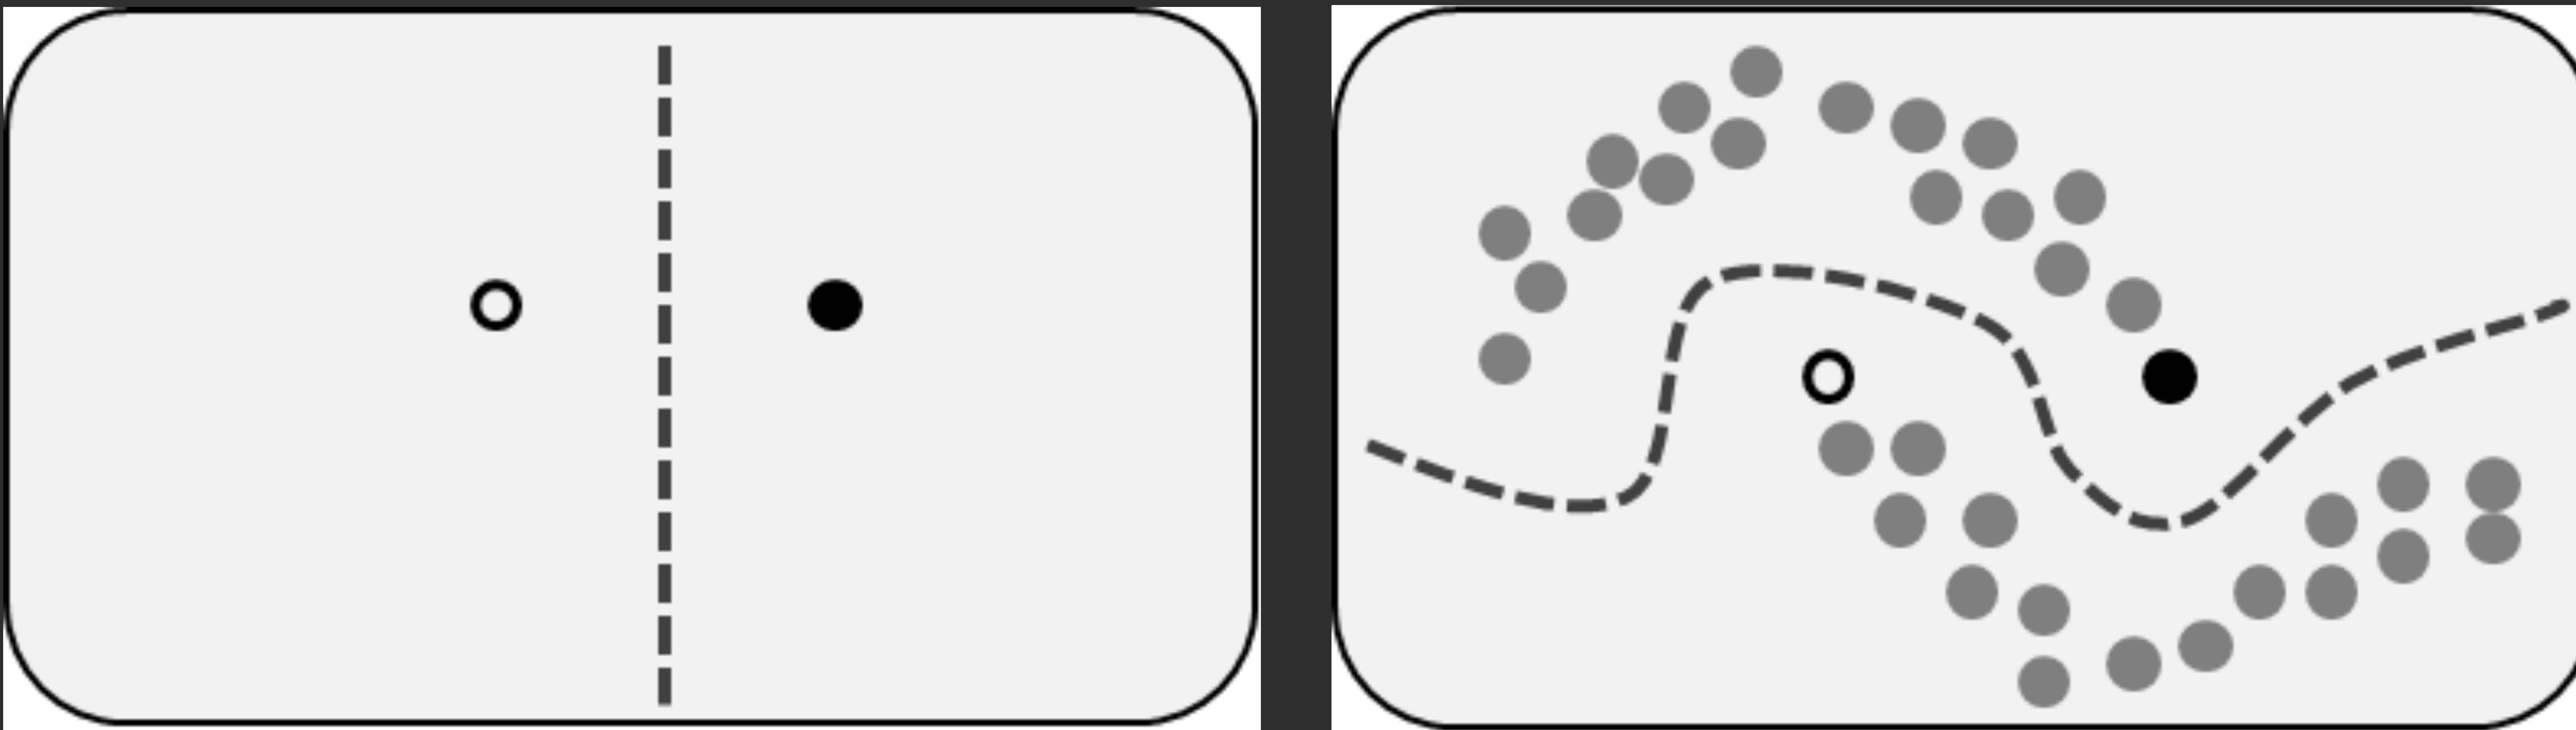


[https://media1.fdncms.com/northcoast/imager/u/zoom/3677369/
field_notes-579c97c184a67189.jpg](https://media1.fdncms.com/northcoast/imager/u/zoom/3677369/field_notes-579c97c184a67189.jpg)

$$p(y|x) = \frac{p(x, y)}{p(x)}$$

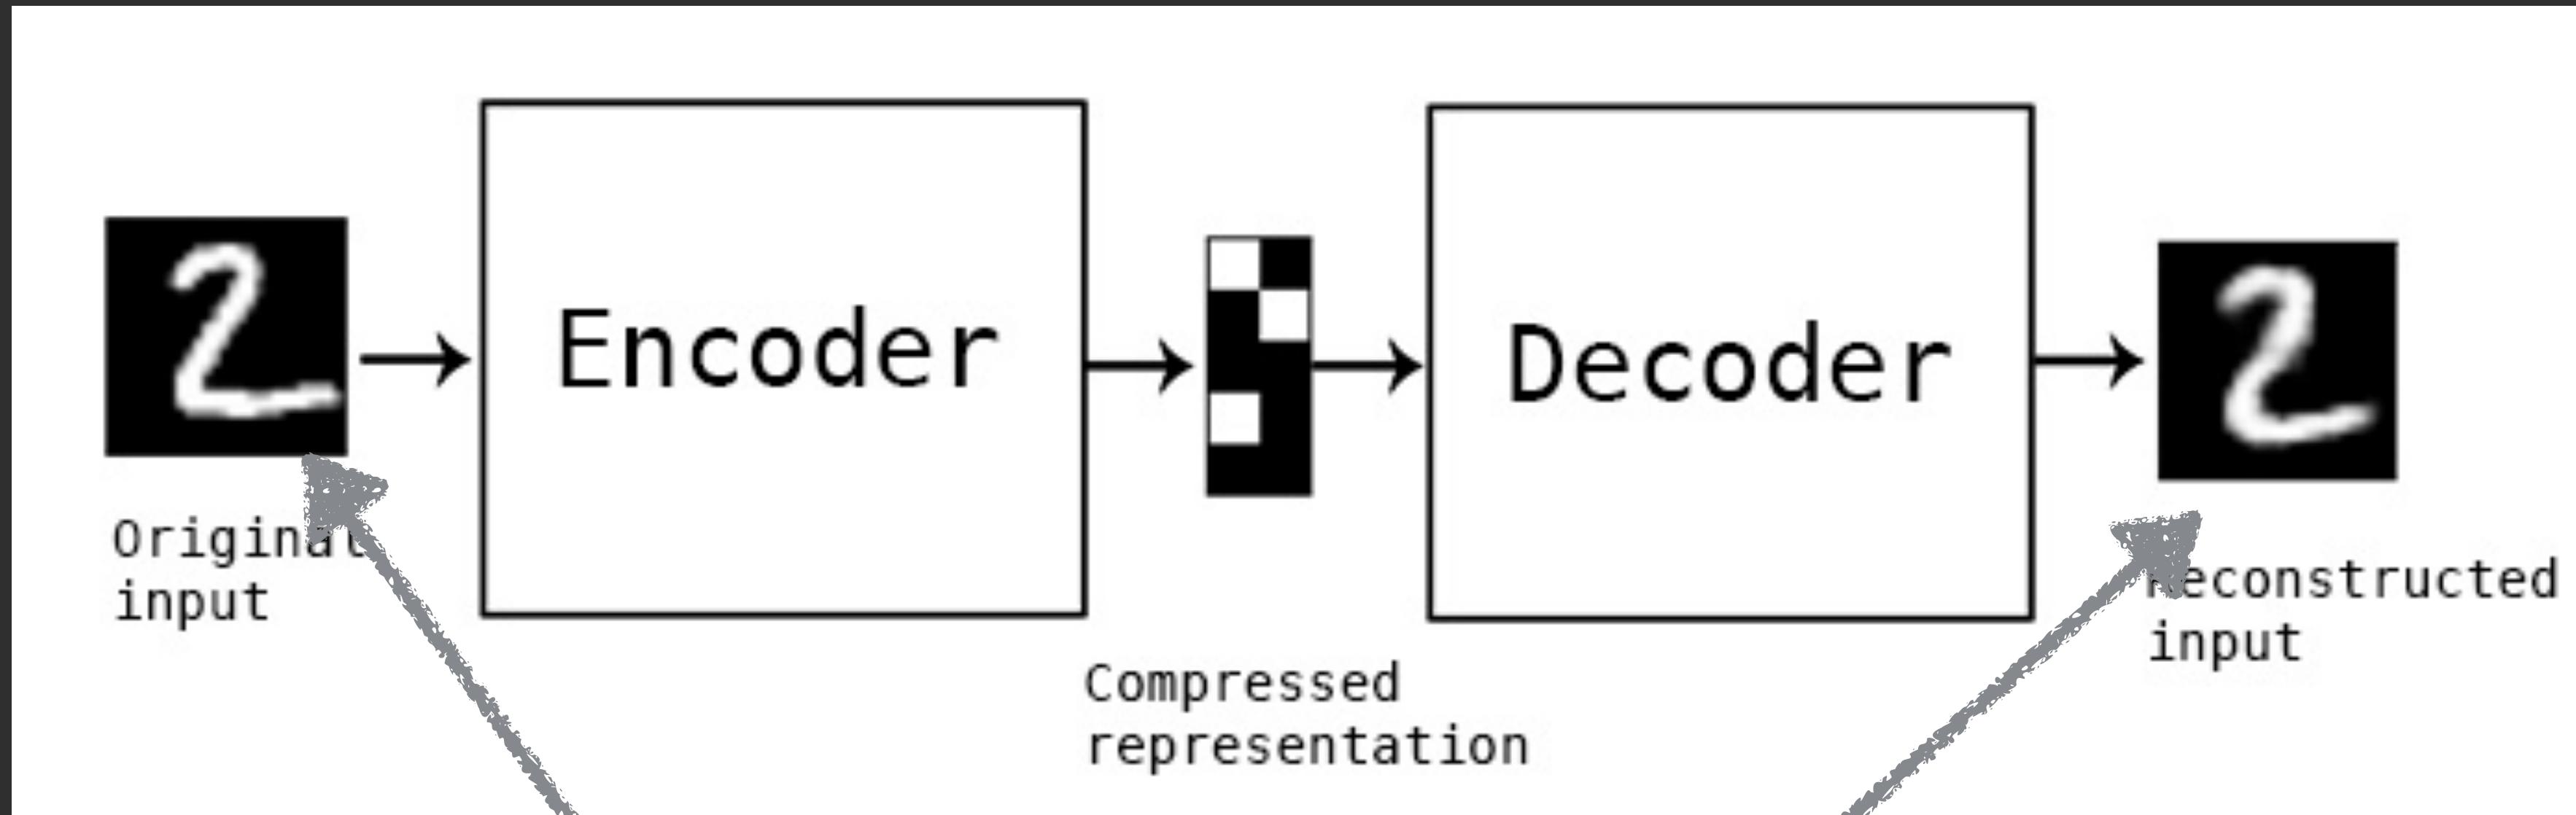


GANs in semi-supervised learning



https://en.wikipedia.org/wiki/Semi-supervised_learning#/media/File:Example_of_unlabeled_data_in_semisupervised_learning.png

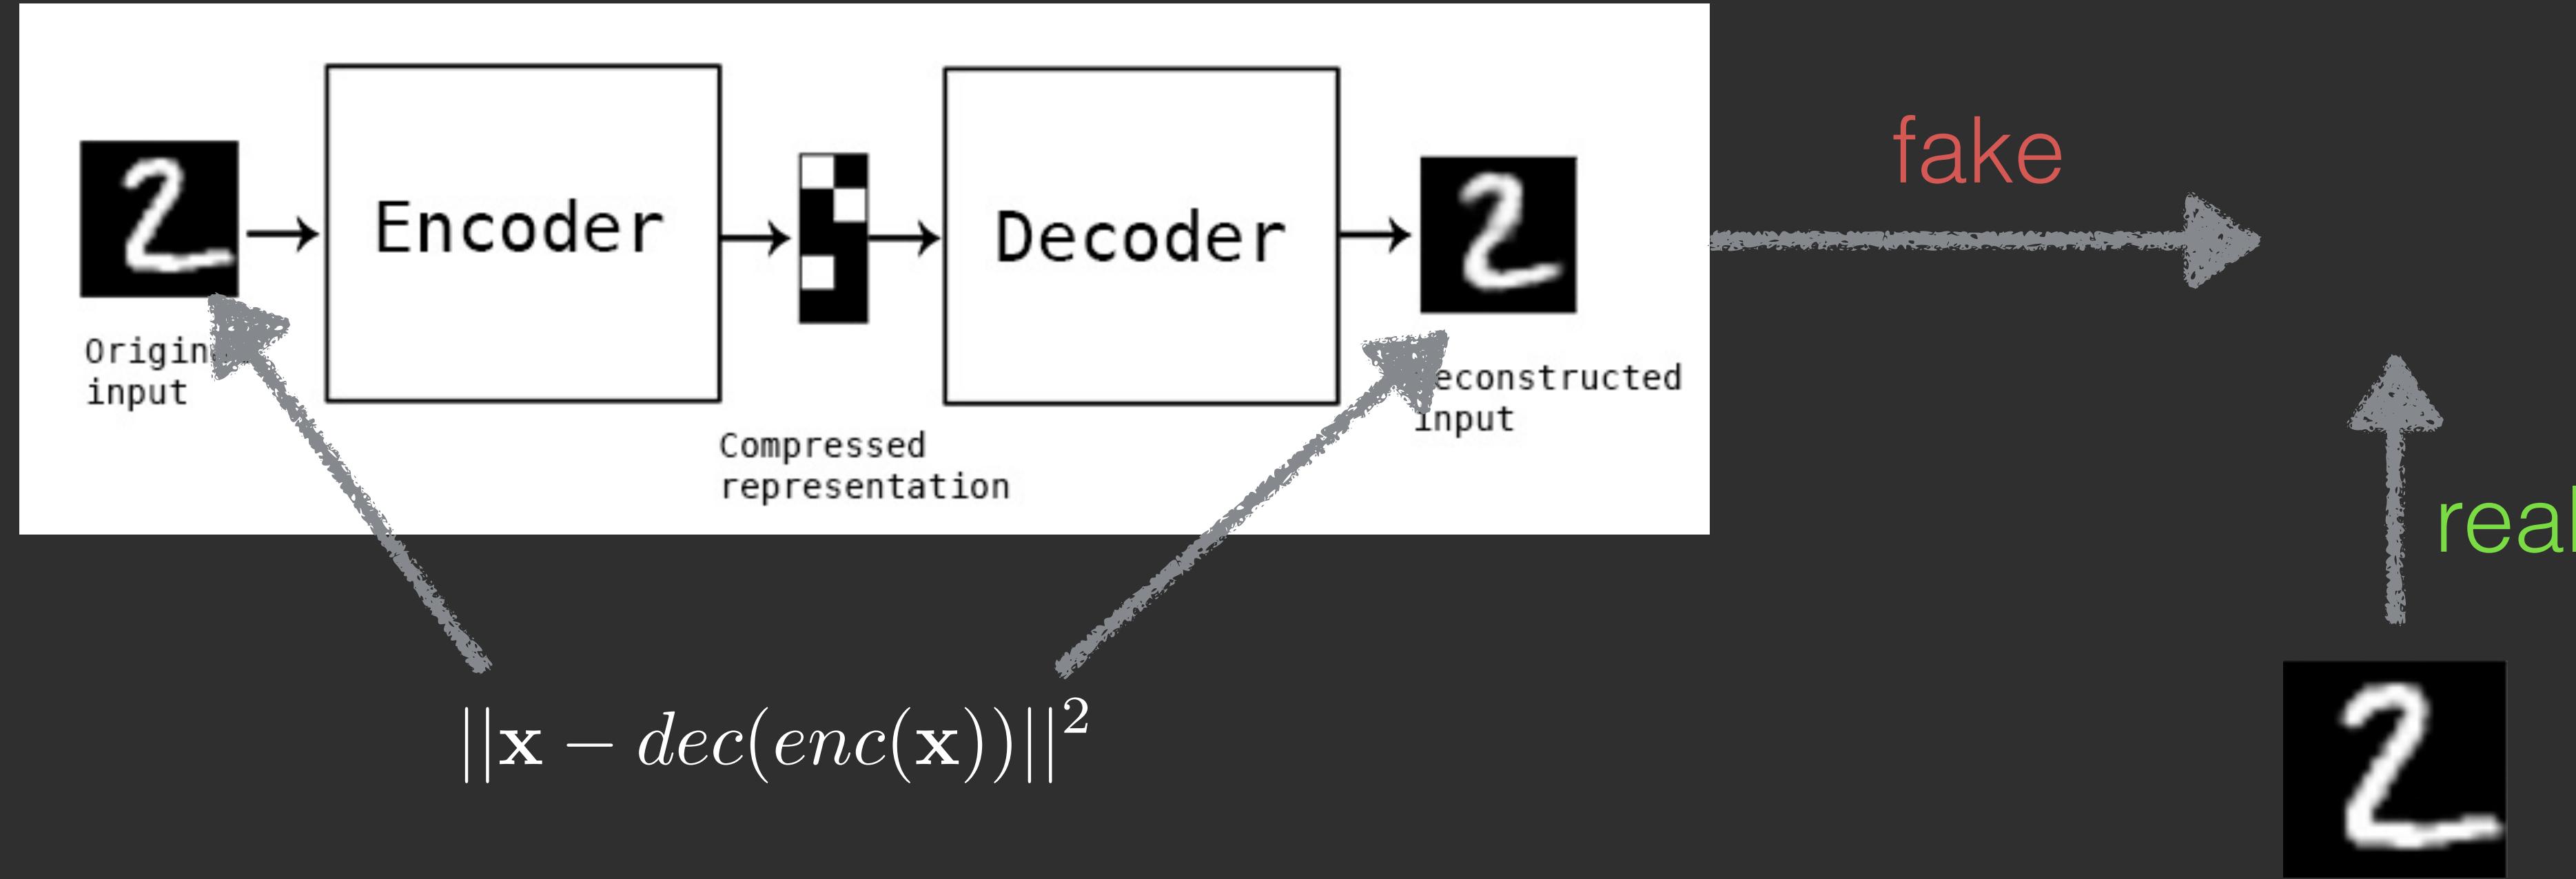
Autoencoders (cooler than PCA)



$$\|\mathbf{x} - dec(enc(\mathbf{x}))\|^2$$

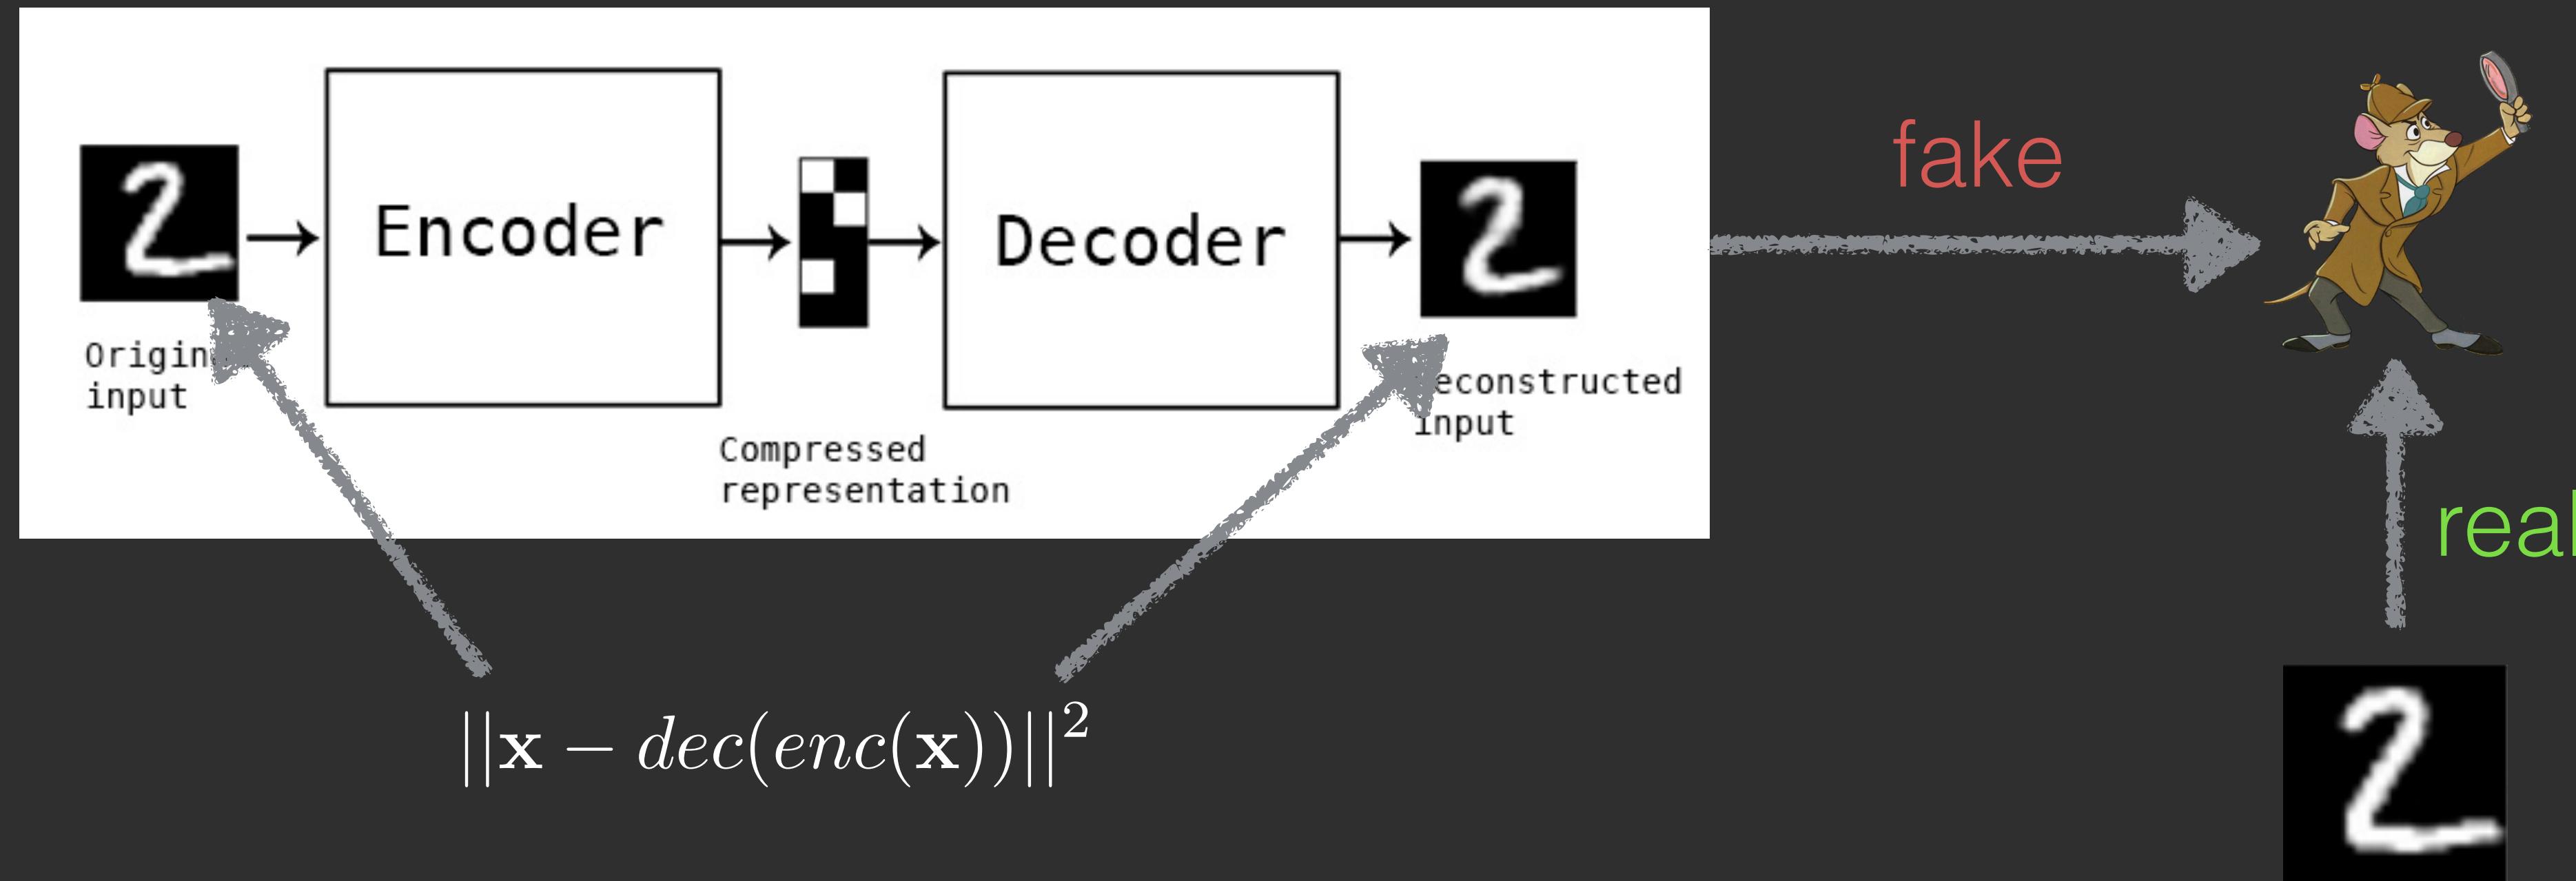
Euclidean distance is a terrible perceptual similarity metric!

Example: *adversarial* autoencoders



Autoencoder minimises Euclidean distance + tries to fool D

Example: *adversarial* autoencoders



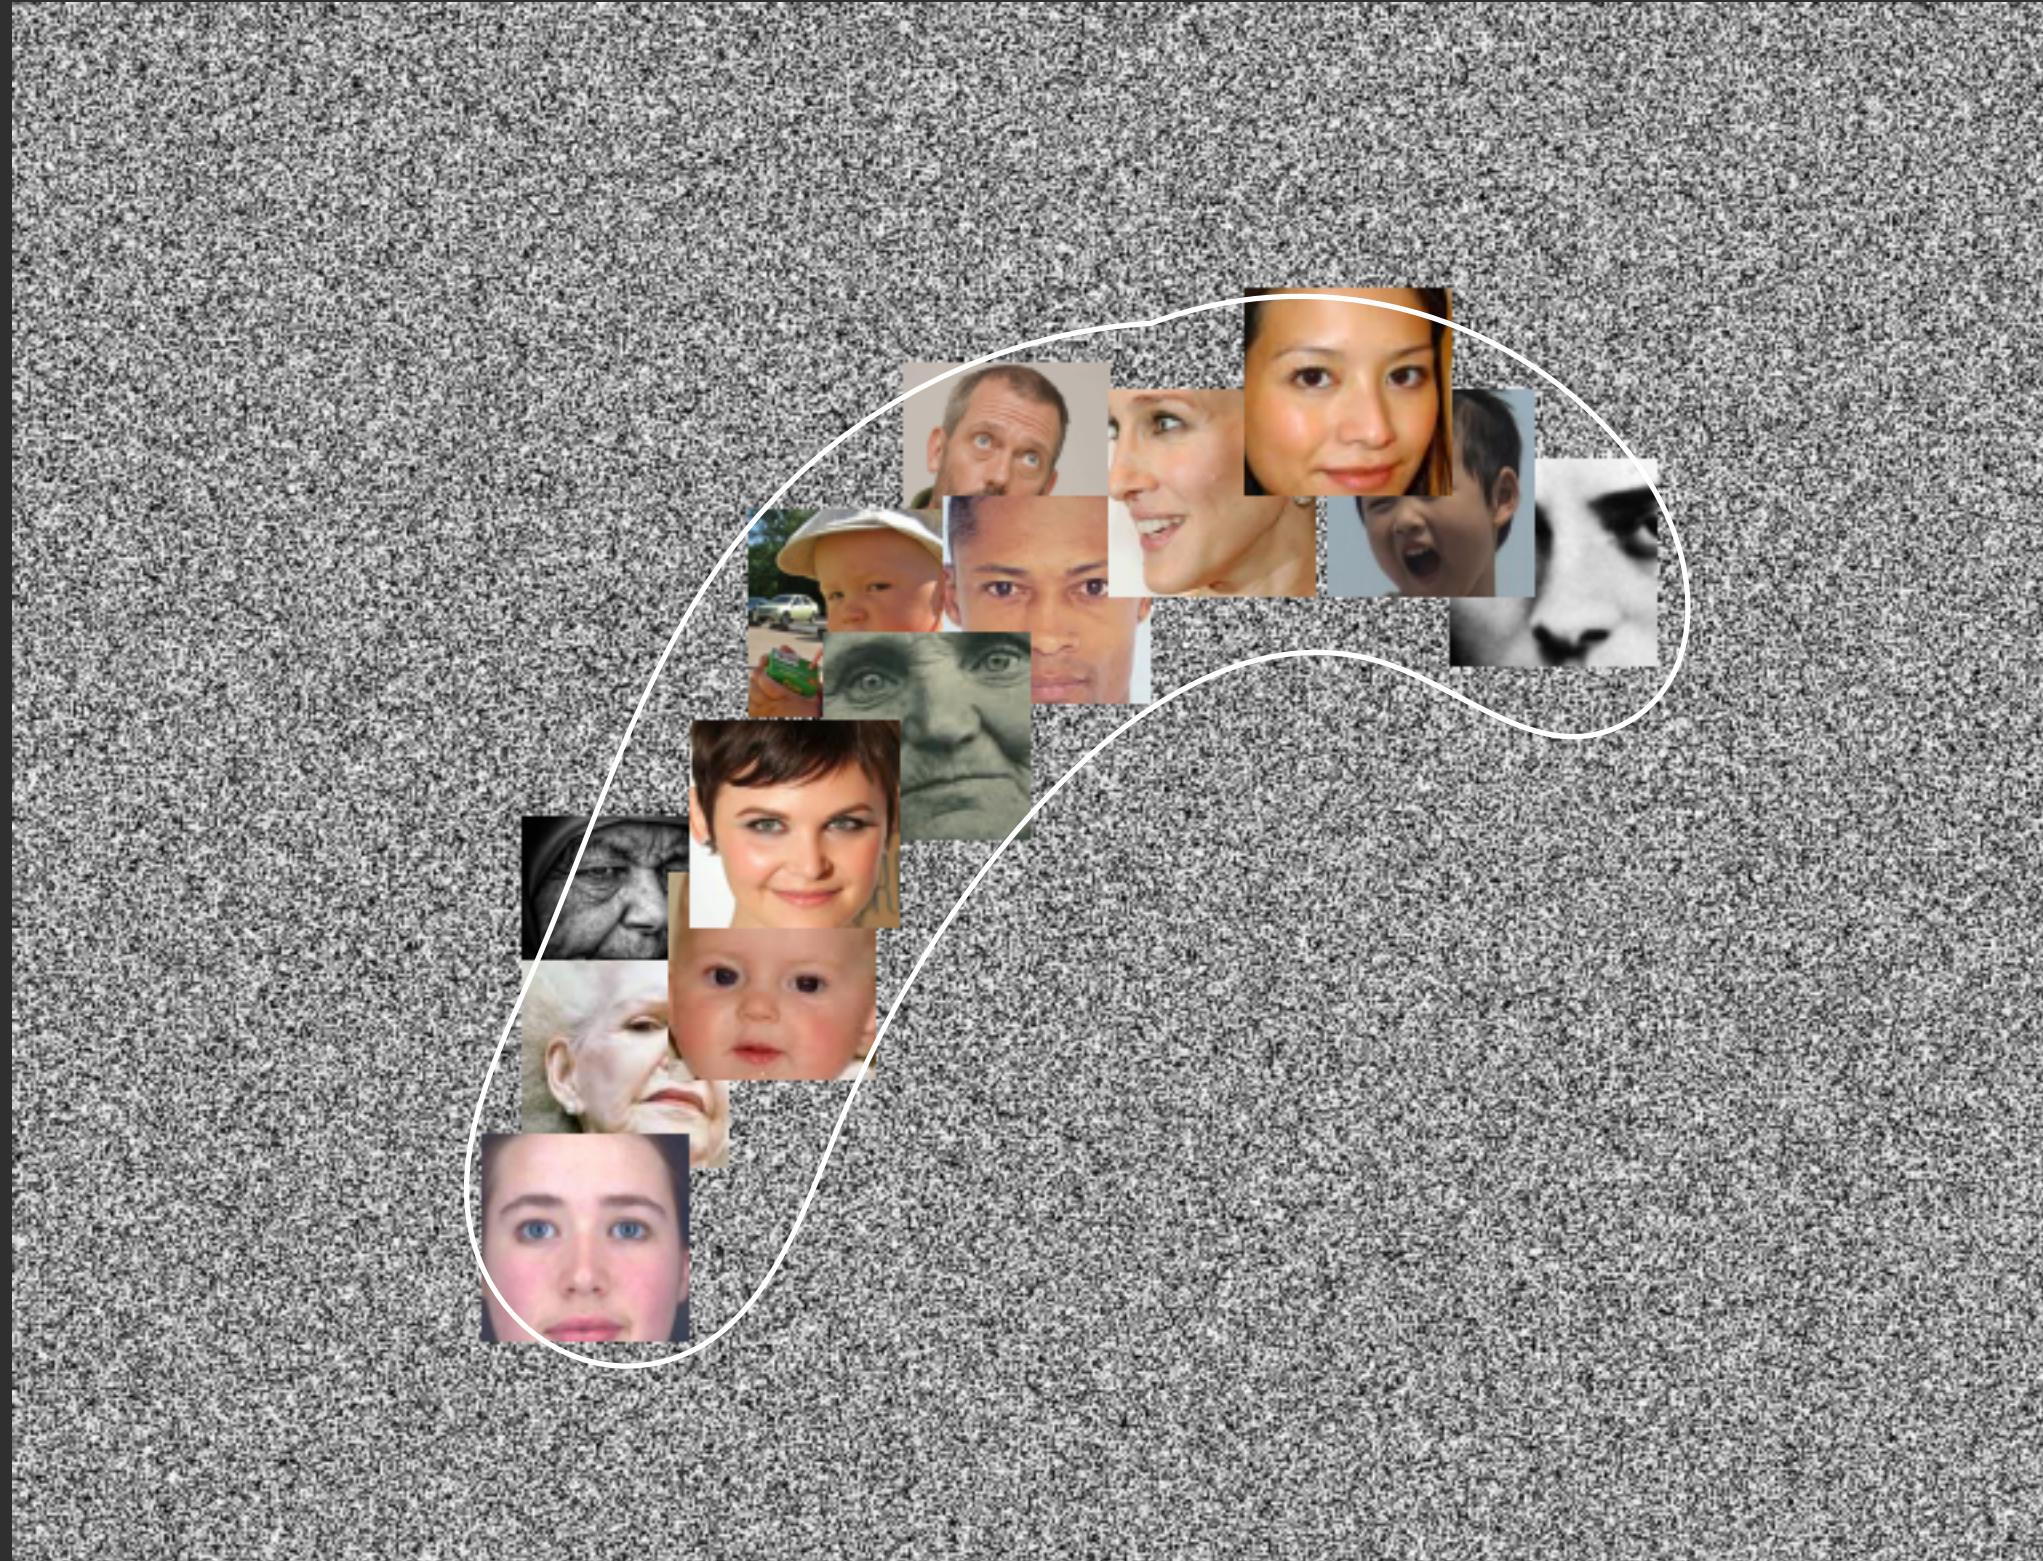
Autoencoder minimises Euclidean distance + tries to fool D

GANs to improve latent representations



Autoencoding beyond pixels using a learned similarity metric (Larsen et al, 2016)

Manifold hypothesis



$$128 \times 128 \times 3 \times 255 = 12.5M$$

Real data lies on a low dimensional manifold in high dimensional space

<https://www.iro.umontreal.ca/~bengioy/talks/gss2012-YB3-algorithms-AE-depth.pdf>

Manifold hypothesis

Z

cGANs with projection discriminator (Miyato & Koyama, 2018)

Manifold hypothesis



\mathbf{z}_1

\mathbf{z}

cGANs with projection discriminator (Miyato & Koyama, 2018)

Manifold hypothesis



\mathbf{z}_1



\mathbf{z}_2

\mathbf{Z}

cGANs with projection discriminator (Miyato & Koyama, 2018)

Manifold hypothesis



\mathbf{z}_1

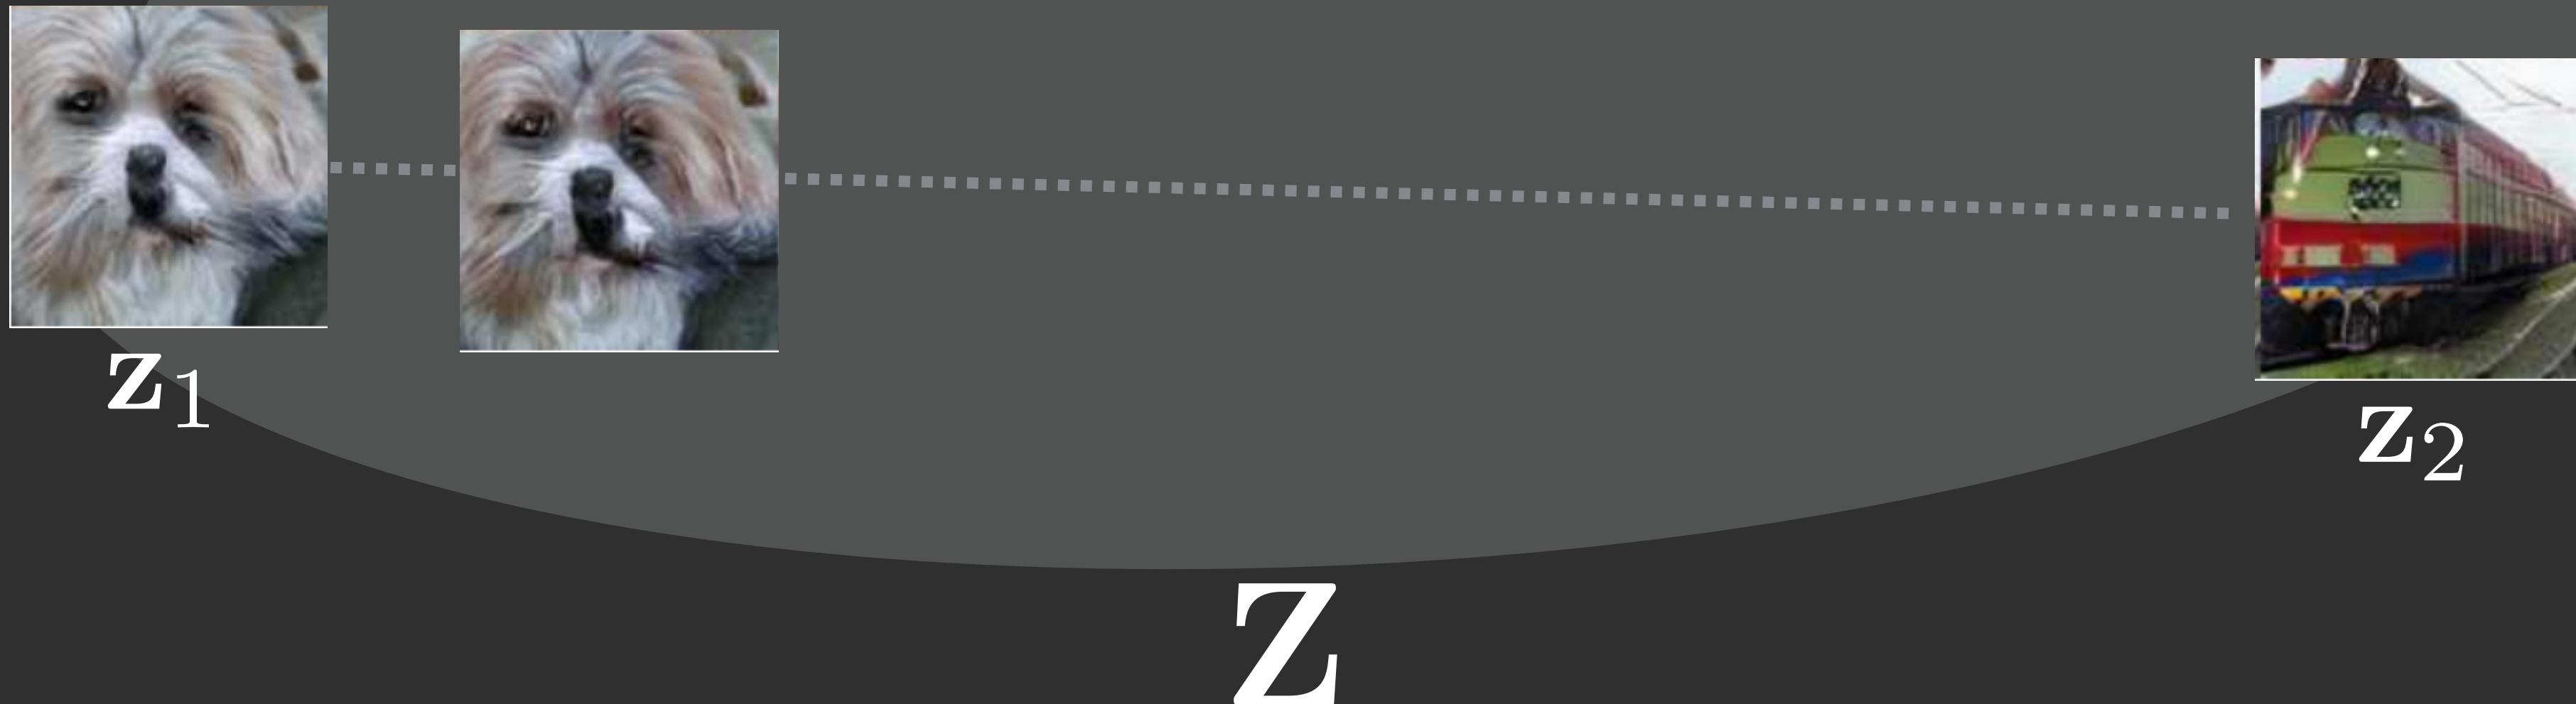


\mathbf{z}_2

\mathbf{Z}

cGANs with projection discriminator (Miyato & Koyama, 2018)

Manifold hypothesis



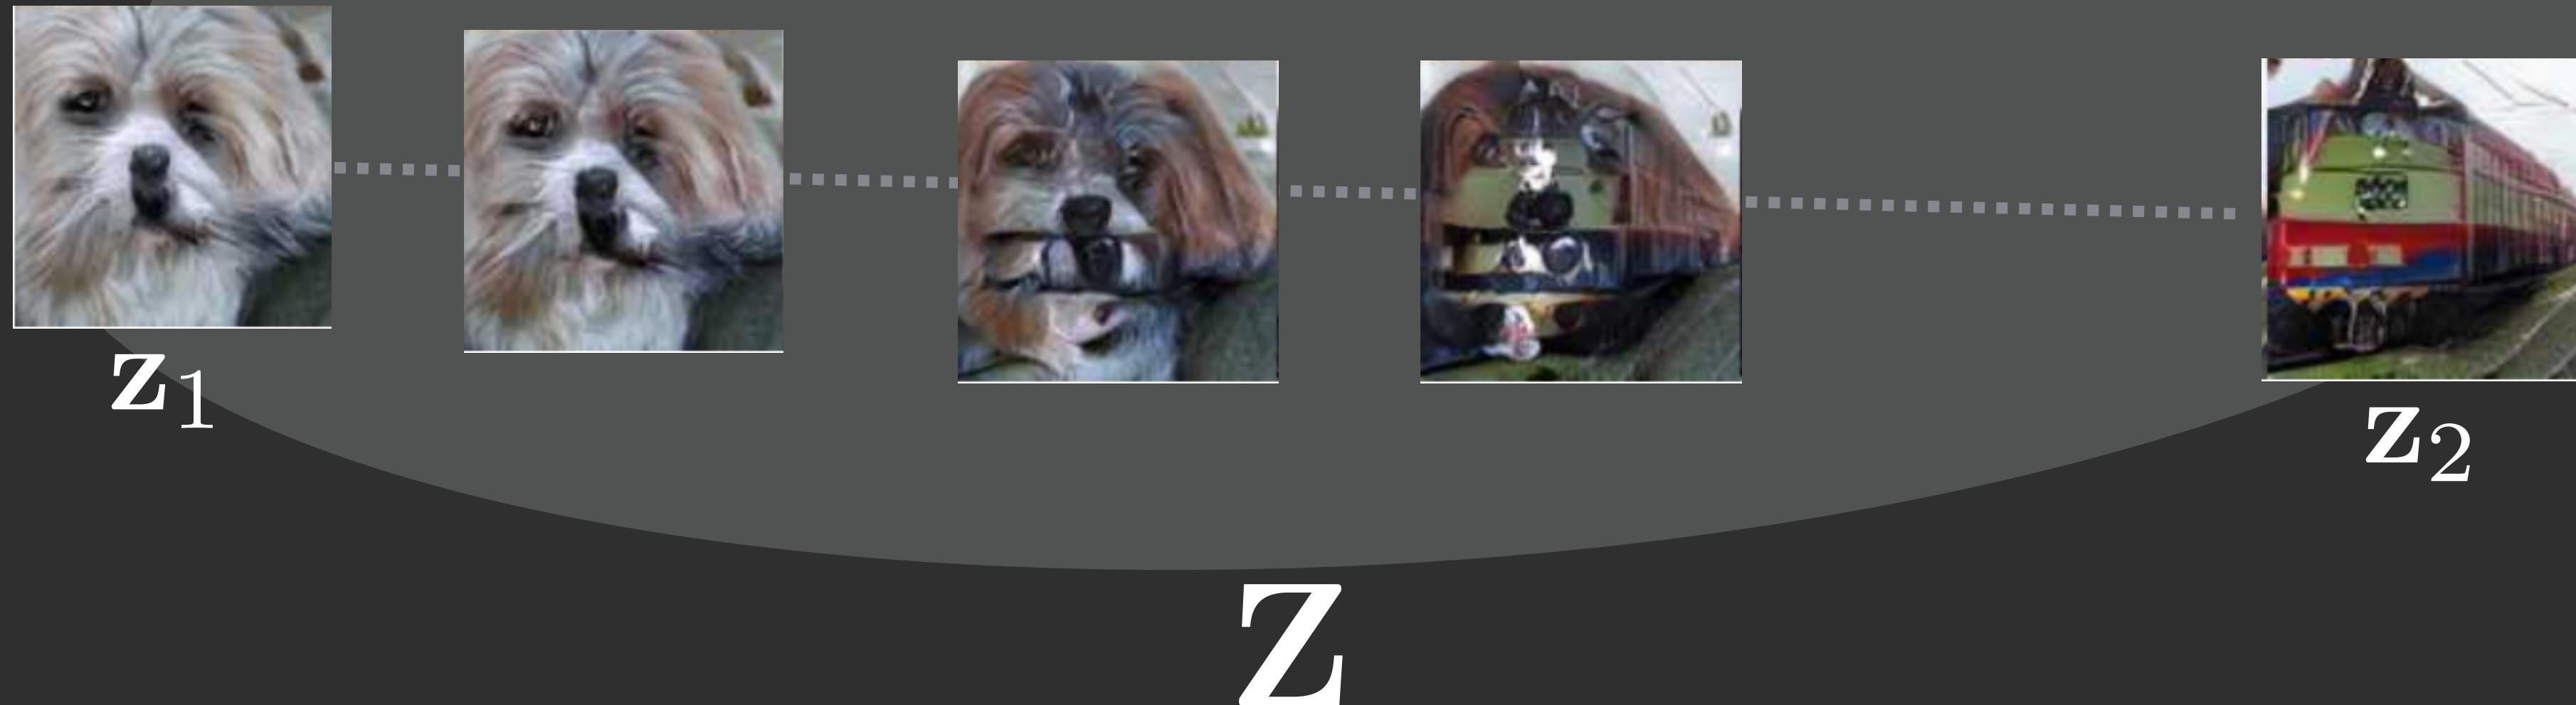
cGANs with projection discriminator (Miyato & Koyama, 2018)

Manifold hypothesis



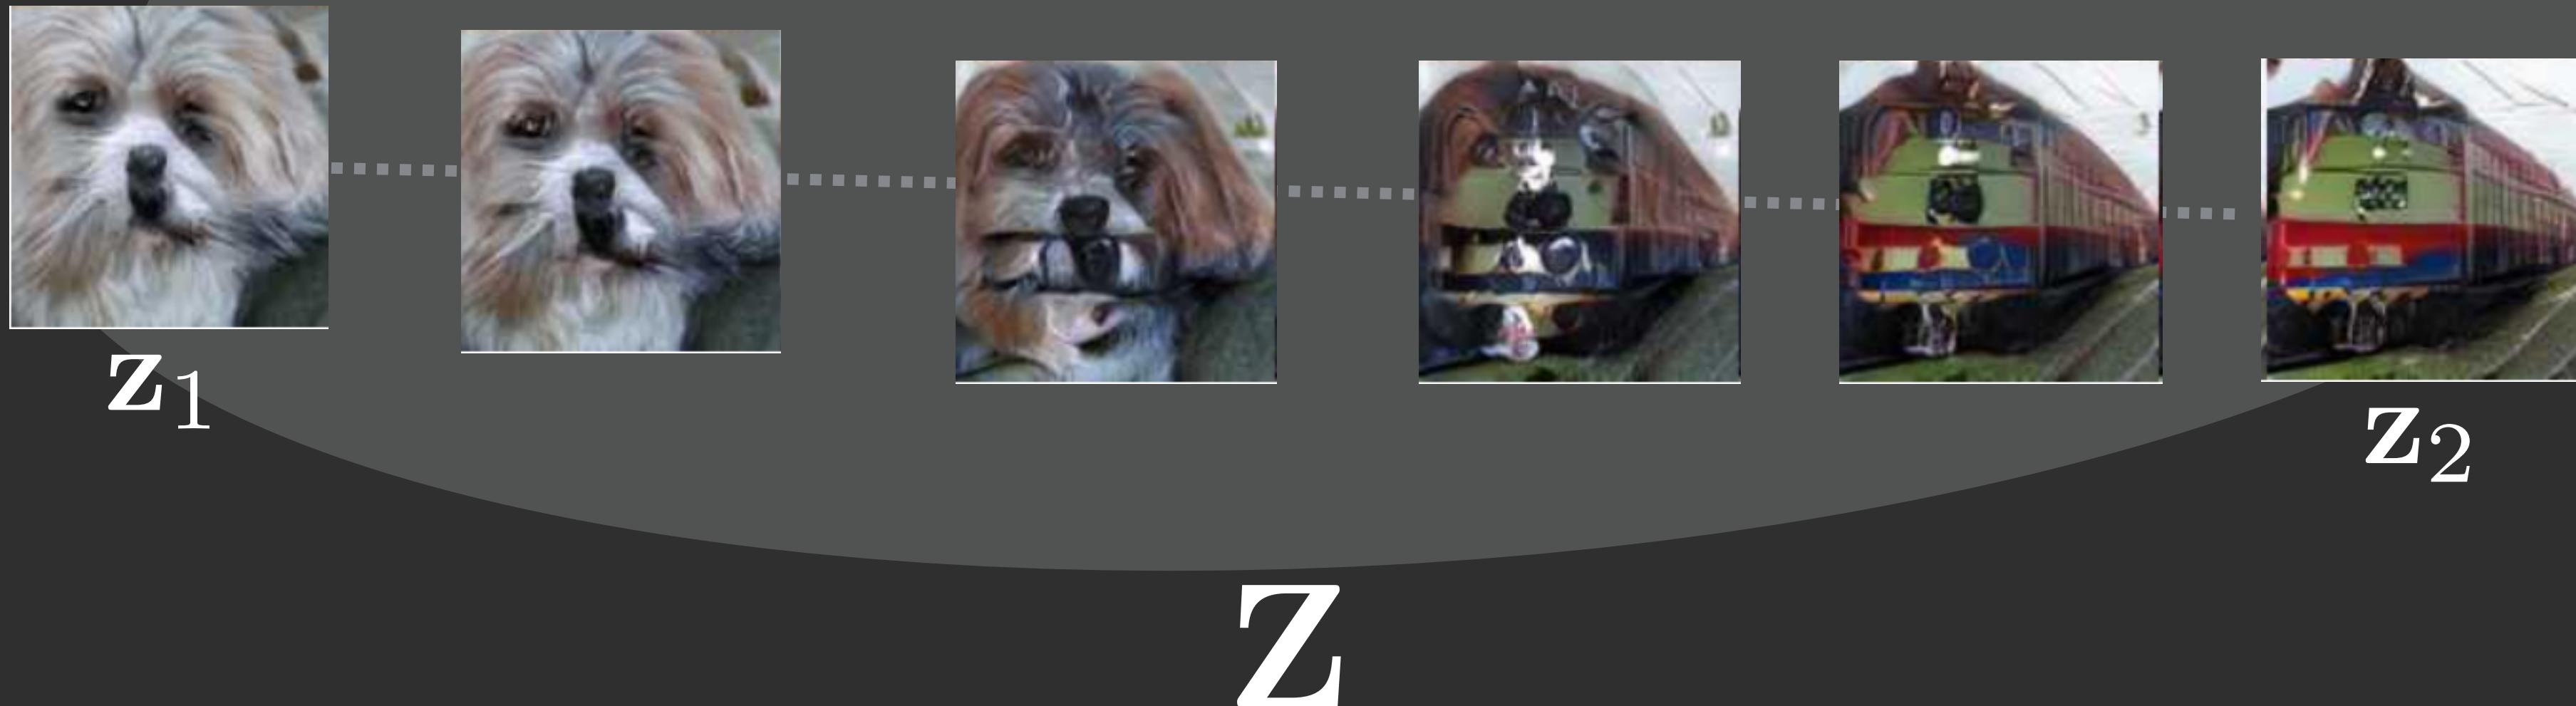
cGANs with projection discriminator (Miyato & Koyama, 2018)

Manifold hypothesis



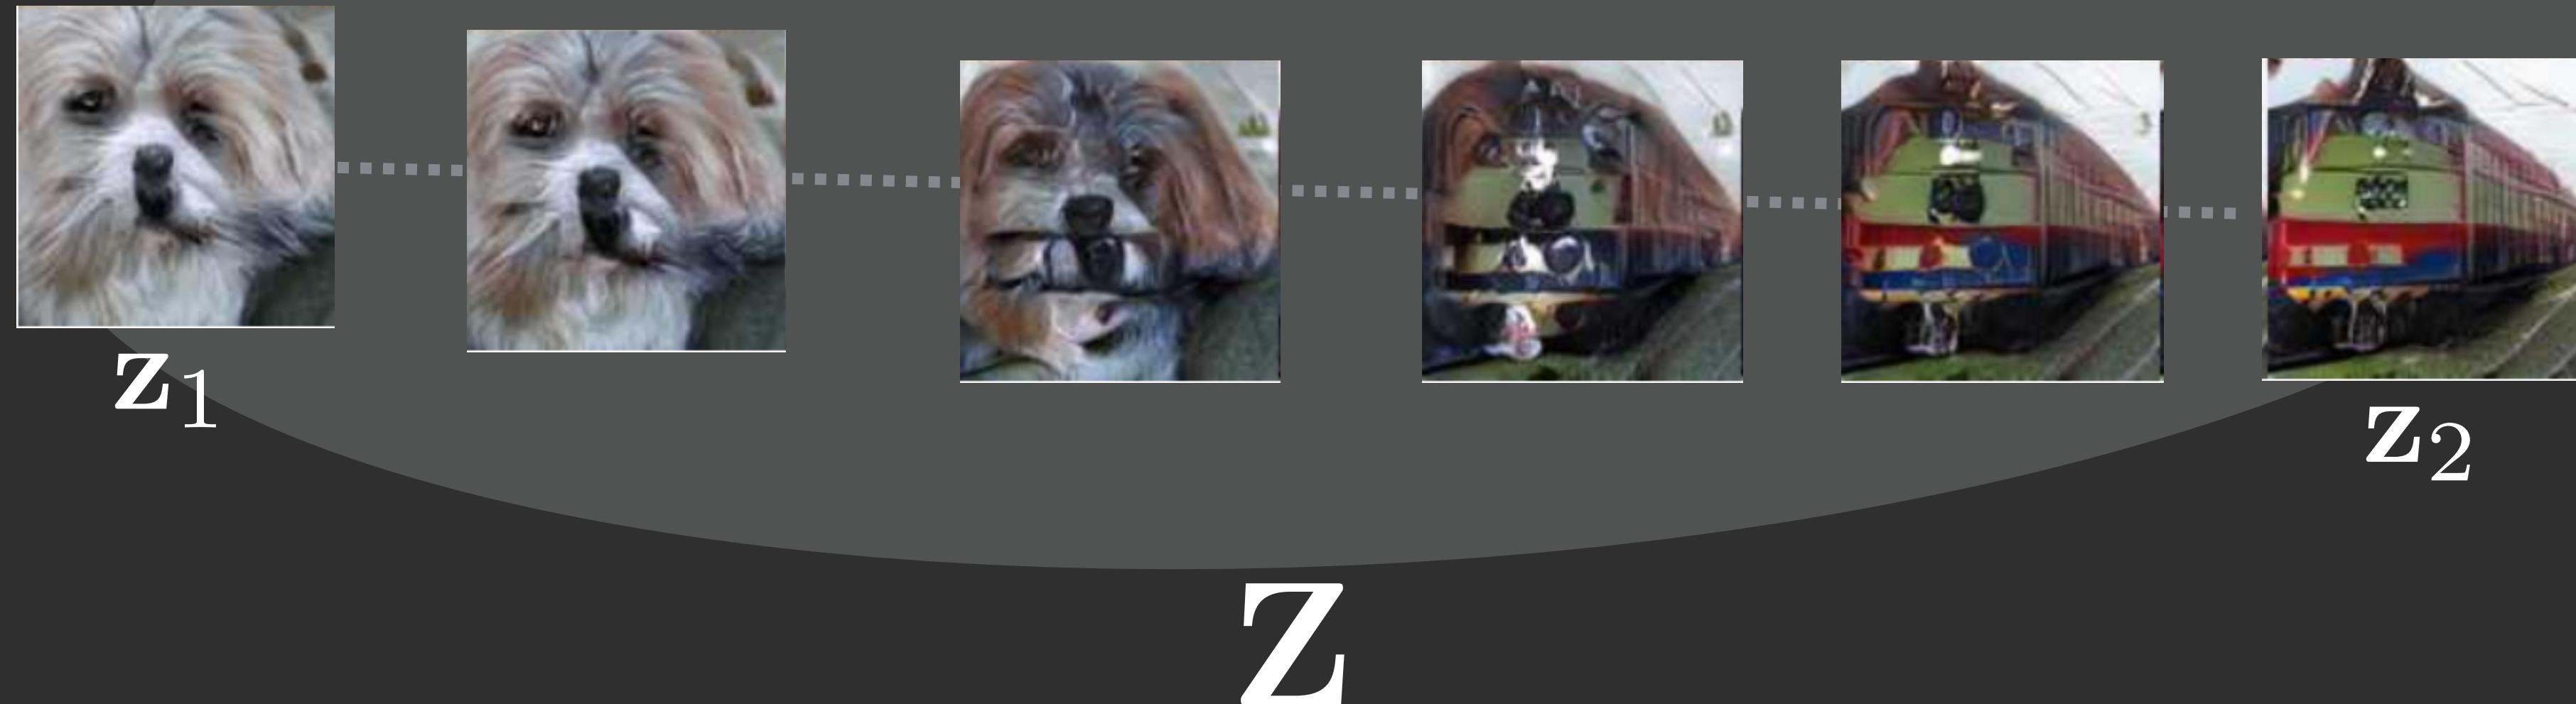
cGANs with projection discriminator (Miyato & Koyama, 2018)

Manifold hypothesis



cGANs with projection discriminator (Miyato & Koyama, 2018)

Manifold hypothesis



cGANs with projection discriminator (Miyato & Koyama, 2018)

Space of the prior becomes semantically meaningful

Thank you!

